

EECS151 : Introduction to Digital Design and ICs

Lecture 2 – Design Process

Bora Nikolić



At HotChips'19 Cerebras announced the largest chip in the world at 8.5 in x 8.5in with 1.2 trillion transistors, and 15kW of power, aimed for training of deep-learning neural networks

At HotChips'21 they showed the next version in 7nm CMOS, with >2x transistor count

- 46,225 mm² silicon
- 2.6 Trillion transistors
- 850,000 AI optimized cores
- 40 Gigabytes on-chip memory
- 20 Petabyte/s memory bandwidth
- 220 Petabit/s fabric bandwidth
- 7nm Process technology at TSMC




Sean Lie,
HotChips'21


Review

- Moore's law is slowing down
 - There are continued improvements in technology, but at a slower pace
- Dennard's scaling has ended a decade ago
 - All designs are now power limited
- Specialization and customization provides added performance
 - Under power constraints and stagnant technology
- Design costs are high
 - Methodology and better reuse to rescue!
 - Abstraction, modularity, regularity are the keys
 - And creativity!

Putting it in Perspective



Performance gains
over the past decade




Lisa Su, HotChips'19 keynote



Digital Logic

Implementing Digital Systems


- Digital systems implement a set of Boolean equations



- How do we implement a digital system?

Modern (Mostly) Digital System-On-A-Chip

- 4x 'Firestorm' Large CPUs
- 4x 'Icestorm' Small CPUs
- GPU
- Neural processing unit (NPU)
- Lots of memory
- DDR memory interfaces



TechInsights

- 5nm CMOS
- Up to 2.49GHz



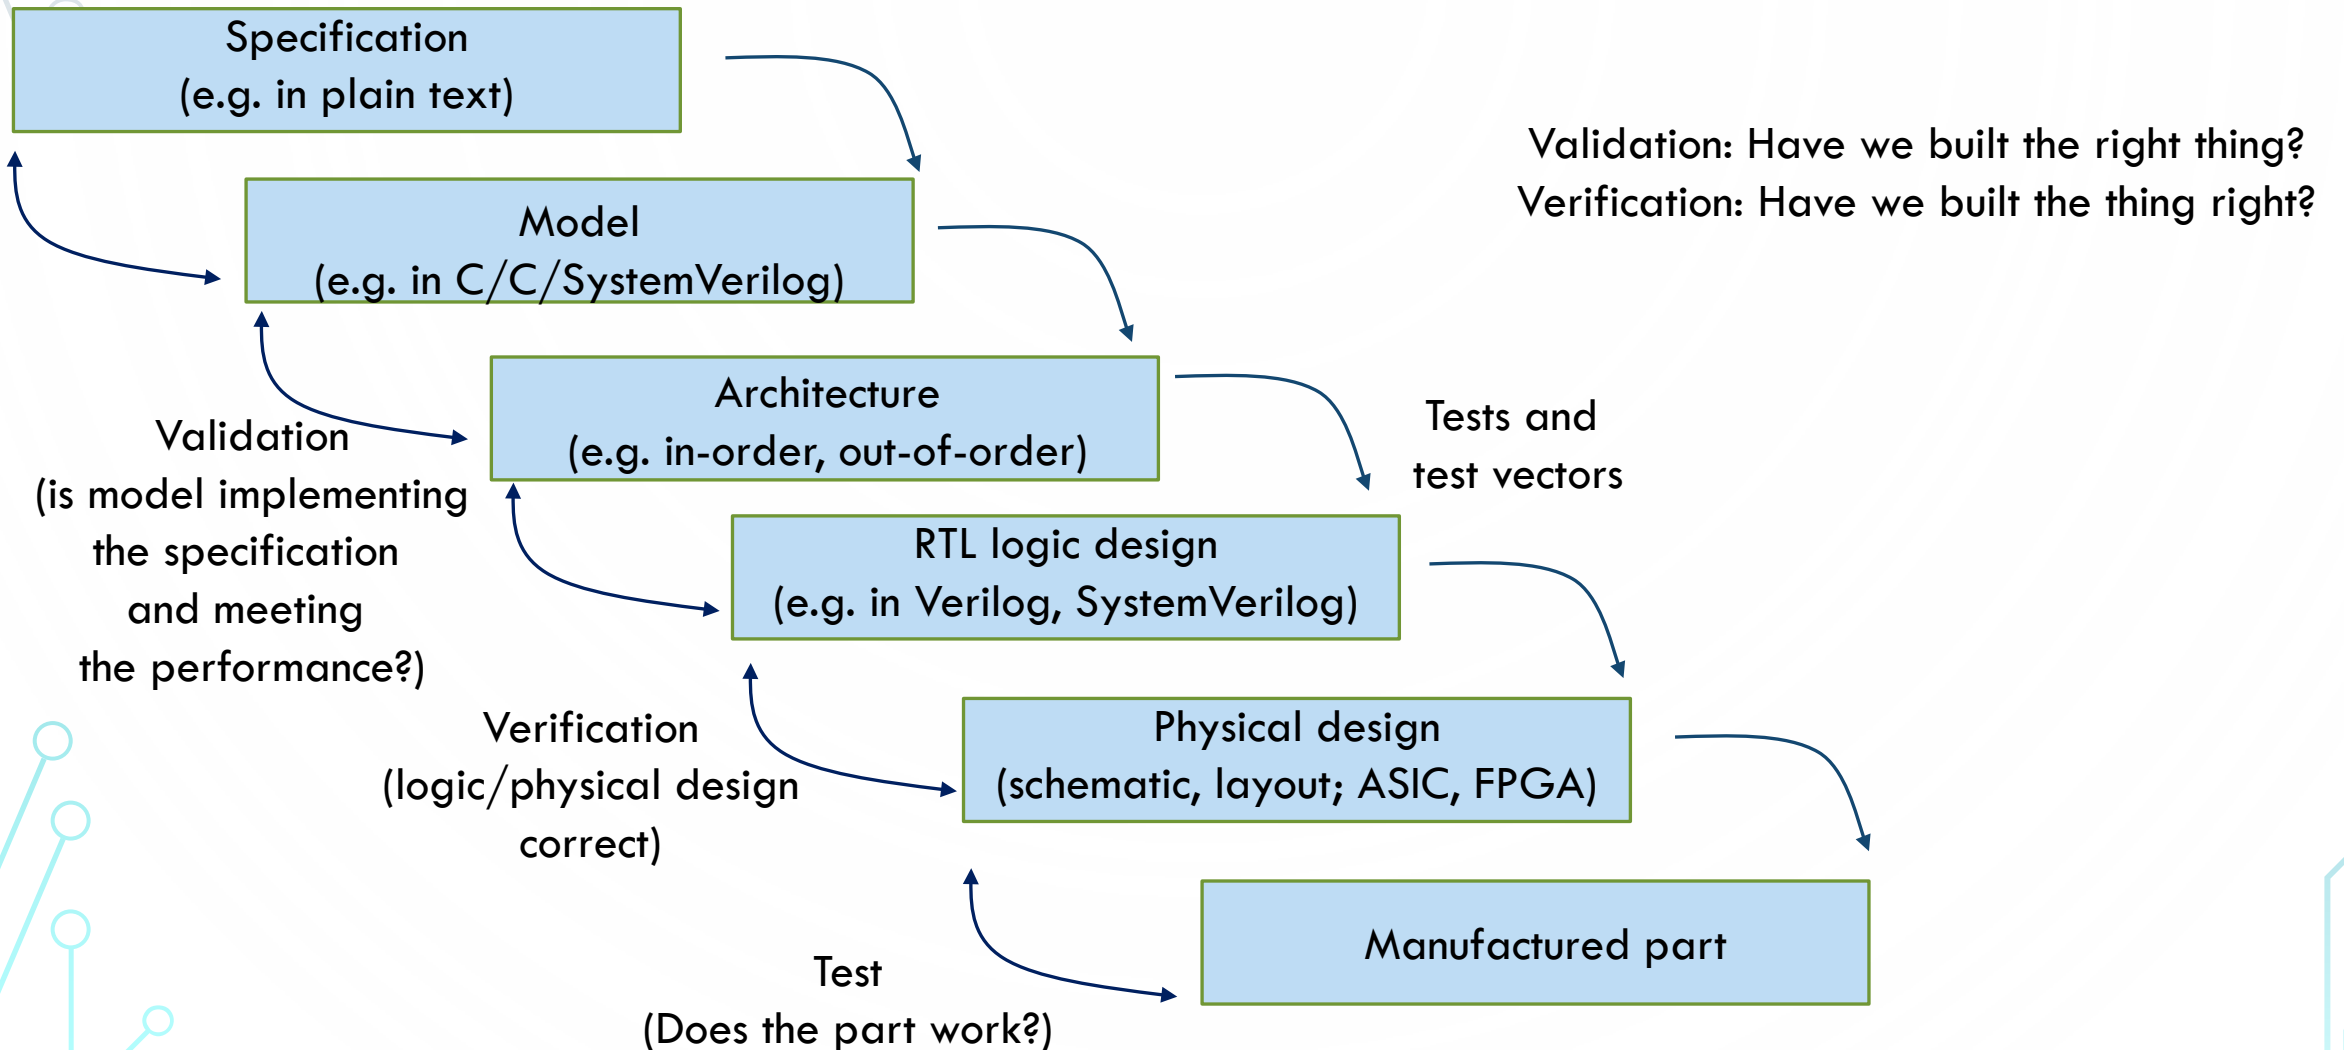
By Henriok

<https://commons.wikimedia.org/w/index.php?curid=96026688>



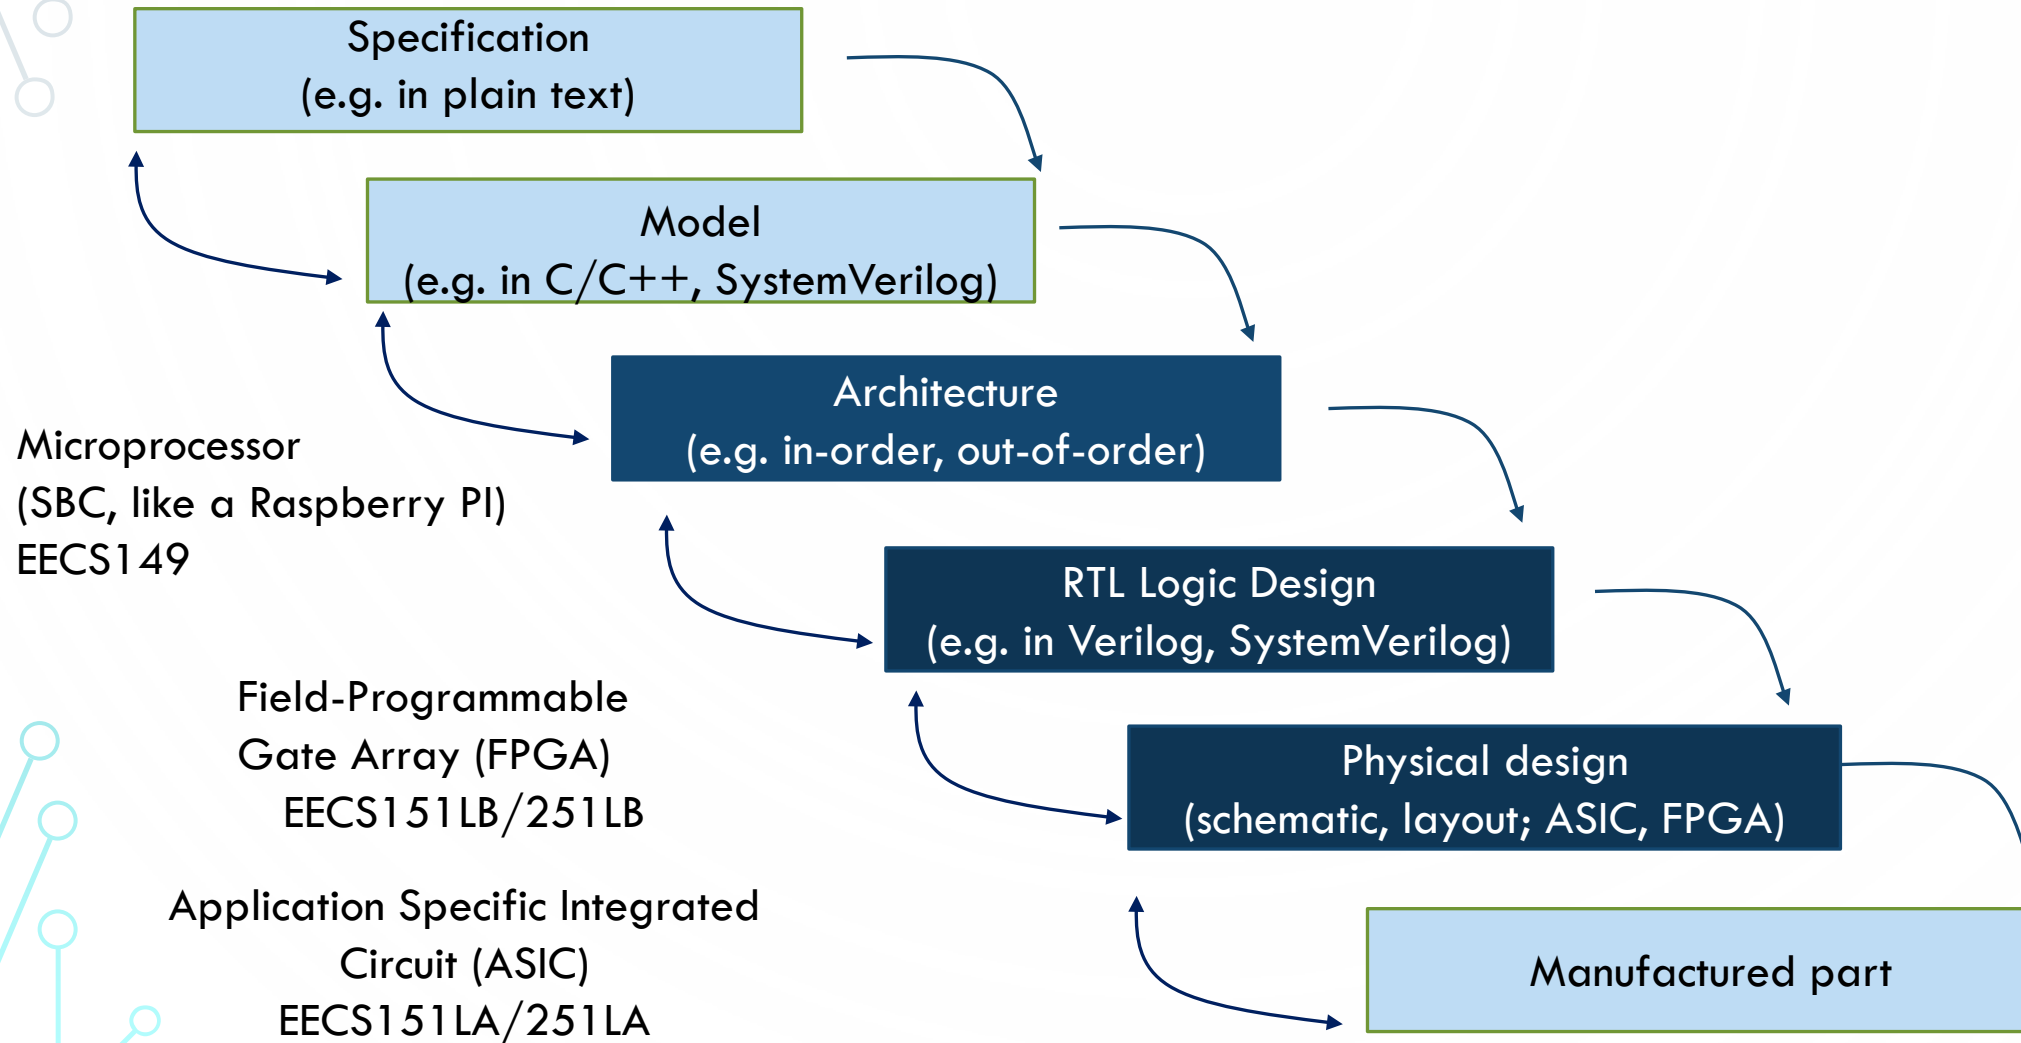
Design Process

- Design through layers of abstractions



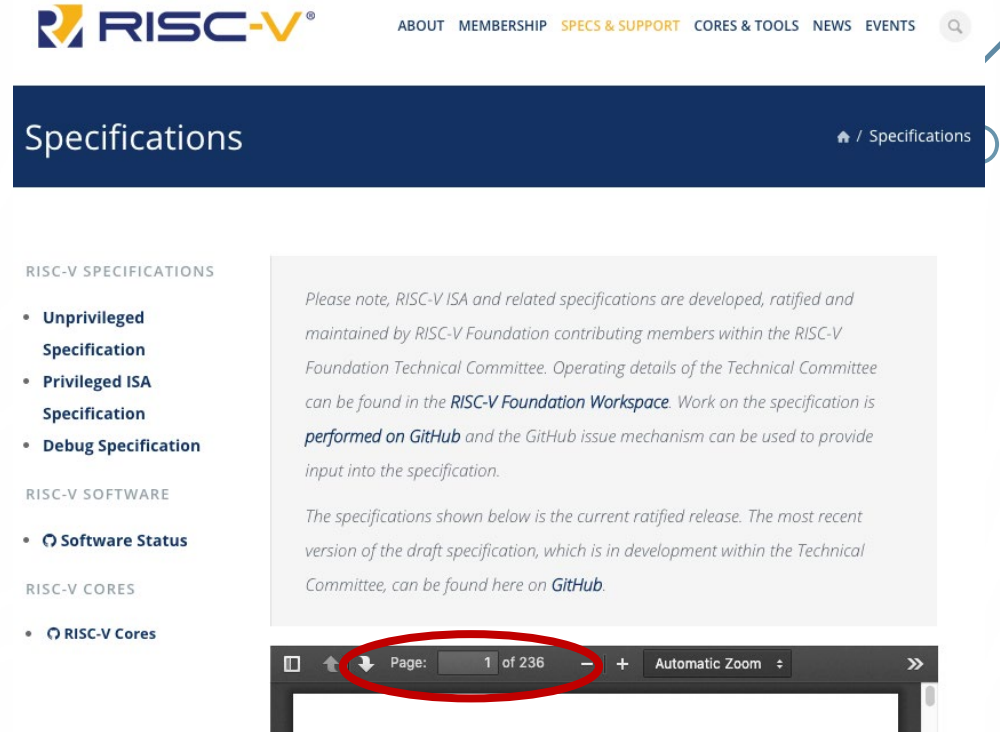
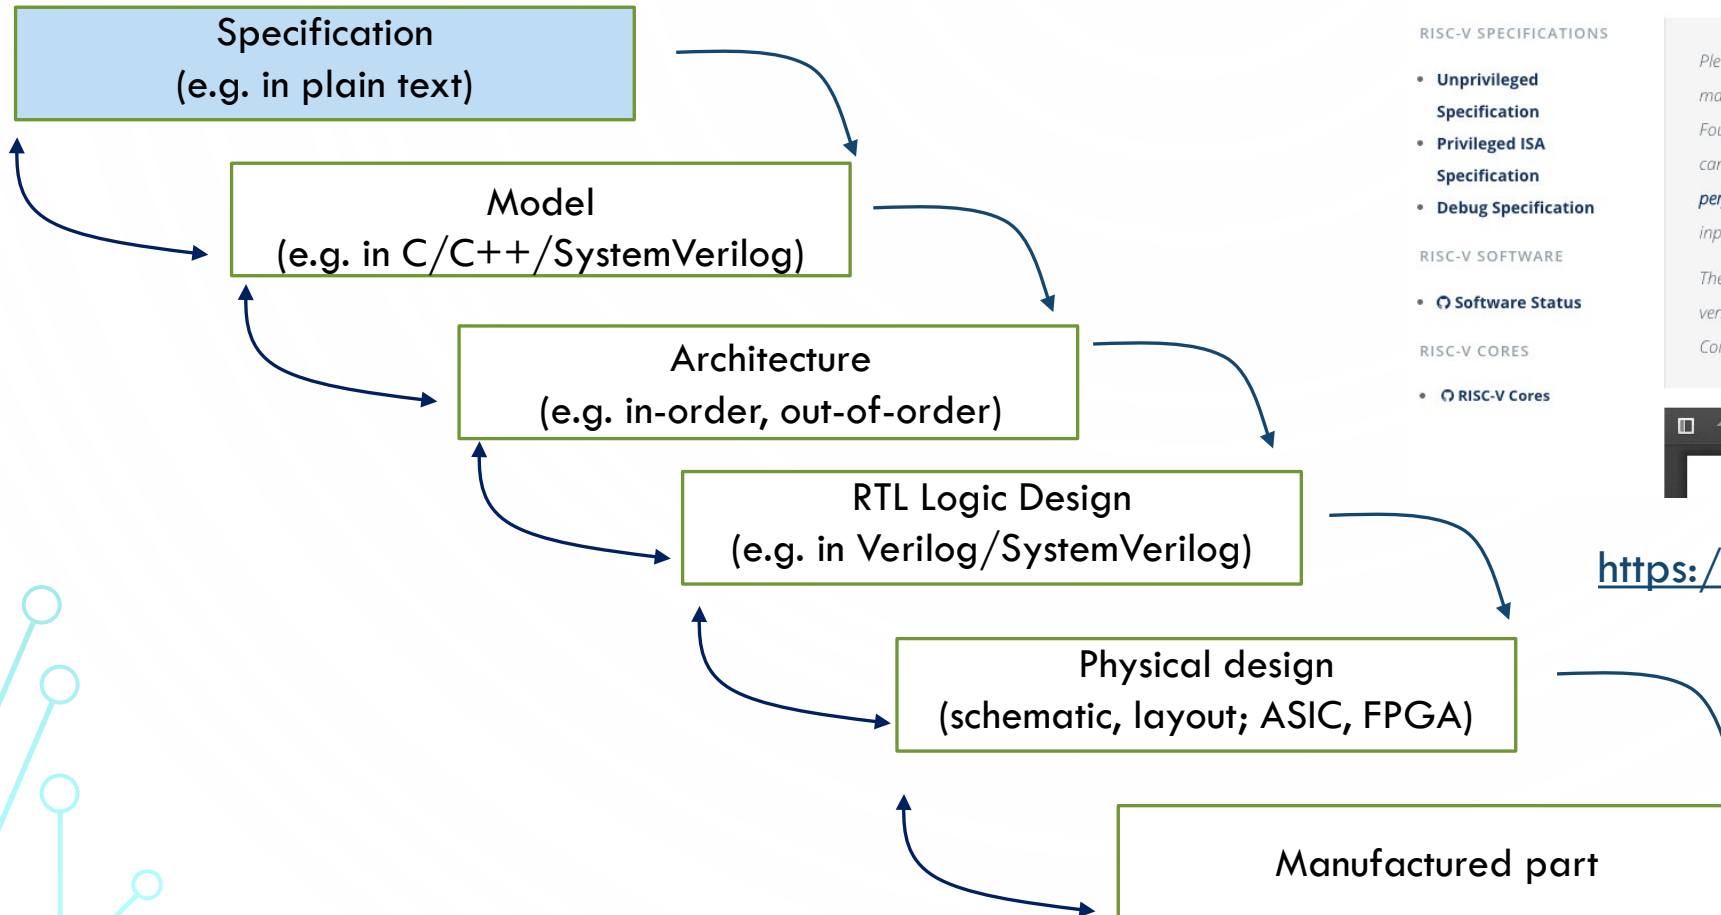
Design Abstractions in EECS151/251A

- Design through layers of abstractions



Example: RISC-V Design Process

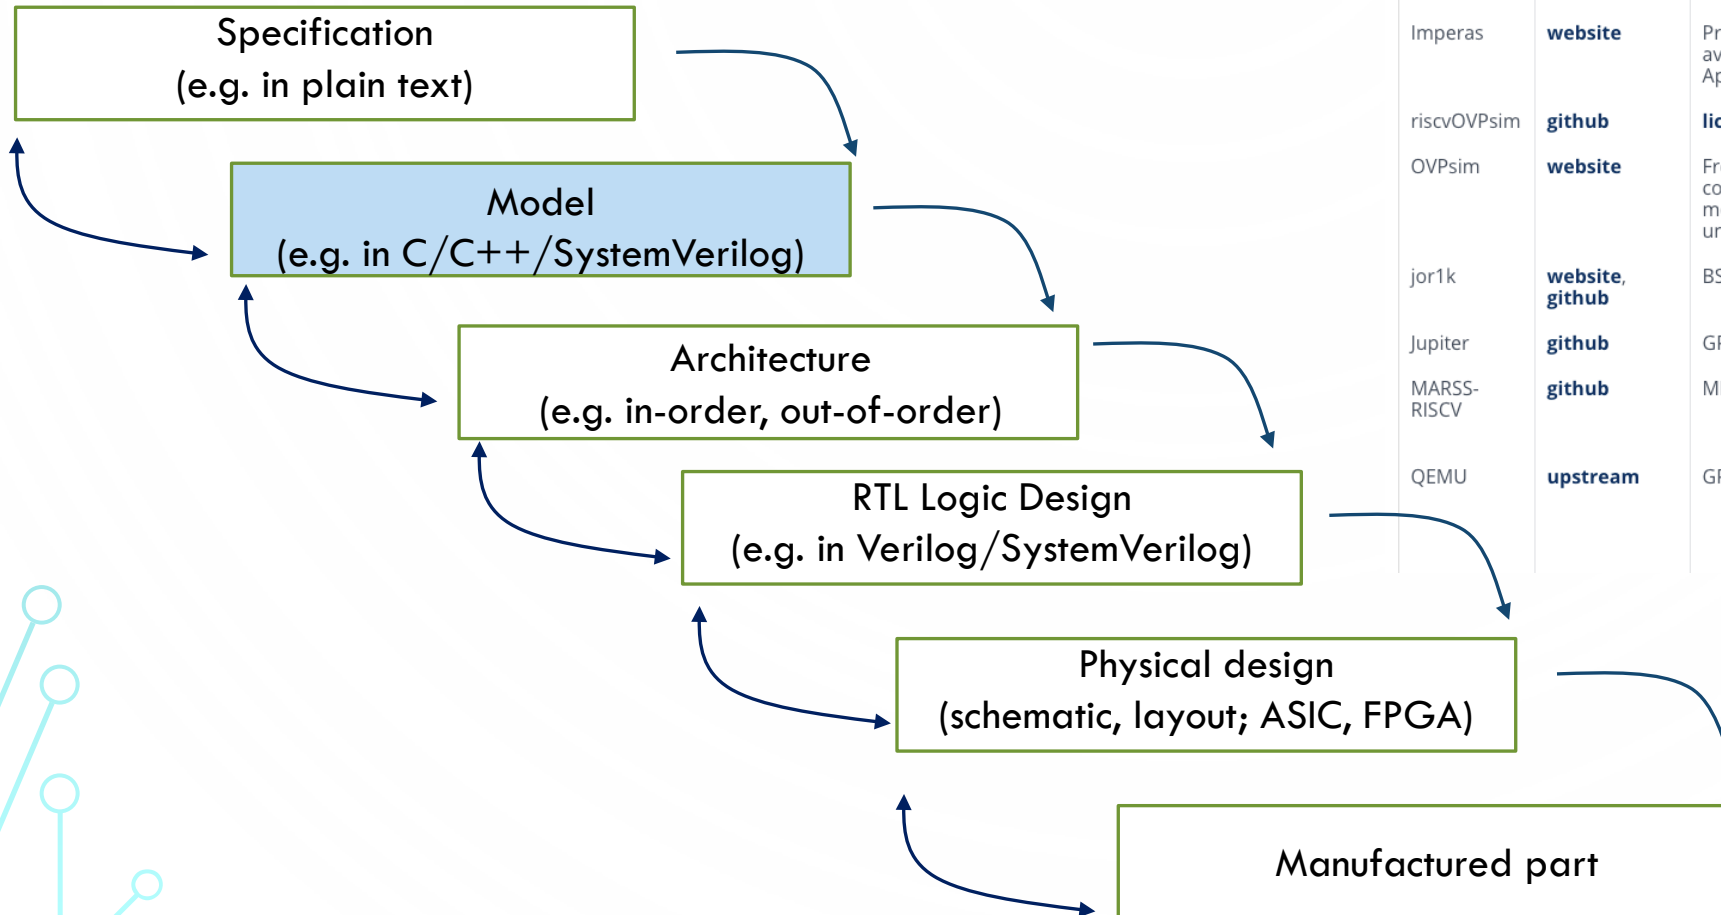
- Design through layers of abstractions



<https://riscv.org/specifications/>

Example: RISC-V Design Process

- Design through layers of abstractions



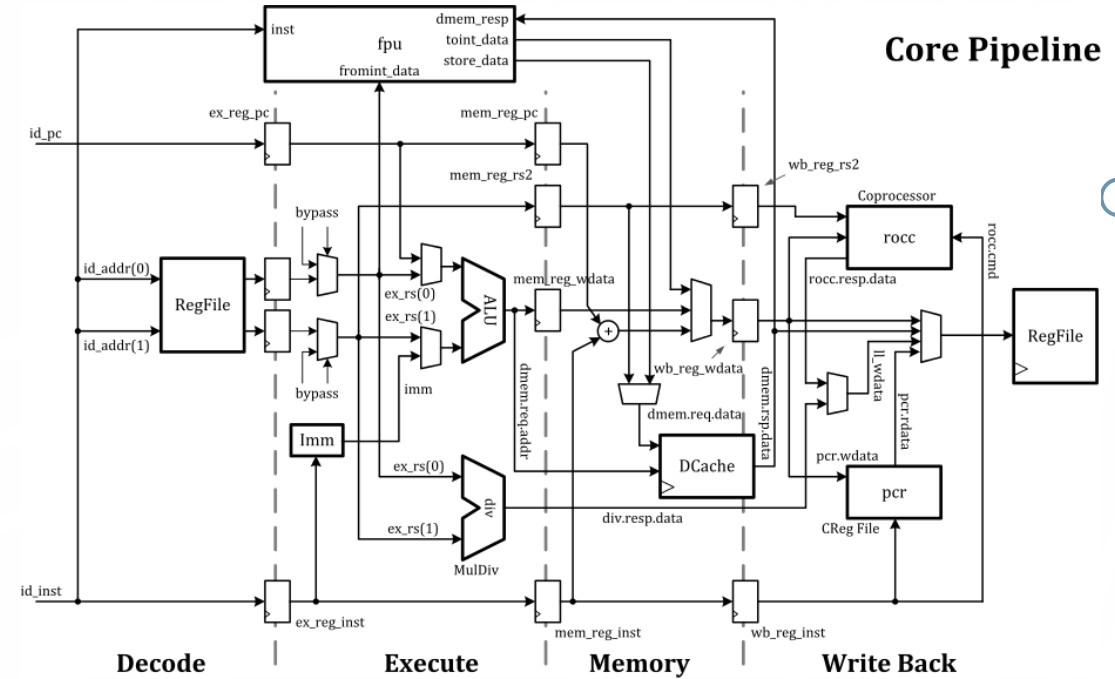
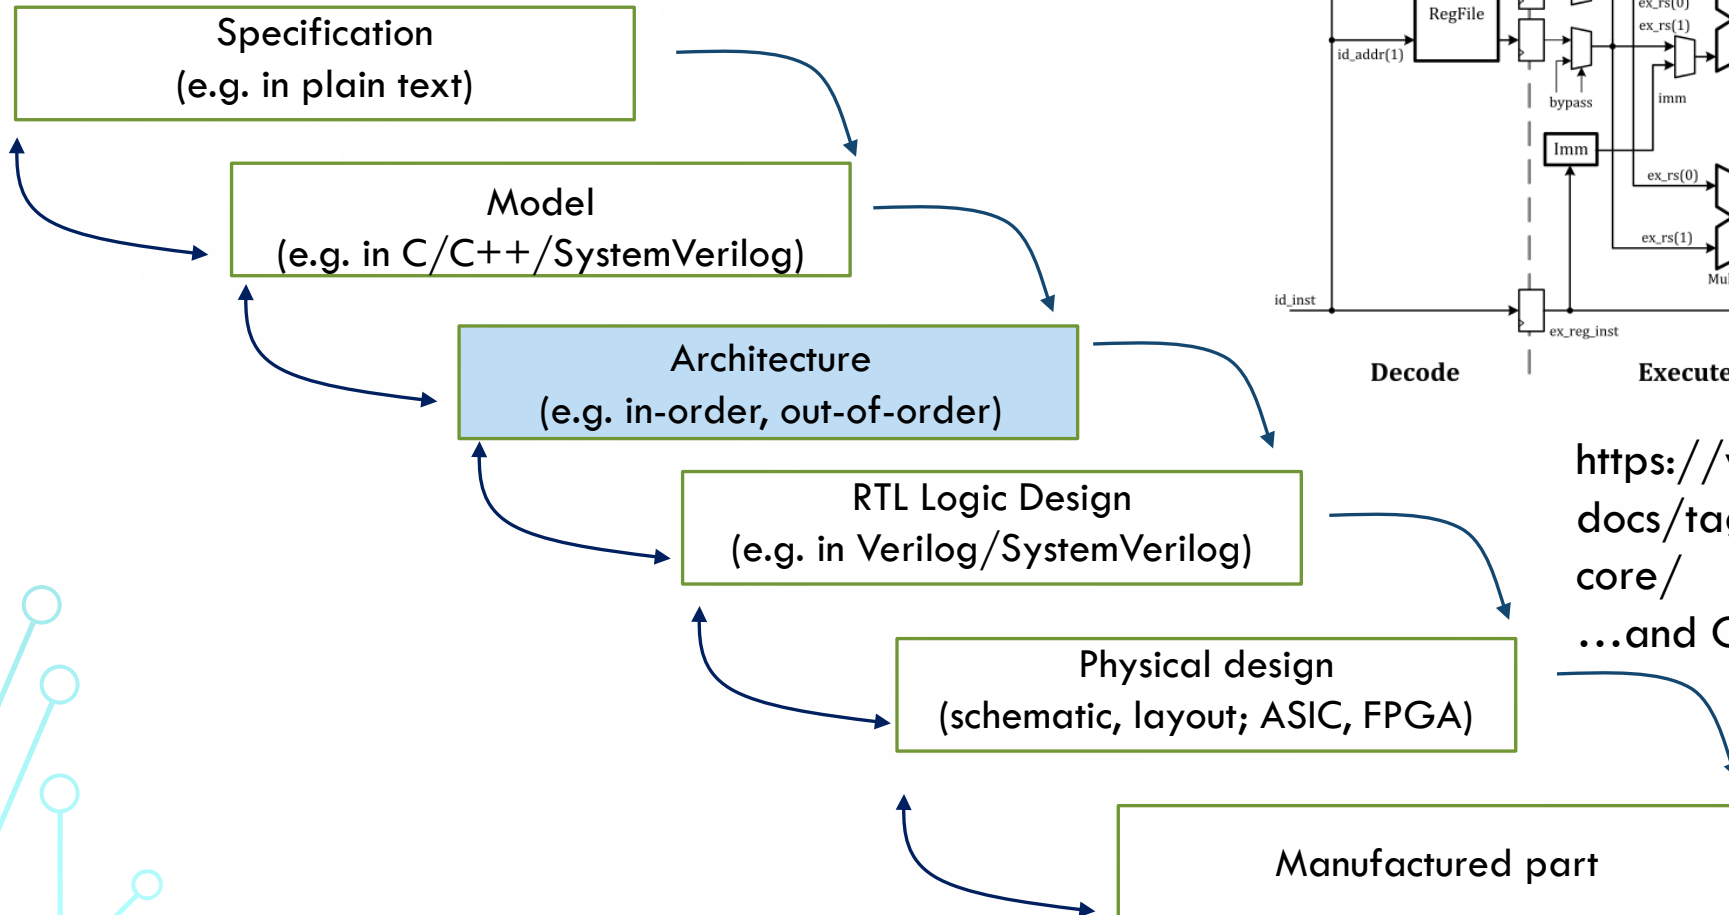
Simulators

| Name | Links | License | Maintainers |
|----------------|---|--|---|
| DBT-RISE-RISCV | github | BSD-3-Clause | MINRES Technologies |
| FireSim | website , mailing list , github , ISCA 2018 Paper | BSD | Sagar Karandikar, Howard Mao, Donggyu Kim, David Biancolin, Alon Amid, Berkeley Architecture Research |
| gem5 | SW-dev thread , repository | BSD-style | Alec Roelke (University of Virginia) |
| Imperas | website | Proprietary, models available under Apache 2.0 | Imperas |
| riscvOVPsim | github | license | Imperas |
| OVPsim | website | Free for non commercial use, models available under Apache 2.0 | Imperas |
| jor1k | website , github | BSD 2-Clause | Sebastian Macke |
| Jupiter | github | GPL-3.0 | Andrés Castellanos |
| MARSS-RISCV | github | MIT | Gaurav N Kothari, Parikshit P Sarnaik, Gokturk Yuksek (State University of New York at Binghamton) |
| QEMU | upstream | GPL | Sagar Karandikar (University of California, Berkeley), Bastian Koppelman (University of Paderborn), Alex Suykov, Stefan O'Rear and Michael Clark (SiFive) |

<https://riscv.org/software-status/#simulators>

Example: RISC-V Design Process

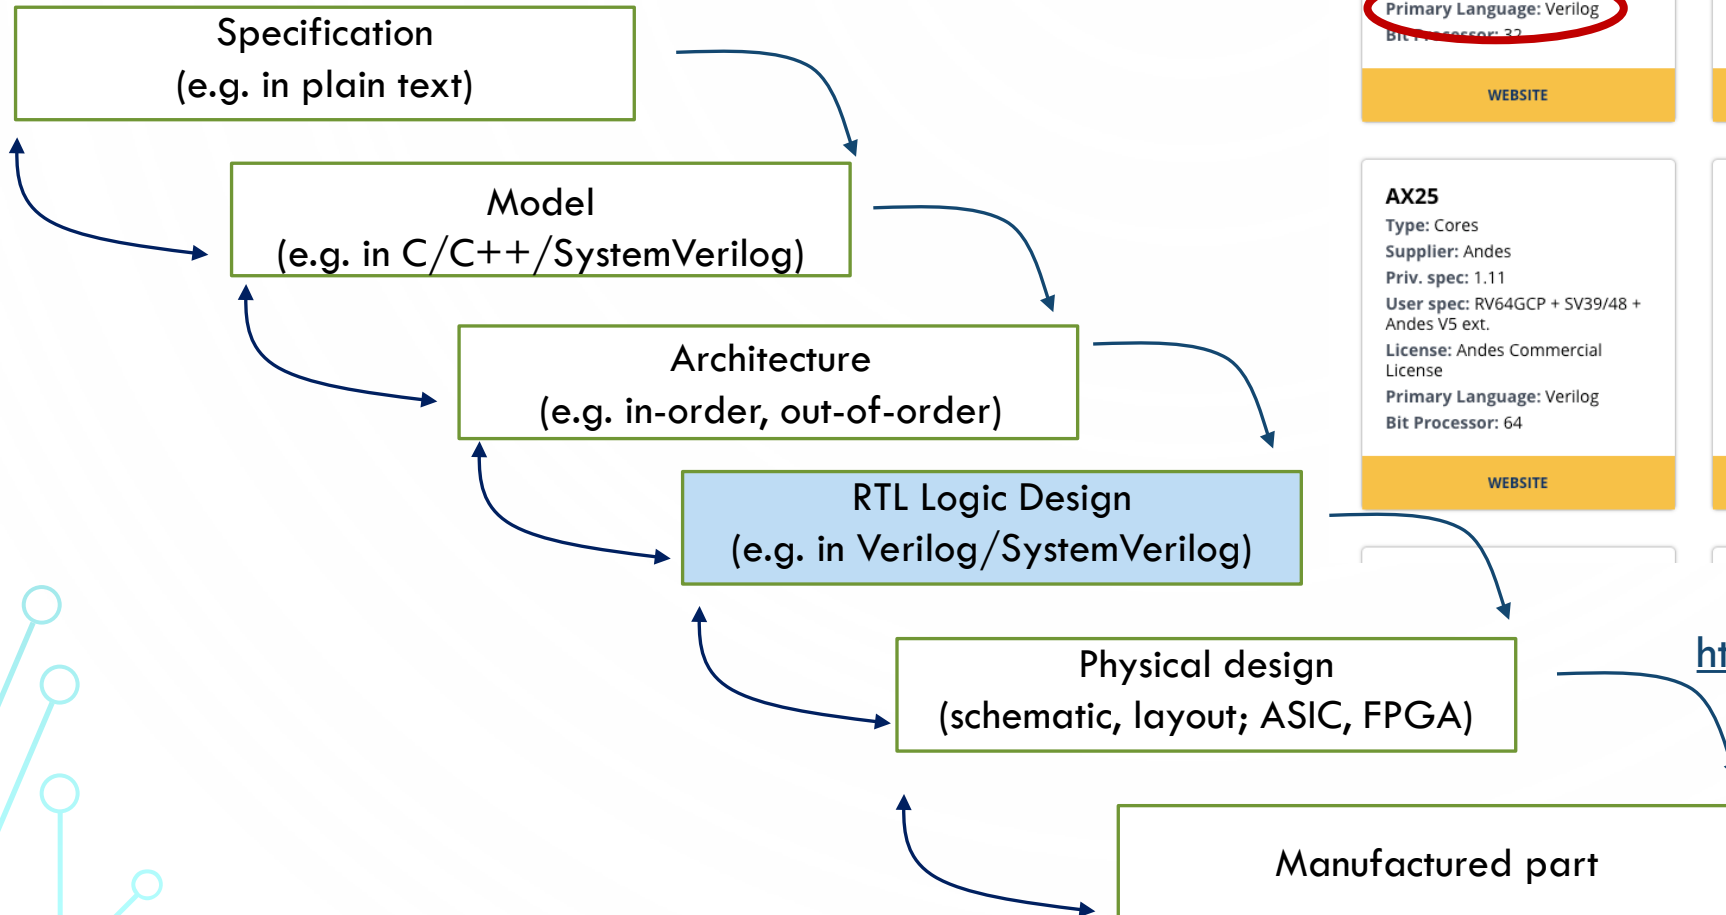
- Design through layers of abstractions



<https://www.lowrisc.org/docs/tagged-memory-v0.1/rocket-core/>
...and CS152

Example: RISC-V Design Process

- Design through layers of abstractions


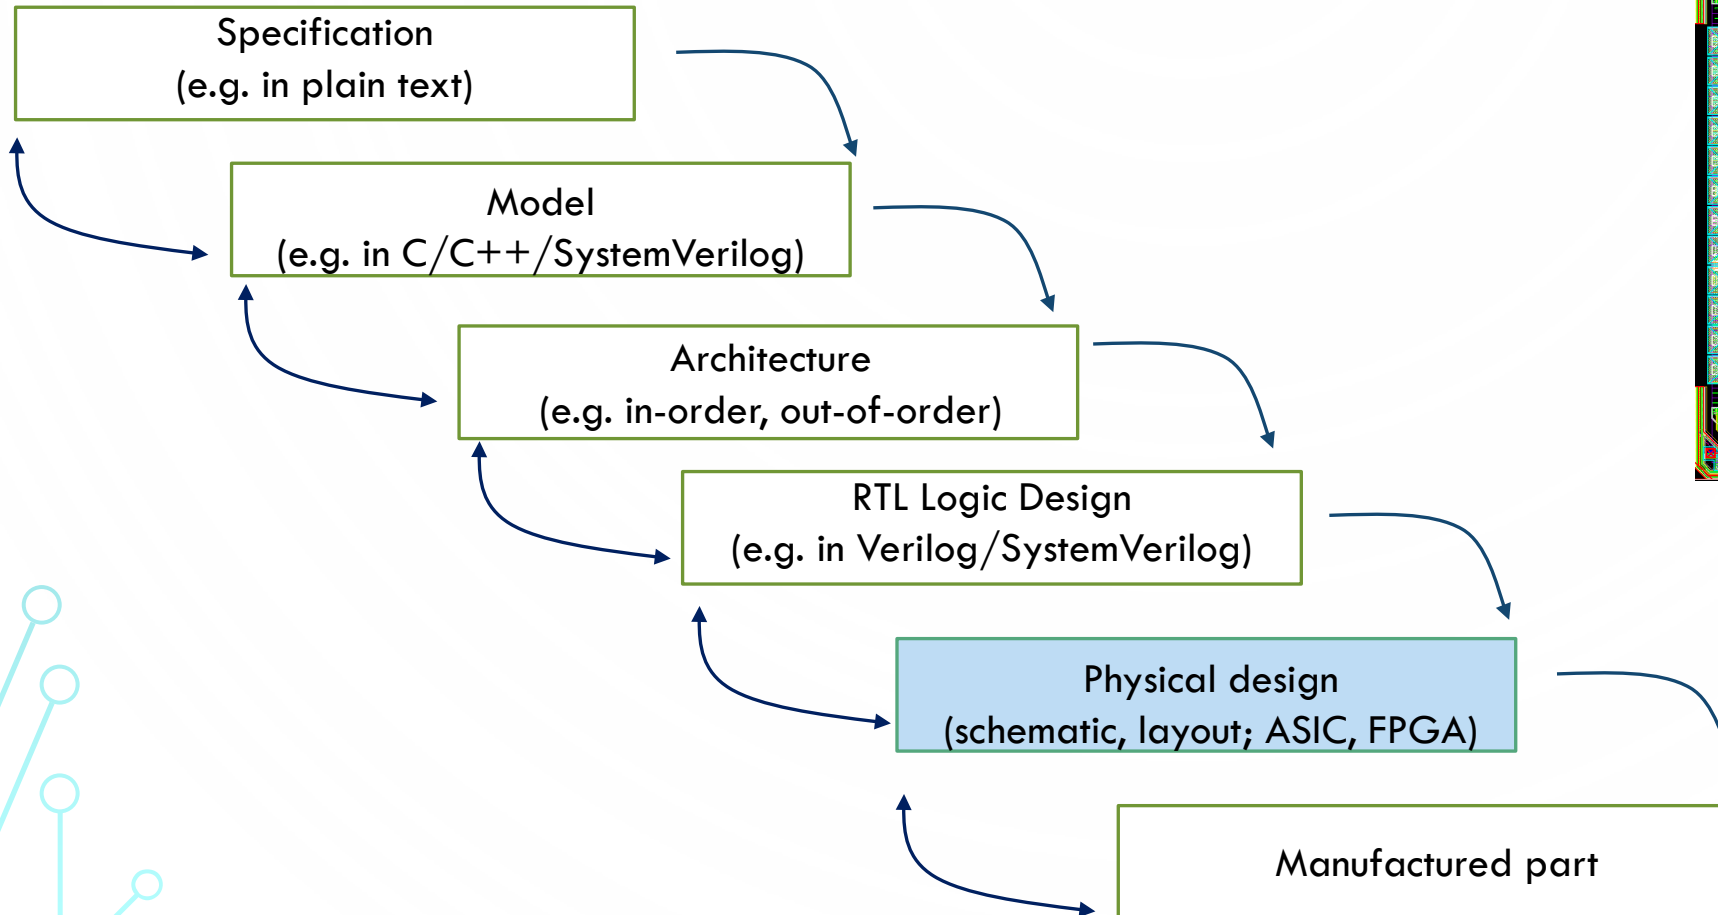


| | | |
|--|---|--|
| A25 Type: Cores Supplier: Andes Priv. spec: 1.11 User spec: RV32GCP + SV32 + Andes V5 ext. License: Andes Commercial License Primary Language: Verilog Bit Processor: 32 WEBSITE | A25MP Type: Cores Supplier: Andes Priv. spec: 1.11 User spec: RV32GCP + SV32 + Andes V5 ext. + Multi-core License: Andes Commercial License Primary Language: Verilog Bit Processor: 32 WEBSITE | Ariane Type: Cores Supplier: ETH Zurich, Università di Bologna Priv. spec: 1.11-draft User spec: RV64GC License: Solderpad Hardware License v. 0.51 Primary Language: SystemVerilog Bit Processor: 64 WEBSITE GITHUB |
| AX25 Type: Cores Supplier: Andes Priv. spec: 1.11 User spec: RV64GCP + SV39/48 + Andes V5 ext. License: Andes Commercial License Primary Language: Verilog Bit Processor: 64 WEBSITE | AX25MP Type: Cores Supplier: Andes Priv. spec: 1.11 User spec: RV64GCP + SV39/48 + Andes V5 ext. + Multi-core License: Andes Commercial License Primary Language: Verilog Bit Processor: 64 WEBSITE | Berkeley Out-of-Order Machine (BOOM) Type: Cores Supplier: Esperanto, UCB Bar Priv. spec: 1.11-draft User spec: 2.3-draft License: BSD Primary Language: Chisel GITHUB |

<https://riscv.org/risc-v-cores/>


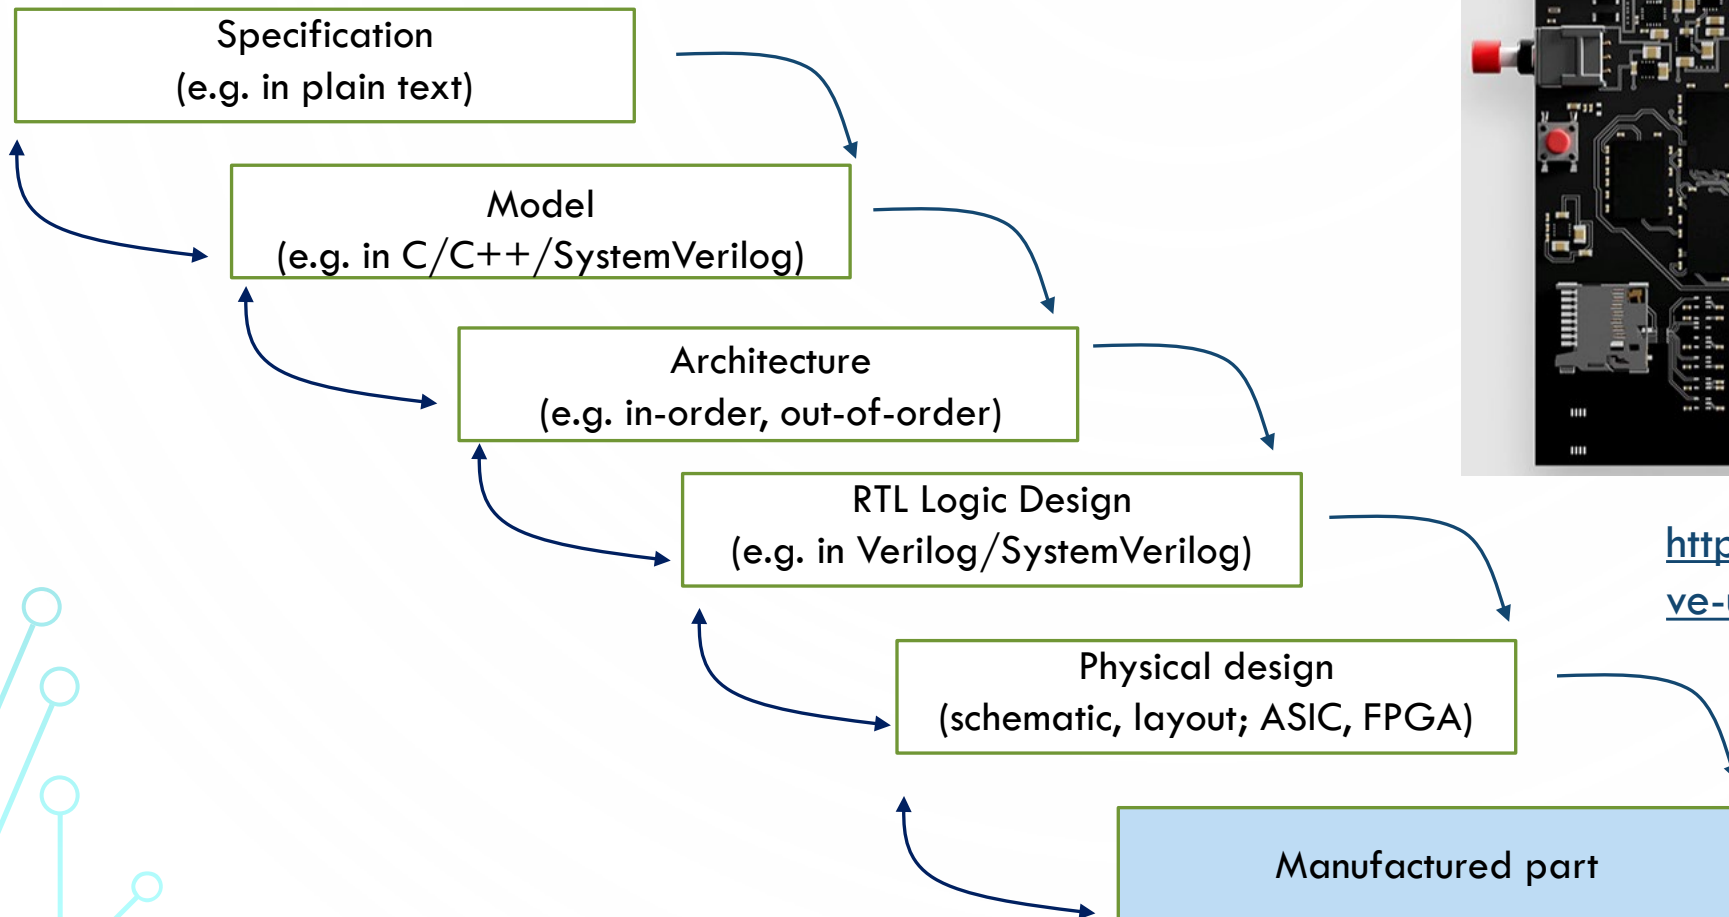
Example: RISC-V Design Process

- Design through layers of abstractions



Example: RISC-V Design Process

- Design through layers of abstractions



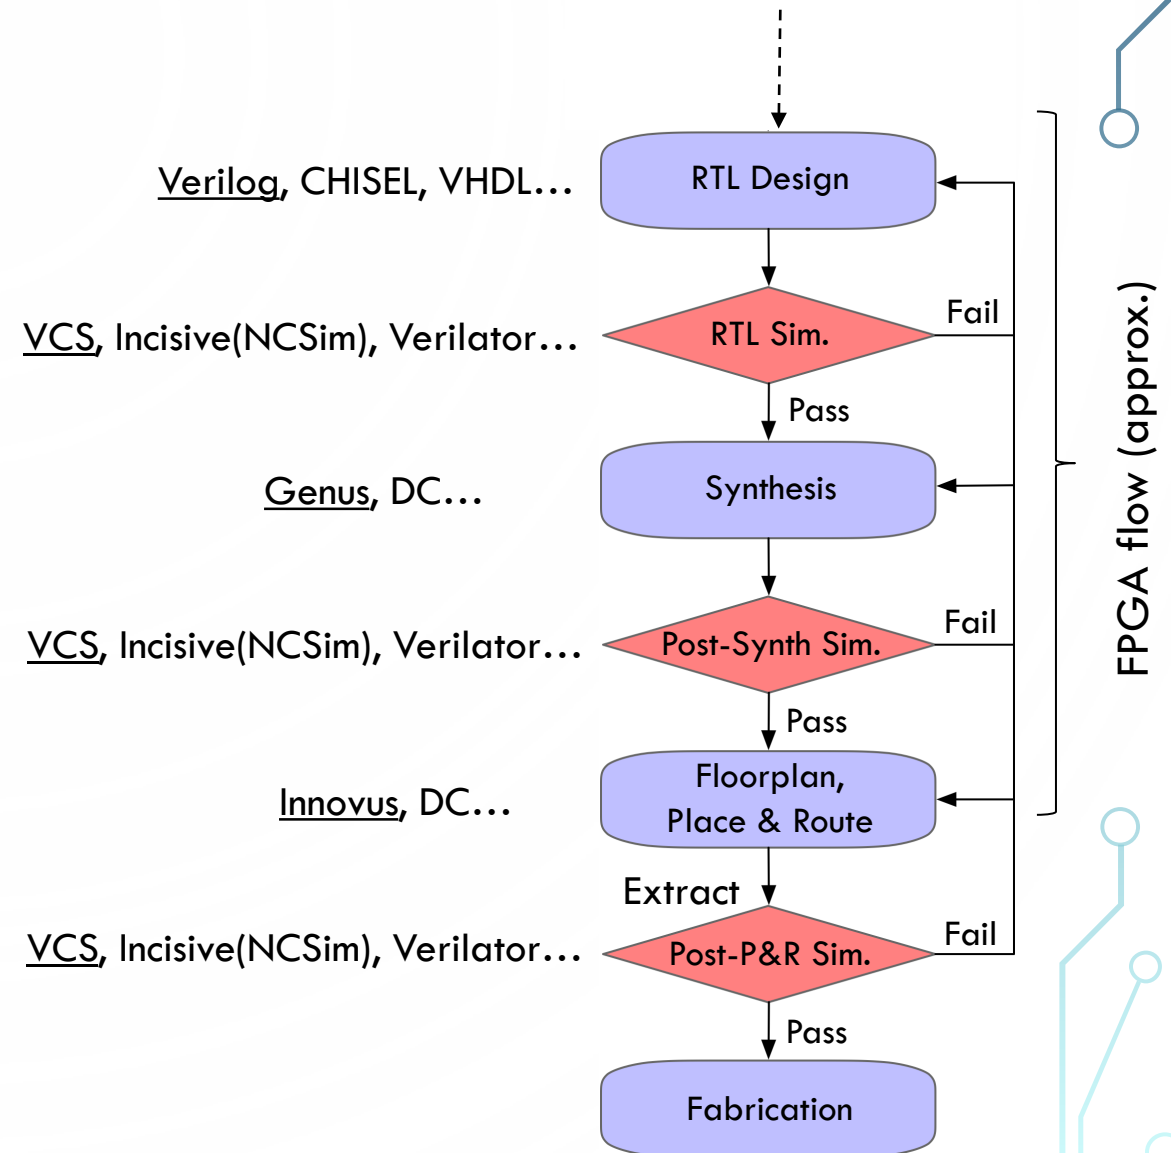
<https://www.sifive.com/boards/hi-five-unleashed>

RTL → Physical Design

RTL Logic Design
(e.g. in Verilog/SystemVerilog)


Physical design
(schematic, layout; ASIC, FPGA)


- Labs focus on a process of translating RTL to physical ASIC or FPGA by using industry-standard tools.
- Explores the entire design stack.



EECS

- <https://efabless.com/projects/35> 16





Boolean Logic in A Nutshell

Boolean Logic and Logic Gates (From CS61C/EE16B)

- Logic gates

Name


Boolean equation

Symbol

Truth table

NOT or Inverter

$$\text{Out} = \overline{A}$$




| A | Out |
|---|-----|
| 0 | 1 |
| 1 | 0 |

Single input

Buffer


$$\text{Out} = A$$



| A | Out |
|---|-----|
| 0 | 0 |
| 1 | 1 |

NAND

$$\text{Out} = \overline{A \cdot B}$$




| A | B | Out |
|---|---|-----|
| 0 | 0 | 1 |
| 0 | 1 | 1 |
| 1 | 0 | 1 |
| 1 | 1 | 0 |

| A | B | Out |
|---|---|-----|
| 0 | 0 | 1 |
| 0 | 1 | 0 |
| 1 | 0 | 0 |
| 1 | 1 | 0 |

NOR

$$\text{Out} = \overline{A + B}$$



- In CMOS, basic logic gates are inverting

More Logic Gates

Name


Boolean equation

Symbol

Truth table


AND

$$\text{Out} = A \cdot B$$




| A | B | Out |
|---|---|-----|
| 0 | 0 | 0 |
| 0 | 1 | 0 |
| 1 | 0 | 0 |
| 1 | 1 | 1 |

In CMOS




OR

$$\text{Out} = A + B$$



| A | B | Out |
|---|---|-----|
| 0 | 0 | 0 |
| 0 | 1 | 1 |
| 1 | 0 | 1 |
| 1 | 1 | 1 |

In CMOS



More Logic Gates

Name


Boolean equation

Symbol

Truth table

Exclusive OR
XOR


$$\text{Out} = A \oplus B$$



| A | B | Out |
|---|---|-----|
| 0 | 0 | 0 |
| 0 | 1 | 1 |
| 1 | 0 | 1 |
| 1 | 1 | 0 |

Exclusive NOR
XNOR

$$\text{Out} = \overline{A \oplus B}$$



| A | B | Out |
|---|---|-----|
| 0 | 0 | 1 |
| 0 | 1 | 0 |
| 1 | 0 | 0 |
| 1 | 1 | 1 |


- XOR and XNOR are both inverting and non-inverting

Multi-Input Gates

3-Input NAND

NAND3 Boolean equation

$$\text{Out} = \overline{A \cdot B \cdot C}$$




| A | B | C | Out |
|---|---|---|-----|
| 0 | 0 | 0 | 1 |
| 0 | 0 | 1 | 1 |
| 0 | 1 | 0 | 1 |
| 0 | 1 | 1 | 1 |
| 1 | 0 | 0 | 1 |
| 1 | 0 | 1 | 1 |
| 1 | 1 | 0 | 1 |
| 1 | 1 | 1 | 0 |

And-Or-Invert

AOI21 Boolean equation

$$\text{Out} = \overline{A \cdot B + C}$$




| A | B | C | Out |
|---|---|---|-----|
| 0 | 0 | 0 | 1 |
| 0 | 0 | 1 | 0 |
| 0 | 1 | 0 | 1 |
| 0 | 1 | 1 | 0 |
| 1 | 0 | 0 | 1 |
| 1 | 0 | 1 | 0 |
| 1 | 1 | 0 | 0 |
| 1 | 1 | 1 | 0 |

- Single gate in modern CMOS usually doesn't have more than 3-4 inputs

Combinational Logic (CL) Blocks

Example four-input function:




Truth Table

| A | B | C | D | Out |
|---|---|---|---|--------------|
| 0 | 0 | 0 | 0 | $F(0,0,0,0)$ |
| 0 | 0 | 0 | 1 | $F(0,0,0,1)$ |
| 0 | 0 | 1 | 0 | $F(0,0,1,0)$ |
| 0 | 0 | 1 | 1 | $F(0,0,1,1)$ |
| 0 | 1 | 0 | 0 | $F(0,1,0,0)$ |
| 0 | 1 | 0 | 1 | $F(0,1,0,1)$ |
| 0 | 1 | 1 | 0 | $F(0,1,1,0)$ |
| 0 | 1 | 1 | 1 | $F(0,1,1,1)$ |
| 1 | 0 | 0 | 0 | $F(1,0,0,0)$ |
| 1 | 0 | 0 | 1 | $F(1,0,0,1)$ |
| 1 | 0 | 1 | 0 | $F(1,0,1,0)$ |
| 1 | 0 | 1 | 1 | $F(1,0,1,1)$ |
| 1 | 1 | 0 | 0 | $F(1,1,0,0)$ |
| 1 | 1 | 0 | 1 | $F(1,1,0,1)$ |
| 1 | 1 | 1 | 0 | $F(1,1,1,0)$ |
| 1 | 1 | 1 | 1 | $F(1,1,1,1)$ |

- Output a function only of the current inputs (no history).
- Truth-table representation of function. Output is explicitly specified for each input combination.
- In general, CL blocks have more than one output signal, in which case, the truth-table will have multiple output columns.

Example CL Block

- 2-bit adder. Takes two 2-bit integers and produces 3-bit result.



| A1 | A0 | B1 | B0 | C2 | C1 | C0 |
|----|----|----|----|----|----|----|
| 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 1 | 0 | 0 | 1 |
| 0 | 0 | 1 | 0 | 0 | 1 | 0 |
| 0 | 0 | 1 | 1 | 0 | 1 | 1 |
| 0 | 1 | 0 | 0 | 0 | 0 | 1 |
| 0 | 1 | 0 | 1 | 0 | 1 | 0 |
| 0 | 1 | 1 | 0 | 0 | 1 | 1 |
| 0 | 1 | 1 | 1 | 1 | 0 | 0 |
| 1 | 0 | 0 | 0 | 0 | 1 | 0 |
| 1 | 0 | 0 | 1 | 0 | 1 | 1 |
| 1 | 0 | 1 | 0 | 1 | 0 | 0 |
| 1 | 0 | 1 | 1 | 1 | 0 | 1 |
| 1 | 1 | 0 | 0 | 0 | 1 | 1 |
| 1 | 1 | 0 | 1 | 1 | 0 | 0 |
| 1 | 1 | 1 | 0 | 1 | 0 | 1 |
| 1 | 1 | 1 | 1 | 1 | 1 | 0 |

- Think about truth table for 32-bit adder. It's possible to write out, but it might take a while!

Theorem:

Any combinational logic function can be implemented as a network of simple logic gates.

Quiz

Total number of possible truth tables with 4 inputs is:

- a) 4
- b) 16
- c) 256
- d) 16,384
- e) 65,536
- f) None of the above

www.yellkey.com/foot

Peer Instruction

Total number of possible truth tables with 4 inputs is:

- a) 4
- b) 16
- c) 256
- d) 16,384
- e) 65,536
- f) None of the above


www.yellkey.com/leg

Logic Circuit

- A logic gate can be implemented in different ways



NAND

$$\text{Out} = \overline{A \cdot B}$$




| A | B | Out |
|---|---|-----|
| 0 | 0 | 1 |
| 0 | 1 | 1 |
| 1 | 0 | 1 |
| 1 | 1 | 0 |

CMOS




Sizing of transistors (W/L) in CMOS changes properties (delay, power, size) of a logic gate



Mechanical LEGO logic gates. A clockwise rotation represents a binary “one” while a counter-clockwise rotation represents a binary “zero.”

Sequential Logic Blocks



- Output is a function of both the current inputs and the state.
- State represents the memory.
- State is a function of previous inputs.
- In synchronous digital systems, state is updated on each clock tick.


Flip-Flop as A Sequential Circuit

- Synchronous state element transfers its input to the output on a rising (or, rarely, falling) clock edge



- Flip-flop

- Rising edge

Signifies 'edge triggered'



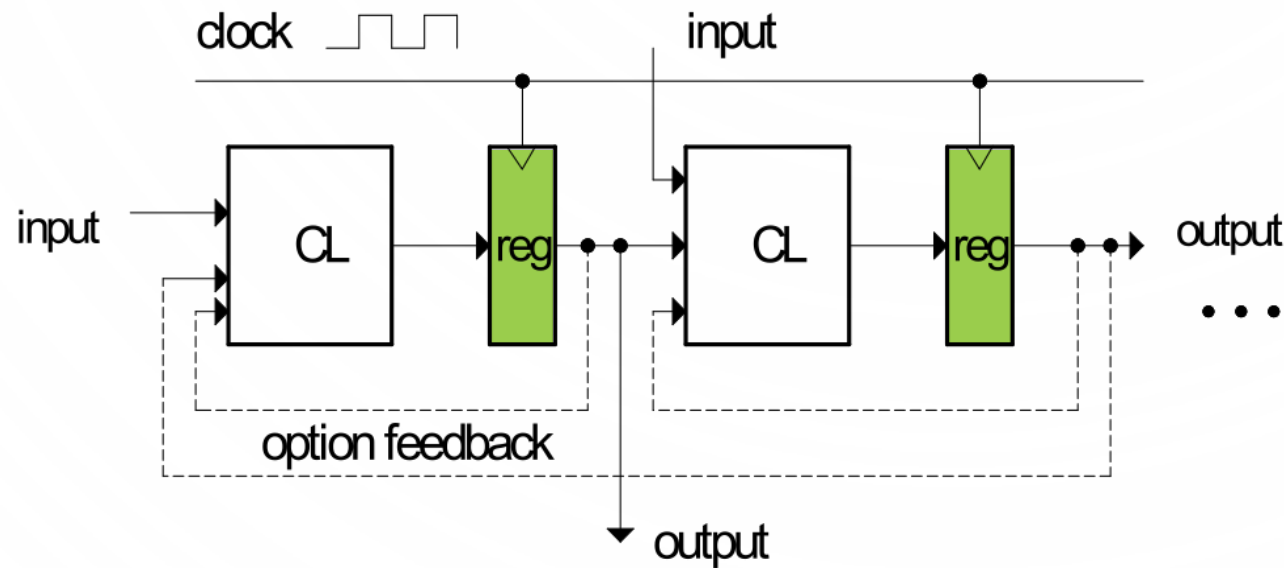
- 4-bit register



Register Transfer Level Abstraction (RTL)

Any synchronous digital circuit can be represented with:

- Combinational Logic (CL) blocks, plus
- State elements (registers or memories)
- Clock orchestrates *sequencing* of CL operations



- State elements are combined with CL blocks to control the flow of data.

Administrivia



- Labs and discussions start this week
- Lab 1 posted, please start it before coming to the lab session
- Lab 2 is more involved
 - Be prepared
 - Verilog primer
- Homework 1 posted this week, due next Friday
 - Start early



Design Metrics: Robustness



What Makes Circuits Digital?

- Chips are noisy
- Supply noise will appear at the output of the logic gate



- The following logic gate should still interpret its inputs as 0s and 1s
- This necessary property is called "Restoration" or "Regeneration"
- A lot of money was spent in the past to unsuccessfully make logic out of non-regenerative gates
 - Some of emerging CMOS replacements don't have gain...


Beneath the Digital Abstraction




- Logic levels:
 - Mapping a continuous voltage onto a discrete binary logic variable
 - Low (0): $[0, V_L]$
 - High (1): $[V_H, V_{DD}]$
 - V_L V_H : nominal voltage levels

Voltage Transfer Characteristic

- A gate should interpret everything that is close to 0V as a logic 0
 - And everything close to V_{DD} as a logic 1




a) Ideal



b) Real

Nominal Voltage Levels





$$\begin{aligned} V_{OH} &= f(V_{OL}) \\ V_{OL} &= f(V_{OH}) \\ V_M &= f(V_M) \end{aligned}$$

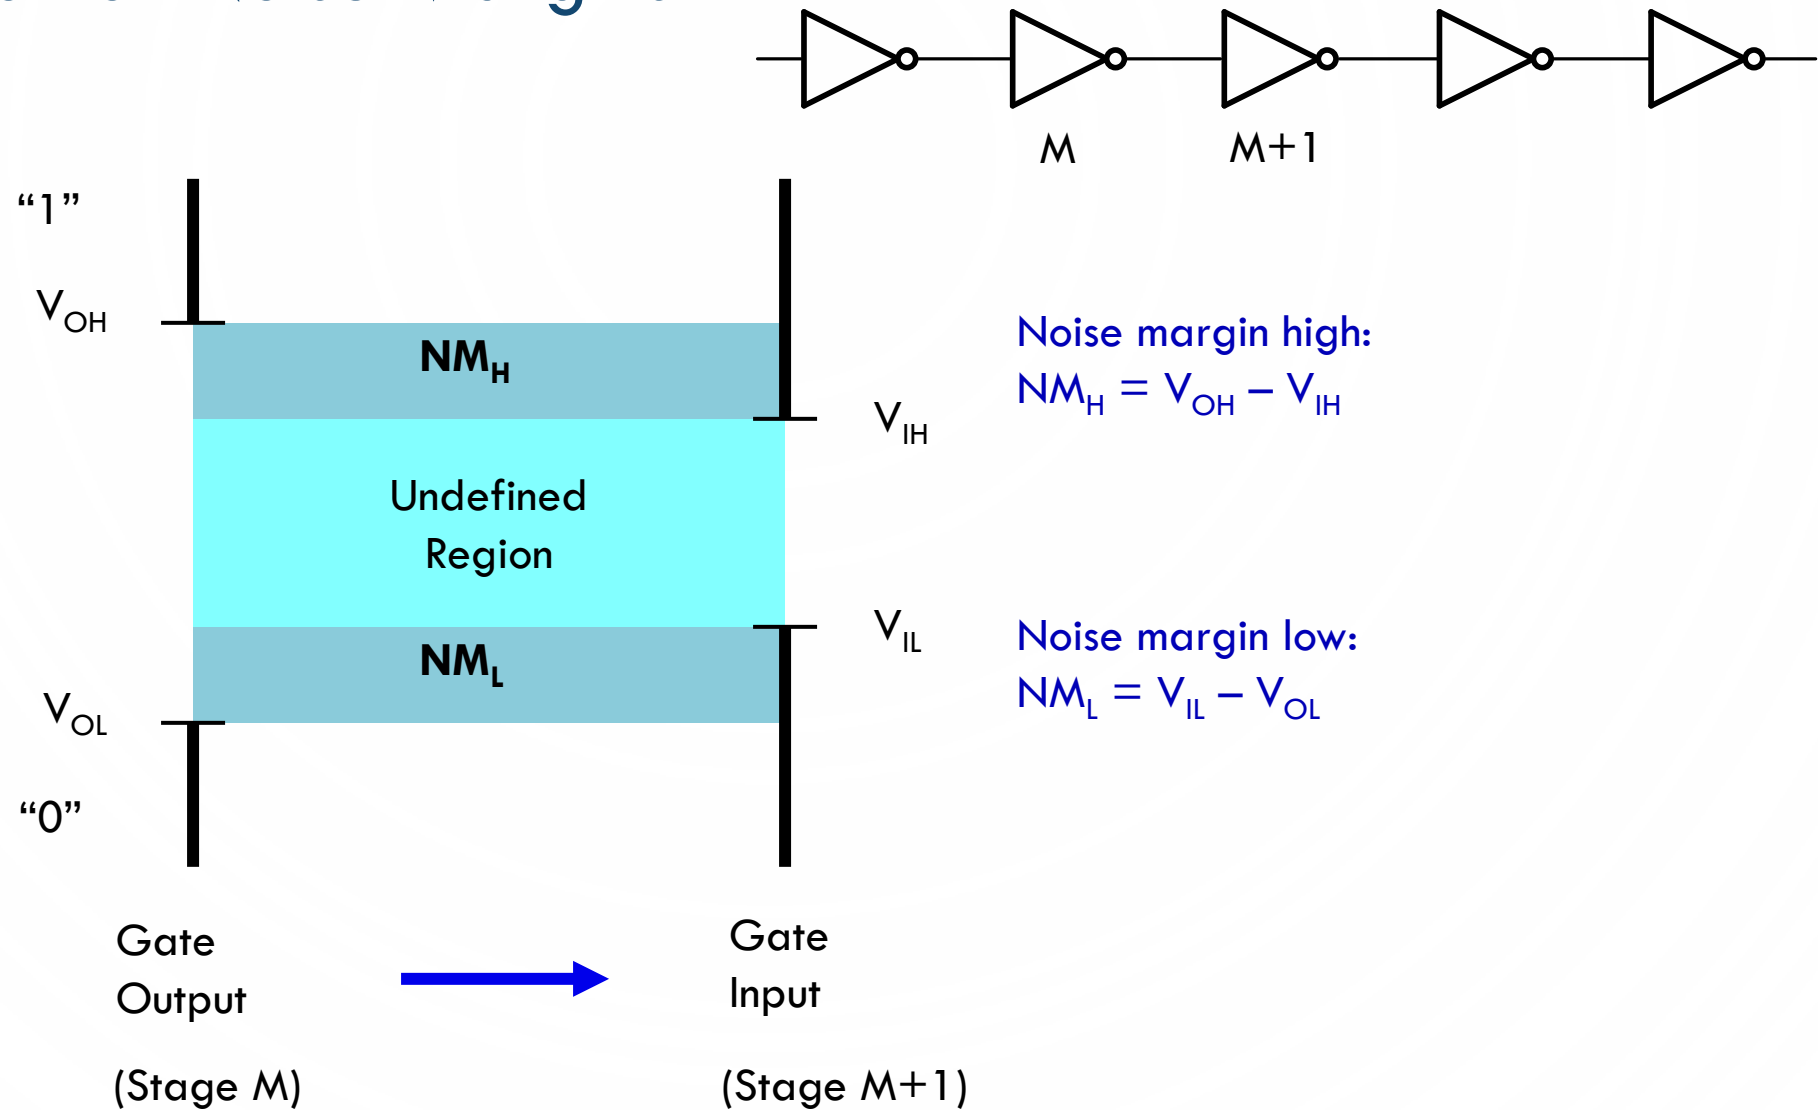
In CMOS:

$$\begin{aligned} V_{OH} &= V_{DD} \\ V_{OL} &= 0 \\ V_M &\sim V_{DD}/2 \end{aligned}$$

Mapping Between Analog Voltages and Digital Signals




Definition of Noise Margins



The amount of **noise** that could be added to a **worst-case output** so that the signal can still be interpreted correctly as a **valid input** to the next gate.

Regenerative Property

- Ensures that a **disturbed** signal gradually **regenerates** one of the **nominal voltage levels** after passing through a few logical stages.
 - Look for a sharp transition in voltage transfer characteristics.






Design Metrics: Performance



Design Tradeoffs

- The desired functionality can be implemented with different performance, power or cost targets




Digital Logic Delay


- Changes at the inputs do not instantaneously appear at the outputs
 - There are finite resistances and capacitances in each gate...




- Propagation through a chain of gates



Delay Definitions

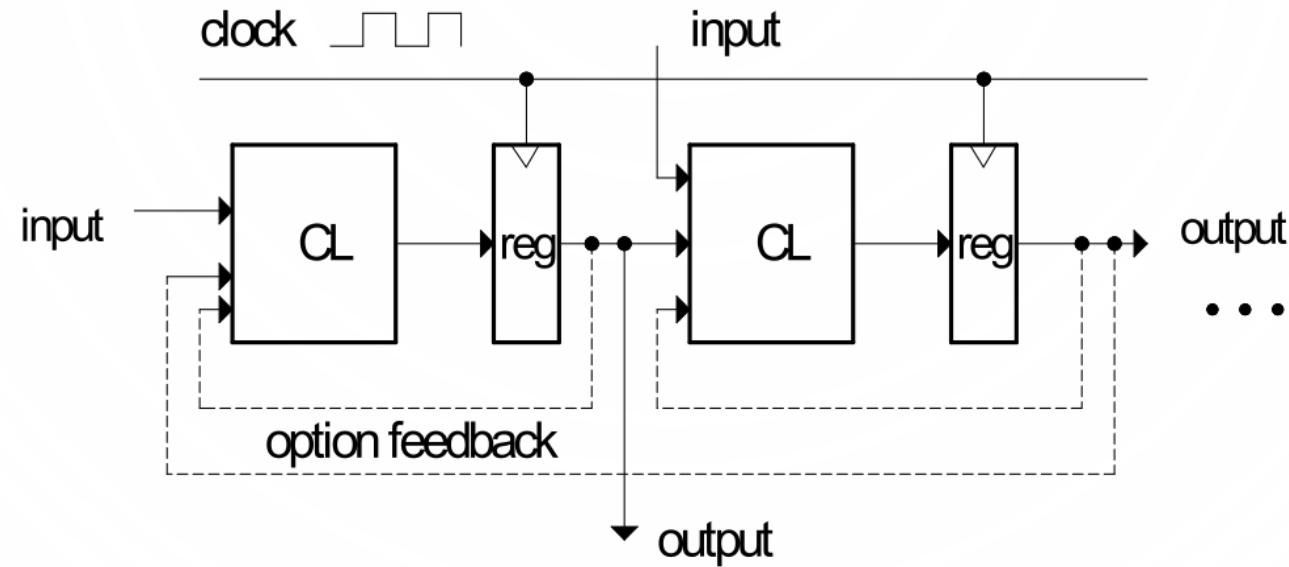


- Delay calculations need to be additive
 - Calculate the delay from the same point in the waveform



Digital Logic Timing

- The longest propagation delay through CL blocks sets the maximum clock frequency



- To increase clock rate:
 - Find the longest path
 - Make it faster


Performance

- Throughput

- Number of tasks performed in a unit of time (operations per second)
- E.g. Google TPUv3 board performs 420 TFLOPS (10^{12} floating-point operations per second, where a floating point operation is BFLOAT16)
- Watch out for 'op' definitions – can be a 1-b ADD or a double-precision FP add (or more complex task)
- Peak vs. average throughput

- Latency

- How long does a task take from start to finish
- E.g. facial recognition on a phone takes 10's of ms
- Sometime expressed in terms of clock cycles
- Average vs. 'tail' latency






Design Metrics: Energy and Power

Energy and Power

- Energy (in joules (J))
 - Needed to perform a task
 - Add two numbers or fetch a datum from memory
 - (or fetch two numbers, add them and store in memory)
 - Active and standby
 - Battery stores certain amount of energy (in $Ws = J$ or Wh)
 - That is what utility charges for (in kWh)
- Power (in watts (W))
 - Energy dissipated in time ($W = J/s$)
 - Sets cooling requirements
 - Heat spreader, size of a heat sink, forced air, liquid, ...





Design Metrics: Cost


Cost

- Non-recurring engineering (NRE) costs
- Cost to develop a design (product)
 - Amortized over all units shipped
 - E.g. \$20M in development adds \$.20 to each of 100M units
- Recurring costs
 - Cost to manufacture, test and package a unit
 - Processed wafer cost is ~10k (around 16nm node) which yields:
 - 1 Cerebras chip
 - 50-100 large FPGAs or GPUs
 - 200 laptop CPUs
 - >1000 cell phone SoCs

$$\text{cost per IC} = \text{variable cost per IC} + \frac{\text{fixed cost}}{\text{volume}}$$

$$\text{variable cost} = \frac{\text{cost of die} + \text{cost of die test} + \text{cost of packaging}}{\text{final test yield}}$$

Die Cost



$$\text{cost of die} = \frac{\text{cost of wafer}}{\text{dies per wafer} * \text{die yield}}$$


From: <http://www.amd.com>

Yield

$$Y = \frac{\text{No. of good chips per wafer}}{\text{Total number of chips per wafer}} \times 100\%$$


$$\text{Die cost} = \frac{\text{Wafer cost}}{\text{Dies per wafer} \times \text{Die yield}}$$

$$\text{Dies per wafer} = \frac{\pi \times (\text{wafer diameter}/2)^2}{\text{die area}} - \frac{\pi \times \text{wafer diameter}}{\sqrt{2} \times \text{die area}}$$




Defects

Yield = 0.25



Yield = 0.76



$$\text{die yield} = \left(1 + \frac{\text{defects per unit area} \times \text{die area}}{\alpha} \right)^{-\alpha}$$

α is approximately 3

$$\text{die cost} = f(\text{die area})^4$$

Summary

- The design process involves traversing the abstraction layers of specification, modeling, architecture, RTL design and physical implementation
- Tests follow the design refinements
- Targets are processors, FPGAs or ASICs
- Automated design flows help manage the complexity
- Optimize for performance, energy and cost