

inst.eecs.berkeley.edu/~eecs151

EECS151 : Introduction to Digital Design and ICs

Lecture 2 – Design Process



Bora Nikolić

At HotChips'19 Cerebras announced the largest chip in the world at 8.5 in x 8.5in with 1.2 trillion transistors, and 15kW of power, aimed for training of deep-learning neural networks

At HotChips'21 they showed the next version in 7nm CMOS, with >2x transistor count

- 46,225 mm² silicon
- 2.6 Trillion transistors
- 850,000 AI optimized cores
- 40 Gigabytes on-chip memory
- 20 Petabyte/s memory bandwidth
- 220 Petabit/s fabric bandwidth
- 7nm Process technology at TSMC

EECS151/251A L02 DESIGN

Sean Lie,
HotChips'21

Berkeley
UNIVERSITY OF CALIFORNIA

CC BY NC SA

1

Review

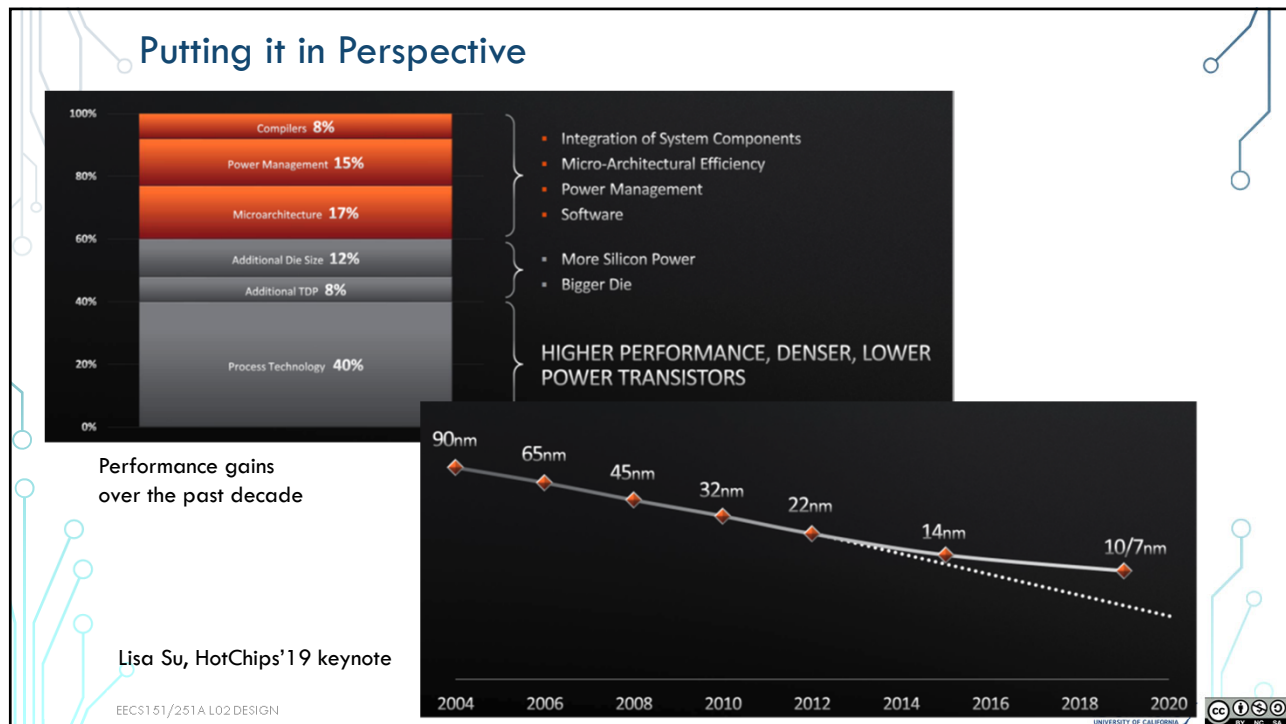
- Moore's law is slowing down
 - There are continued improvements in technology, but at a slower pace
- Dennard's scaling has ended a decade ago
 - All designs are now power limited
- Specialization and customization provides added performance
 - Under power constraints and stagnant technology
- Design costs are high
 - Methodology and better reuse to rescue!
 - Abstraction, modularity, regularity are the keys
 - And creativity!

EECS151/251A L02 DESIGN

2 Berkeley
UNIVERSITY OF CALIFORNIA

CC BY NC SA

2



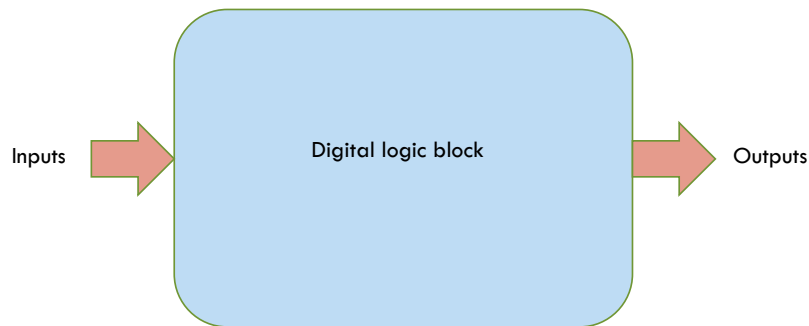
3



4

Implementing Digital Systems

- Digital systems implement a set of Boolean equations



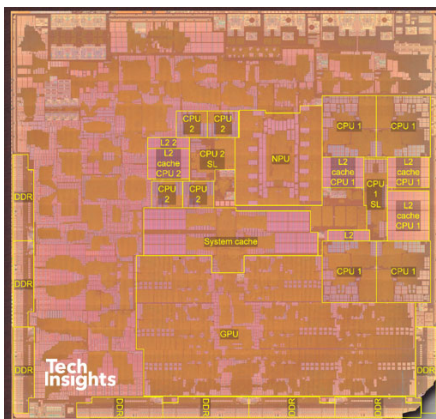
- How do we implement a digital system?

EECS151/251A L02 DESIGN

5 Berkeley
UNIVERSITY OF CALIFORNIA

5

Modern (Mostly) Digital System-On-A-Chip



TechInsights

- 5nm CMOS
- Up to 2.49GHz

- 4x 'Firestorm' Large CPUs
- 4x 'Icestorm' Small CPUs
- GPU
- Neural processing unit (NPU)
- Lots of memory
- DDR memory interfaces

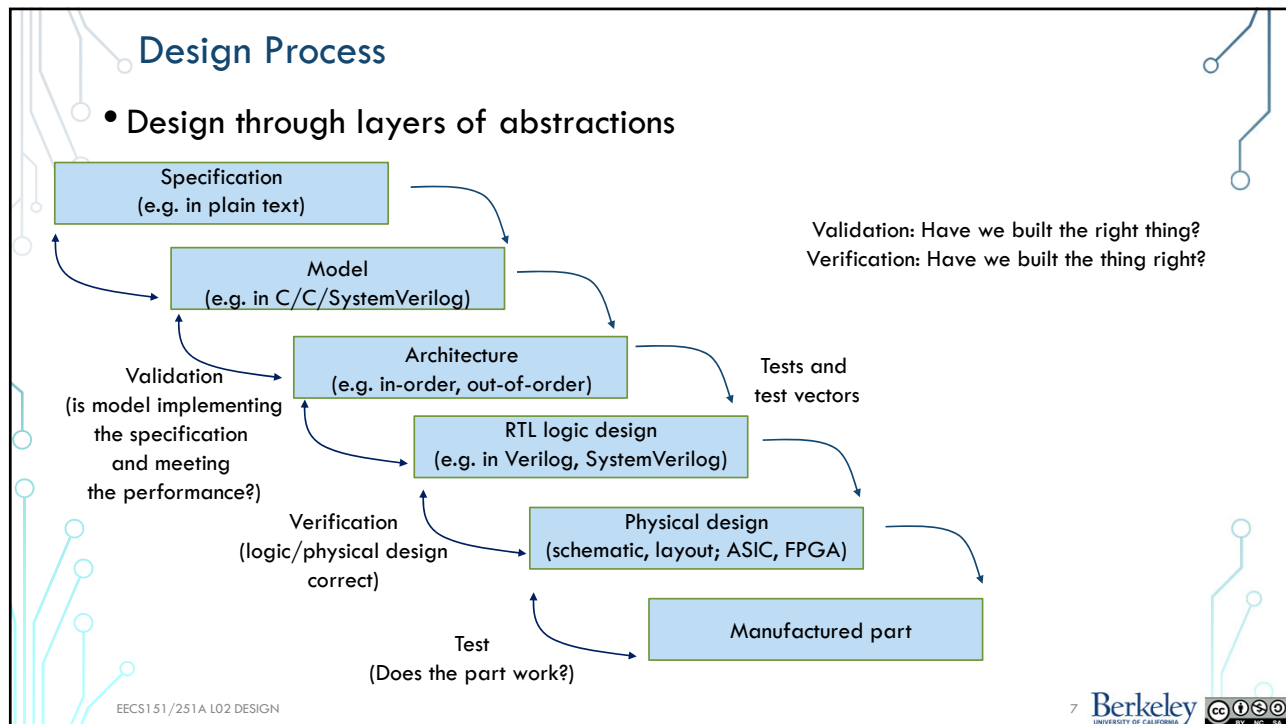


By Henriok

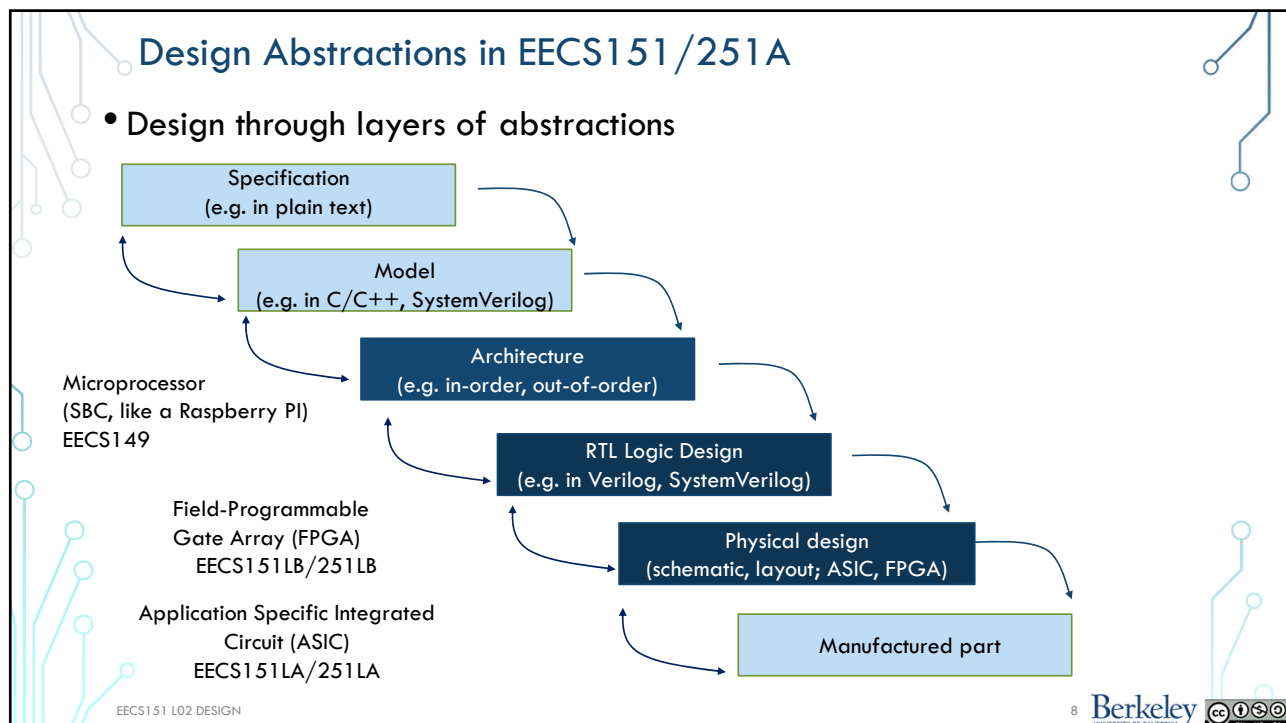
<https://commons.wikimedia.org/w/index.php?curid=96026688>


6 Berkeley
UNIVERSITY OF CALIFORNIA

6



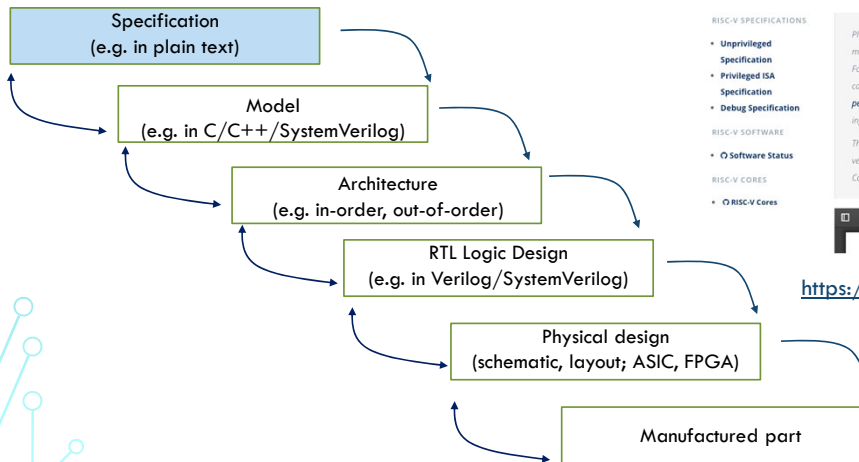
7



8

Example: RISC-V Design Process

- Design through layers of abstractions



ABOUT MEMBERSHIP SPECS & SUPPORT CORES & TOOLS NEWS EVENTS

Specifications

Specifications

RISC-V SPECIFICATIONS

- Unprivileged Specification
- Privileged ISA Specification
- Debug Specification
- RISC-V SOFTWARE
- Software Status
- RISC-V CORES
- RISC-V Cores

Please note, RISC-V ISA and related specifications are developed, ratified and maintained by RISC-V Foundation contributing members within the RISC-V Foundation Technical Committee. Operating details of the Technical Committee can be found in the [RISC-V Foundation Workspace](#). Work on the specification is performed on [GitHub](#) and the GitHub issue mechanism can be used to provide input into the specification.

The specifications shown below is the current ratified release. The most recent version of the draft specification, which is in development within the Technical Committee, can be found here on [GitHub](#).



<https://riscv.org/specifications/>

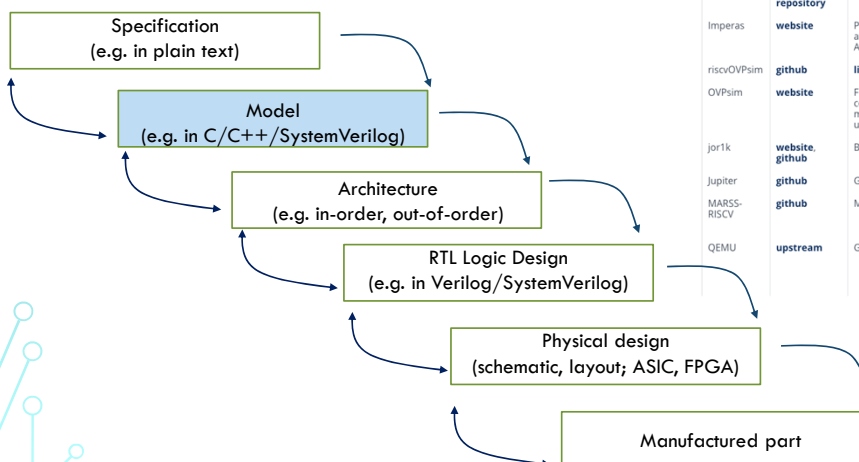
EECS151 L02 DESIGN



9

Example: RISC-V Design Process

- Design through layers of abstractions



Simulators

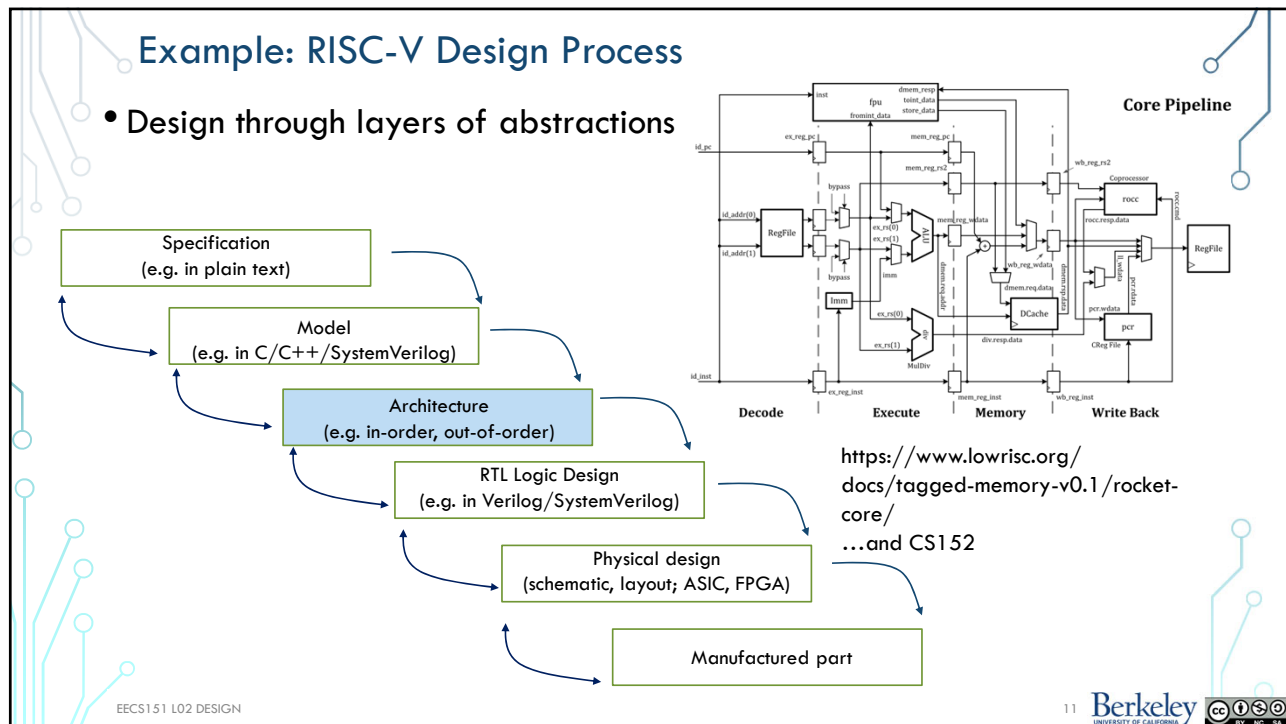
Name	Links	License	Maintainers
DBT-RISE-RISC-V	github	BSD-3-Clause	MINRES Technologies
FireSim	website , mailing list , github , ISCA 2018 Paper	BSD	Sagar Karandikar, Howard Mao, Donggyu Kim, David Biancolin, Alon Amid, Berkeley Architecture Research
gem5	SW-dev thread repository	BSD-style	Alec Roelke (University of Virginia)
Imperas	website	Proprietary, models available under Apache 2.0	Imperas
riscvOVPsim	github	license	Imperas
OVPsim	website	Free for non commercial use, models available under Apache 2.0	Imperas
jor1k	website , github	BSD 2-Clause	Sebastian Macke
Jupiter	github	GPL-3.0	Andrés Castellanos
MARSS-RISC-V	github	MIT	Gaurav N Kothari, Parikshit P Sarnaik, Gokturk Yuksek (State University of New York at Binghamton)
QEMU	upstream	GPL	Sagar Karandikar (University of California, Berkeley), Bastian Koppelman (University of Paderborn), Alex Suykov, Stefan O'Rear and Michael Clark (SiFive)

<https://riscv.org/software-status/#simulators>

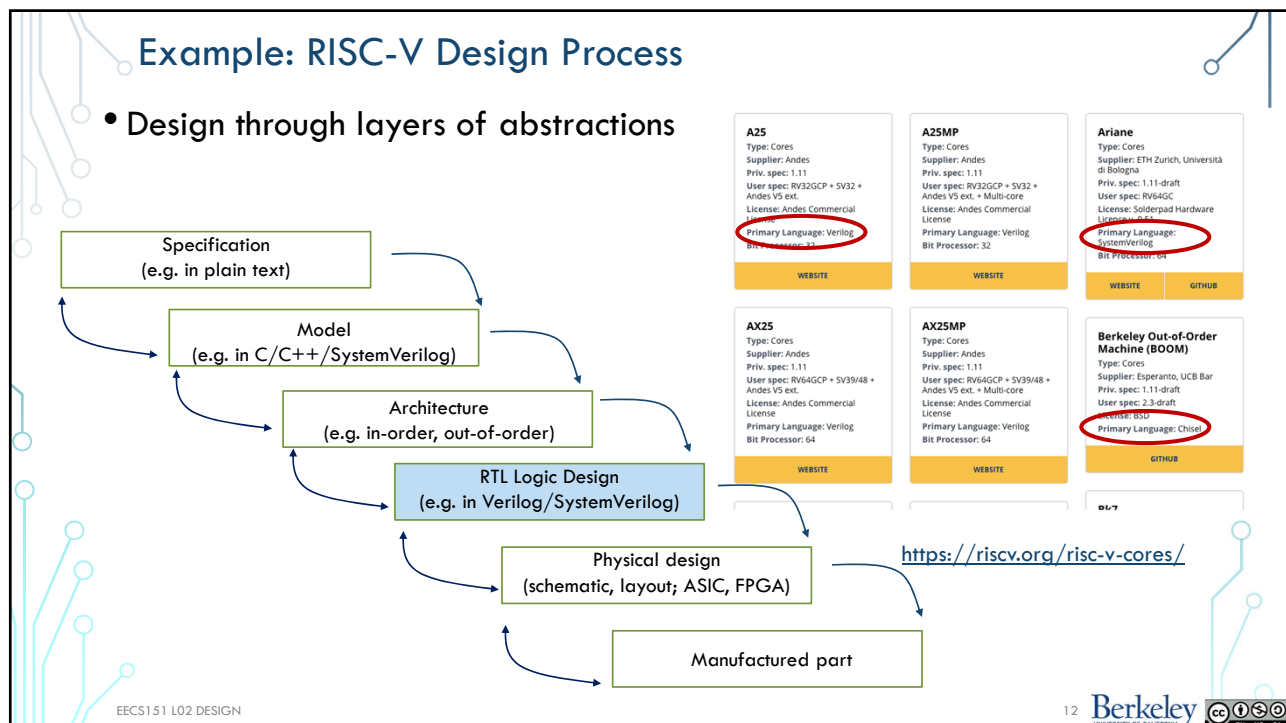
EECS151 L02 DESIGN



10



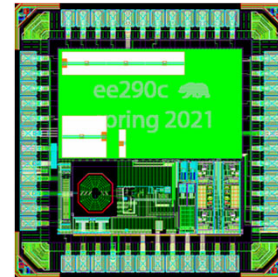
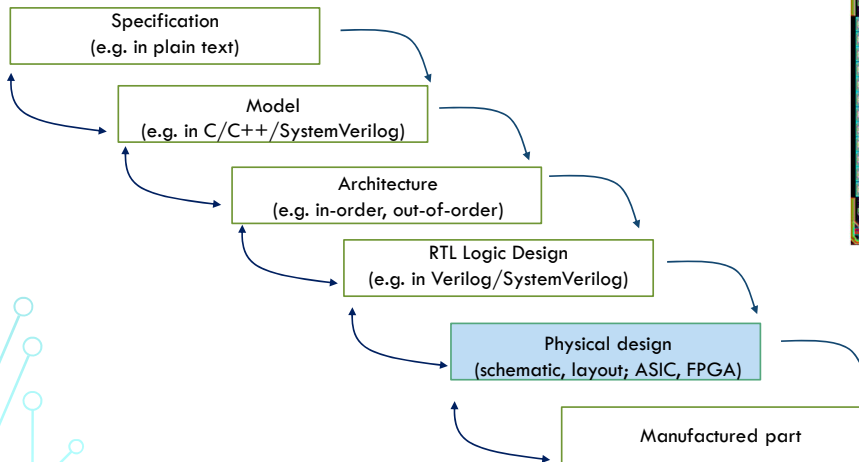
11



12

Example: RISC-V Design Process

- Design through layers of abstractions



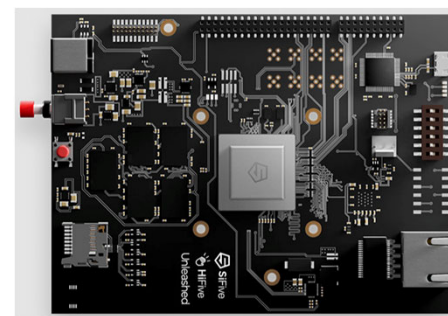
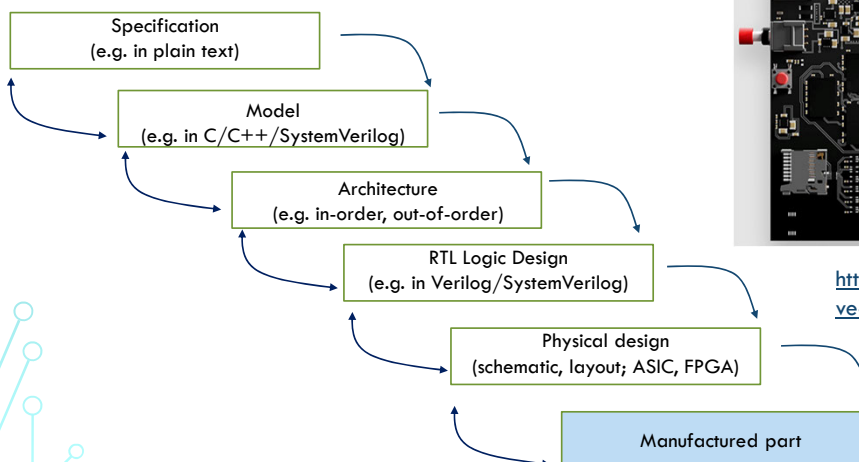
EECS151 L02 DESIGN

13 Berkeley UNIVERSITY OF CALIFORNIA

13

Example: RISC-V Design Process

- Design through layers of abstractions



<https://www.sifive.com/boards/hifive-unleashed>

EECS151 L02 DESIGN

14 Berkeley UNIVERSITY OF CALIFORNIA

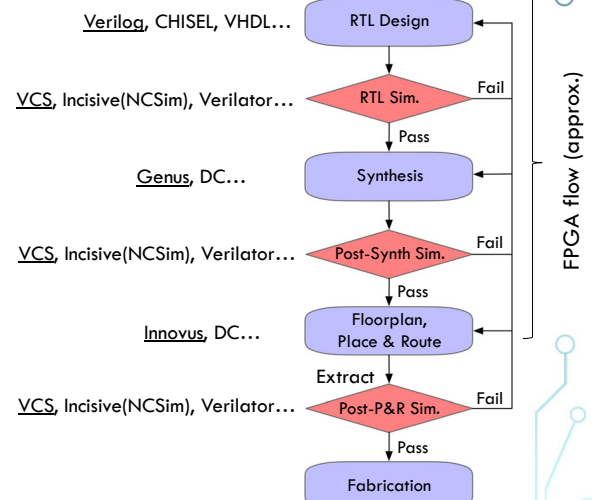
14

RTL → Physical Design

RTL Logic Design
(e.g. in Verilog/SystemVerilog)

Physical design
(schematic, layout; ASIC, FPGA)

- Labs focus on a process of translating RTL to physical ASIC or FPGA by using industry-standard tools.
- Explores the entire design stack.



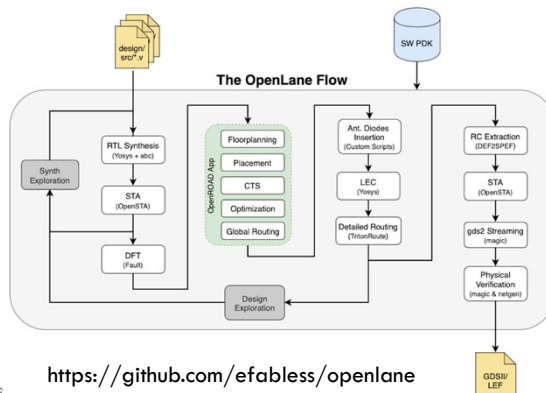
EECS151 L02 DESIGN

15 Berkeley UNIVERSITY OF CALIFORNIA

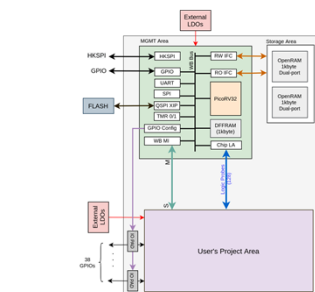
15

Open-Source Flows

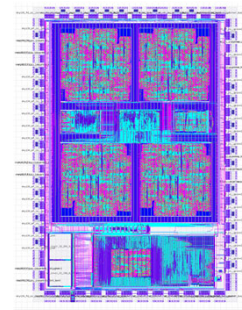
- Skywater 130nm is an open-source design kit
- OpenROAD (UCSD) and OpenLane (eFabless) are open-source design flows
 - Work with Sky130
 - A version of ASIC labs can target Sky130nm



<https://github.com/efabless/openlane>


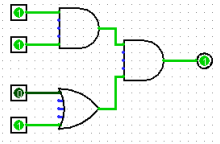


<https://github.com/efabless/caravel>





<https://efabless.com/projects/35> 16 Berkeley UNIVERSITY OF CALIFORNIA

16

Boolean Logic in A Nutshell




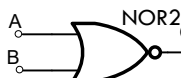
EECS151/251A L02 DESIGN

17  

17

Boolean Logic and Logic Gates (From CS61C/EE16B)



- Logic gates

Name	Boolean equation	Symbol	Truth table															
NOT or Inverter	$\text{Out} = \bar{A}$	 <p style="text-align: center;">NOT/INV</p>	<table border="1"> <thead> <tr> <th>A</th> <th>Out</th> </tr> </thead> <tbody> <tr> <td>0</td> <td>1</td> </tr> <tr> <td>1</td> <td>0</td> </tr> </tbody> </table>	A	Out	0	1	1	0									
A	Out																	
0	1																	
1	0																	
Buffer	$\text{Out} = A$	 <p style="text-align: center;">BUF</p>	<table border="1"> <thead> <tr> <th>A</th> <th>Out</th> </tr> </thead> <tbody> <tr> <td>0</td> <td>0</td> </tr> <tr> <td>1</td> <td>1</td> </tr> </tbody> </table>	A	Out	0	0	1	1									
A	Out																	
0	0																	
1	1																	
NAND	$\text{Out} = \overline{A \cdot B}$	 <p style="text-align: center;">NAND2</p>	<table border="1"> <thead> <tr> <th>A</th> <th>B</th> <th>Out</th> </tr> </thead> <tbody> <tr> <td>0</td> <td>0</td> <td>1</td> </tr> <tr> <td>0</td> <td>1</td> <td>1</td> </tr> <tr> <td>1</td> <td>0</td> <td>1</td> </tr> <tr> <td>1</td> <td>1</td> <td>0</td> </tr> </tbody> </table>	A	B	Out	0	0	1	0	1	1	1	0	1	1	1	0
A	B	Out																
0	0	1																
0	1	1																
1	0	1																
1	1	0																
NOR	$\text{Out} = \overline{A + B}$	 <p style="text-align: center;">NOR2</p>	<table border="1"> <thead> <tr> <th>A</th> <th>B</th> <th>Out</th> </tr> </thead> <tbody> <tr> <td>0</td> <td>0</td> <td>1</td> </tr> <tr> <td>0</td> <td>1</td> <td>0</td> </tr> <tr> <td>1</td> <td>0</td> <td>0</td> </tr> <tr> <td>1</td> <td>1</td> <td>0</td> </tr> </tbody> </table>	A	B	Out	0	0	1	0	1	0	1	0	0	1	1	0
A	B	Out																
0	0	1																
0	1	0																
1	0	0																
1	1	0																

Single input

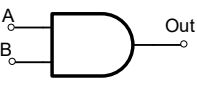
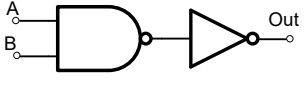
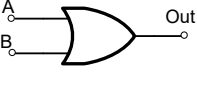
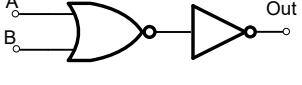
- In CMOS, basic logic gates are inverting

EECS151/251A L02 DESIGN



18  

18

More Logic Gates

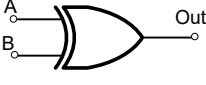
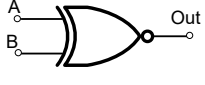
Name	Boolean equation	Symbol	Truth table															
AND	$Out = A \cdot B$		<table border="1"> <thead> <tr> <th>A</th> <th>B</th> <th>Out</th> </tr> </thead> <tbody> <tr><td>0</td><td>0</td><td>0</td></tr> <tr><td>0</td><td>1</td><td>0</td></tr> <tr><td>1</td><td>0</td><td>0</td></tr> <tr><td>1</td><td>1</td><td>1</td></tr> </tbody> </table>	A	B	Out	0	0	0	0	1	0	1	0	0	1	1	1
A	B	Out																
0	0	0																
0	1	0																
1	0	0																
1	1	1																
In CMOS																		
OR	$Out = A + B$		<table border="1"> <thead> <tr> <th>A</th> <th>B</th> <th>Out</th> </tr> </thead> <tbody> <tr><td>0</td><td>0</td><td>0</td></tr> <tr><td>0</td><td>1</td><td>1</td></tr> <tr><td>1</td><td>0</td><td>1</td></tr> <tr><td>1</td><td>1</td><td>1</td></tr> </tbody> </table>	A	B	Out	0	0	0	0	1	1	1	0	1	1	1	1
A	B	Out																
0	0	0																
0	1	1																
1	0	1																
1	1	1																
In CMOS																		

EECS151/251A L02 DESIGN

19  



19

More Logic Gates

Name	Boolean equation	Symbol	Truth table															
Exclusive OR XOR	$Out = A \oplus B$		<table border="1"> <thead> <tr> <th>A</th> <th>B</th> <th>Out</th> </tr> </thead> <tbody> <tr><td>0</td><td>0</td><td>0</td></tr> <tr><td>0</td><td>1</td><td>1</td></tr> <tr><td>1</td><td>0</td><td>1</td></tr> <tr><td>1</td><td>1</td><td>0</td></tr> </tbody> </table>	A	B	Out	0	0	0	0	1	1	1	0	1	1	1	0
A	B	Out																
0	0	0																
0	1	1																
1	0	1																
1	1	0																
Exclusive NOR XNOR	$Out = \overline{A \oplus B}$		<table border="1"> <thead> <tr> <th>A</th> <th>B</th> <th>Out</th> </tr> </thead> <tbody> <tr><td>0</td><td>0</td><td>1</td></tr> <tr><td>0</td><td>1</td><td>0</td></tr> <tr><td>1</td><td>0</td><td>0</td></tr> <tr><td>1</td><td>1</td><td>1</td></tr> </tbody> </table>	A	B	Out	0	0	1	0	1	0	1	0	0	1	1	1
A	B	Out																
0	0	1																
0	1	0																
1	0	0																
1	1	1																

• XOR and XNOR are both inverting and non-inverting

EECS151/251A L02 DESIGN

20  

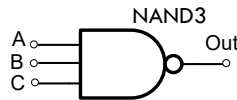
20

Multi-Input Gates

3-Input NAND

NAND3 Boolean equation

$$\text{Out} = \overline{A \cdot B \cdot C}$$



A	B	C	Out
0	0	0	1
0	0	1	1
0	1	0	1
0	1	1	1
1	0	0	1
1	0	1	1
1	1	0	1
1	1	1	0

And-Or-Invert

AOI21 Boolean equation

$$\text{Out} = \overline{A \cdot B + C}$$



A	B	C	Out
0	0	0	1
0	0	1	0
0	1	0	1
0	1	1	0
1	0	0	1
1	0	1	0
1	1	0	0
1	1	1	0

- Single gate in modern CMOS usually doesn't have more than 3-4 inputs

EECS151/251A L02 DESIGN

21

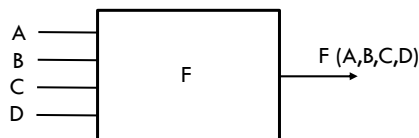
Berkeley



21

Combinational Logic (CL) Blocks

Example four-input function:



- Output a function only of the current inputs (no history).
- Truth-table representation of function. Output is explicitly specified for each input combination.
- In general, CL blocks have more than one output signal, in which case, the truth-table will have multiple output columns.

Truth Table

A	B	C	D	Out
0	0	0	0	F(0,0,0,0)
0	0	0	1	F(0,0,0,1)
0	0	1	0	F(0,0,1,0)
0	0	1	1	F(0,0,1,1)
0	1	0	0	F(0,1,0,0)
0	1	0	1	F(0,1,0,1)
0	1	1	0	F(0,1,1,0)
0	1	1	1	F(0,1,1,1)
1	0	0	0	F(1,0,0,0)
1	0	0	1	F(1,0,0,1)
1	0	1	0	F(1,0,1,0)
1	0	1	1	F(1,0,1,1)
1	1	0	0	F(1,1,0,0)
1	1	0	1	F(1,1,0,1)
1	1	1	0	F(1,1,1,0)
1	1	1	1	F(1,1,1,1)

EECS151/251A L02 DESIGN

22

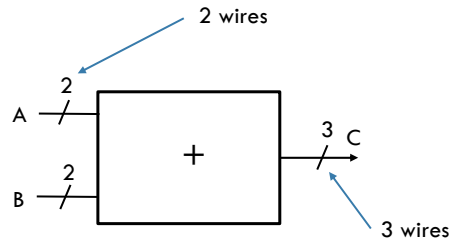
Berkeley



22

Example CL Block

- 2-bit adder. Takes two 2-bit integers and produces 3-bit result.



A1	A0	B1	B0	C2	C1	C0
0	0	0	0	0	0	0
0	0	0	1	0	0	1
0	0	1	0	0	1	0
0	0	1	1	0	1	1
0	1	0	0	0	0	1
0	1	0	1	0	1	0
0	1	1	0	0	1	1
0	1	1	1	1	0	0
1	0	0	0	0	1	0
1	0	0	1	0	1	1
1	0	1	0	1	0	0
1	0	1	1	1	0	1
1	1	0	0	0	1	1
1	1	0	1	1	0	0
1	1	1	0	1	0	1
1	1	1	1	1	1	0

- Think about truth table for 32-bit adder. It's possible to write out, but it might take a while!

Theorem:

Any combinational logic function can be implemented as a network of simple logic gates.

EECS151/251A L02 DESIGN

23

23

Quiz

Total number of possible truth tables with 4 inputs is:

- 4
- 16
- 256
- 16,384
- 65,536
- None of the above

www.yellkey.com/foot

EECS151/251A L02 DESIGN

24

24

Peer Instruction

Total number of possible truth tables with 4 inputs is:

- a) 4
- b) 16
- c) 256
- d) 16,384
- e) 65,536
- f) None of the above

www.yellkey.com/leg

EECS151/251A L02 DESIGN

25 Berkeley UNIVERSITY OF CALIFORNIA

25

Logic Circuit

- A logic gate can be implemented in different ways

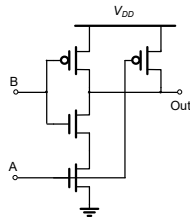
NAND

$$\text{Out} = \overline{A \cdot B}$$

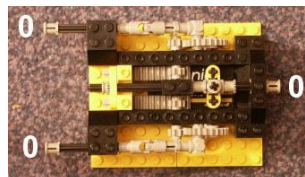
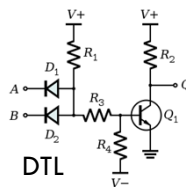


A	B	Out
0	0	1
0	1	1
1	0	1
1	1	0

CMOS



Sizing of transistors (W/L) in CMOS changes properties (delay, power, size) of a logic gate



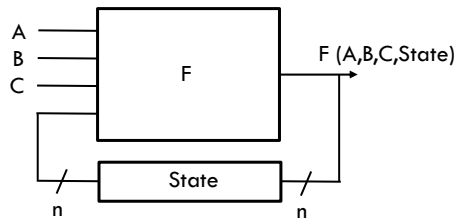
Mechanical LEGO logic gates. A clockwise rotation represents a binary "one" while a counter-clockwise rotation represents a binary "zero."

EECS151/251A L02 DESIGN

26 Berkeley UNIVERSITY OF CALIFORNIA

26

Sequential Logic Blocks



- Output is a function of both the current inputs and the state.
- State represents the memory.
- State is a function of previous inputs.
- In synchronous digital systems, state is updated on each clock tick.

EECS151/251A L02 DESIGN

27 Berkeley UNIVERSITY OF CALIFORNIA

27

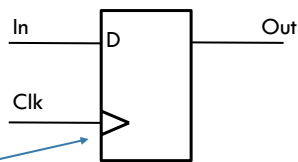
Flip-Flop as A Sequential Circuit

- Synchronous state element transfers its input to the output on a rising (or, rarely, falling) clock edge

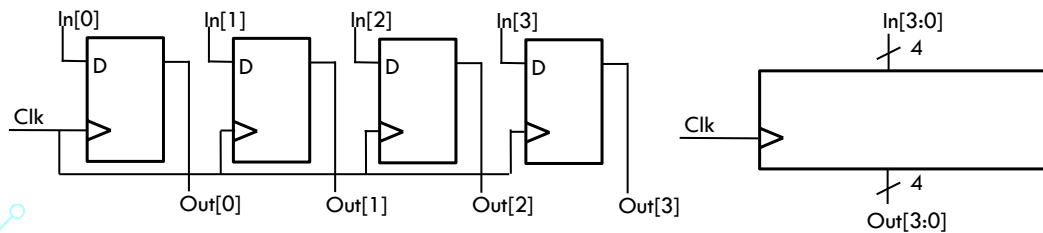
- Flip-flop

- Rising edge

Signifies 'edge triggered'



- 4-bit register



EECS151/251A L02 DESIGN

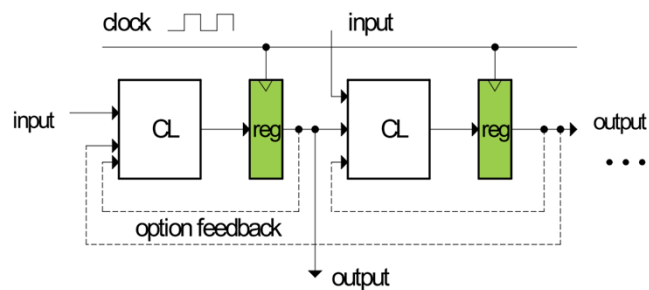
28 Berkeley UNIVERSITY OF CALIFORNIA

28

Register Transfer Level Abstraction (RTL)

Any synchronous digital circuit can be represented with:

- Combinational Logic (CL) blocks, plus
- State elements (registers or memories)
- Clock orchestrates *sequencing* of CL operations



- State elements are combined with CL blocks to control the flow of data.

EECS151/251A L02 DESIGN

29 Berkeley UNIVERSITY OF CALIFORNIA

29

Administrivia

- Labs and discussions start this week
- Lab 1 posted, please start it before coming to the lab session
- Lab 2 is more involved
 - Be prepared
 - Verilog primer
- Homework 1 posted this week, due next Friday
 - Start early

EECS151/251A L02 DESIGN

30 Berkeley UNIVERSITY OF CALIFORNIA

30



Design Metrics: Robustness

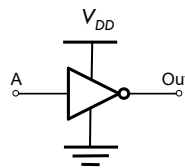
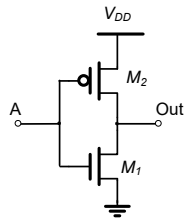
EECS151/251A L02 DESIGN

31 Berkeley UNIVERSITY OF CALIFORNIA

31

What Makes Circuits Digital?

- Chips are noisy
- Supply noise will appear at the output of the logic gate



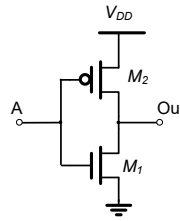
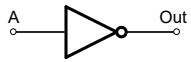
- The following logic gate should still interpret its inputs as 0s and 1s
- This necessary property is called "Restoration" or "Regeneration"
- A lot of money was spent in the past to unsuccessfully make logic out of non-regenerative gates
 - Some of emerging CMOS replacements don't have gain...

EECS151/251A L02 DESIGN

32 Berkeley UNIVERSITY OF CALIFORNIA

32

Beneath the Digital Abstraction



- Logic levels:
 - Mapping a continuous voltage onto a discrete binary logic variable
 - Low (0): $[0, V_L]$
 - High (1): $[V_H, V_{DD}]$
 - V_L, V_H : nominal voltage levels

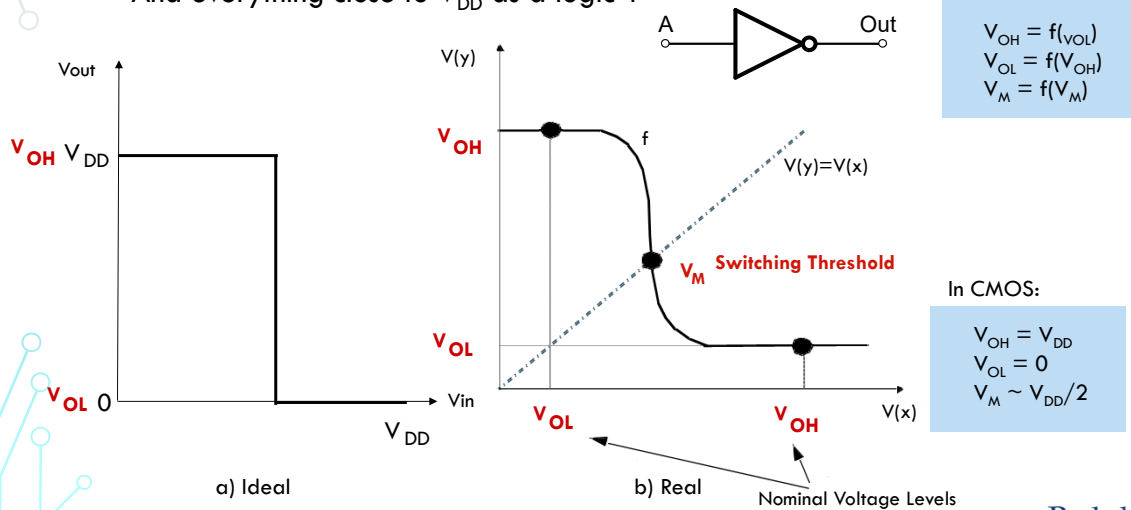
EECS151 L02 DESIGN

33 Berkeley UNIVERSITY OF CALIFORNIA

33

Voltage Transfer Characteristic

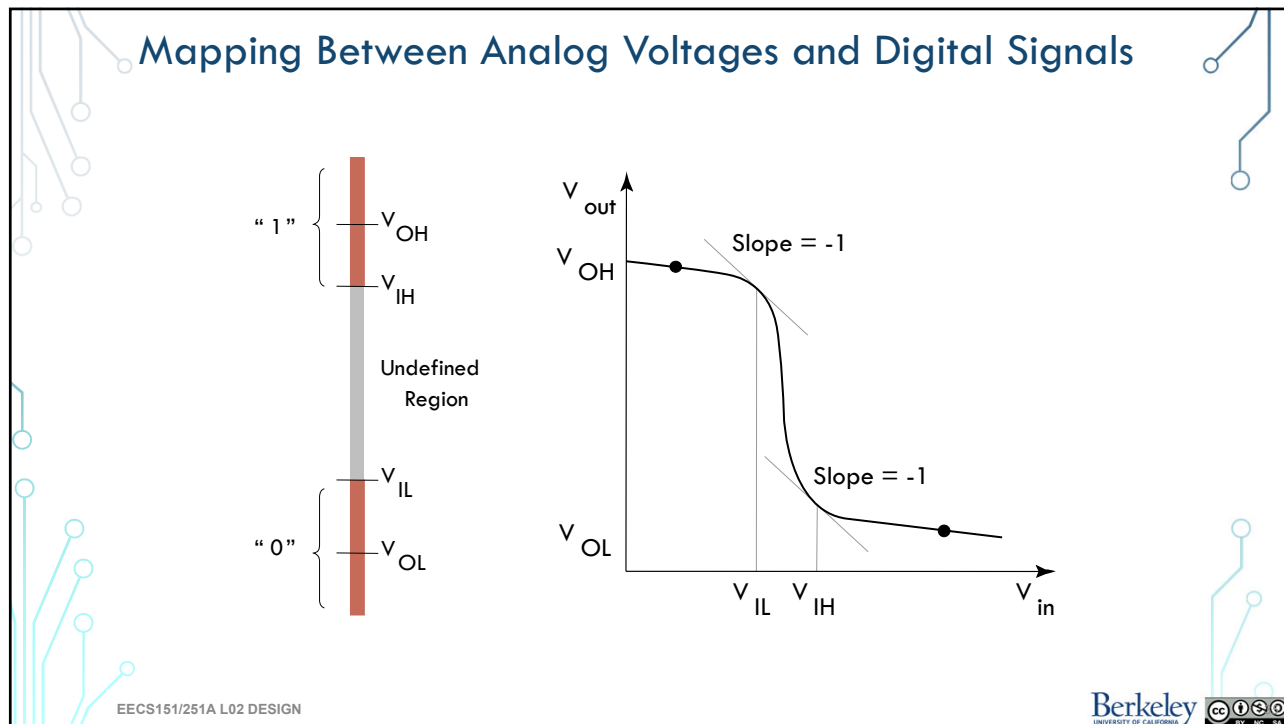
- A gate should interpret everything that is close to 0V as a logic 0
 - And everything close to V_{DD} as a logic 1



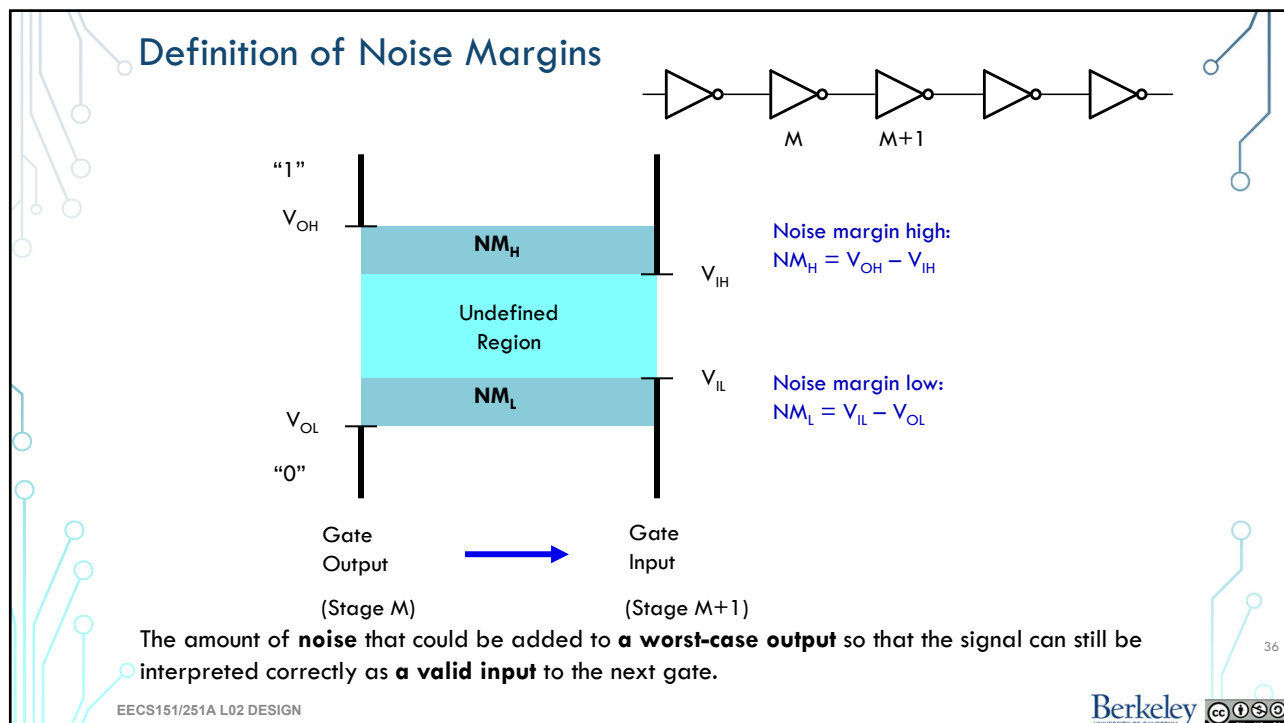
EECS151/251A L02 DESIGN

34 Berkeley UNIVERSITY OF CALIFORNIA

34



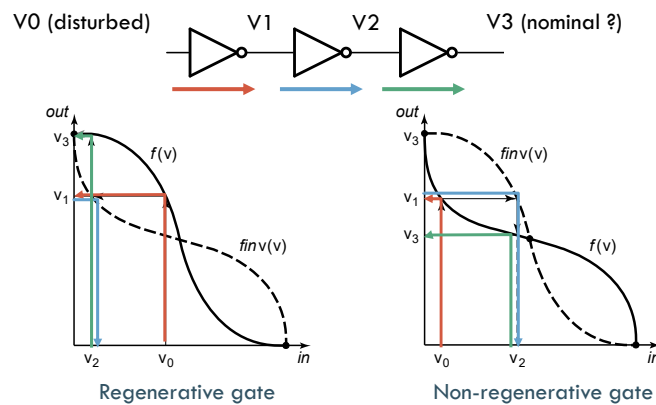
35



36

Regenerative Property

- Ensures that a **disturbed** signal gradually **regenerates** one of the **nominal voltage levels** after passing through a few logical stages.
 - Look for a sharp transition in voltage transfer characteristics.



EECS151/251A L02 DESIGN

 Berkeley
 UNIVERSITY OF CALIFORNIA

37

37

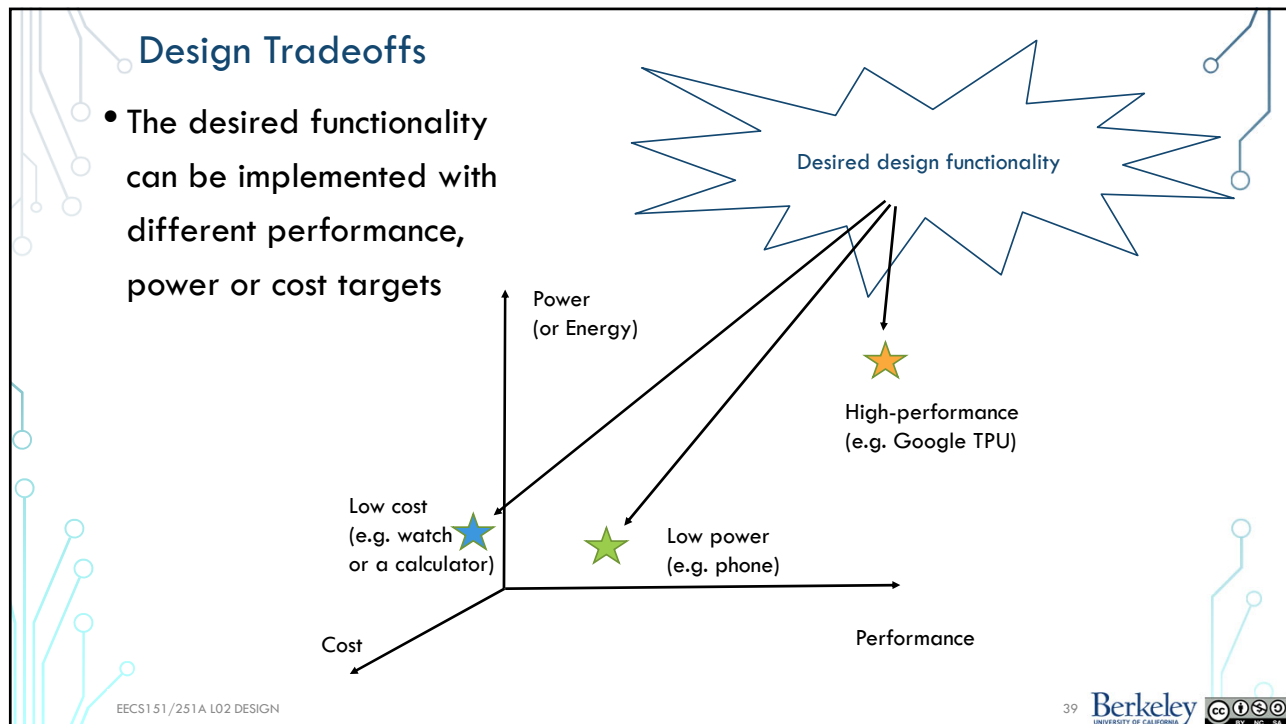


Design Metrics: Performance

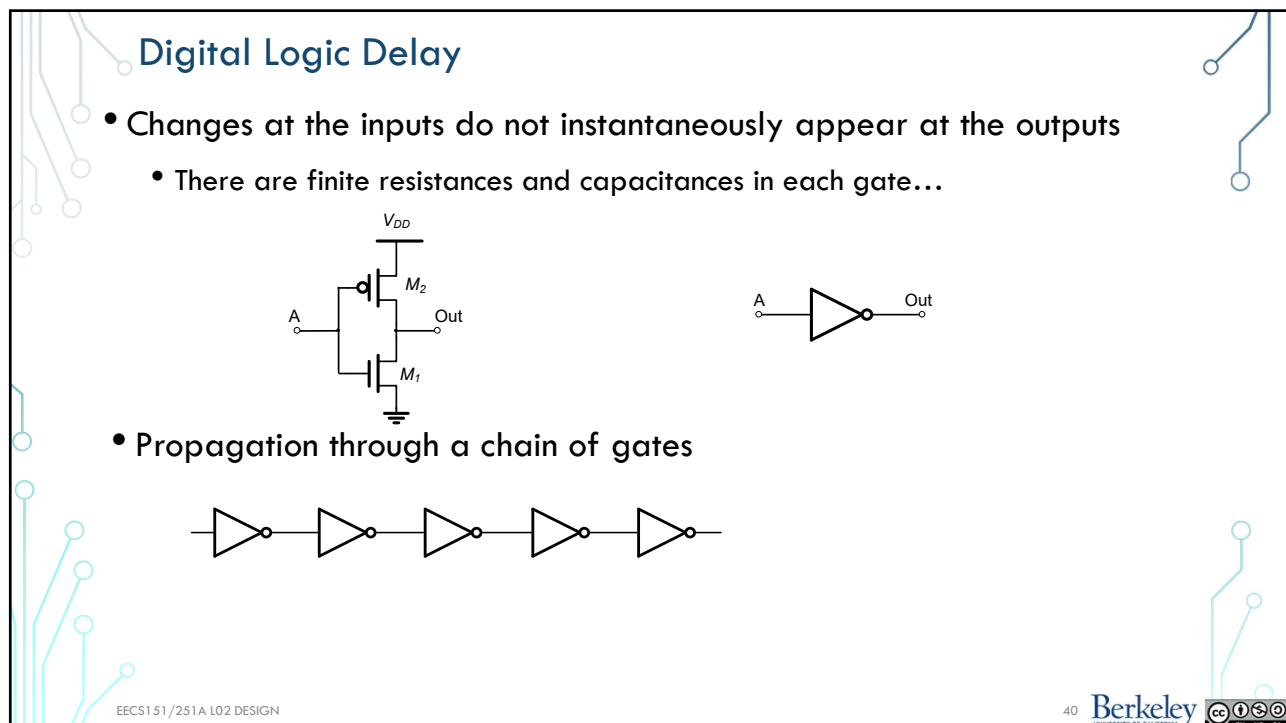
EECS151/251A L02 DESIGN

 38 Berkeley
 UNIVERSITY OF CALIFORNIA

38

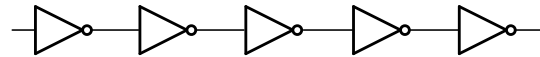


39

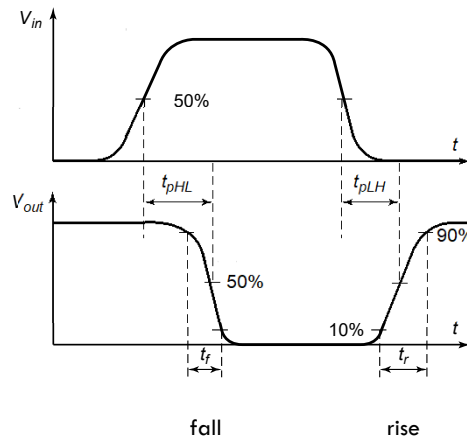


40

Delay Definitions



- Delay calculations need to be additive
 - Calculate the delay from the same point in the waveform



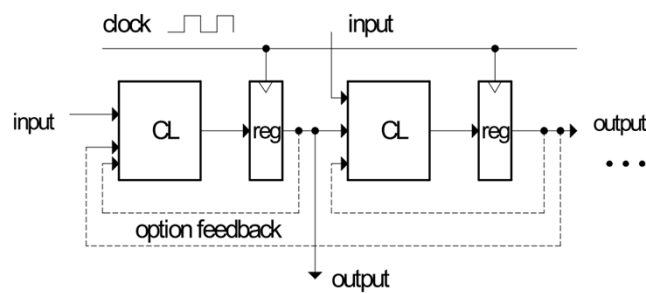
EECS151/251A L02 DESIGN

41

41

Digital Logic Timing

- The longest propagation delay through CL blocks sets the maximum clock frequency



- To increase clock rate:
 - Find the longest path
 - Make it faster

EECS151/251A L02 DESIGN

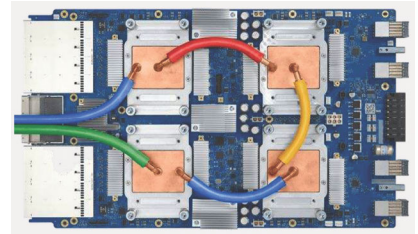
42

42

Performance

- Throughput

- Number of tasks performed in a unit of time (operations per second)
- E.g. Google TPUv3 board performs 420 TFLOPS (10^{12} floating-point operations per second, where a floating point operation is BFLOAT16)
- Watch out for 'op' definitions – can be a 1-b ADD or a double-precision FP add (or more complex task)
- Peak vs. average throughput



- Latency

- How long does a task take from start to finish
- E.g. facial recognition on a phone takes 10's of ms
- Sometime expressed in terms of clock cycles
- Average vs. 'tail' latency

EECS151/251A L02 DESIGN

43 Berkeley UNIVERSITY OF CALIFORNIA

43



Design Metrics: Energy and Power

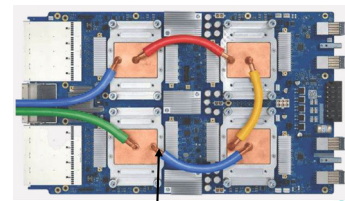
EECS151/251A L02 DESIGN

44 Berkeley UNIVERSITY OF CALIFORNIA

44

Energy and Power

- Energy (in joules (J))
 - Needed to perform a task
 - Add two numbers or fetch a datum from memory
 - (or fetch two numbers, add them and store in memory)
 - Active and standby
 - Battery stores certain amount of energy (in $Ws = J$ or Wh)
 - That is what utility charges for (in kWh)
- Power (in watts (W))
 - Energy dissipated in time ($W = J/s$)
 - Sets cooling requirements
 - Heat spreader, size of a heat sink, forced air, liquid, ...



Liquid

EECS151/251A L02 DESIGN

45 Berkeley UNIVERSITY OF CALIFORNIA

45



Design Metrics:
Cost

EECS151/251A L02 DESIGN

46 Berkeley UNIVERSITY OF CALIFORNIA

46

Cost

- Non-recurring engineering (NRE) costs
- Cost to develop a design (product)
 - Amortized over all units shipped
 - E.g. \$20M in development adds \$.20 to each of 100M units
- Recurring costs
 - Cost to manufacture, test and package a unit
 - Processed wafer cost is ~10k (around 16nm node) which yields:
 - 1 Cerebras chip
 - 50-100 large FPGAs or GPUs
 - 200 laptop CPUs
 - >1000 cell phone SoCs

$$\text{cost per IC} = \text{variable cost per IC} + \frac{\text{fixed cost}}{\text{volume}}$$

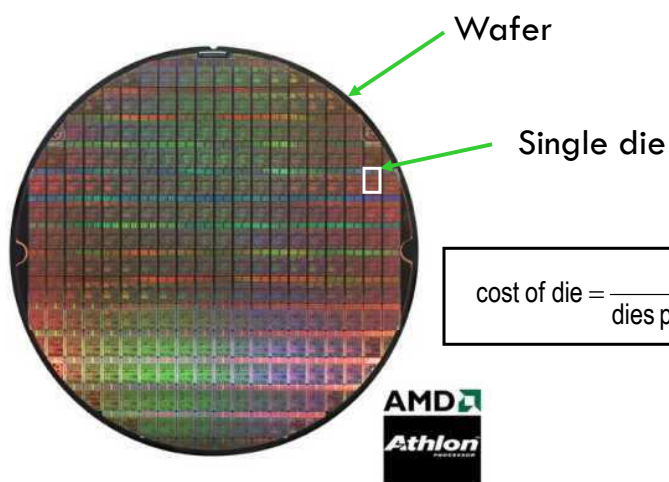
$$\text{variable cost} = \frac{\text{cost of die} + \text{cost of die test} + \text{cost of packaging}}{\text{final test yield}}$$

EECS151/251A L02 DESIGN

47 Berkeley UNIVERSITY OF CALIFORNIA

47

Die Cost



$$\text{cost of die} = \frac{\text{cost of wafer}}{\text{dies per wafer} * \text{die yield}}$$

From: <http://www.amd.com>

EECS151/251A L02 DESIGN

Berkeley UNIVERSITY OF CALIFORNIA

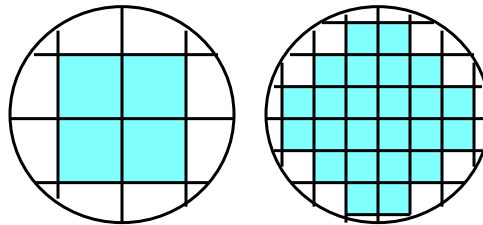
48

Yield

$$Y = \frac{\text{No. of good chips per wafer}}{\text{Total number of chips per wafer}} \times 100\%$$

$$\text{Die cost} = \frac{\text{Wafer cost}}{\text{Dies per wafer} \times \text{Die yield}}$$

$$\text{Dies per wafer} = \frac{\pi \times (\text{wafer diameter}/2)^2}{\text{die area}} - \frac{\pi \times \text{wafer diameter}}{\sqrt{2} \times \text{die area}}$$



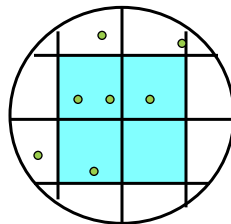
EECS151/251A L02 DESIGN

 Berkeley
 UNIVERSITY OF CALIFORNIA

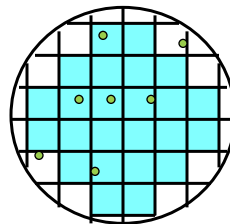

49

Defects

Yield = 0.25



Yield = 0.76



$$\text{die yield} = \left(1 + \frac{\text{defects per unit area} \times \text{die area}}{\alpha} \right)^{-\alpha}$$

 α is approximately 3

$$\text{die cost} = f(\text{die area})^4$$

EECS151/251A L02 DESIGN

 Berkeley
 UNIVERSITY OF CALIFORNIA


50

Summary

- The design process involves traversing the abstraction layers of specification, modeling, architecture, RTL design and physical implementation
- Tests follow the design refinements
- Targets are processors, FPGAs or ASICs
- Automated design flows help manage the complexity
- Optimize for performance, energy and cost