



# EECS151 : Introduction to Digital Design and ICs

## Lecture 2 – Design Process

Bora Nikolić

At HotChips'19 Cerebras announced the largest chip in the world at 8.5 in x 8.5in with 1.2 trillion transistors, and 15kW of power, aimed for training of deep-learning neural networks

At HotChips'21 they showed the next version in 7nm CMOS, with >2x transistor count

46,225 mm<sup>2</sup> silicon  
2.6 Trillion transistors  
850,000 AI optimized cores  
40 Gigabytes on-chip memory  
20 Petabyte/s memory bandwidth  
220 Petabit/s fabric bandwidth  
7nm Process technology at TSMC

EECS151/251A L02 DESIGN



Sean Lie,  
HotChips'21




## Review

- Moore's law is slowing down
  - There are continued improvements in technology, but at a slower pace
- Dennard's scaling has ended a decade ago
  - All designs are now power limited
- Specialization and customization provides added performance
  - Under power constraints and stagnant technology
- Design costs are high
  - Methodology and better reuse to rescue!
  - Abstraction, modularity, regularity are the keys
    - And creativity!

EECS151/251A L02 DESIGN




## Putting it in Perspective



Performance gains over the past decade

Lisa Su, HotChips'19 keynote

EECS151/251A L02 DESIGN




## Digital Logic

EECS151/251A L02 DESIGN



## Implementing Digital Systems

- Digital systems implement a set of Boolean equations




- How do we implement a digital system?

EECS151/251A L02 DESIGN

## Modern (Mostly) Digital System-On-A-Chip

- 4x 'Firestorm' Large CPUs
- 4x 'Icestorm' Small CPUs
- GPU
- Neural processing unit (NPU)
- Lots of memory
- DDR memory interfaces



Technisights

- 5nm CMOS
- Up to 2.49GHz



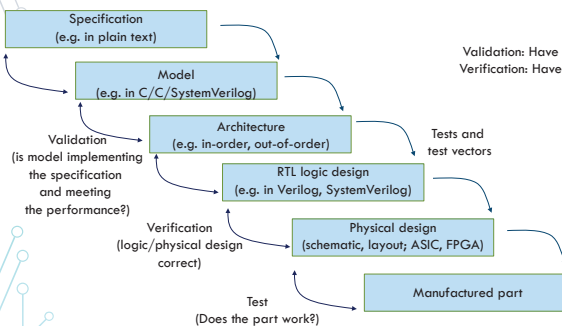
By Haxivik

<https://commons.wikimedia.org/wiki/File:Index.php?curid=96026688>



## Design Process

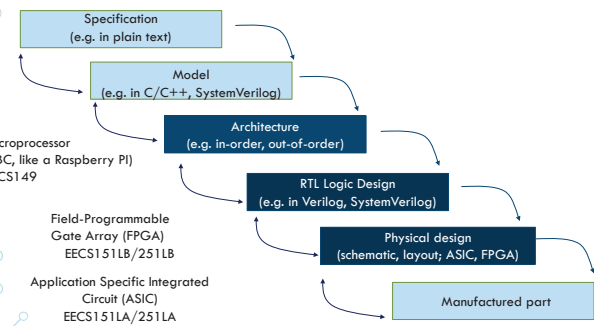
- Design through layers of abstractions



EECS151/251A L02 DESIGN

## Design Abstractions in EECS151/251A

- Design through layers of abstractions

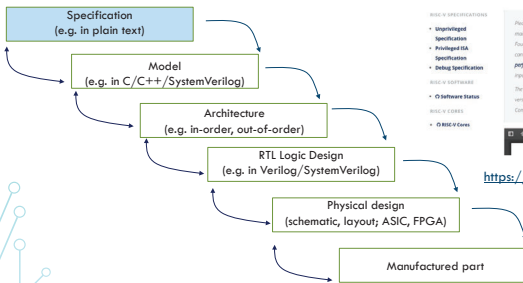


EECS151 L02 DESIGN



## Example: RISC-V Design Process

- Design through layers of abstractions

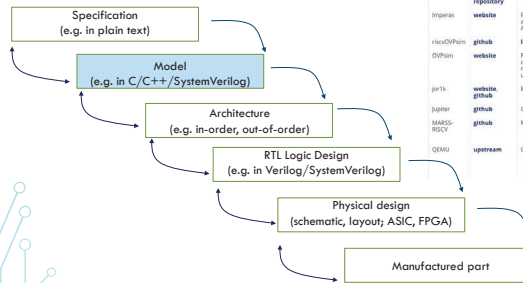


BEC5151 102 DESIGN



## Example: RISC-V Design Process

- Design through layers of abstractions



BEC5151 102 DESIGN



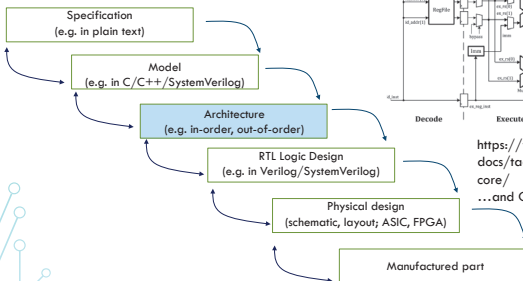
### Simulators

Name	Links	License	Maintainers
BEST-BISE-RISC-V	github	BSD-3-Clause	MHIRES Technologies
FireSim	website, mailing list, github, SCA 2017 Paper	BSD	Sagar Karandikar, Howard Miao, Christopher Kim, David Burch, Alan Alami, Berkeley Architecture Research
gem5	github	BSD-style	Alec Rose (University of Virginia)
Impress	website	Proprietary, models available under Apache 2.0	Impress
riscvO3sim	github	License	Impress
O3sim	github	Free for non-commercial use, models available under Apache 2.0	Impress
perls	website, github	BSD-3-Clause	Selvakumar Marthi
Juster	github	GPL-3.0	Andreas Carlsson
ARMV8-RISC-V	github	MIT	Ben Baer, University of California, Berkeley
QEMU	upstream	GPL	Sagar Karandikar (University of California, Berkeley), Srinivas Aravamudan (University of Pennsylvania), Alex Avdoyan, Stefan Oefner and Michael Clark (QEMU)

<https://riscv.org/software-status/#simulators>

## Example: RISC-V Design Process

- Design through layers of abstractions

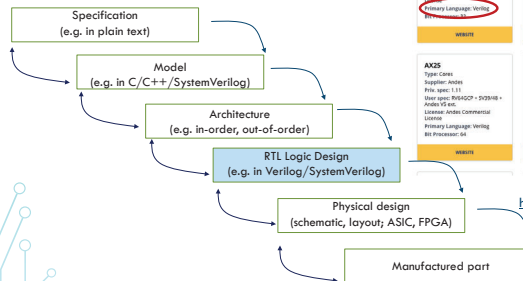


BEC5151 102 DESIGN



## Example: RISC-V Design Process

- Design through layers of abstractions



BEC5151 102 DESIGN

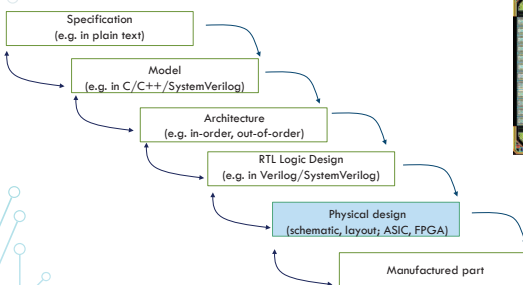


<b>AZ5</b> Type: Core Supplier: Andes Price: \$1.11 User spec: RV32GCP + RV32M + Andes VLE ext + Multi-core License: Andes Commercial License Primary Language: Verilog Bit Processor: 32	<b>AZ5MP</b> Type: Core Supplier: Andes Price: \$1.11 User spec: RV32GCP + RV32M + Andes VLE ext + Multi-core License: Andes Commercial License Primary Language: Verilog Bit Processor: 32	<b>Ariane</b> Type: Core Supplier: Armada, University of Bologna Price: \$1.11 User spec: RV32GCP + RV32M + Andes VLE ext + Multi-core License: Andes Commercial License Primary Language: Verilog Bit Processor: 32
<b>AX25</b> Type: Core Supplier: Andes Price: \$1.11 User spec: RV32GCP + RV32M + Andes VLE ext + Multi-core License: Andes Commercial License Primary Language: Verilog Bit Processor: 32	<b>AX25MP</b> Type: Core Supplier: Andes Price: \$1.11 User spec: RV32GCP + RV32M + Andes VLE ext + Multi-core License: Andes Commercial License Primary Language: Verilog Bit Processor: 32	<b>Berkeley Out-of-Order Machine (BOOM)</b> Type: Core Supplier: Berkeley, OpenMata, USB Bar Price: \$1.11 User spec: RV32GCP + RV32M + Andes VLE ext + Multi-core License: Andes Commercial License Primary Language: Verilog Bit Processor: 32

<https://riscv.org/risc-v-cores/>

## Example: RISC-V Design Process

- Design through layers of abstractions

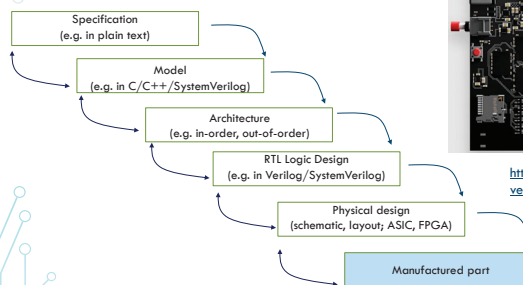


BEC5151 102 DESIGN




## Example: RISC-V Design Process

- Design through layers of abstractions




BEC5151 102 DESIGN

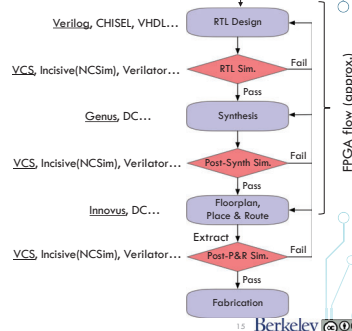


<https://www.sifive.com/boards/hifi-ve-unleashed>

## RTL → Physical Design



- Labs focus on a process of translating RTL to physical ASIC or FPGA by using industry-standard tools.
- Explores the entire design stack.




BEC5151 102 DESIGN



## Open-Source Flows


- Skywater 130nm is an open-source design kit
- OpenROAD (UCSD) and OpenLane (eFabless) are open-source design flows
  - Work with Sky130
  - A version of ASIC labs can target Sky130nm



<https://github.com/efabless/openlane>

BEC5151 102 DESIGN

<https://github.com/efabless/caravel>







<https://efabless.com/projects/35>




## Boolean Logic and Logic Gates (From CS61C/EE16B)

### Logic gates

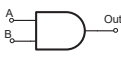
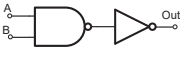
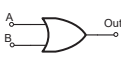
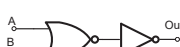
Name	Boolean equation	Symbol	NOT/INV
NOT or Inverter	$Out = \bar{A}$		Out
Buffer	$Out = A$		Out
NAND	$Out = \overline{A \cdot B}$		Out
NOR	$Out = \overline{A + B}$		Out

• In CMOS, basic logic gates are inverting



## Boolean Logic in A Nutshell



## More Logic Gates

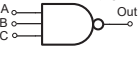

Name	Boolean equation	Symbol	Truth table															
AND	$Out = A \cdot B$		<table><tr><th>A</th><th>B</th><th>Out</th></tr><tr><td>0</td><td>0</td><td>0</td></tr><tr><td>0</td><td>1</td><td>0</td></tr><tr><td>1</td><td>0</td><td>0</td></tr><tr><td>1</td><td>1</td><td>1</td></tr></table>	A	B	Out	0	0	0	0	1	0	1	0	0	1	1	1
A	B	Out																
0	0	0																
0	1	0																
1	0	0																
1	1	1																
In CMOS																		
OR	$Out = A + B$		<table><tr><th>A</th><th>B</th><th>Out</th></tr><tr><td>0</td><td>0</td><td>0</td></tr><tr><td>0</td><td>1</td><td>1</td></tr><tr><td>1</td><td>0</td><td>1</td></tr><tr><td>1</td><td>1</td><td>1</td></tr></table>	A	B	Out	0	0	0	0	1	1	1	0	1	1	1	1
A	B	Out																
0	0	0																
0	1	1																
1	0	1																
1	1	1																
In CMOS																		

## More Logic Gates

Name	Boolean equation	Symbol	Truth table															
Exclusive OR XOR	$Out = A \oplus B$		<table><tr><th>A</th><th>B</th><th>Out</th></tr><tr><td>0</td><td>0</td><td>0</td></tr><tr><td>0</td><td>1</td><td>1</td></tr><tr><td>1</td><td>0</td><td>1</td></tr><tr><td>1</td><td>1</td><td>0</td></tr></table>	A	B	Out	0	0	0	0	1	1	1	0	1	1	1	0
A	B	Out																
0	0	0																
0	1	1																
1	0	1																
1	1	0																
Exclusive NOR XNOR	$Out = \overline{A \oplus B}$		<table><tr><th>A</th><th>B</th><th>Out</th></tr><tr><td>0</td><td>0</td><td>1</td></tr><tr><td>0</td><td>1</td><td>0</td></tr><tr><td>1</td><td>0</td><td>0</td></tr><tr><td>1</td><td>1</td><td>1</td></tr></table>	A	B	Out	0	0	1	0	1	0	1	0	0	1	1	1
A	B	Out																
0	0	1																
0	1	0																
1	0	0																
1	1	1																

• XOR and XNOR are both inverting and non-inverting


## Multi-Input Gates

3-Input NAND NAND3	Boolean equation $Out = \overline{A \cdot B \cdot C}$	
And-Or-Invert AOI21	Boolean equation $Out = \overline{A \cdot B + C}$	

• Single gate in modern CMOS usually doesn't have more than 3-4 inputs

## Combinational Logic (CL) Blocks

Example four-input function:




- Output a function only of the current inputs (no history).
- Truth-table representation of function. Output is explicitly specified for each input combination.
- In general, CL blocks have more than one output signal, in which case, the truth-table will have multiple output columns.

A	B	C	D	Out
0	0	0	0	F(0,0,0,0)
0	0	0	1	F(0,0,0,1)
0	0	1	0	F(0,0,1,0)
0	0	1	1	F(0,0,1,1)
0	1	0	0	F(0,1,0,0)
0	1	0	1	F(0,1,0,1)
0	1	1	0	F(0,1,1,0)
0	1	1	1	F(0,1,1,1)
1	0	0	0	F(1,0,0,0)
1	0	0	1	F(1,0,0,1)
1	0	1	0	F(1,0,1,0)
1	0	1	1	F(1,0,1,1)
1	1	0	0	F(1,1,0,0)
1	1	0	1	F(1,1,0,1)
1	1	1	0	F(1,1,1,0)
1	1	1	1	F(1,1,1,1)

## Example CL Block

• 2-bit adder. Takes two 2-bit integers and produces 3-bit result.



• Think about truth table for 32-bit adder. It's possible to write out, but it might take a while!

A1	A0	B1	B0	C2	C1	C0
0	0	0	0	0	0	0
0	0	0	1	0	0	1
0	0	1	0	0	1	0
0	0	1	1	0	1	1
0	1	0	0	0	0	1
0	1	0	1	0	1	0
0	1	1	0	0	1	1
0	1	1	1	1	0	0
1	0	0	0	0	1	0
1	0	0	1	0	1	1
1	0	1	0	1	0	0
1	0	1	1	1	0	1
1	1	0	0	0	1	1
1	1	0	1	1	0	0
1	1	1	0	1	0	1
1	1	1	1	1	1	0

## Quiz

Total number of possible truth tables with 4 inputs is:

- 4
- 16
- 256
- 16,384
- 65,536
- None of the above

[www.yellkey.com/foot](http://www.yellkey.com/foot)

Theorem:  
Any combinational logic function can be implemented as a network of simple logic gates.

## Peer Instruction


Total number of possible truth tables with 4 inputs is:

- a) 4
- b) 16
- c) 256
- d) 16,384
- e) 65,536
- f) None of the above

[www.yellkey.com/leg](http://www.yellkey.com/leg)

## Logic Circuit

- A logic gate can be implemented in different ways




Sizing of transistors (W/L) in CMOS changes properties (delay, power, size) of a logic gate



Mechanical LEGO logic gates. A clockwise rotation represents a binary "one" while a counter-clockwise rotation represents a binary "zero."


## Sequential Logic Blocks




- Output is a function of both the current inputs and the state.
- State represents the memory.
- State is a function of previous inputs.
- In synchronous digital systems, state is updated on each clock tick.

## Flip-Flop as A Sequential Circuit

- Synchronous state element transfers its input to the output on a rising (or, rarely, falling) clock edge




- 4-bit register



## Register Transfer Level Abstraction (RTL)

Any synchronous digital circuit can be represented with:

- Combinational Logic (CL) blocks, plus
- State elements (registers or memories)
- Clock orchestrates sequencing of CL operations



- State elements are combined with CL blocks to control the flow of data.

## Administrivia

- Labs and discussions start this week
- Lab 1 posted, please start it before coming to the lab session
- Lab 2 is more involved
  - Be prepared
  - Verilog primer
- Homework 1 posted this week, due next Friday
  - Start early

## Design Metrics: Robustness




## What Makes Circuits Digital?

- Chips are noisy
- Supply noise will appear at the output of the logic gate



- The following logic gate should still interpret its inputs as 0s and 1s
- This necessary property is called "Restoration" or "Regeneration"
- A lot of money was spent in the past to unsuccessfully make logic out of non-regenerative gates
  - Some of emerging CMOS replacements don't have gain...

## Beneath the Digital Abstraction




### Logic levels:

- Mapping a continuous voltage onto a discrete binary logic variable
- Low (0):  $[0, V_L]$
- High (1):  $[V_H, V_{DD}]$
- $V_L, V_H$ : nominal voltage levels

## Voltage Transfer Characteristic


- A gate should interpret everything that is close to 0V as a logic 0
- And everything close to  $V_{DD}$  as a logic 1



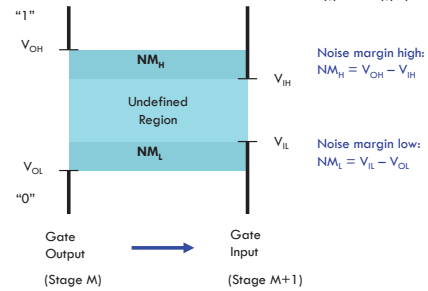
In CMOS:

$$\begin{aligned} V_{OH} &= V_{DD} \\ V_{OL} &= 0 \\ V_M &= V_{DD}/2 \end{aligned}$$

## Mapping Between Analog Voltages and Digital Signals




## Definition of Noise Margins



The amount of **noise** that could be added to a **worst-case output** so that the signal can still be interpreted correctly as a **valid input** to the next gate.

## Regenerative Property

- Ensures that a **disturbed** signal gradually **regenerates** one of the **nominal voltage levels** after passing through a few logical stages.
- Look for a sharp transition in voltage transfer characteristics.

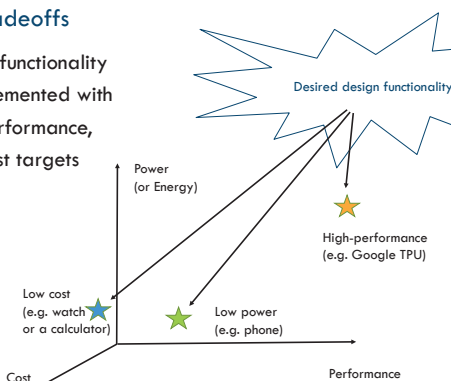


## Design Metrics: Performance




## Design Tradeoffs

- The desired functionality can be implemented with different performance, power or cost targets



## Digital Logic Delay

- Changes at the inputs do not instantaneously appear at the outputs
- There are finite resistances and capacitances in each gate...




- Propagation through a chain of gates





## Delay Definitions


- Delay calculations need to be additive
  - Calculate the delay from the same point in the waveform



EECS151/251A L02 DESIGN

## Digital Logic Timing

- The longest propagation delay through CL blocks sets the maximum clock frequency




- To increase clock rate:
  - Find the longest path
  - Make it faster

EECS151/251A L02 DESIGN

## Performance


- Throughput
  - Number of tasks performed in a unit of time (operations per second)
  - E.g. Google TPUv3 board performs 420 TFLOPS ( $10^{12}$  floating-point operations per second, where a floating point operation is BFLOAT16)
  - Watch out for 'op' definitions – can be a 1-b ADD or a double-precision FP add (or more complex task)
  - Peak vs. average throughput
- Latency
  - How long does a task take from start to finish
  - E.g. facial recognition on a phone takes 10's of ms
  - Sometime expressed in terms of clock cycles
  - Average vs. 'tail' latency



Design Metrics:  
Energy and Power

## Energy and Power

- Energy (in joules (J))
  - Needed to perform a task
  - Add two numbers or fetch a datum from memory
    - (or fetch two numbers, add them and store in memory)
  - Active and standby
  - Battery stores certain amount of energy (in  $Ws = J$  or  $Wh$ )
  - That is what utility charges for (in  $kWh$ )
- Power (in watts (W))
  - Energy dissipated in time ( $W = J/s$ )
  - Sets cooling requirements
    - Heat spreader, size of a heat sink, forced air, liquid, ...



Liquid



Design Metrics:  
Cost


## Cost

- Non-recurring engineering (NRE) costs
- Cost to develop a design (product)
  - Amortized over all units shipped
  - E.g. \$20M in development adds \$.20 to each of 100M units
- Recurring costs
  - Cost to manufacture, test and package a unit
  - Processed wafer cost is ~10k (around 16nm node) which yields:
    - 1 Cerebras chip
    - 50-100 large FPGAs or GPUs
    - 200 laptop CPUs
    - >1000 cell phone SoCs

$$\text{cost per IC} = \text{variable cost per IC} + \frac{\text{fixed cost}}{\text{volume}}$$

$$\text{variable cost} = \frac{\text{cost of die} + \text{cost of die test} + \text{cost of packaging}}{\text{final test yield}}$$

## Die Cost



$$\text{cost of die} = \frac{\text{cost of wafer}}{\text{dies per wafer} * \text{die yield}}$$



From: <http://www.amd.com>

EECS151/251A L02 DESIGN

Berkeley

EECS151/251A L02 DESIGN


Berkeley

## Yield

$$Y = \frac{\text{No. of good chips per wafer}}{\text{Total number of chips per wafer}} \times 100\%$$


$$\text{Die cost} = \frac{\text{Wafer cost}}{\text{Dies per wafer} \times \text{Die yield}}$$

$$\text{Dies per wafer} = \frac{\pi \times (\text{wafer diameter}/2)^2}{\text{die area}} = \frac{\pi \times \text{wafer diameter}}{\sqrt{2} \times \text{die area}}$$




## Defects

Yield = 0.25



Yield = 0.76



$$\text{die yield} = \left( 1 + \frac{\text{defects per unit area} \times \text{die area}}{\alpha} \right)^{-\alpha}$$

$\alpha$  is approximately 3

$$\text{die cost} = f(\text{die area})^4$$

## Summary

- The design process involves traversing the abstraction layers of specification, modeling, architecture, RTL design and physical implementation
- Tests follow the design refinements
- Targets are processors, FPGAs or ASICs
- Automated design flows help manage the complexity
- Optimize for performance, energy and cost