

EECS151 : Introduction to Digital Design and ICs


Lecture 8 – RISC-V ISA

Bora Nikolić



September 21, 2021, EET Asia
RISC-V to Shake Up \$8.6B Semiconductor IP Market

RISC-V is now a rising star in the industry, largely due to its open-source advantage, better power consumption performance promise, reliable security functions and lower political risk impact yet.



<https://www.eetasia.com/risc-v-to-shake-up-8-6b-semiconductor-ip-market/>

EECS151 L08 RISC-V

Nikolić, Fall 2021



Review

- Finite state machines: Common example of sequential logic
 - Moore's machine: Output depends only on the current state
 - Mealy's machine: Output depends on the current state and the input
- Large state machines can be factored
- Common Verilog patterns for FSMs
- Common job interview questions ☺

Nikolić, Fall 2021



Building a RISC-V Processor



Berkeley RISC-V ISA

www.riscv.org

- An open, license-free ISA
 - Runs GCC, LLVM, Linux distributions, ...
 - RV32, RV64, and RV128 variants for 32b, 64b, and 128b address spaces
- Originally developed for teaching classes at Berkeley, now widely adopted
- Base ISA only ~40 integer instructions
- Extensions provide full general-purpose ISA, including IEEE-754/2008 floating-point
- Designed for extension, customization
- Developed at UC Berkeley, now maintained by RISC-V Foundation
- Open and commercial implementations
- RISC-V ISA, datapath, and control covered in CS61C; summarized here



EECS151 L08 RISC-V

Nikolić, Fall 2021




EECS151 L08 RISC-V

Nikolić, Fall 2021



RISC-V Processor Design

- Spec: Unprivileged ISA, RV32I (and a look at RV64I)



The RISC-V Instruction Set Manual
Volume 1: Unprivileged ISA
Document Version 20190808, Hao-Ran He
Editors: Andrew Waterman*, Kate Asanasi*
*CS Division, EECS Department, University of California, Berkeley
andrew@cs.berkeley.edu, kate@cs.berkeley.edu
June 5, 2020


- Tests provided as a part of the project
- Architecture: Single-cycle and pipelined in-order processor
 - Expanded from CS61C

EECS151 L08 RISC-V

Nikolić, Fall 2021



One-Instruction-Per-Cycle RISC-V Machine



- On every tick of the clock, the computer executes one instruction
- Current state outputs drive the inputs to the combinational logic, whose outputs settle at the values of the state before the next clock edge
- At the rising clock edge, all the state elements are updated with the combinational logic outputs, and execution moves to the next clock cycle

EECS151 L08 RISC-V

Nikolić, Fall 2021



State Required by RV32I ISA

Each instruction reads and updates this state during execution:

- Registers ($x0 \dots x31$)
 - Register file (regfile) **Reg** holds 32 registers x 32 bits/register: **Reg**[0] . . **Reg**[31]
 - First register read specified by *rs1* field in instruction
 - Second register read specified by *rs2* field in instruction
 - Write register (destination) specified by *rd* field in instruction
 - $x0$ is always 0 (writes to **Reg**[0] are ignored)
- Program counter (**PC**)
 - Holds address of current instruction
- Memory (**MEM**)
 - Holds both instructions & data, in one 32-bit byte-addressed memory space
 - We'll use separate memories for instructions (**IMEM**) and data (**DMEM**)
 - These are placeholders for instruction and data caches
 - Instructions are read (fetched) from instruction memory
 - Load/store instructions access data memory

EECS151 L08 RISC-V

Nikolić, Fall 2021



Stages of the Datapath : Overview

- Problem: A single, "monolithic" CL block that "executes an instruction" (performs all necessary operations beginning with fetching the instruction and completing with the register access) is too bulky and inefficient
- Solution: Break up the process of "executing an instruction" into stages, and then connect the stages to create the whole datapath
 - smaller stages are easier to design
 - easy to optimize (change) one stage without touching the others (modularity)

EECS151 L08 RISC-V

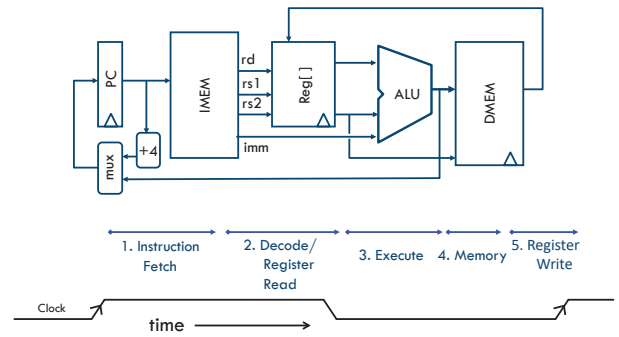
Nikolić, Fall 2021



Five Stages of the Datapath


- Stage 1: *Instruction Fetch (IF)*
- Stage 2: *Instruction Decode (ID)*
- Stage 3: *Execute (EX) - ALU (Arithmetic-Logic Unit)*
- Stage 4: *Memory Access (MEM)*
- Stage 5: *Write Back to Register (WB)*

Basic Phases of Instruction Execution



Datapath Components: Combinational

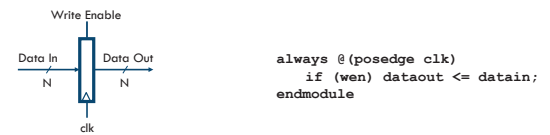
- Combinational Elements



- Storage Elements + Clocking Methodology
- Building Blocks

Datapath Elements: State and Sequencing (1/4)


- Register



- Write Enable:
 - Negated (or deasserted) (0): Data Out will not change
 - Asserted (1): Data Out will become Data In on positive edge of clock

Datapath Elements: State and Sequencing (2/4)

- Register file (regfile, RF) consists of 32 registers:
 - Two 32-bit output busses: busA and busB
 - One 32-bit input bus: busW
 - x0 is wired to 0
- Register is selected by:
 - RA (number) selects the register to put on busA (data)
 - RB (number) selects the register to put on busB (data)
 - RV (number) selects the register to be written via busW (data) when Write Enable is 1
- Clock input (clk)
 - CLK input is a factor ONLY during write operation
 - During read operation, behaves as a combinational logic block:
 - RA or RB valid \Rightarrow busA or busB valid after "access time."




Datapath Elements: State and Sequencing (3/4)

- Reg file in Verilog

```
module rv32i_regs (
    input clk, wen,
    input [4:0] rw,
    input [4:0] ra,
    input [4:0] rb,
    input [31:0] busw,
    output [31:0] busa,
    output [31:0] busb
);
    reg [31:0] regs [0:30];
    always @(posedge clk)
        if (wen) regs[rw] <= busw;
    assign busa = (ra == 5'd0) ? 32'd0 : regs[ra];
    assign busb = (rb == 5'd0) ? 32'd0 : regs[rb];
endmodule
```


x0 is zero



- How does RV64I register file look like?

Datapath Elements: State and Sequencing (4/4)

- "Magic" memory
 - One input bus: Data In
 - One output bus: Data Out
- Memory word is found by:
 - For Read: Address selects the word to put on Data Out
 - For Write: Set Write Enable = 1: address selects the memory word to be written via the Data In bus
- Clock input (CLK)
 - CLK input is a factor ONLY during write operation
 - During read operation, behaves as a combinational logic block:
 - Address valid \Rightarrow Data Out valid after "access time"
- Real memory later in the class



Review: Complete RV32I ISA

Open Reference Card									
Base Integer Instructions: RV32I									
Category	Name	Fast	RV32I Base	Category	Name	Fast	RV32I Base	Category	Name
Shifts	Shift Left Logical	R	SLL rd,rs1,rs2	Loads	Load Byte	I	LB rd,rs1,imm	Branches	Branch =
	Shift Left Logical Imm.	I	SLLI rd,rs1,shamt		Load Halfword	I	LH rd,rs1,imm		Branch <
	Shift Right Logical	R	SRL rd,rs1,rs2		Load Byte Unaligned	I	LBU rd,rs1,imm		Branch >
	Shift Right Logical Imm.	I	SRLI rd,rs1,shamt		Load Halfword Unaligned	I	LHU rd,rs1,imm		Branch <=
Arithmetic	Shift Right Arithmetic	R	SRA rd,rs1,rs2	Stores	Store Word	S	SW rs1,rs2,imm	Logical	XOR
	Shift Right Arithmetic Imm.	I	SRAI rd,rs1,shamt		Store Byte	S	SB rs1,rs2,imm		XOR Immediate
	ADD Immediate	I	ADDI rd,rs1,imm		Store Halfword	S	SH rs1,rs2,imm		OR
	SUB Immediate	I	SUBI rd,rs1,imm		Store Word	S	SW rs1,rs2,imm		OR Immediate
Logical	Load Upper Imm to PC	U	LUI PC,imm	Jump & Link	Jump	J	JAL rd,imm	Compare	Set <
	ADD Upper Imm to PC	U	AUIPC PC,imm		Jump & Link Imm.	J	JALR rd,rs1,imm		Set <=
	XOR	R	XOR rd,rs1,rs2		Jump & Link Imm.	J	JALR rd,rs1,imm		Set <= Imm Unaligned
	XOR Immediate	I	XORI rd,rs1,imm		Jump & Link Imm.	J	JALR rd,rs1,imm		Set < Imm Unaligned
AND	AND	R	AND rd,rs1,rs2	Synch	Environment Break	I	ECALL	Synch thread	Environment Break
	AND Immediate	I	ANDI rd,rs1,imm		Environment Break	I	ECALL		Environment Break
	Set <	R	SLT rd,rs1,rs2		Environment Break	I	ECALL		Environment Break
	Set <=	R	SLTI rd,rs1,imm		Environment Break	I	ECALL		Environment Break

- Need datapath and control to implement these instructions

Quiz

- 1) We should use the main ALU to compute $PC=PC+4$ in order to save some gates
- 2) The ALU is a sequential element
- 3) Program counter is a register

www.yellkey.com/picture

123
FFF
FFT
FTF
FTT
TFF
TFT
TTF
TTT




R-Format Instructions: Datapath

Summary of RISC-V Instruction Formats

31	30	25	24	21	20	19	15	14	12	11	8	7	6	0	
funct7		rs2		rs1		funct3		rd		opcode		R-type			
imm[11:0]				rs1		funct3		rd		opcode		I-type			
imm[11:5]		rs2		rs1		funct3		imm[4:0]		opcode		S-type			
imm[12:10:5]		rs2		rs1		funct3		imm[4:1:11]		opcode		B-type			
imm[31:12]								rd		opcode		U-type			
imm[20:10:1:11]]				imm[19:12]				rd		opcode		J-type			

R-Format Instruction Layout



- 32-bit instruction word divided into six fields of varying numbers of bits each: $7+5+5+3+5+7 = 32$

Examples

- **opcode** is a 7-bit field that lives in bits 6-0 of the instruction
- **rs2** is a 5-bit field that lives in bits 24-20 of the instruction

R-Format Instructions opcode/funct fields

31	25	24	20	19	15	14	12	11	7	6	0		
funct7				rs2		rs1		funct3		rd		opcode	
7				5		5		3		5		7	

- **opcode**: partially specifies what instruction it is
 - Note: This field is equal to **0110011** for all R-Format register-register arithmetic instructions
- **funct7+funct3**: combined with **opcode**, these two fields describe what operation to perform
- **Question: You have been professing simplicity, so why aren't opcode and funct7 and funct3 a single 17-bit field?**
 - Simpler implementation is more important than simpler spec

R-Format Instructions register specifiers

31	25	24	20	19	15	14	12	11	7	6	0			
funct7					rs2		rs1		funct3		rd		opcode	
7					5		5		3		5		7	

- **rs1** (Source Register #1): specifies register containing first operand
- **rs2**: specifies second register operand
- **rd** (Destination Register): specifies register which will receive result of computation
- Each register field holds a 5-bit unsigned integer (0-31) corresponding to a register number (**x0-x31**)

R-Format Example

- RISC-V Assembly Instruction:
add x18,x19,x10

31							25 24		20 19		15 14		12 11		7 6		0
funct7							rs2		rs1		funct3		rd		opcode		
7							5		5		3		5		7		

0000000	01010	10011	000	10010	0110011
add	rs2=10	rs1=19	add	rd=18	Reg-Reg OP

Implementing the add instruction

31	25	24	20	19	15	14	12	11	7	6	0			
funct7					rs2		rs1		funct3		rd		opcode	
7					5		5		3		5		7	

0000000	rs2	rs1	000	rd	0110011
---------	-----	-----	-----	----	---------

add rs2 rs1 add rd Reg-Reg OP

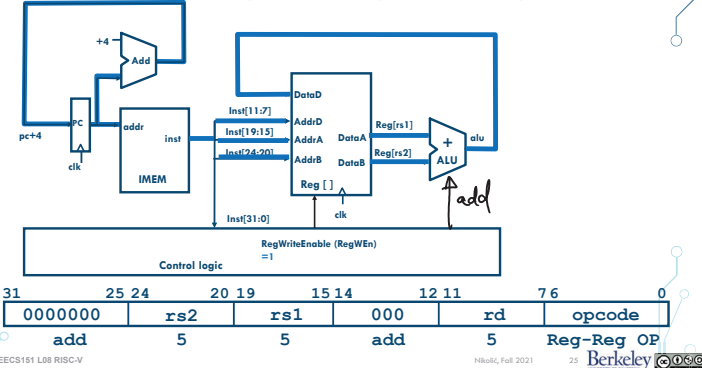
add rd, rs1, rs2

- Instruction makes two changes to machine's state:

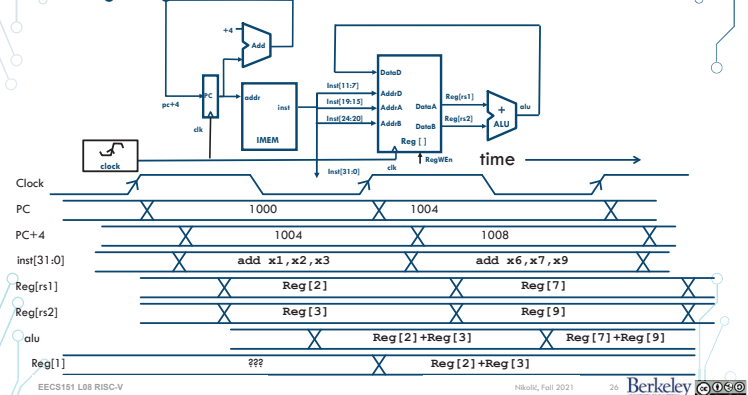
- $Reg[rd] = Reg[rs1] + Reg[rs2]$
- $PC = PC + 4$

Datapath for add

$$PC = PC + 4 \quad \text{Reg}[rd] = \text{Reg}[rs1] + \text{Reg}[rs2]$$



Timing Diagram for add



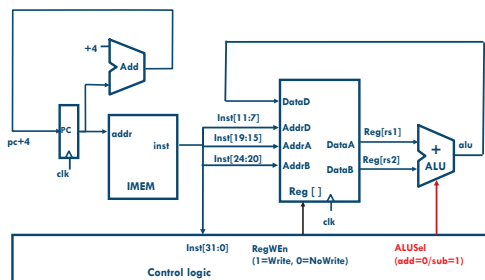
Implementing the sub instruction

31	25 24	20 19	15 14	12 11	7 6	0	
0000000	rs2	rs1	000	rd	0110011		add
0100000	rs2	rs1	000	rd	0110011		sub

sub rd, rs1, rs2

- Almost the same as add, except now have to subtract operands instead of adding them
- inst[30] selects between add and subtract

Datapath for add/sub



Implementing other R-Format instructions

0000000	rs2	rs1	000	rd	0110011	add
0100000	rs2	rs1	000	rd	0110011	sub
0000000	rs2	rs1	001	rd	0110011	sll
0000000	rs2	rs1	010	rd	0110011	slt
0000000	rs2	rs1	011	rd	0110011	sltu
0000000	rs2	rs1	100	rd	0110011	xor
0000000	rs2	rs1	101	rd	0110011	srl
0100000	rs2	rs1	101	rd	0110011	sra
0000000	rs2	rs1	110	rd	0110011	or
0000000	rs2	rs1	111	rd	0110011	and

- All implemented by decoding funct3 and funct7 fields and selecting appropriate ALU function

Administrivia

- Homework 3 is due next Monday
 - Homework 4 will be posted this week, due before midterm 1
- Lab 4 this week
- Lab 5 next week
- Midterm 1 on October 7, 7-8:30pm



I-Format Instructions: Datapath

Instruction Encoding


- Instructions are encoded to simplify logic
 - sub and sra differ in Inst[30] from add and srl
- RV64I widens registers (XLEN=64)
- Additional instructions manipulate 32-bit values, identified by a suffix W
 - ADDW, SUBW
 - RV64I opcode field for 'W' instructions is 0111011 (0110011 for RV32I)

0000000	rs2	rs1	000	rd	0110011	addw
0000000	rs2	rs1	000	rd	0111011	add

32b

64b

I-Format Instruction Layout



- Only one field is different from R-format, rs2 and funct7 replaced by 12-bit signed immediate, $\text{imm}[11:0]$
- Remaining fields (rs1, funct3, rd, opcode) same as before
- $\text{imm}[11:0]$ can hold values in range $[-2048_{\text{ten}}, +2047_{\text{ten}}]$
- Immediate is always sign-extended to 32-bits before use in an arithmetic operation
- Other instructions handle immediates > 12 bits

All RV32 I-format Arithmetic Instructions

imm[11:0]		rs1	000	rd	0010011	addi
imm[11:0]		rs1	010	rd	0010011	slti
imm[11:0]		rs1	011	rd	0010011	sltiu
imm[11:0]		rs1	100	rd	0010011	xori
imm[11:0]		rs1	110	rd	0010011	ori
imm[11:0]		rs1	111	rd	0010011	andi
0000000	shamt	rs1	001	rd	0010011	slli
0000000	shamt	rs1	101	rd	0010011	srlr
0100000	shamt	rs1	101	rd	0010011	srai



The same $\text{Inst}[30]$ immediate bit is used to distinguish "shift right logical" (SRL) from "shift right arithmetic" (SRAI)

"Shift-by-immediate" instructions only use lower 5 bits of the immediate value for shift amount (can only shift by 0-31 bit positions)

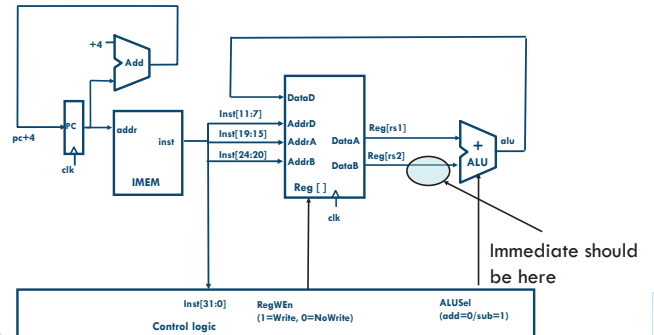
Implementing I-Format - addi instruction

- RISC-V Assembly Instruction – add immediate:

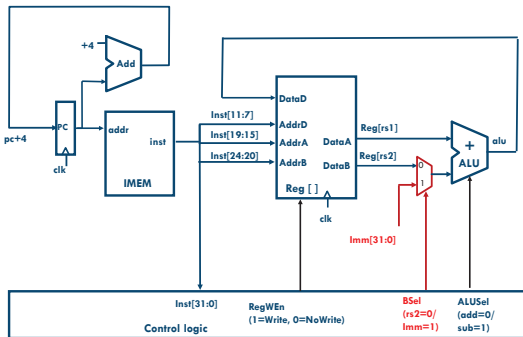
`addi x15, x1, -50`



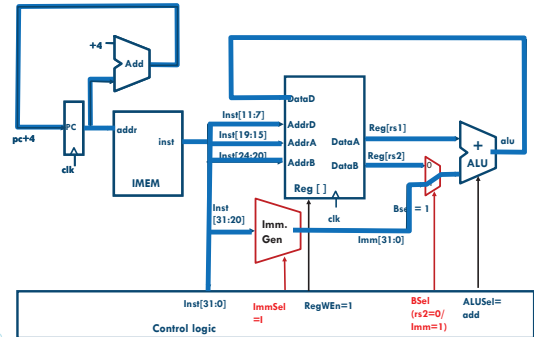
Datapath for add/sub




Adding addi to Datapath



Adding addi to Datapath

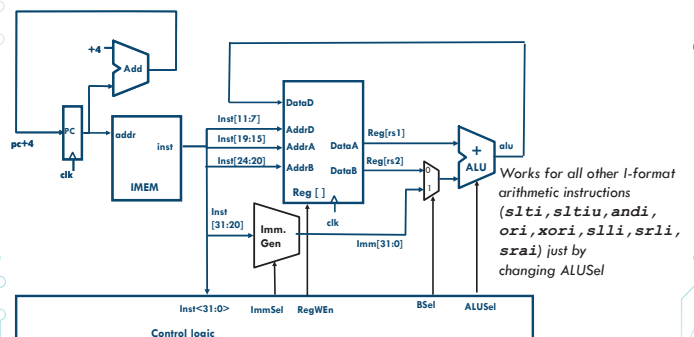


I-Format immediates



- High 12 bits of instruction ($\text{inst}[31:20]$) copied to low 12 bits of immediate ($\text{imm}[11:0]$)
- Immediate is sign-extended by copying value of $\text{inst}[31]$ to fill the upper 20 bits of the immediate value ($\text{imm}[31:12]$)
- Sign extension often in critical path

R+I Datapath




Works for all other I-format arithmetic instructions (`slti`, `sltiu`, `andi`, `ori`, `xori`, `slli`, `srlr`, `srai`) just by changing ALUSel

Add 1w to Datapath

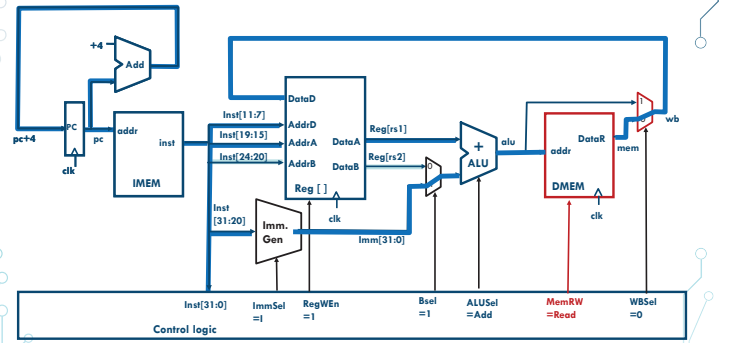
- RISC-V Assembly Instruction (l-type):

lw x14, 8(x2)



- The 12-bit signed immediate is added to the base address in register rs1 to form the **memory** address
- This is very similar to the add-immediate operation but used to create address not to create final result
- The value loaded from **memory** is stored in register **rd**

Adding 1w to Datapath




All RV32 Load Instructions

imm[11:0]	rs1	000	rd	0000011	1b
imm[11:0]	rs1	001	rd	0000011	1h
imm[11:0]	rs1	010	rd	0000011	1w
imm[11:0]	rs1	100	rd	0000011	1bu
imm[11:0]	rs1	101	rd	0000011	1hu


funct3 field encodes size and 'signedness' of load data

- Supporting the narrower loads requires additional logic to extract the correct byte/halfword from the value loaded from memory, and sign- or zero-extend the result to 32 bits before writing back to register file.
- It is just a mux for load extend, similar to sign extension for immediates



S-Format Instructions: Datapath



S-Format Used for Stores



- Store needs to read two registers, rs1 for base memory address, and rs2 for data to be stored, as well immediate offset!
- Can't have both rs2 and immediate in same place as other instructions!
- Note that stores don't write a value to the register file, **no rd**!
- RISC-V design decision is move low 5 bits of immediate to where rd field was in other instructions – keep rs1/rs2 fields in same place
- register names more critical than immediate bits in hardware design

Adding sw Instruction

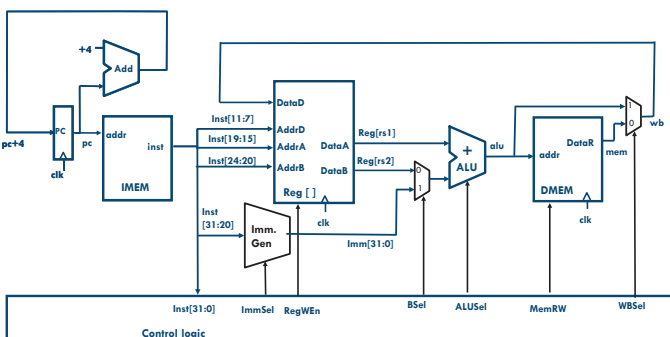
- sw**: Reads two registers, rs1 for base memory address, and rs2 for data to be stored, as well immediate offset!



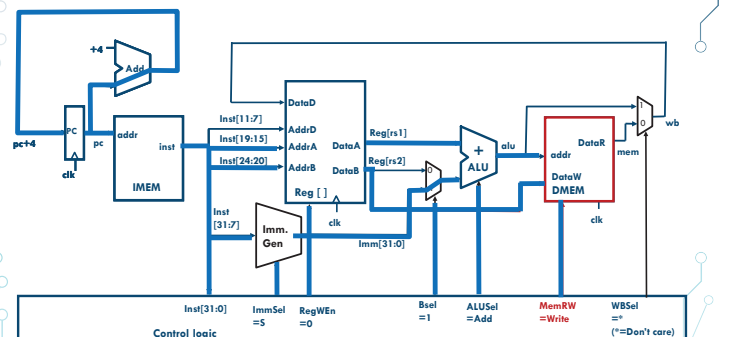
offset[11:5] rs2=14 rs1=2 SW offset[4:0] STORE
=0 =8

0000000 01000 combined 12-bit offset = 8

Datapath with 1w



Adding sw to Datapath




All RV32 Store Instructions

Imm[11:5]	rs2	rs1	000	imm[4:0]	0100011	sb
Imm[11:5]	rs2	rs1	001	imm[4:0]	0100011	sh
Imm[11:5]	rs2	rs1	010	imm[4:0]	0100011	sw

width

- Store byte, halfword, word

I+S Immediate Generation



- Just need a 5-bit mux to select between two positions where low five bits of immediate can reside in instruction
- Other bits in immediate are wired to fixed positions in instruction

Summary

- RISC-V ISA
 - Open, with increasing adoption
- RISC-V processor
 - A large state machine
 - Datapath + control
 - Reviewed R-, I-, S-format instructions and corresponding datapath elements