

EECS151 : Introduction to Digital Design and ICs

Lecture 9 – RISC-V ISA, Pipelining

Bora Nikolić



August 20, 2021, Tom's hardware
**Tesla Packs 50 Billion Transistors Onto D1 Dojo Chip
Designed to Conquer Artificial Intelligence Training**

D1 delivers 362 TeraFLOPs of power.
Called the D1, the chip resembles a part of the Dojo supercomputer used to train AI models inside Tesla HQ, which are later deployed in various applications. The D1 chip is a product of TSMC's manufacturing efforts, forged in a 7nm semiconductor node. Packing over 50 billion transistors, the chip boasts a huge die size of 645mm^2 .




Image credit: Dennis Hong / Twitter

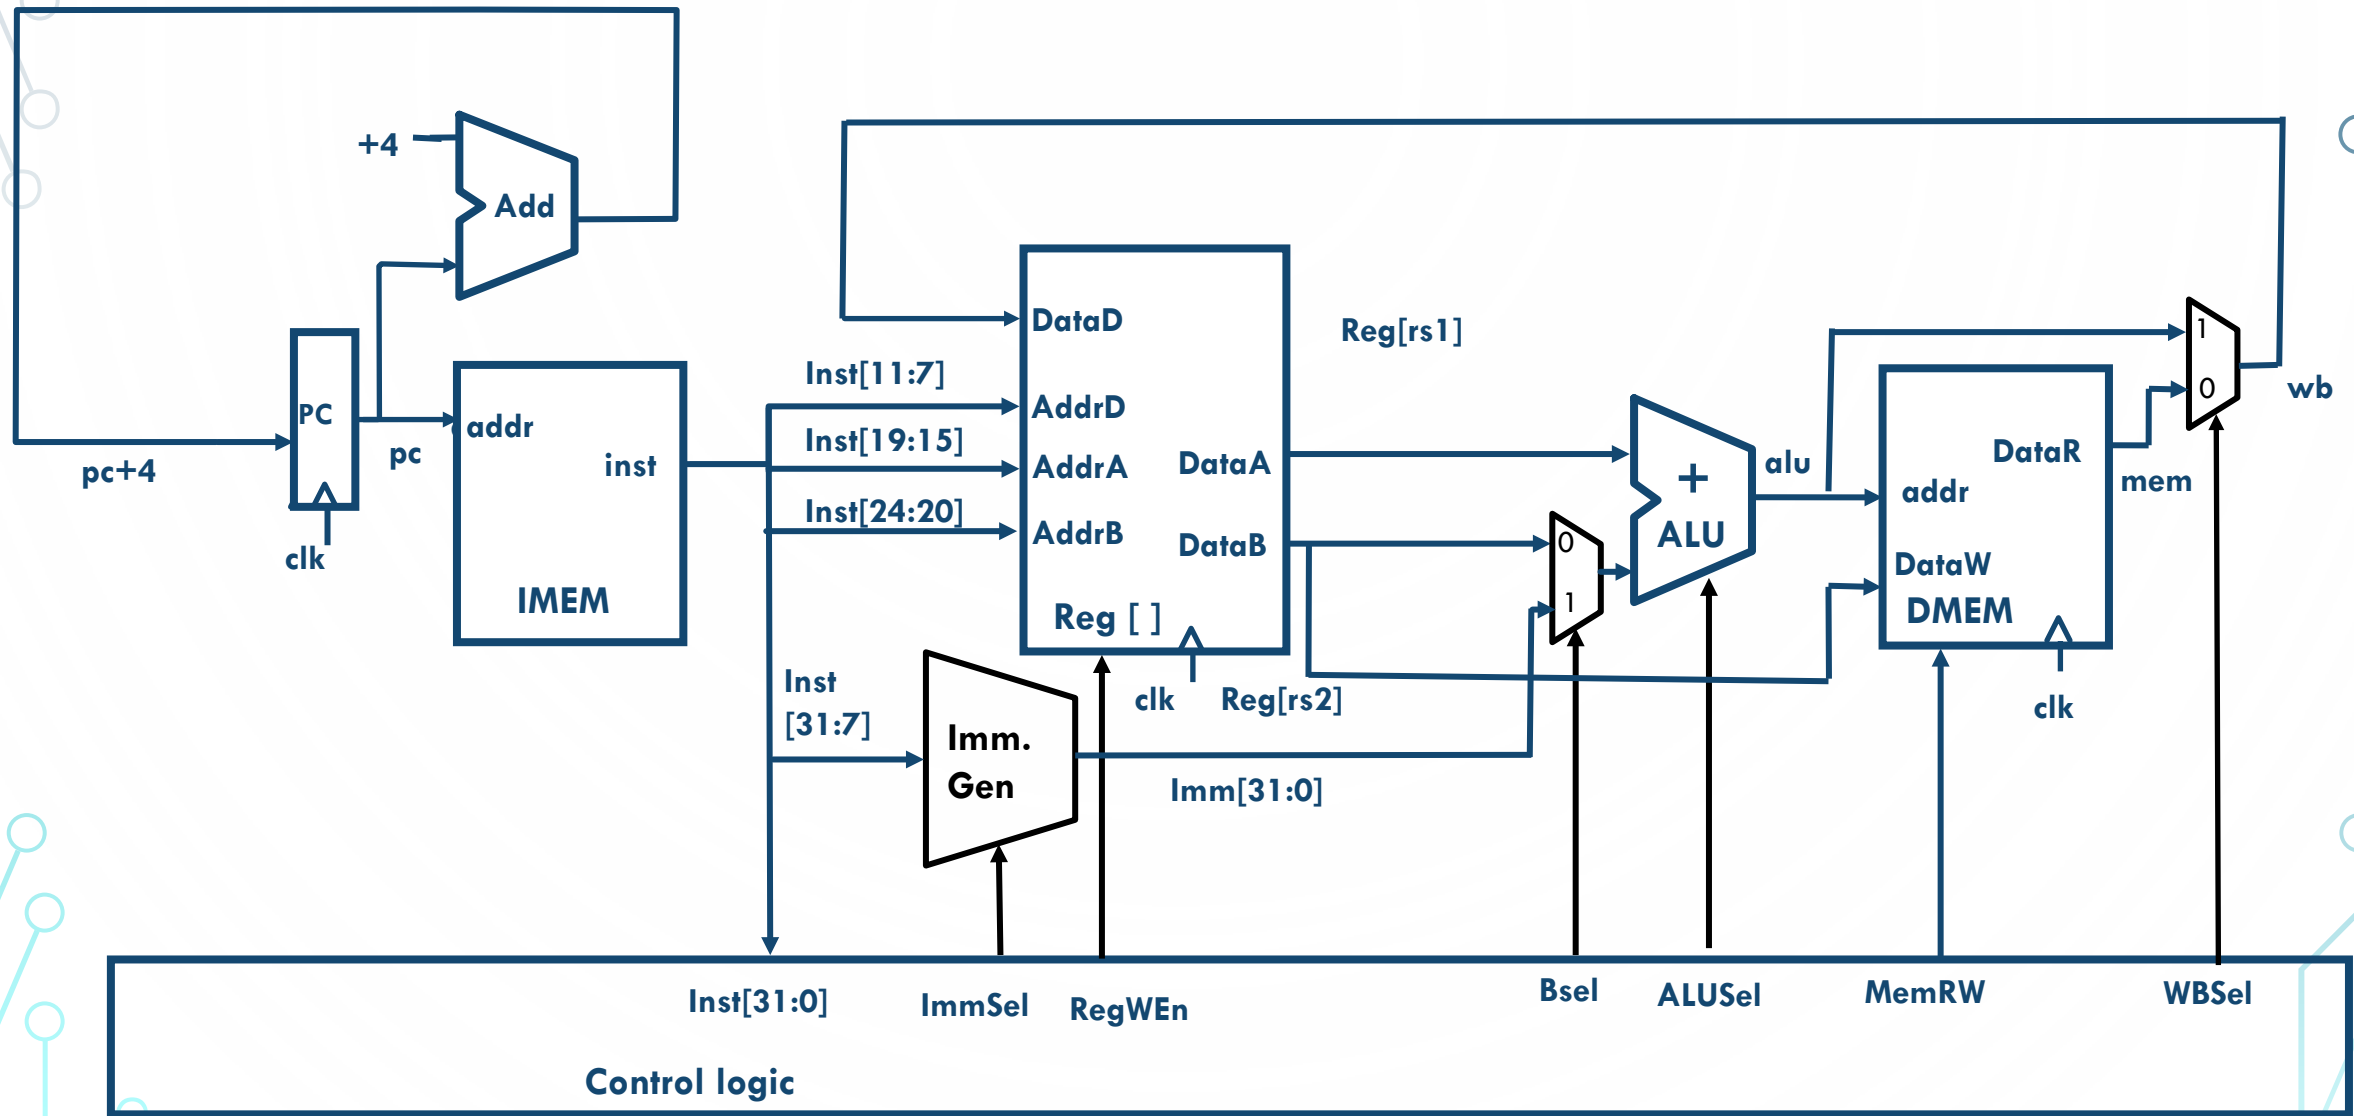
Review

- RISC-V ISA
 - Open, with increasing adoption
- RISC-V processor
 - A large state machine
 - Datapath + control
 - Reviewed R-, I-, S-format instructions and corresponding datapath elements



RISC-V B-Format Instructions

Datapath So Far (R-, I-, S Instruction Types)



B-Format - RISC-V Conditional Branches

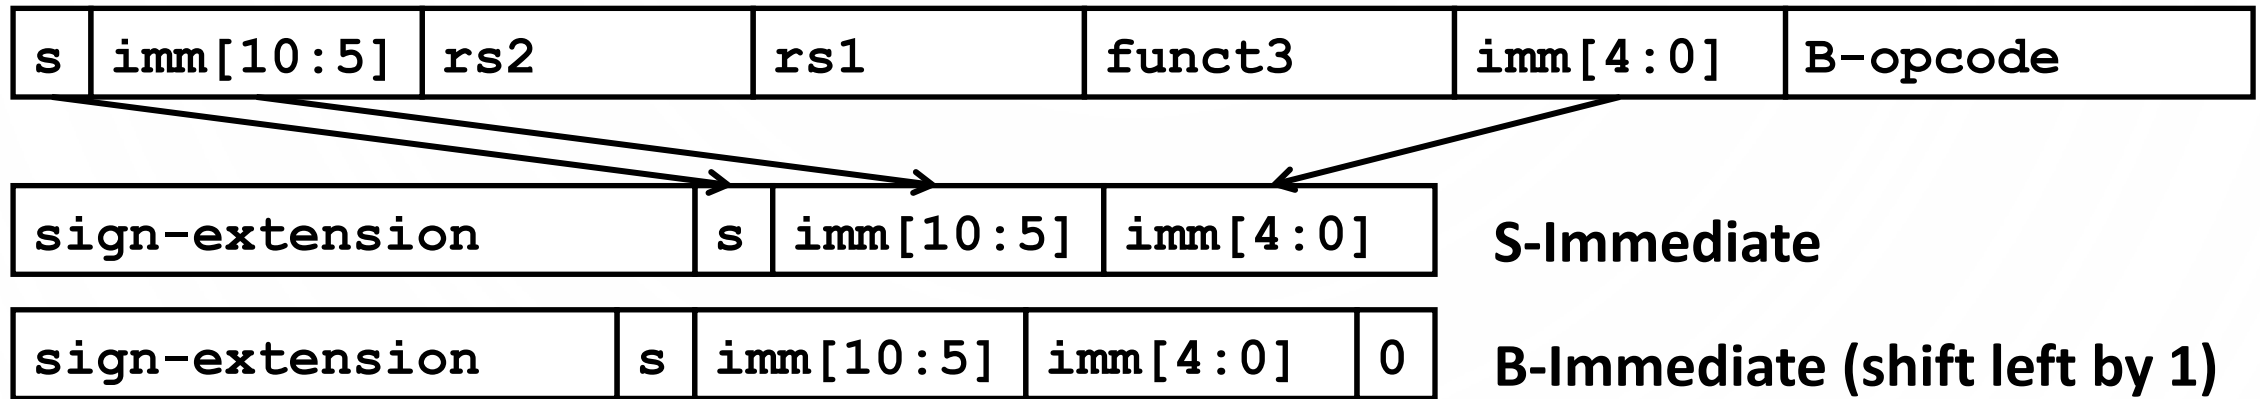
- E.g., **BEQ x1, x2, Label**
- Branches read two registers but don't write a register (similar to stores)
- How to encode label, i.e., where to branch to?

RISC-V Feature, $n \times 16$ -bit instructions

- Extensions to RISC-V base ISA support 16-bit compressed instructions and also variable-length instructions that are multiples of 16-bits in length
- To enable this, RISC-V scales the branch offset by 2 bytes even when there are no 16-bit instructions
- Reduces branch reach by half and means that $\frac{1}{2}$ of possible targets will be errors on RISC-V processors that only support 32-bit instructions (as used in this class)
- RISC-V conditional branches can only reach $\pm 2^{10} \times 32$ -bit instructions on either side of PC

RISC-V Branch Immediates

- 12-bit immediate encodes PC-relative offset of -4096 to +4094 bytes in multiples of 2 bytes
- RISC-V approach: keep 11 immediate bits in fixed position in output value, and rotate LSB of S-format to be bit 12 of B-format



Only one bit changes position between S and B, so only need a single-bit 2-way mux

RISC-V Immediate Encoding

Instruction encodings, inst[31:0]

31	30	25	24	20	19	15	14	12	11	8	7	6	0					
funct7					rs2			rs1			funct3			rd		opcode		R-type
imm[11:0]						rs1			funct3			rd		opcode		I-type		
imm[11:5]				rs2			rs1			funct3			imm[4:0]			opcode		S-type
imm[12 10:5]				rs2			rs1			funct3			imm[4:1 11]			opcode		B-type

32-bit immediates produced, imm[31:0]

31	25	24	12	11	10	5	4	1	0				
-inst[31]-					inst[30:25]			inst[24:21]			inst[20]	I-imm.	
-inst[31]-					inst[30:25]			inst[11:8]			inst[7]	S-imm.	
-inst[31]-					inst[7]			inst[30:25]			inst[11:8]	0	B-imm.

Upper bits sign-extended from inst[31]
always

Only bit 7 of instruction changes role in
immediate between S and B

Branch Example, complete encoding

beq **x19,x10,** **offset = 16 bytes**

13-bit immediate, imm[12:0], with value 16

0000000010000

imm[0] discarded,
→ always zero

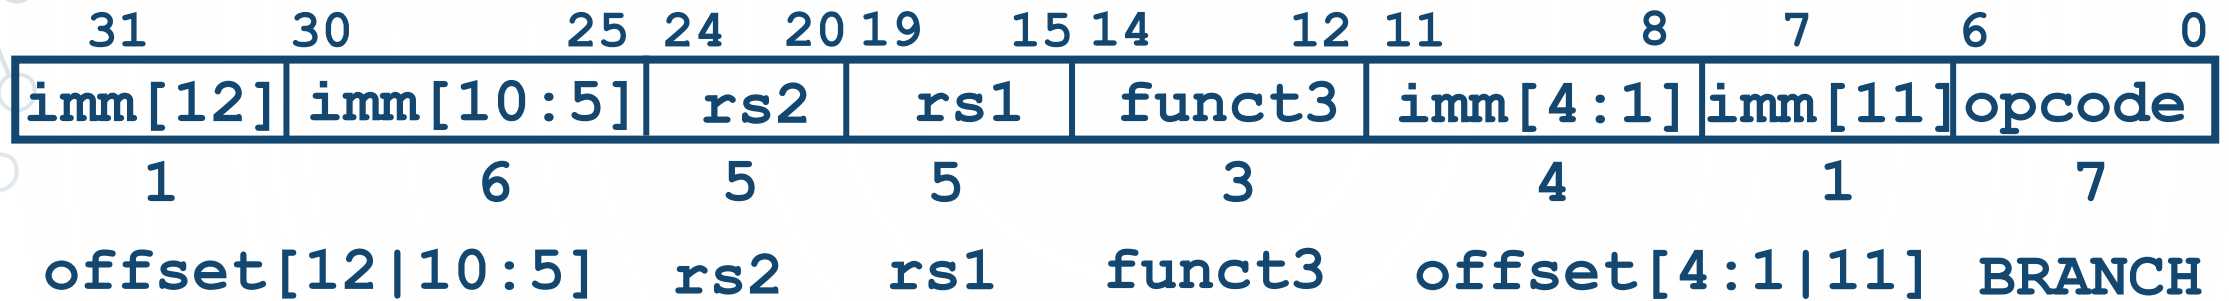
imm[12]

imm[11]

0	000000	01010	10011	000	1000	0	1100011
---	--------	-------	-------	-----	------	---	---------

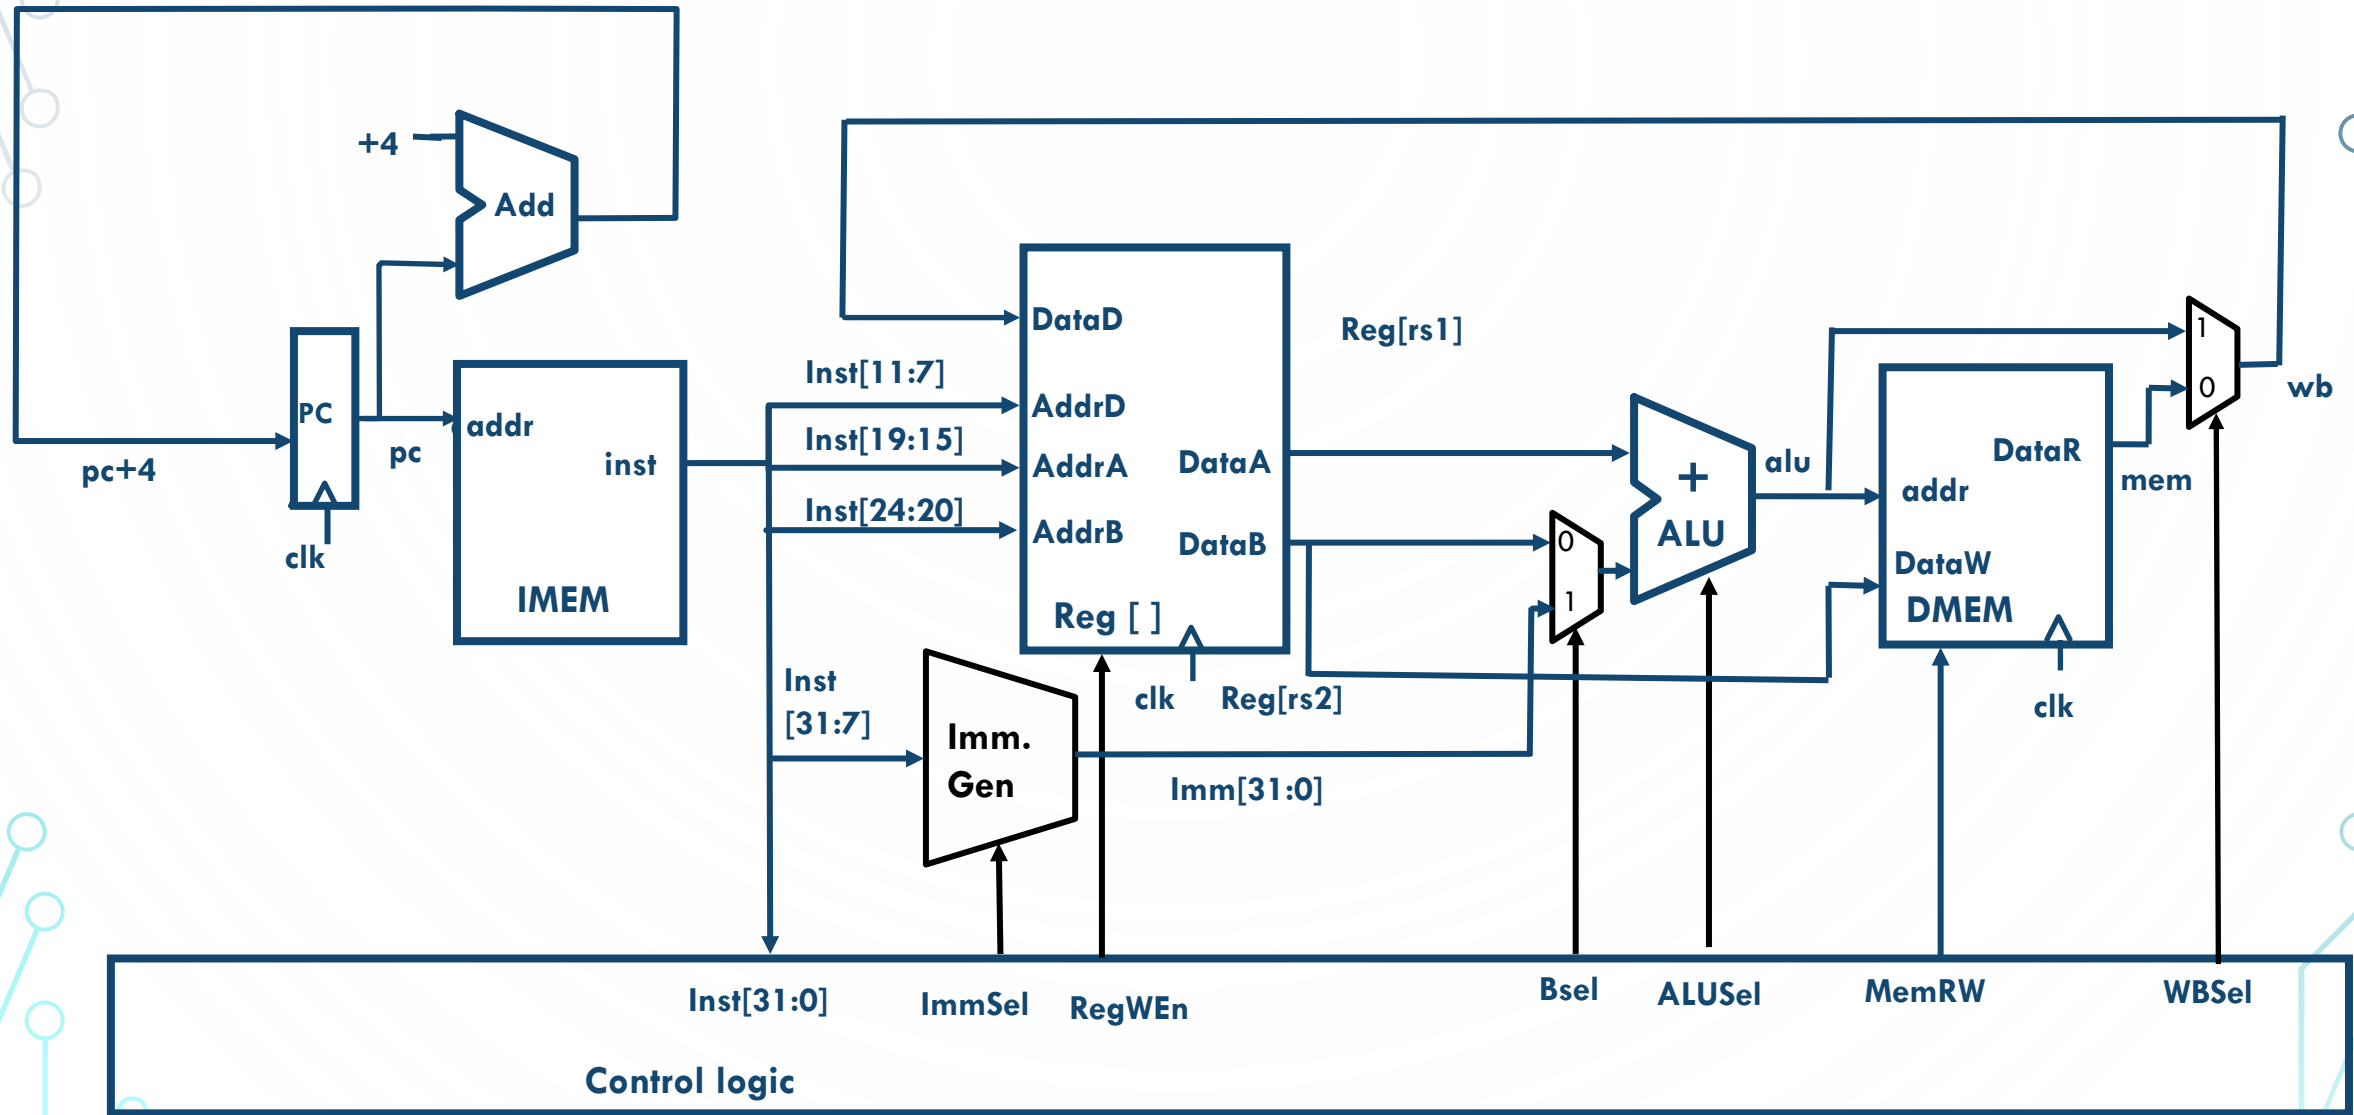
imm[10:5] rs2=10 rs1=19 BEQ imm[4:1] BRANCH

Implementing Branches



- B-format is mostly same as S-format, with two register sources (rs1/rs2) and a 12-bit immediate
- But now immediate represents values -4096 to +4094 in 2-byte increments
- The 12 immediate bits encode even 13-bit signed byte offsets (lowest bit of offset is always zero, so no need to store it)

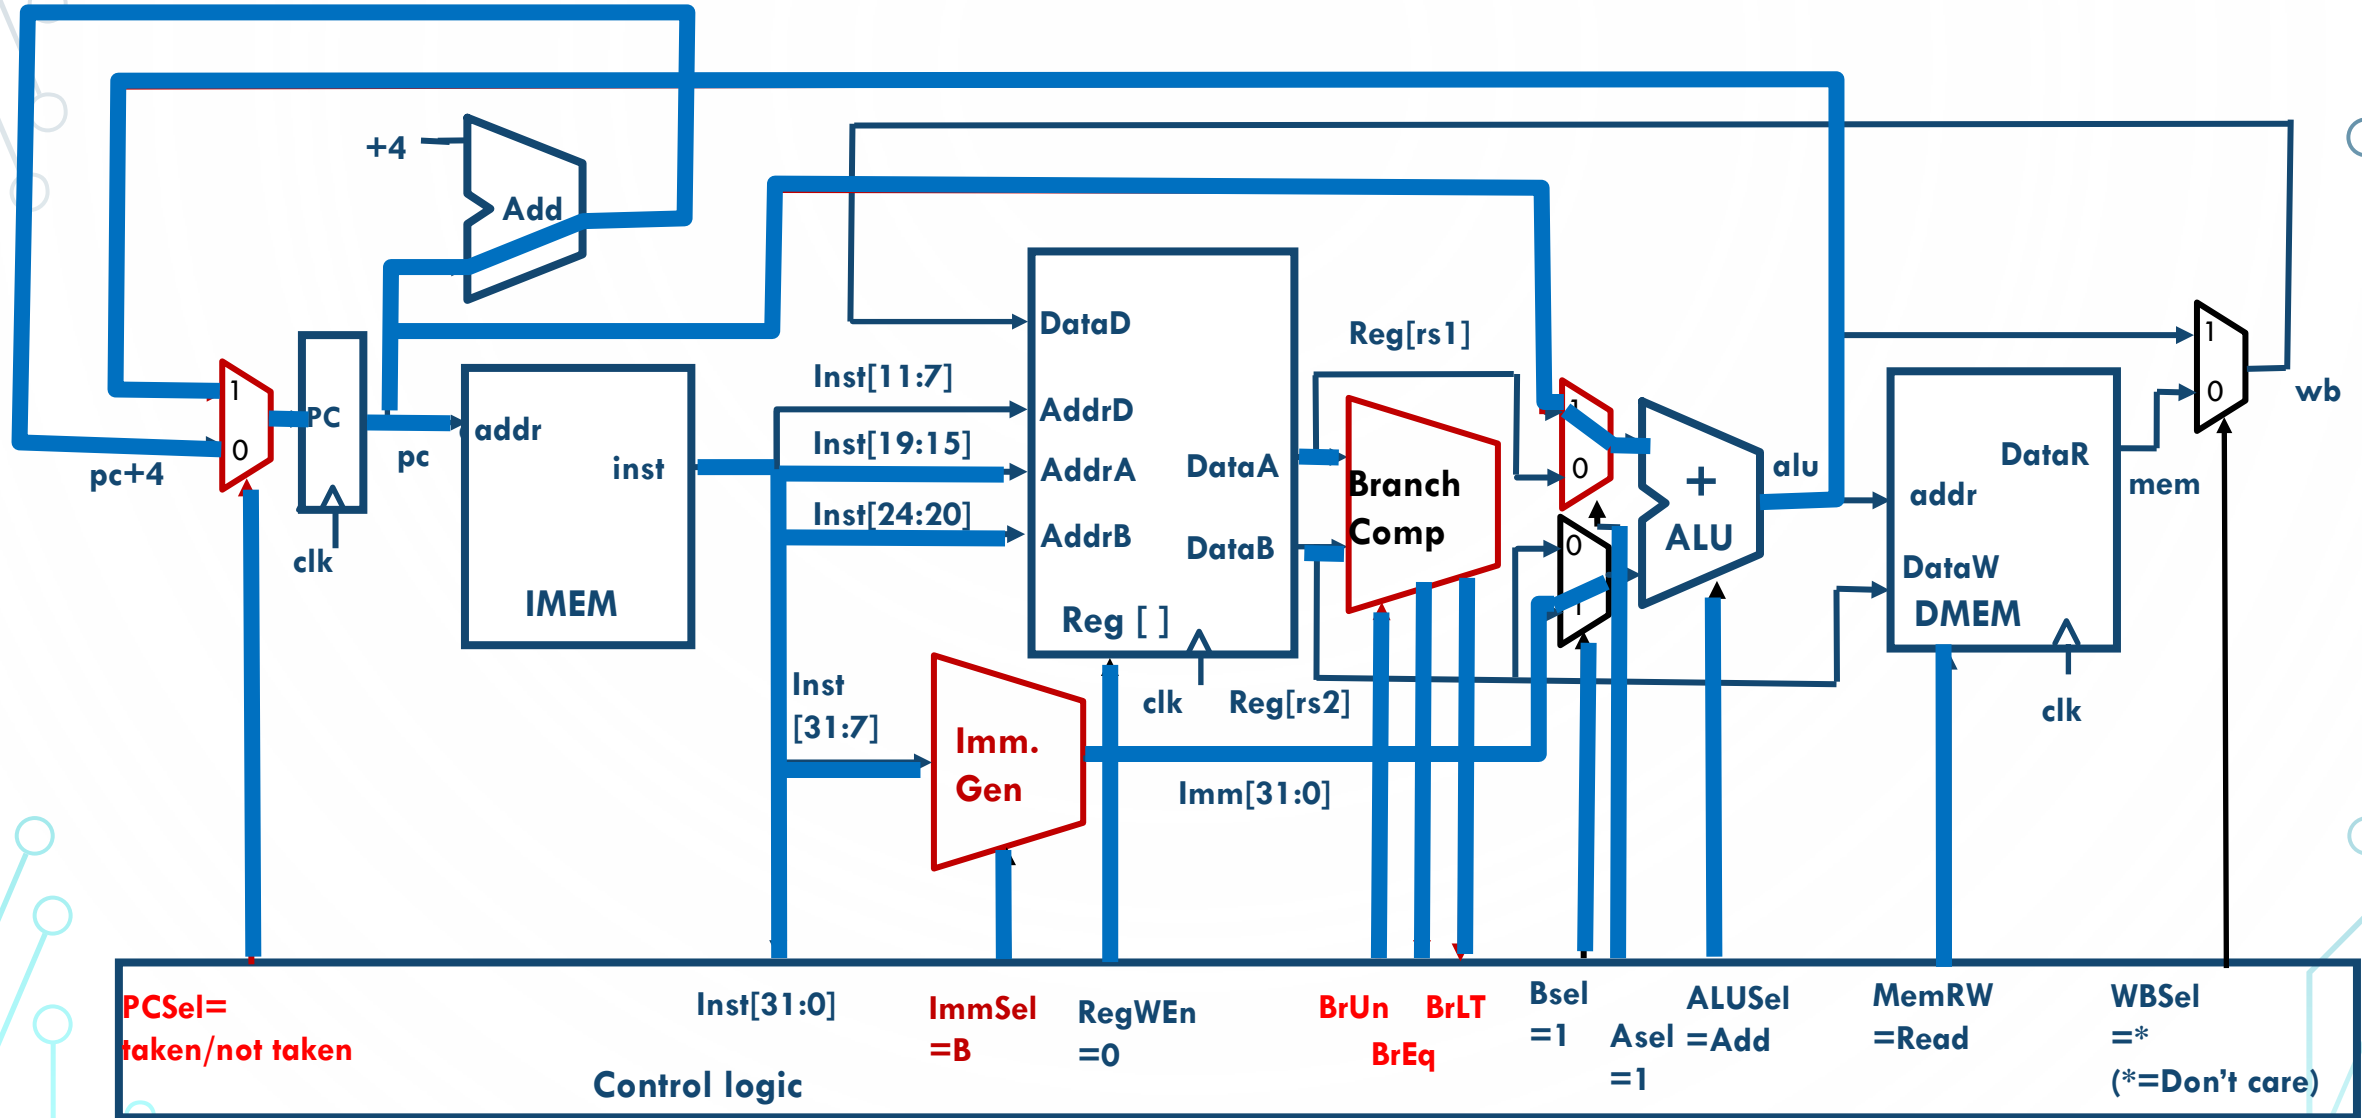
Datapath So Far




To Add Branches

- Different change to the state:
 - $PC = \begin{cases} PC + 4, & \text{branch not taken} \\ PC + \text{immediate}, & \text{branch taken} \end{cases}$
- Six branch instructions: **BEQ**, **BNE**, **BLT**, **BGE**, **BLTU**, **BGEU**
- Need to compute **PC + immediate** and to compare values of **rs1** and **rs2**
 - Need another add/sub unit

Adding Branches



Branch Comparator



- $\text{BrEq} = 1$, if $A = B$
- $\text{BrLT} = 1$, if $A < B$
- $\text{BrUn} = 1$ selects unsigned comparison for BrLT , 0=signed
- BGE branch: $A \geq B$, if $\overline{A < B}$

All RISC-V Branch Instructions

<code>imm[12 10:5]</code>	<code>rs2</code>	<code>rs1</code>	<code>000</code>	<code>imm[4:1 11]</code>	<code>1100011</code>
<code>imm[12 10:5]</code>	<code>rs2</code>	<code>rs1</code>	<code>001</code>	<code>imm[4:1 11]</code>	<code>1100011</code>
<code>imm[12 10:5]</code>	<code>rs2</code>	<code>rs1</code>	<code>100</code>	<code>imm[4:1 11]</code>	<code>1100011</code>
<code>imm[12 10:5]</code>	<code>rs2</code>	<code>rs1</code>	<code>101</code>	<code>imm[4:1 11]</code>	<code>1100011</code>
<code>imm[12 10:5]</code>	<code>rs2</code>	<code>rs1</code>	<code>110</code>	<code>imm[4:1 11]</code>	<code>1100011</code>
<code>imm[12 10:5]</code>	<code>rs2</code>	<code>rs1</code>	<code>111</code>	<code>imm[4:1 11]</code>	<code>1100011</code>

BEQ

BNE

BLT

BGE

BLTU

BGEU

Administrivia


- Homework 4 is due next Monday
 - No new homework this week
 - Homework 5 will be posted next week, due after the midterm
- Lab 5 this week
 - No lab next week
 - Lab 6 (last) after the midterm
- Midterm 1 on October 7, 7-8:30pm



RISC-V


J-Format Instructions

J-Format for Jump Instructions



- JAL saves PC+4 in register rd (the return address)
 - Assembler “j” jump is pseudo-instruction, uses JAL but sets rd=x0 to discard return address
- Set PC = PC + offset (PC-relative jump)
- Target somewhere within $\pm 2^{19}$ locations, 2 bytes apart
 - $\pm 2^{18}$ 32-bit instructions
- Immediate encoding optimized similarly to branch instruction to reduce hardware cost


JALR Instruction (I-Format)



- JALR rd, rs, immediate
 - Writes PC+4 to rd (return address)
 - Sets $PC = rs1 + \text{immediate}$ (and sets the LSB to 0)
 - Uses same immediates as arithmetic and loads
 - **no** multiplication by 2 bytes
 - In contrast to branches and JAL

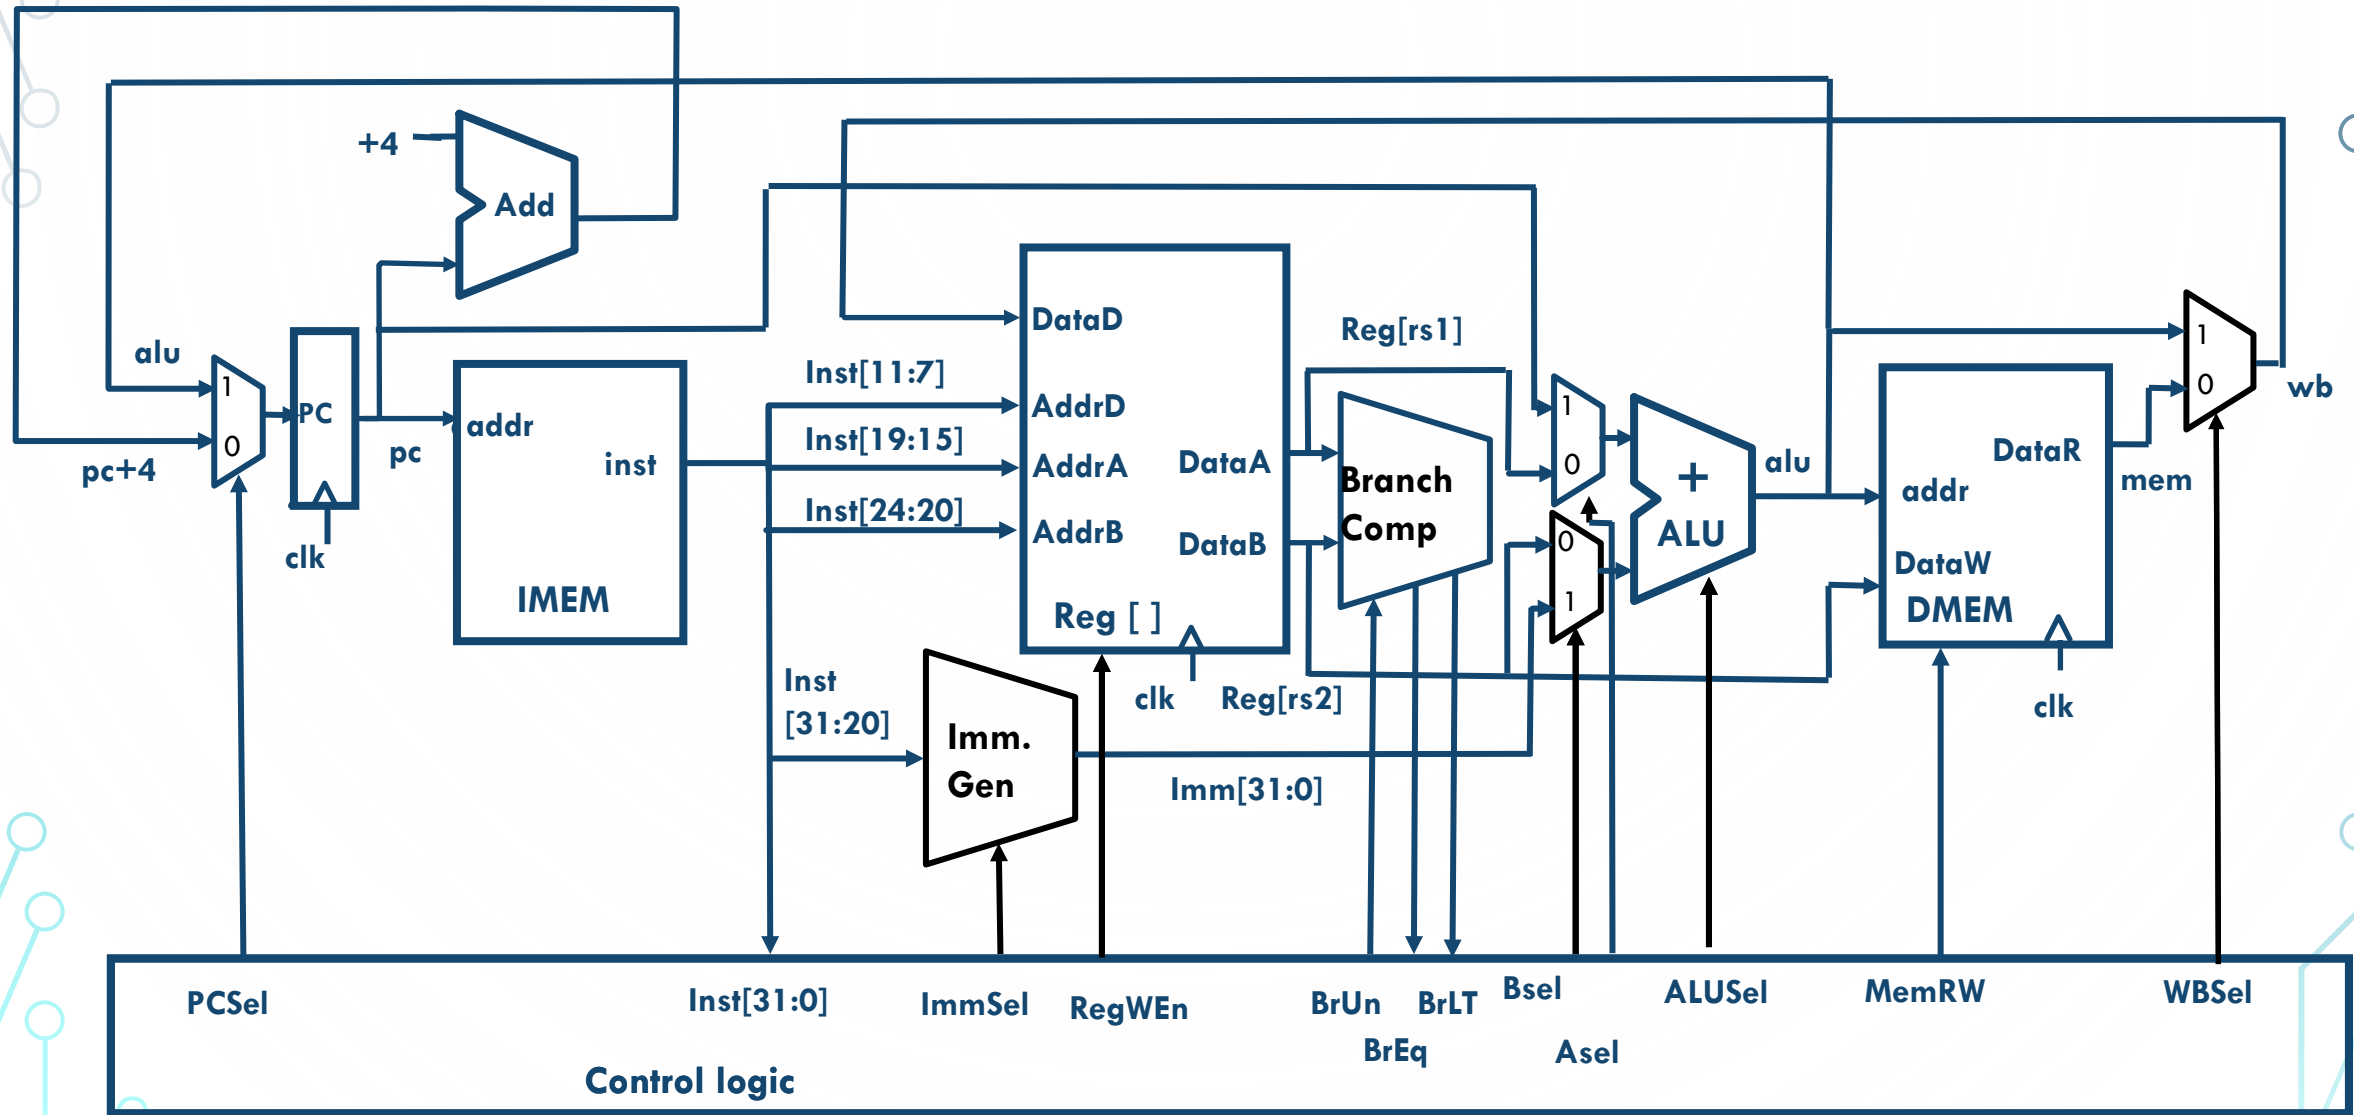
Otherwise, we would have yet another new encoding

Let's Add JALR (I-Format)

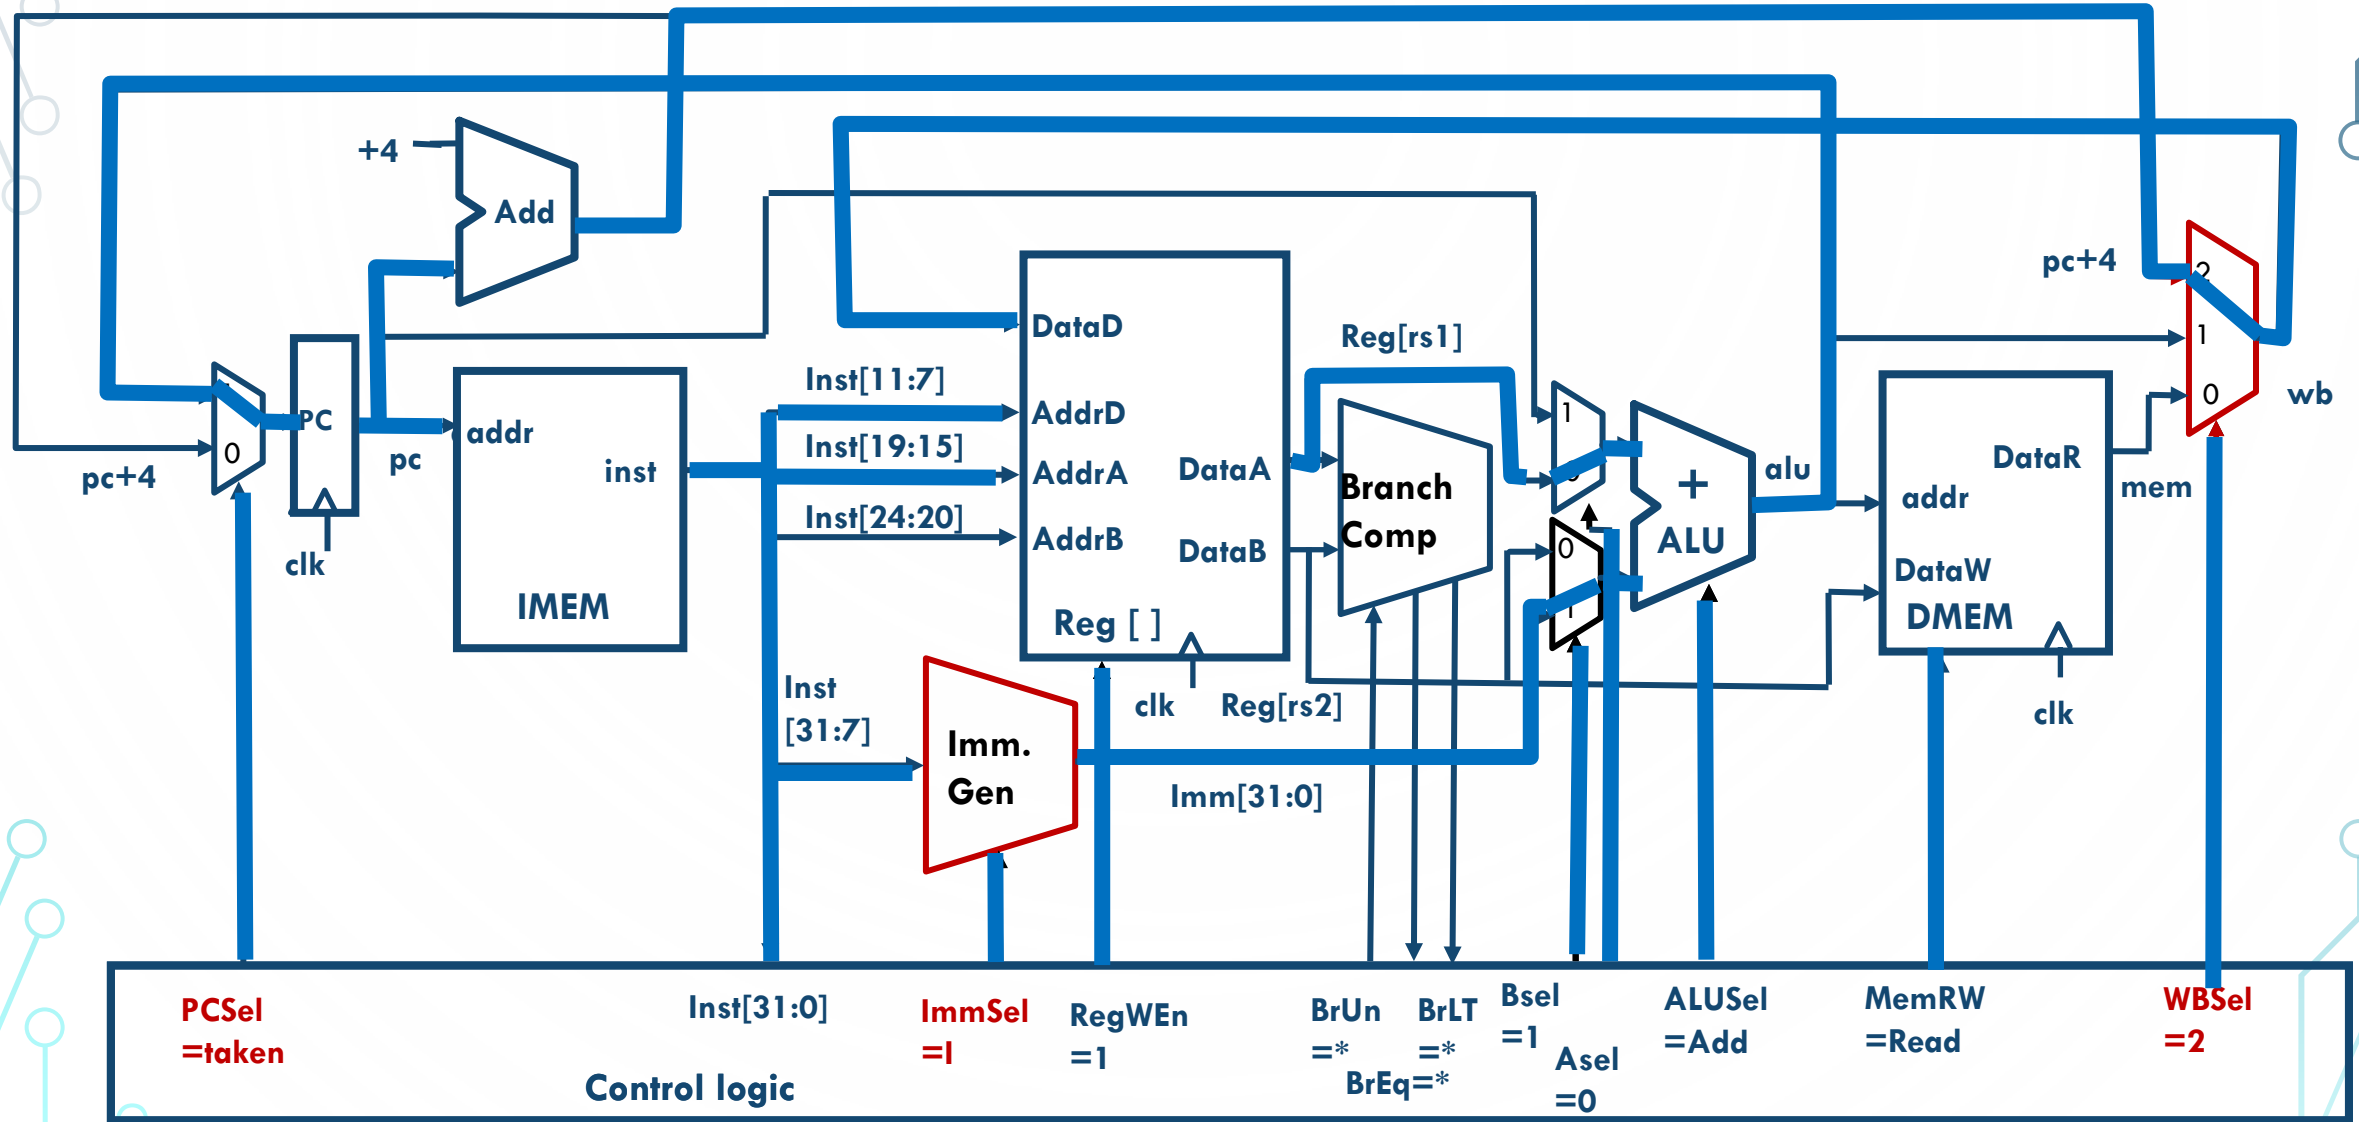


- JALR rd, rs, immediate
- Two changes to the state
 - Writes PC+4 to rd (return address)
 - Sets PC = rs + immediate
 - Uses same immediates as arithmetic and loads
 - **no** multiplication by 2 bytes
 - LSB is ignored


Datapath So Far, with Branches



Adding JALR

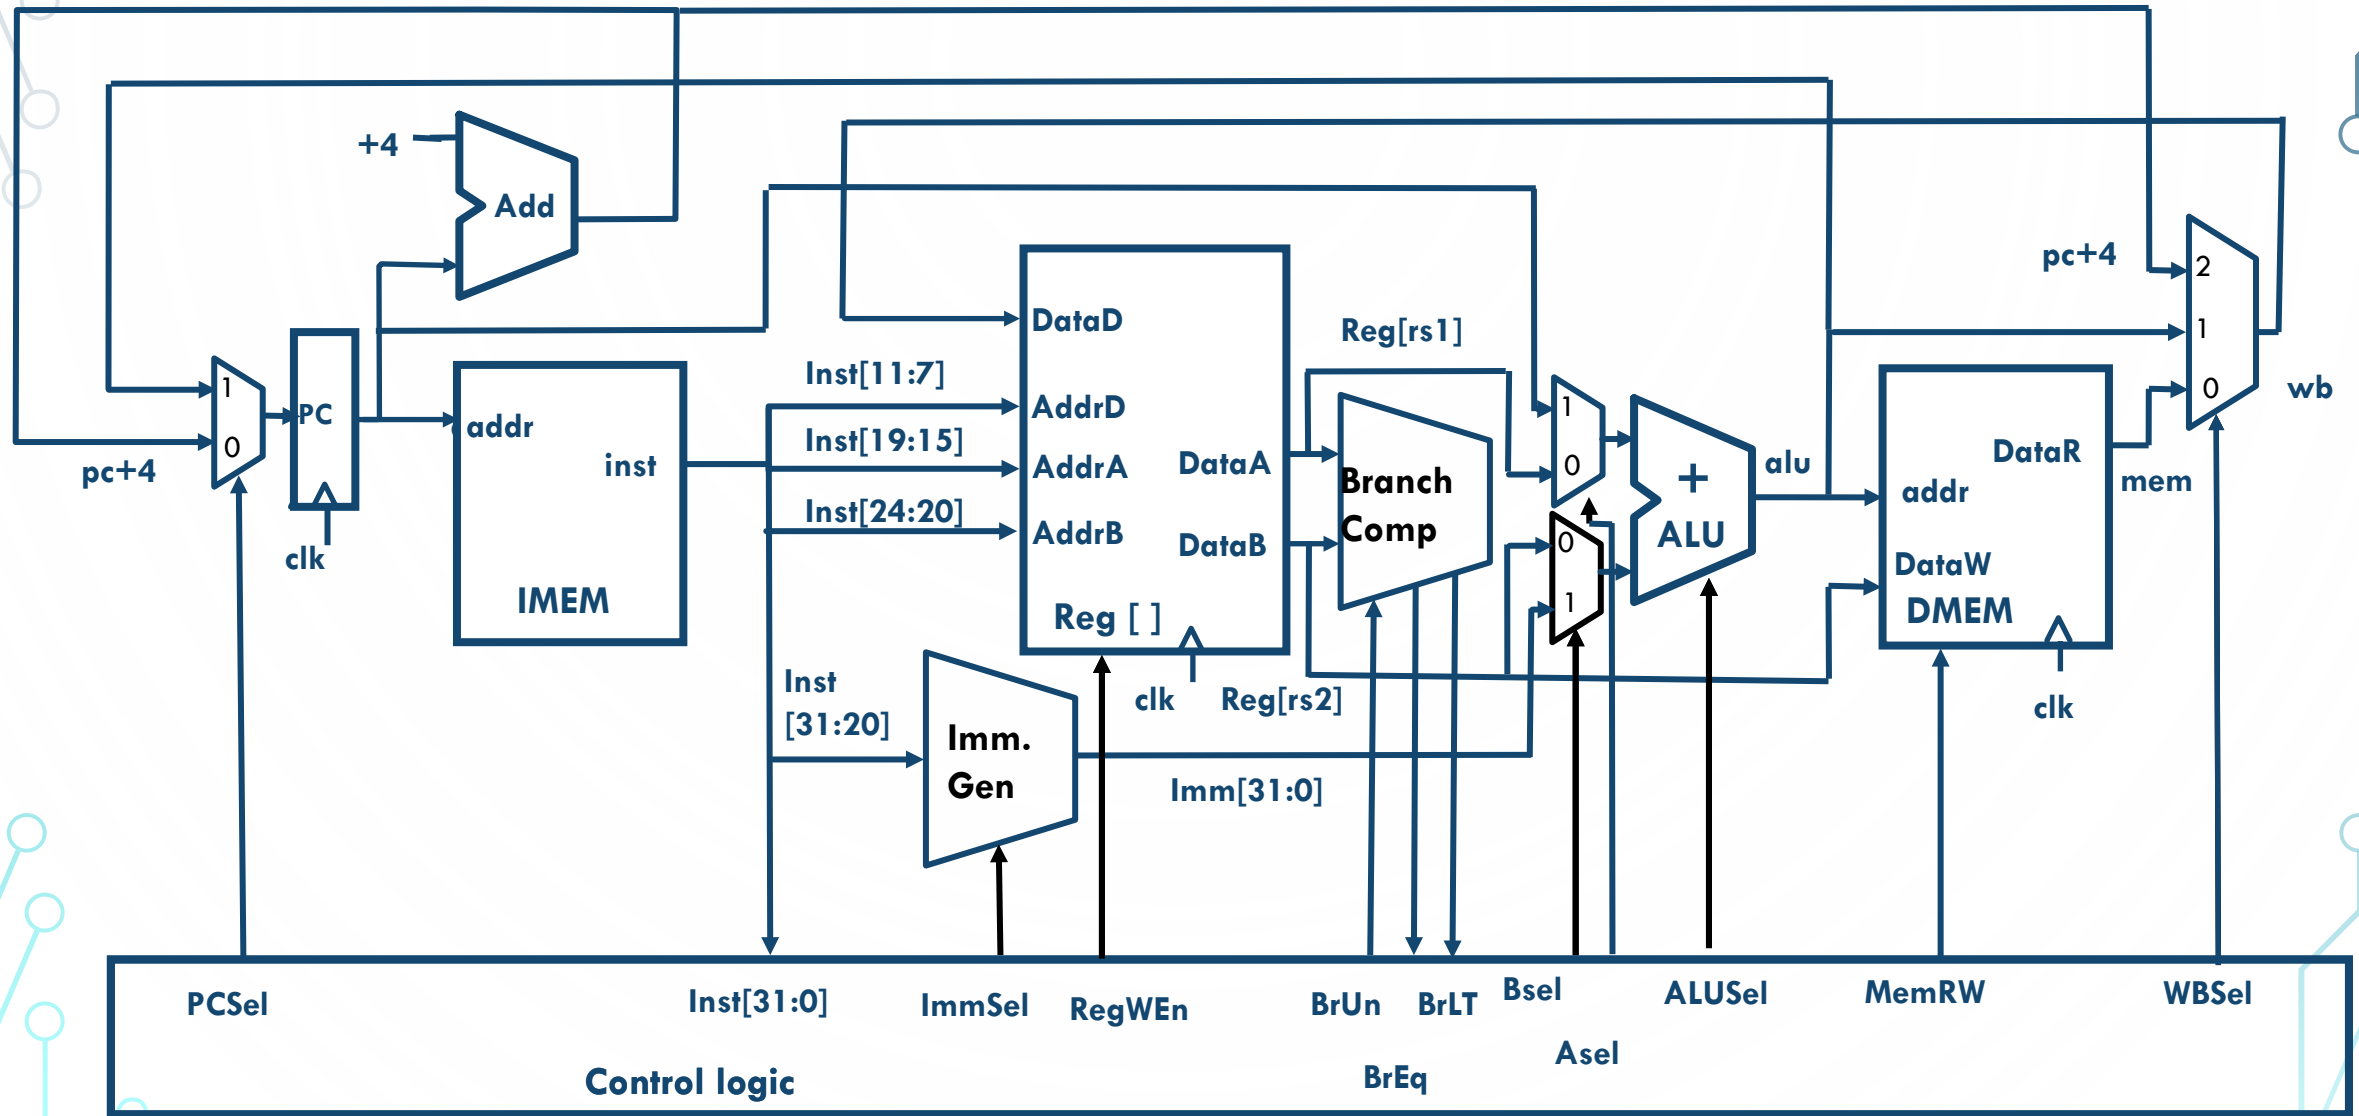


Adding JAL

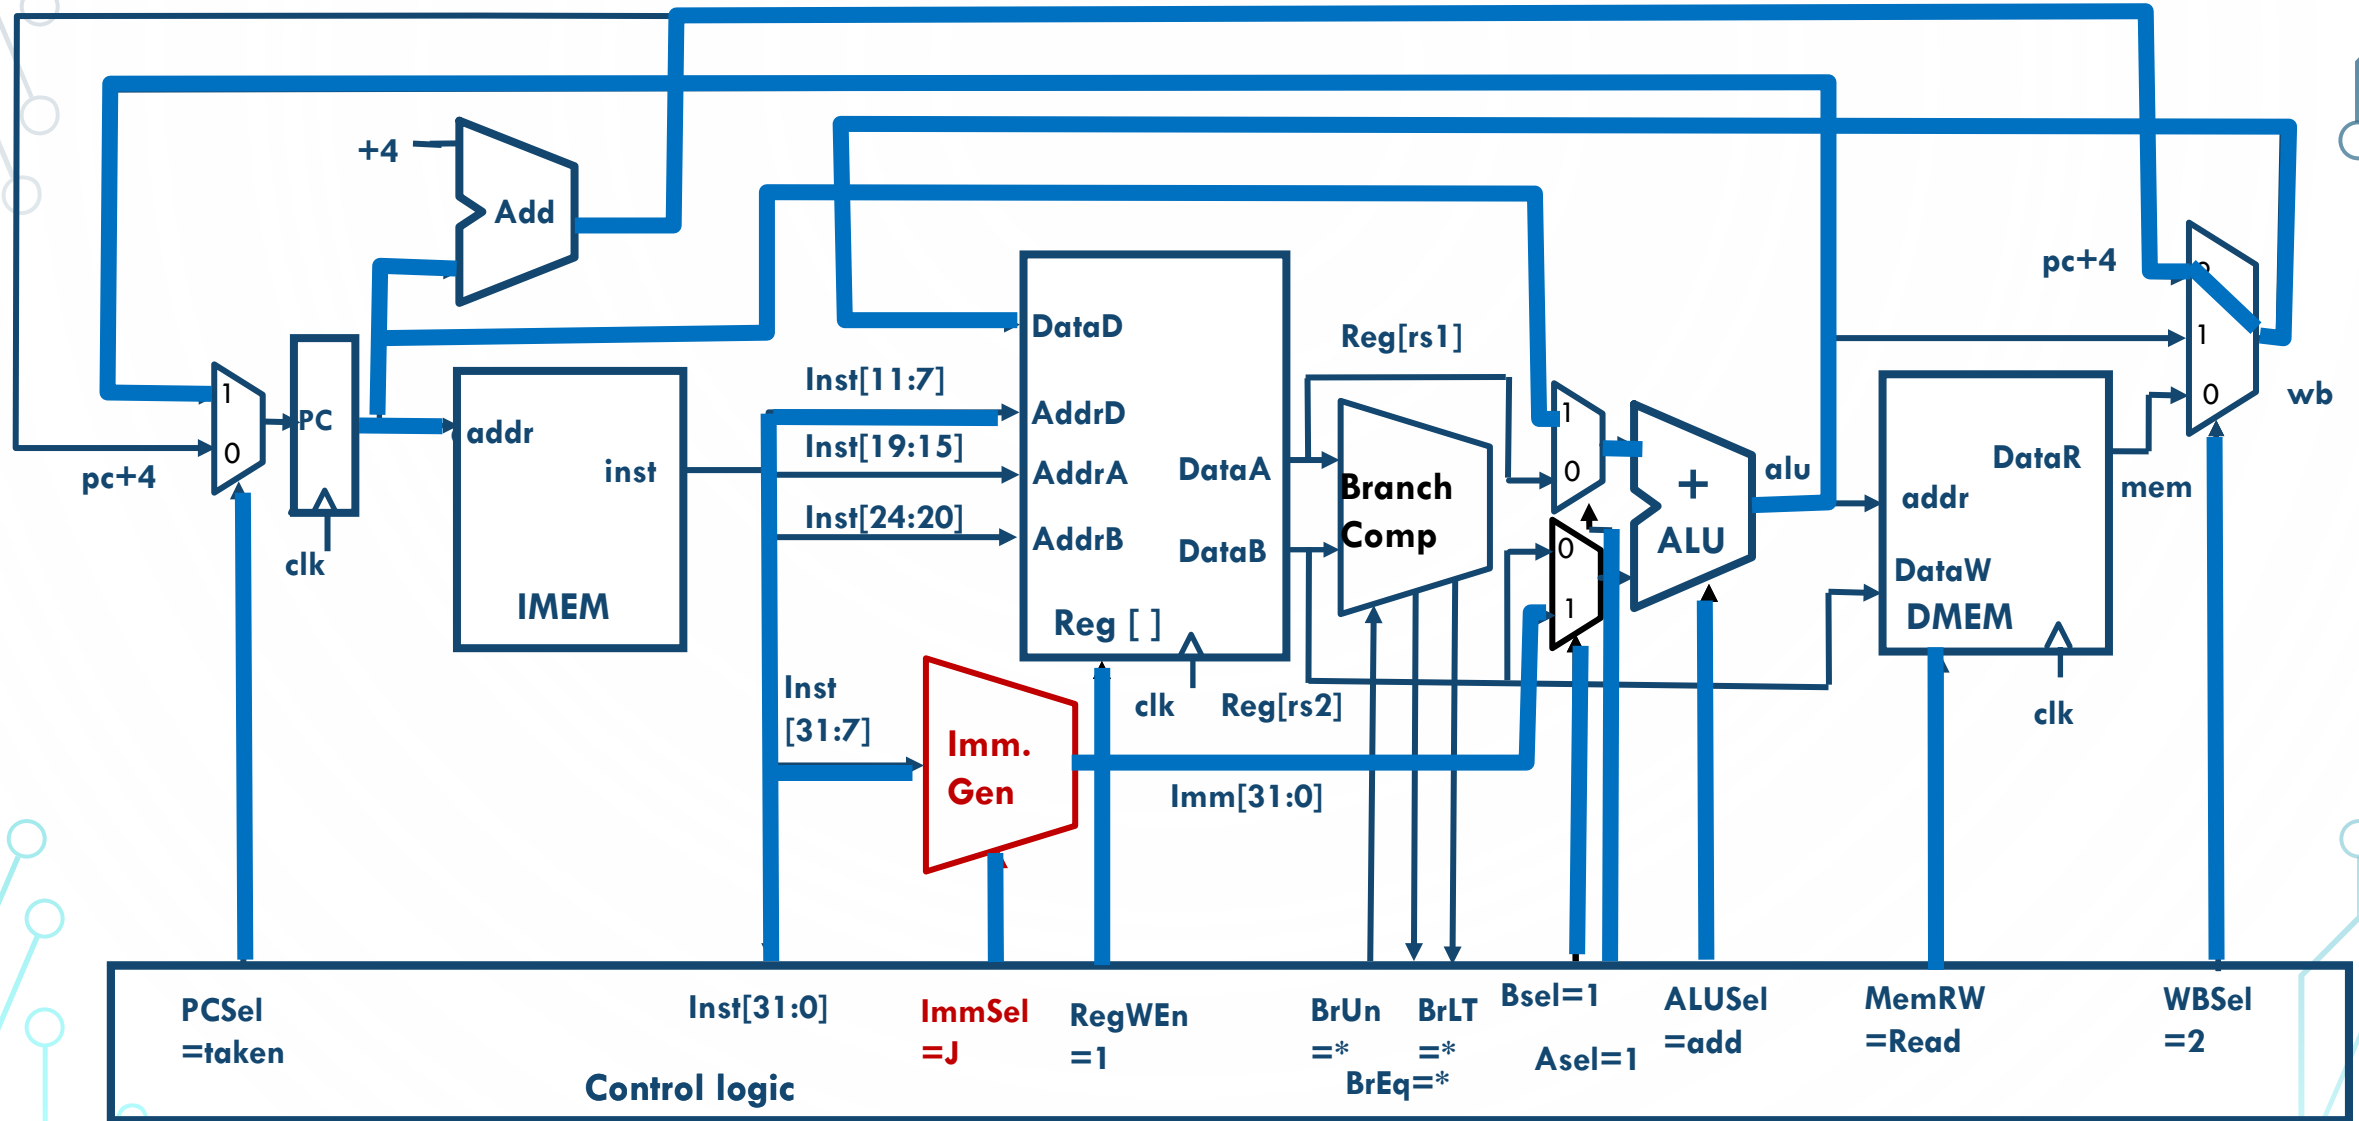


- JAL saves PC+4 in register rd (the return address)
- Set PC = PC + offset (PC-relative jump)
- Target somewhere within $\pm 2^{19}$ locations, 2 bytes apart
 - $\pm 2^{18}$ 32-bit instructions
- Immediate encoding optimized similarly to branch instruction to reduce hardware cost

Datapath with JALR



Adding JAL






RISC-V

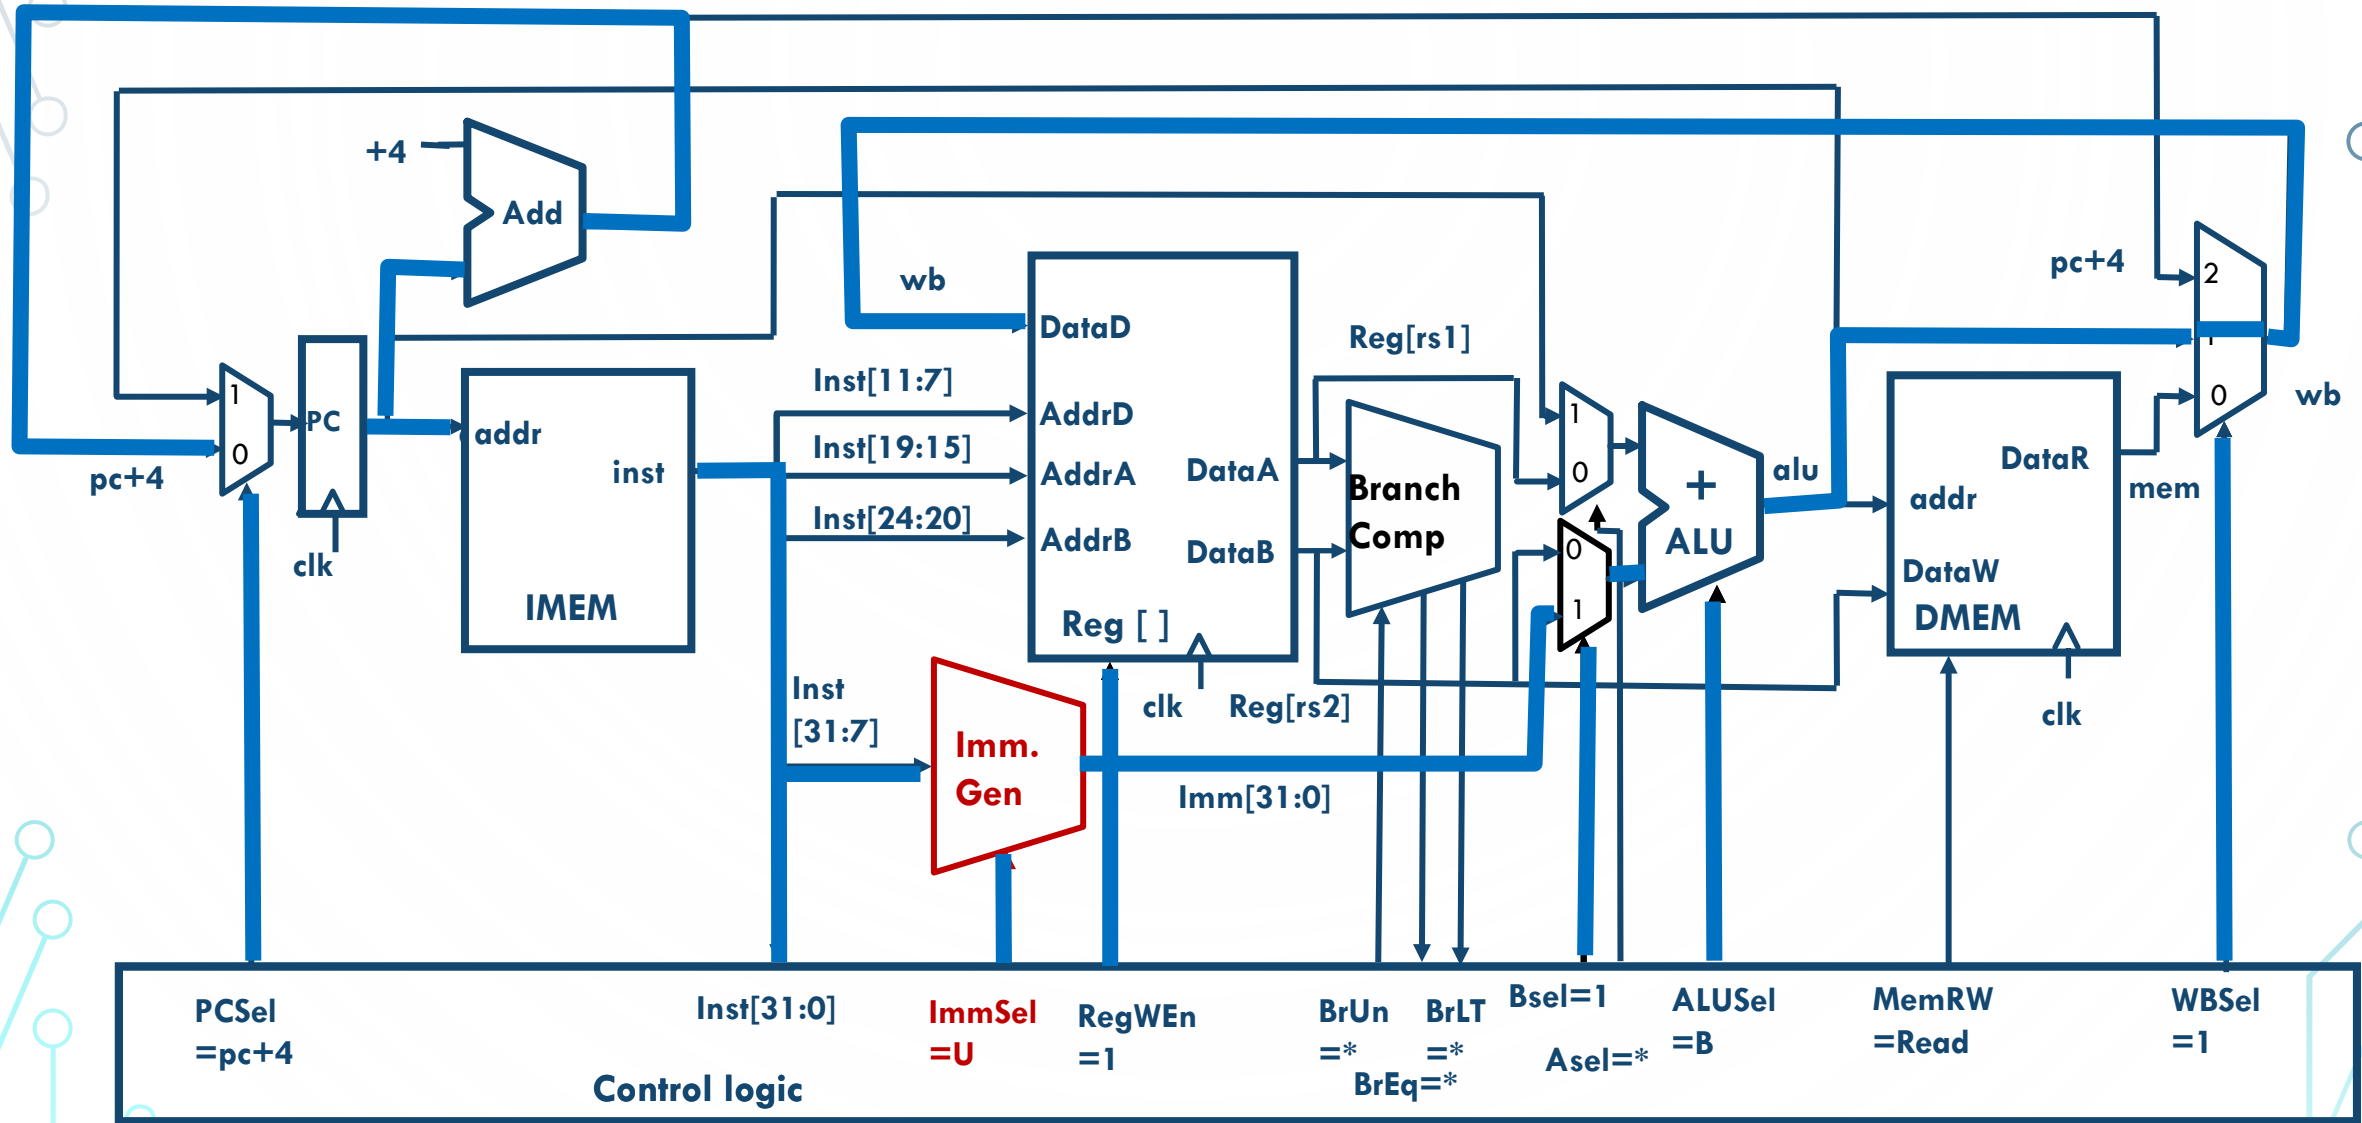
U-Format Instructions

U-Format for “Upper Immediate” Instructions

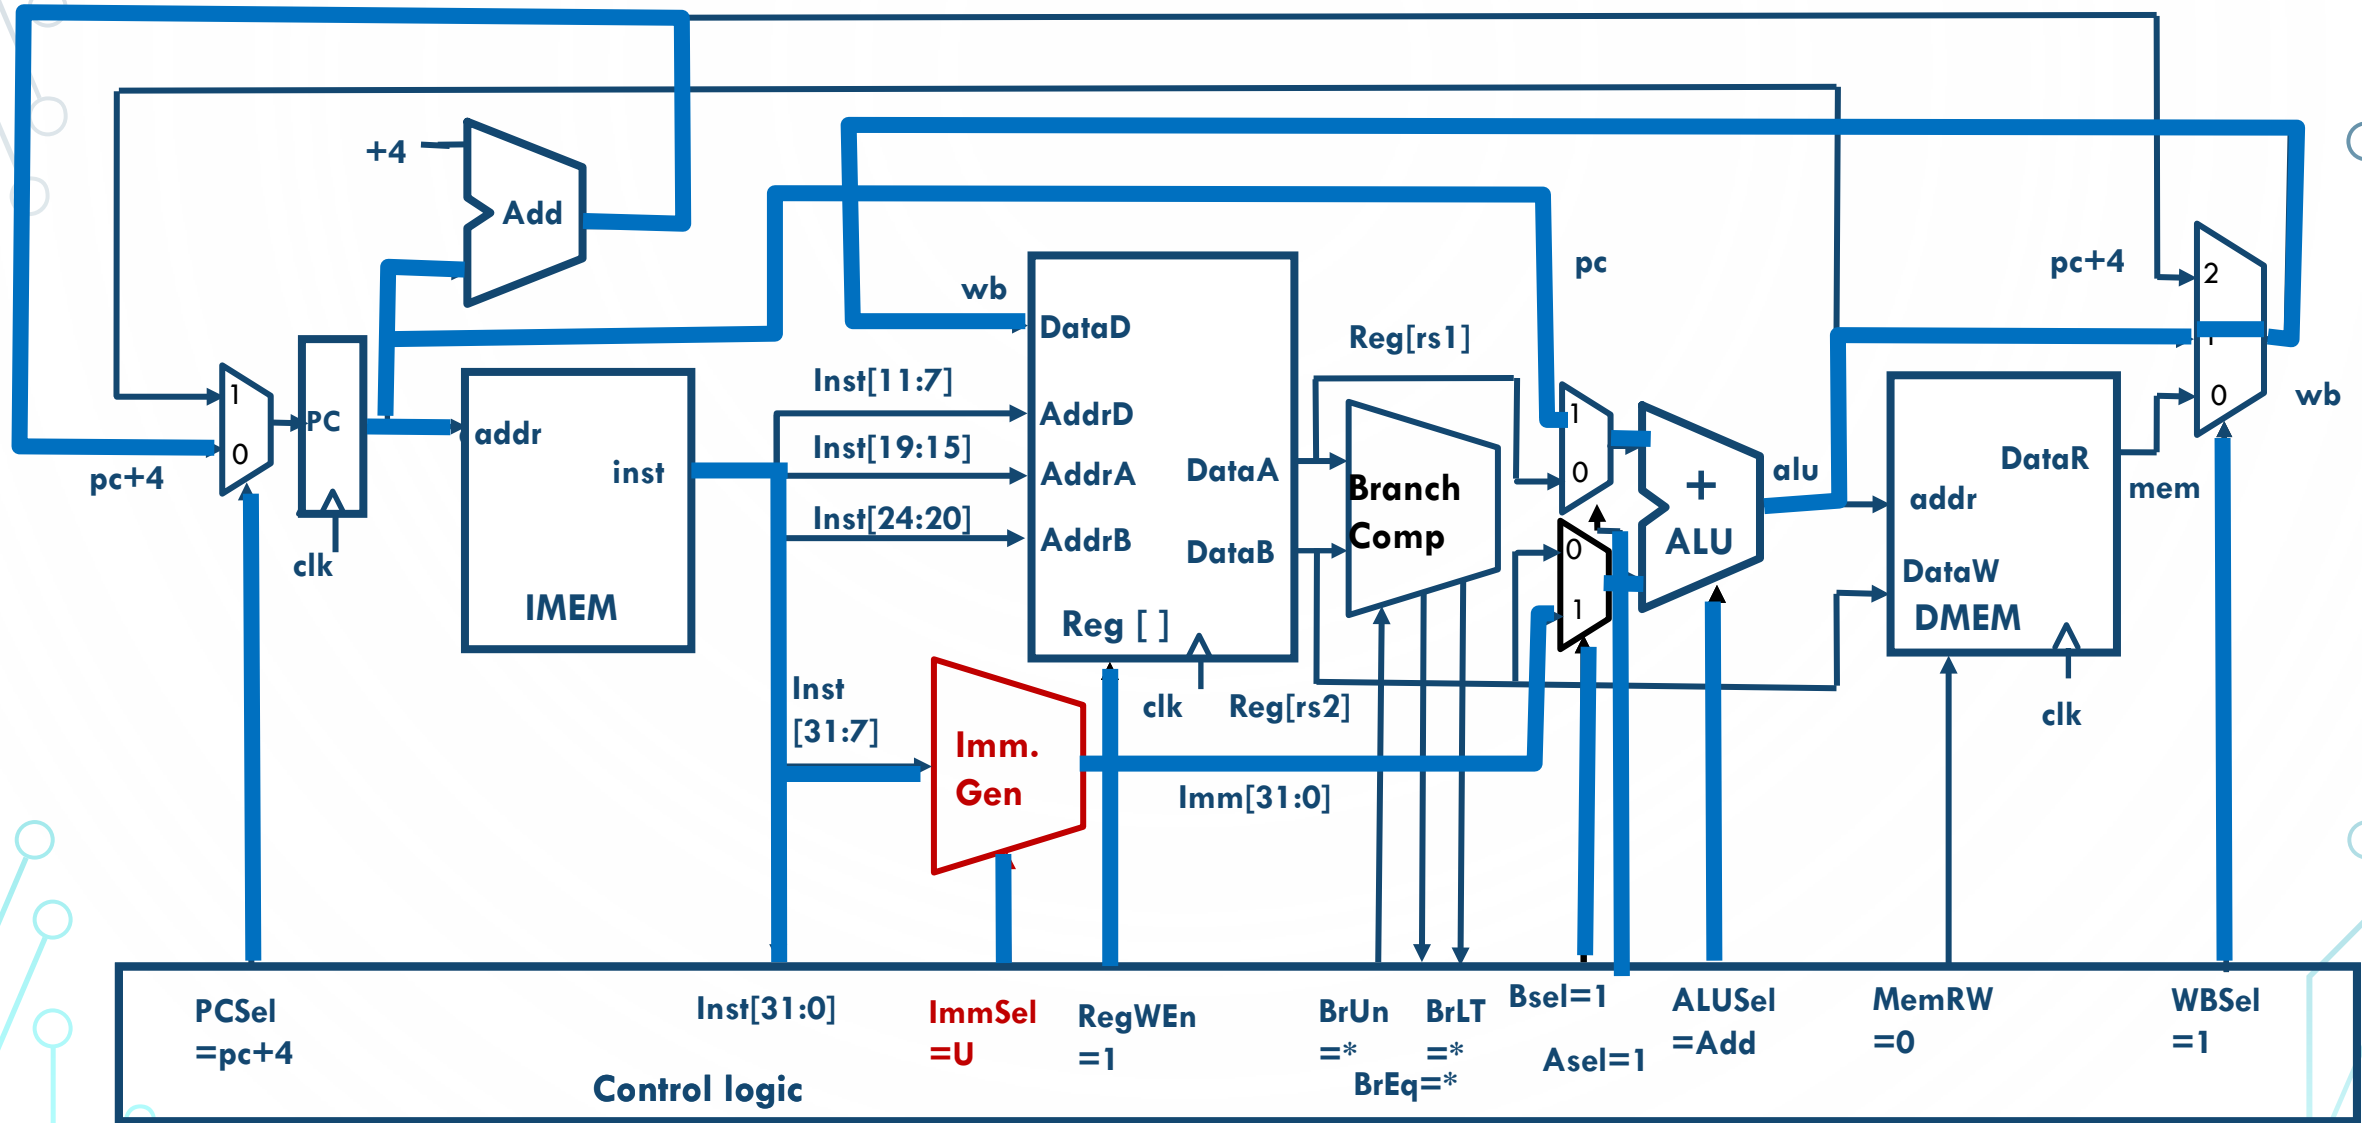


- Has 20-bit immediate in upper 20 bits of 32-bit instruction word
- One destination register, **rd**
- Used for two instructions
 - **lui** – Load Upper Immediate
 - **auipc** – Add Upper Immediate to PC

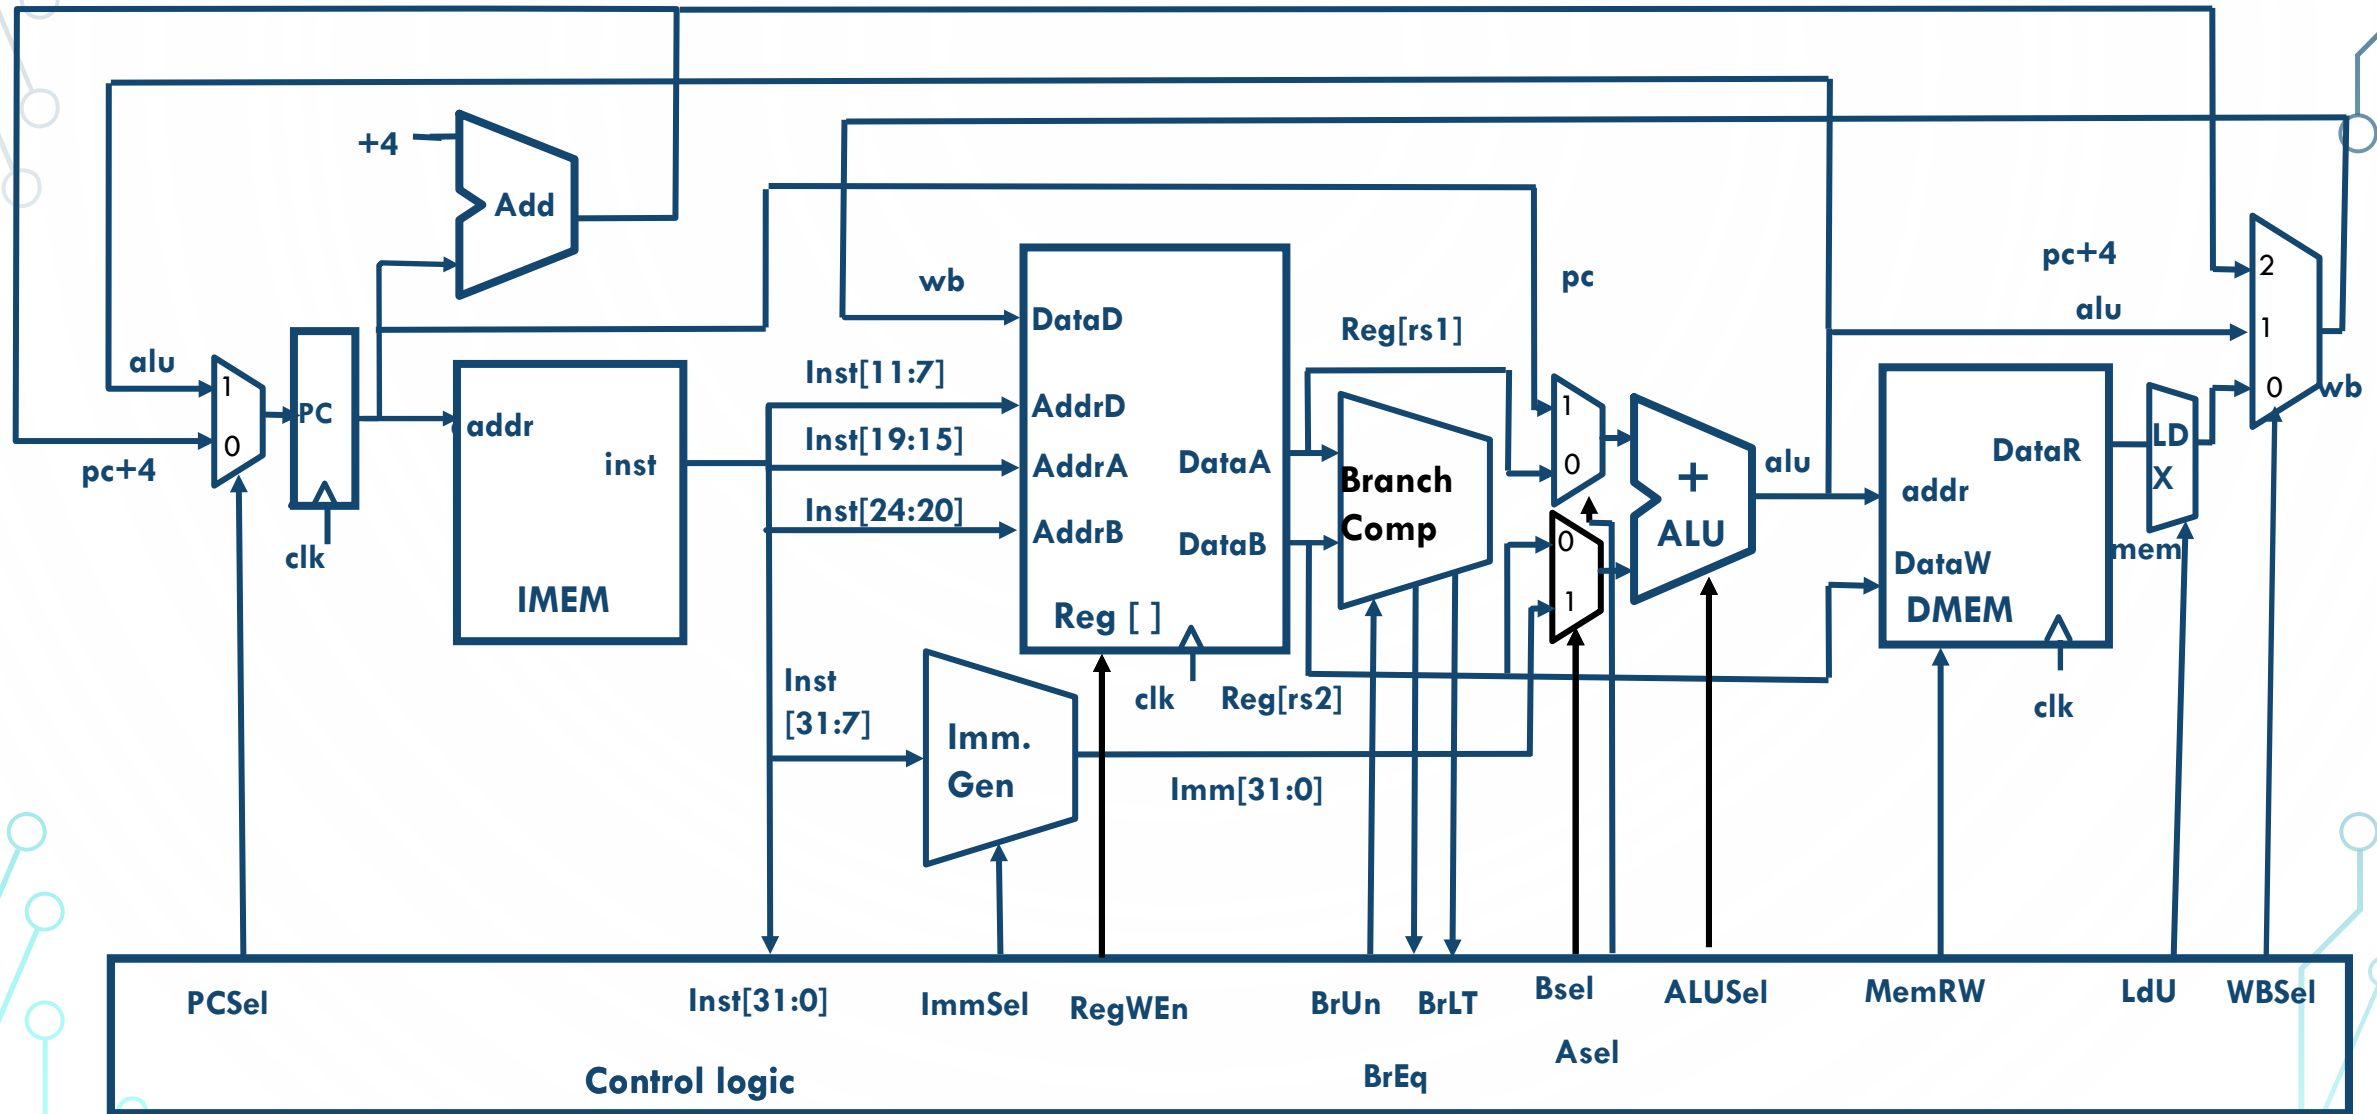
Implementing lui



Implementing `auipc`



Complete RV32I Datapath!



Recap: Complete RV32I ISA

Open

Reference Card

Base Integer Instructions: RV32I							
Category	Name	Fmt	RV32I Base	Category	Name	Fmt	RV32I Base
Shifts	Shift Left Logical	R	SLL rd,rs1,rs2	Loads	Load Byte	I	LB rd,rs1,imm
	Shift Left Log. Imm.	I	SLLI rd,rs1,shamt		Load Halfword	I	LH rd,rs1,imm
	Shift Right Logical	R	SRL rd,rs1,rs2		Load Byte Unsigned	I	LBU rd,rs1,imm
	Shift Right Log. Imm.	I	SRLI rd,rs1,shamt		Load Half Unsigned	I	LHU rd,rs1,imm
	Shift Right Arithmetic	R	SRA rd,rs1,rs2		Load Word	I	LW rd,rs1,imm
				Stores	Store Byte	S	SB rs1,rs2,imm
	Shift Right Arith. Imm.	I	SRAI rd,rs1,shamt		Store Halfword	S	SH rs1,rs2,imm
Arithmetic	ADD	R	ADD rd,rs1,rs2		Store Word	S	SW rs1,rs2,imm
	ADD Immediate	I	ADDI rd,rs1,imm	Branches	Branch =	B	BEQ rs1,rs2,imm
	SUBtract	R	SUB rd,rs1,rs2		Branch ≠	B	BNE rs1,rs2,imm
Logical	Load Upper Imm	U	LUI rd,imm		Branch <	B	BLT rs1,rs2,imm
	Add Upper Imm to PC	U	AUIPC rd,imm		Branch ≥	B	BGE rs1,rs2,imm
	XOR	R	XOR rd,rs1,rs2		Branch < Unsigned	B	BLTU rs1,rs2,imm
	XOR Immediate	I	XORI rd,rs1,imm		Branch ≥ Unsigned	B	BGEU rs1,rs2,imm
	OR	R	OR rd,rs1,rs2	Jump & Link	J&L	J	JAL rd,imm
	OR Immediate	I	ORI rd,rs1,imm		Jump & Link Register	I	JALR rd,rs1,imm
	AND	R	AND rd,rs1,rs2	Synch	Synch thread	I	FENCE
	AND Immediate	I	ANDI rd,rs1,imm				
	Compare Set <	R	SLT rd,rs1,rs2				
	Set < Immediate	I	SLTI rd,rs1,imm	Environment	CALL	I	ECALL
Compare	Set < Unsigned	R	SLTU rd,rs1,rs2		BREAK	I	EBREAK
	Set < Imm Unsigned	I	SLTIU rd,rs1,imm				


- 40 instructions are enough to run any C program

Summary of RISC-V Instruction Formats

31	30	25	24	21	20	19	15	14	12	11	8	7	6	0					
funct7					rs2			rs1			funct3			rd			opcode		R-type
imm[11:0]							rs1			funct3			rd			opcode		I-type	
imm[11:5]				rs2			rs1			funct3			imm[4:0]			opcode		S-type	
imm[12 10:5]				rs2			rs1			funct3			imm[4:1 11]			opcode		B-type	
imm[31:12]										rd			opcode		U-type				
imm[20 10:1 11]]							imm[19:12]					rd			opcode		J-type		

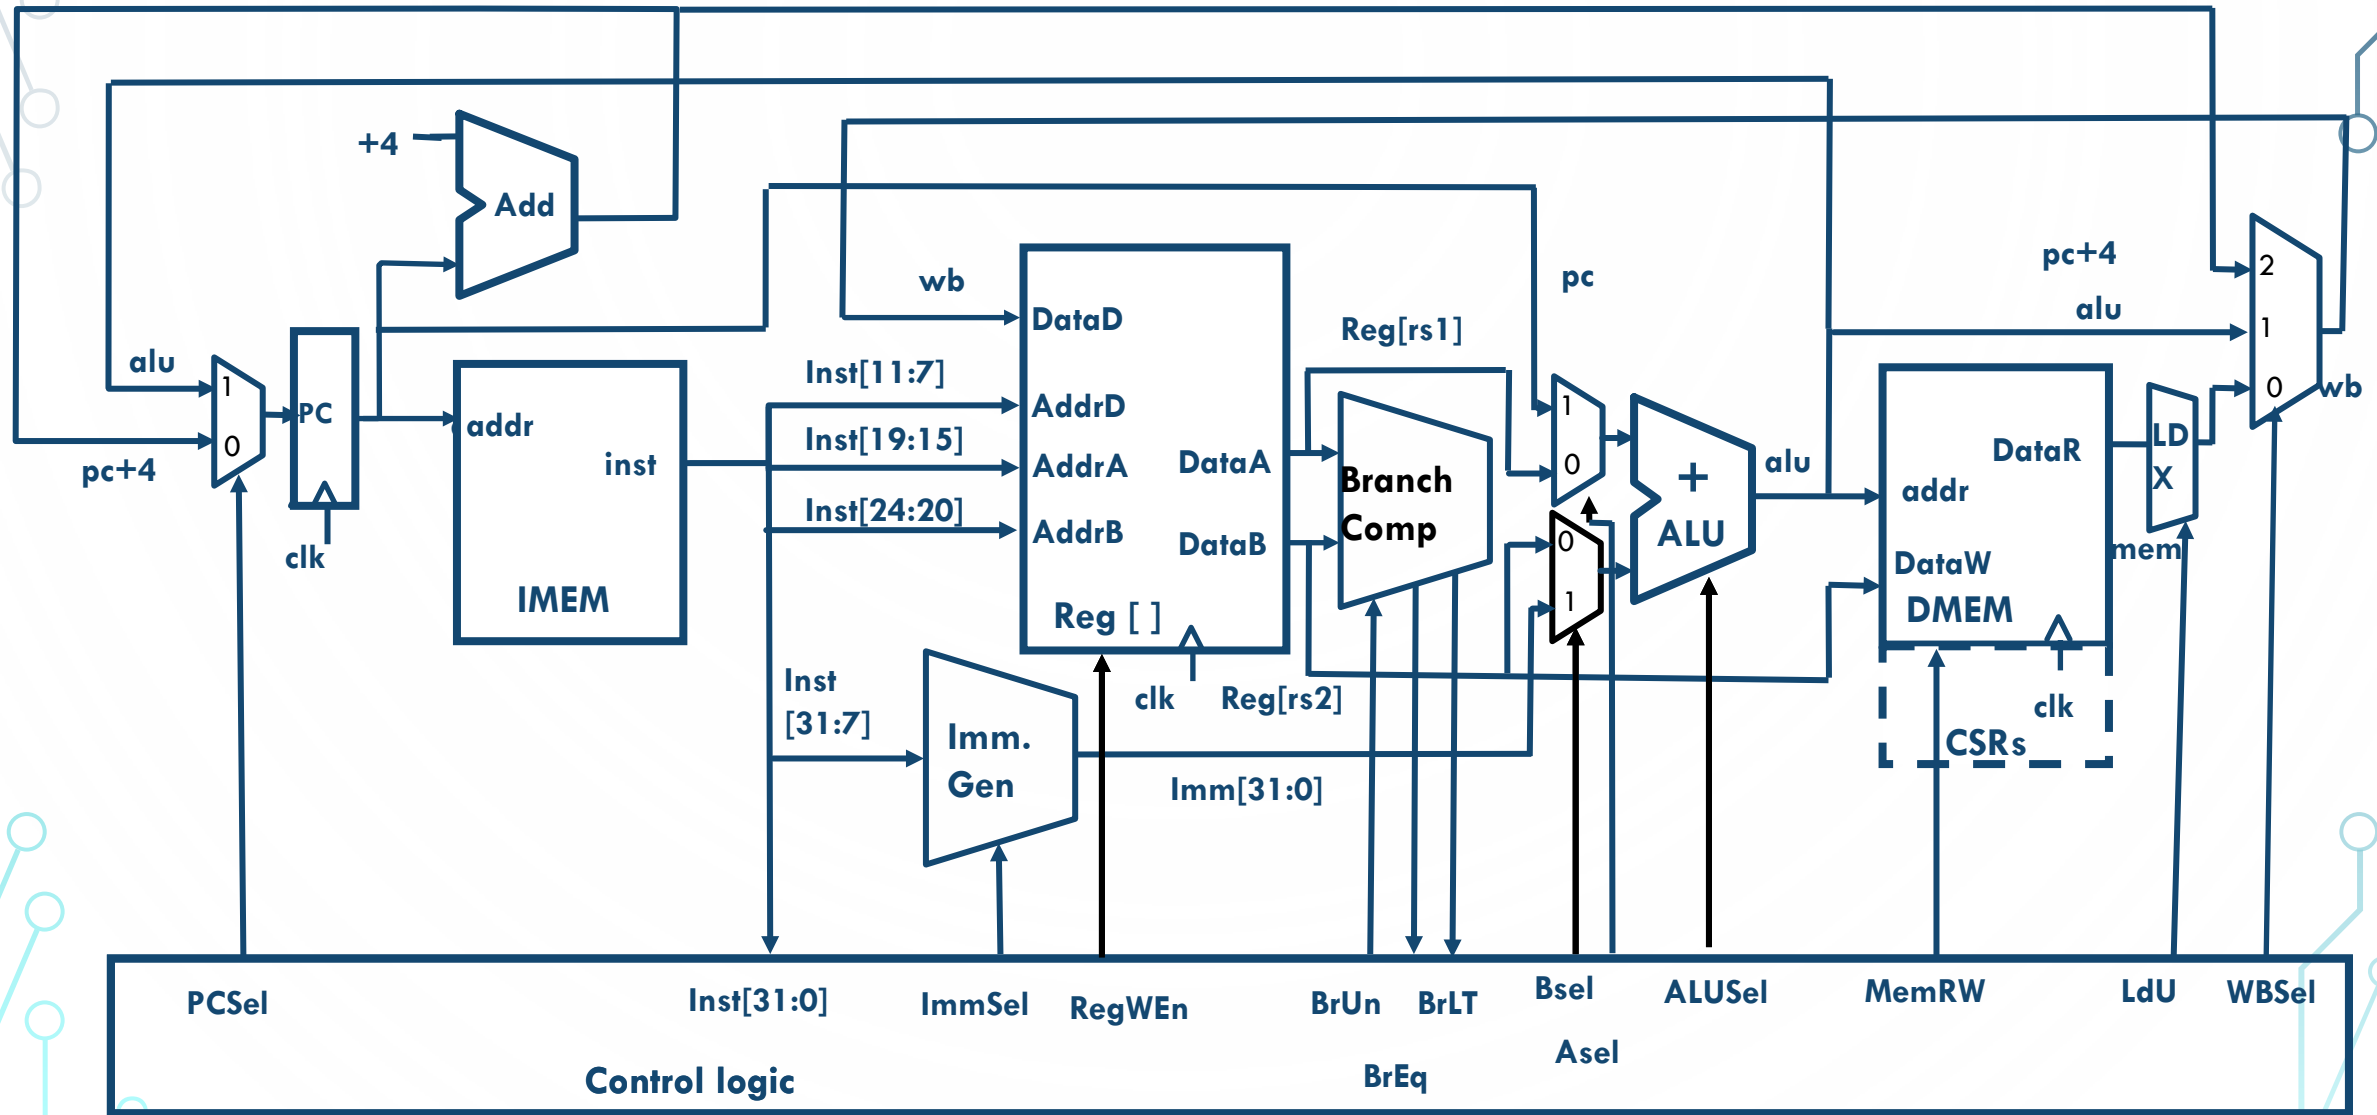
Control and Status Registers (CSRs)

- 4096 CSRs in a separate address space



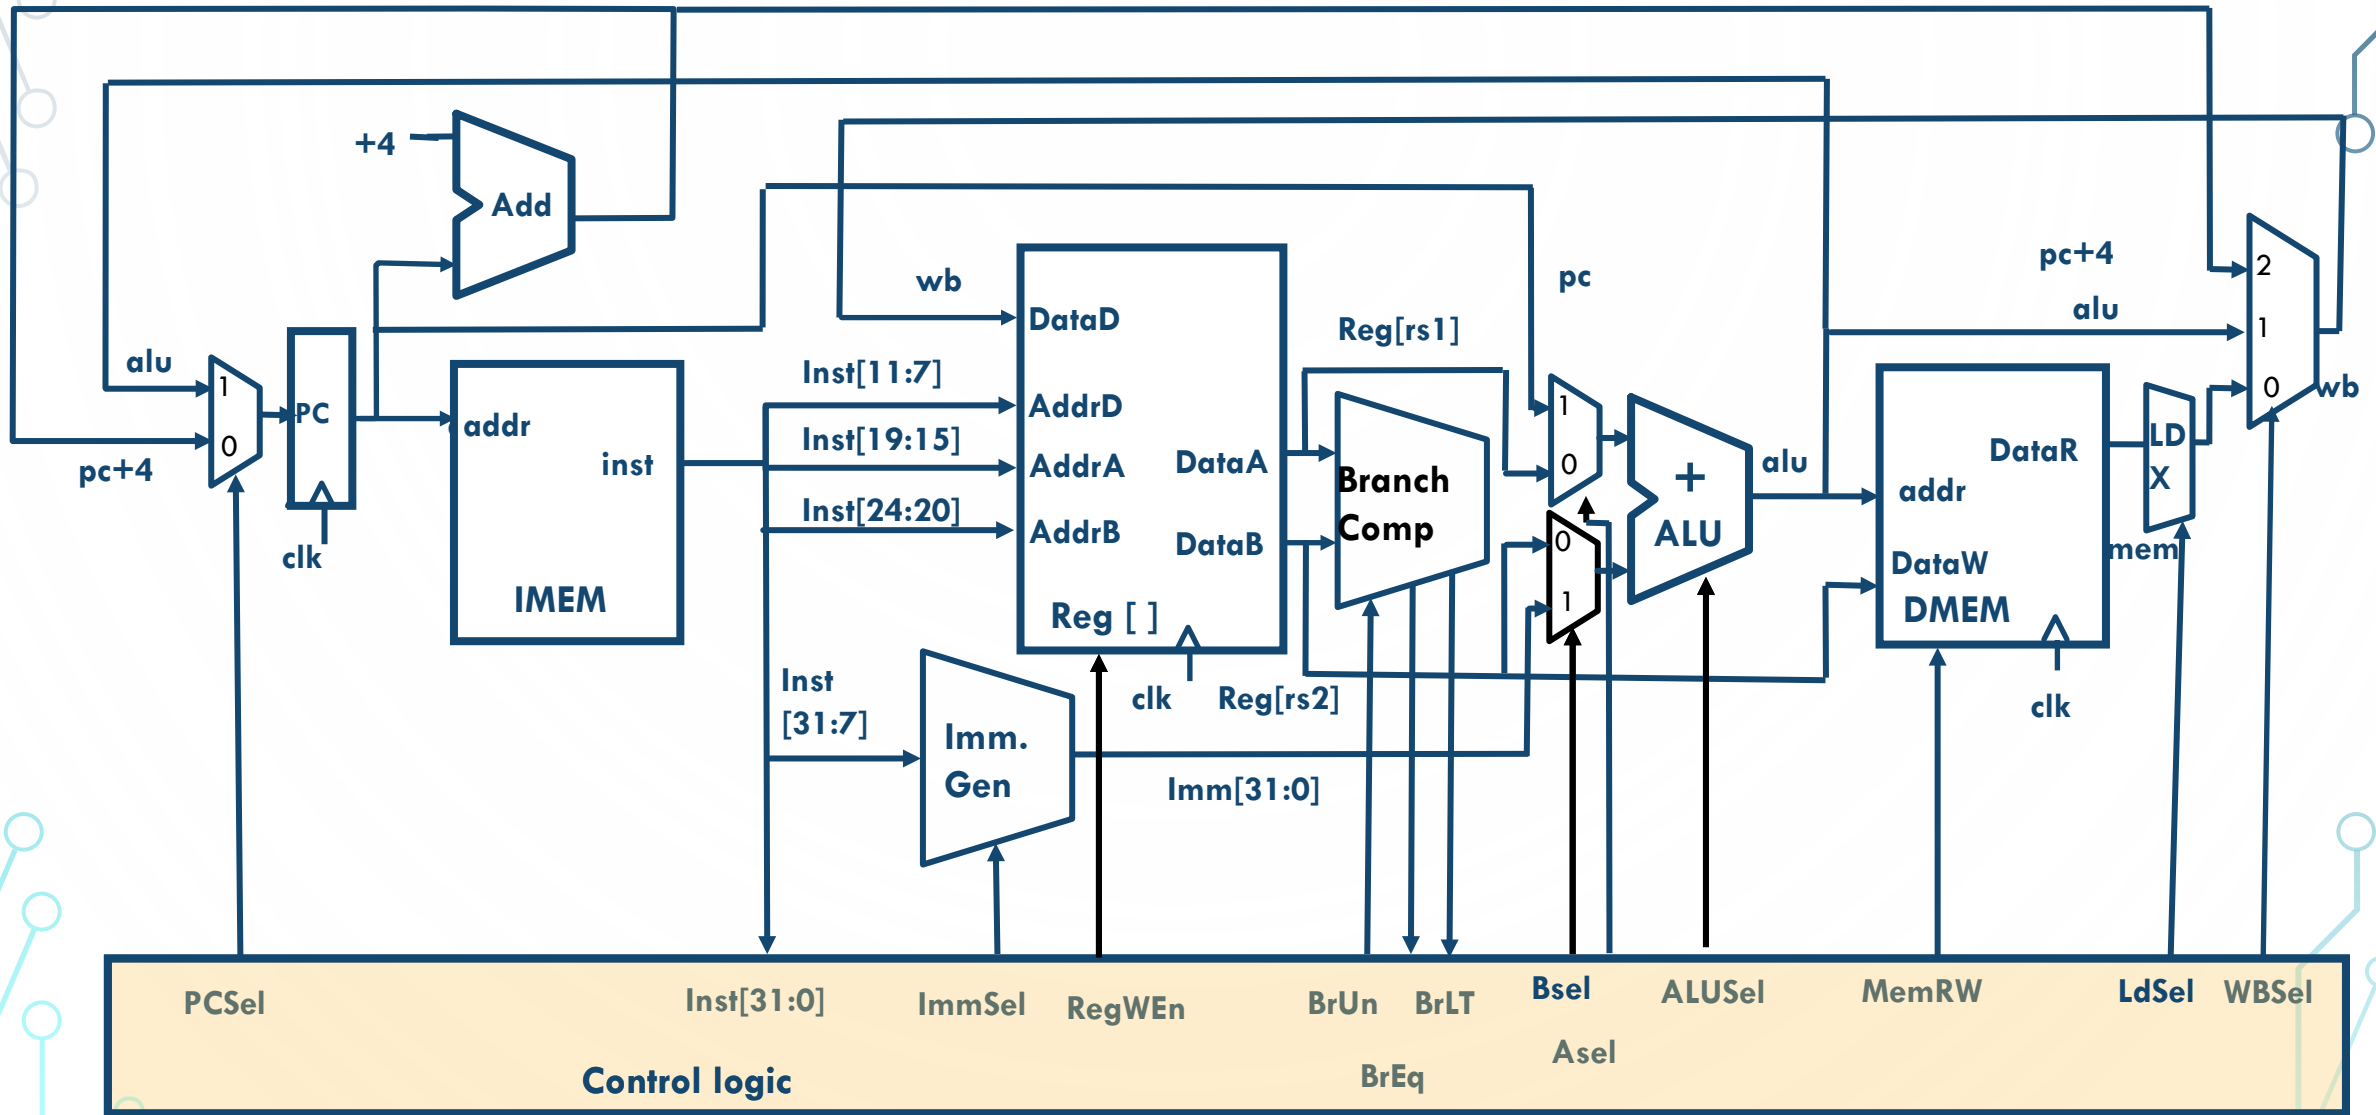
- **csrrw** reads the old value of CSR, zero-extends and writes to **rd**
- Initial value of **rs1** is written to CSR
- Pseudo-instructions: **csrr**, **csrw** (**csrrw** **x0**, **csr**, **rs1**)

Add the `csrrw`!

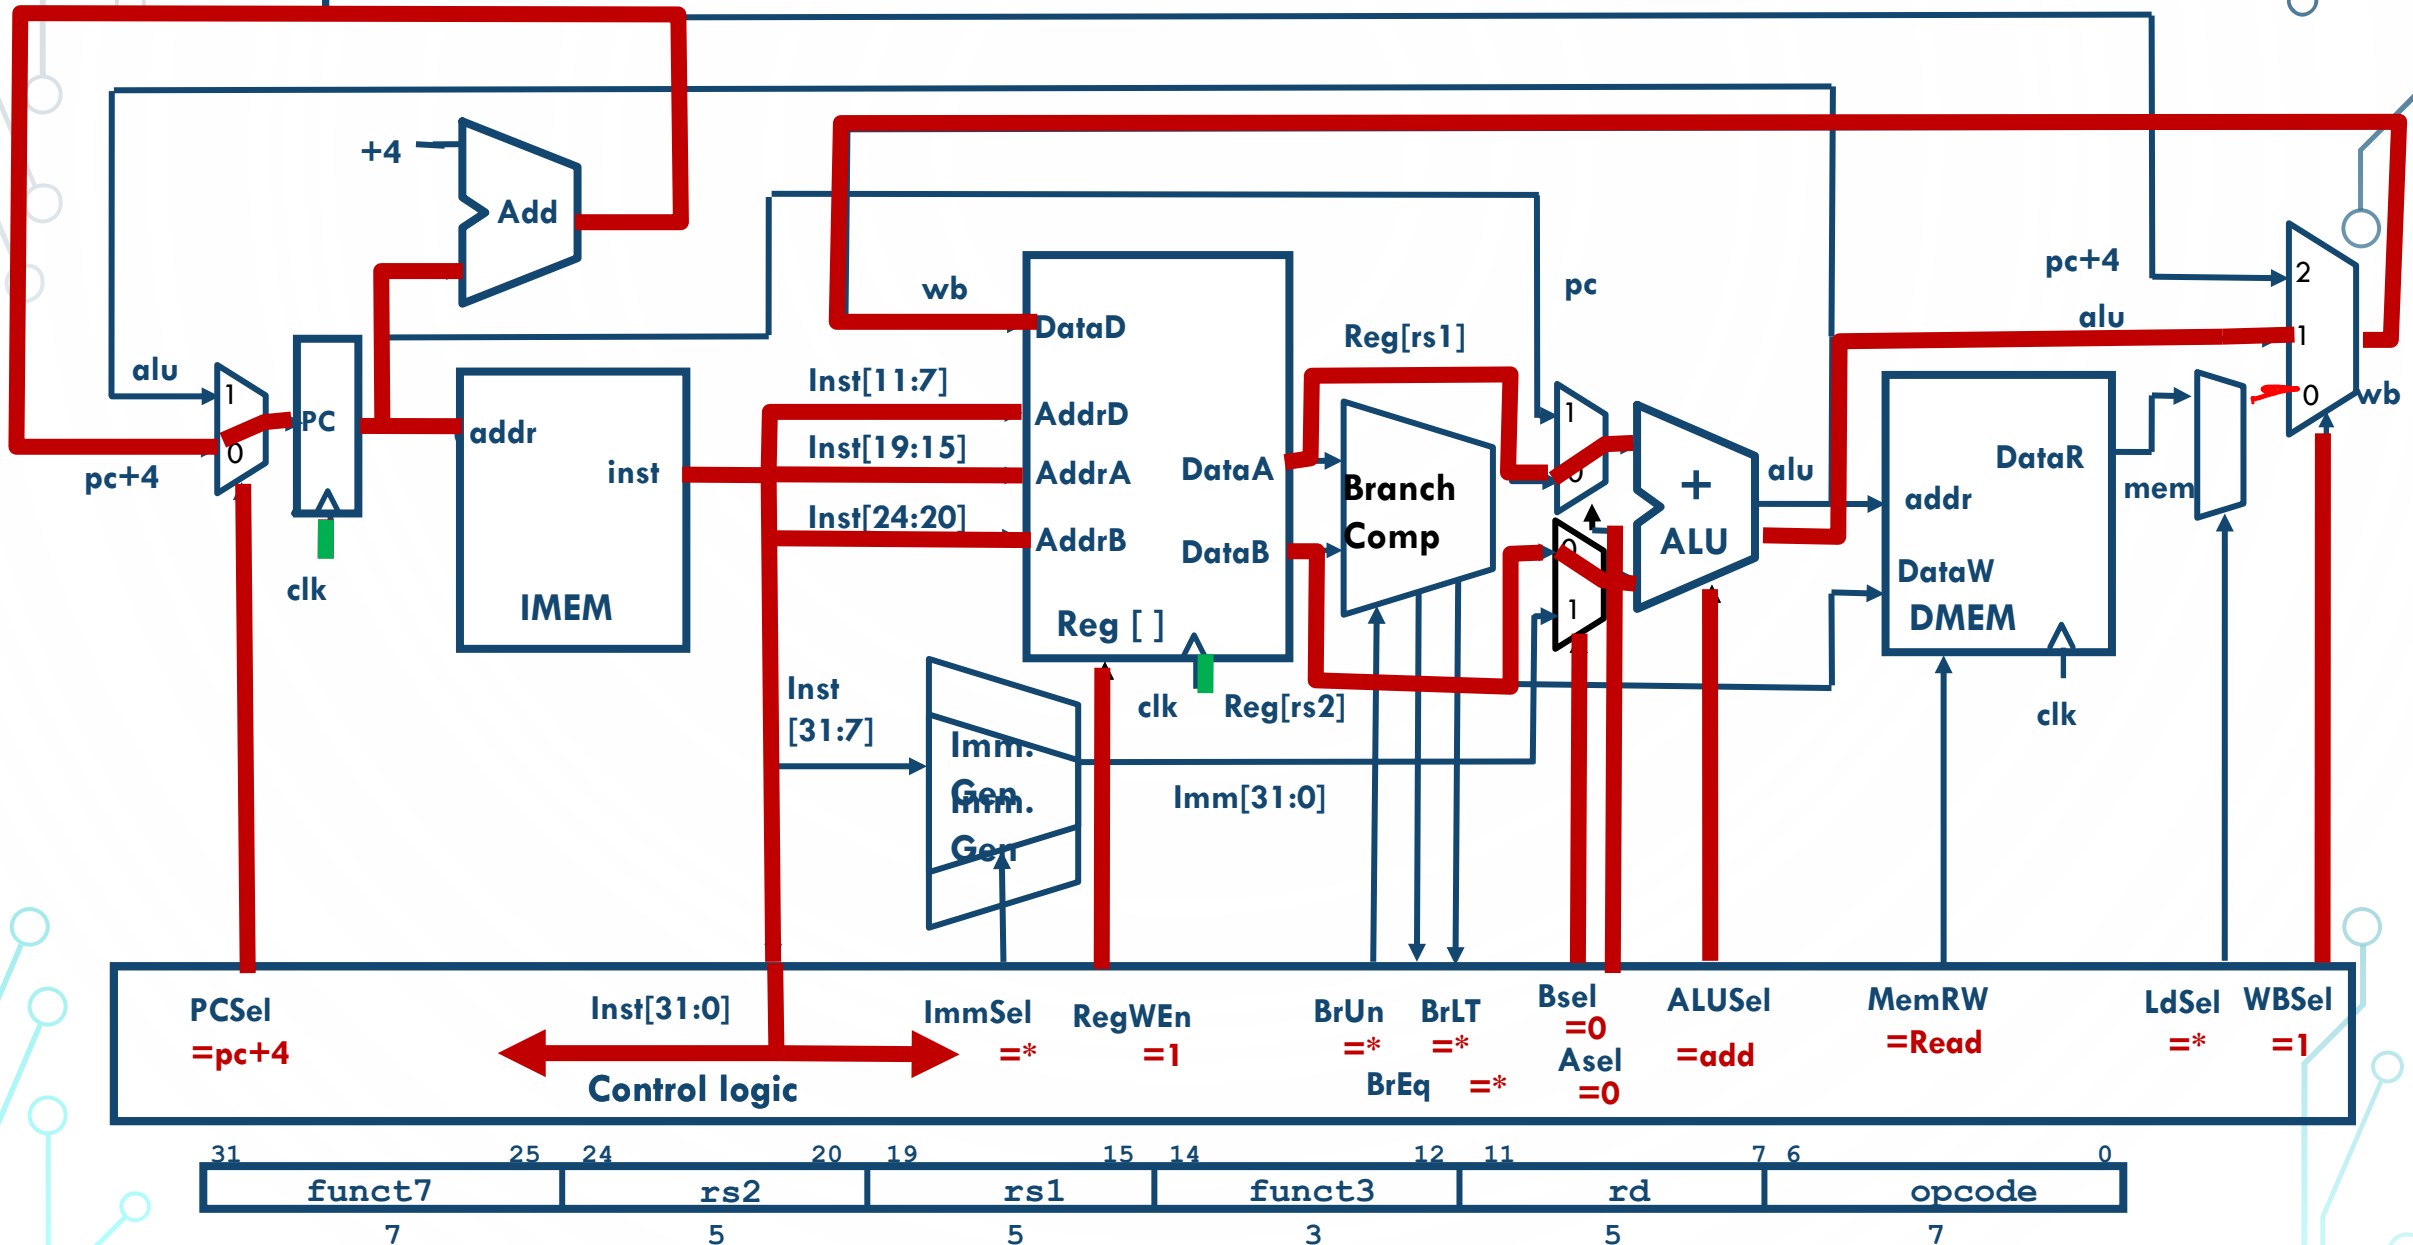




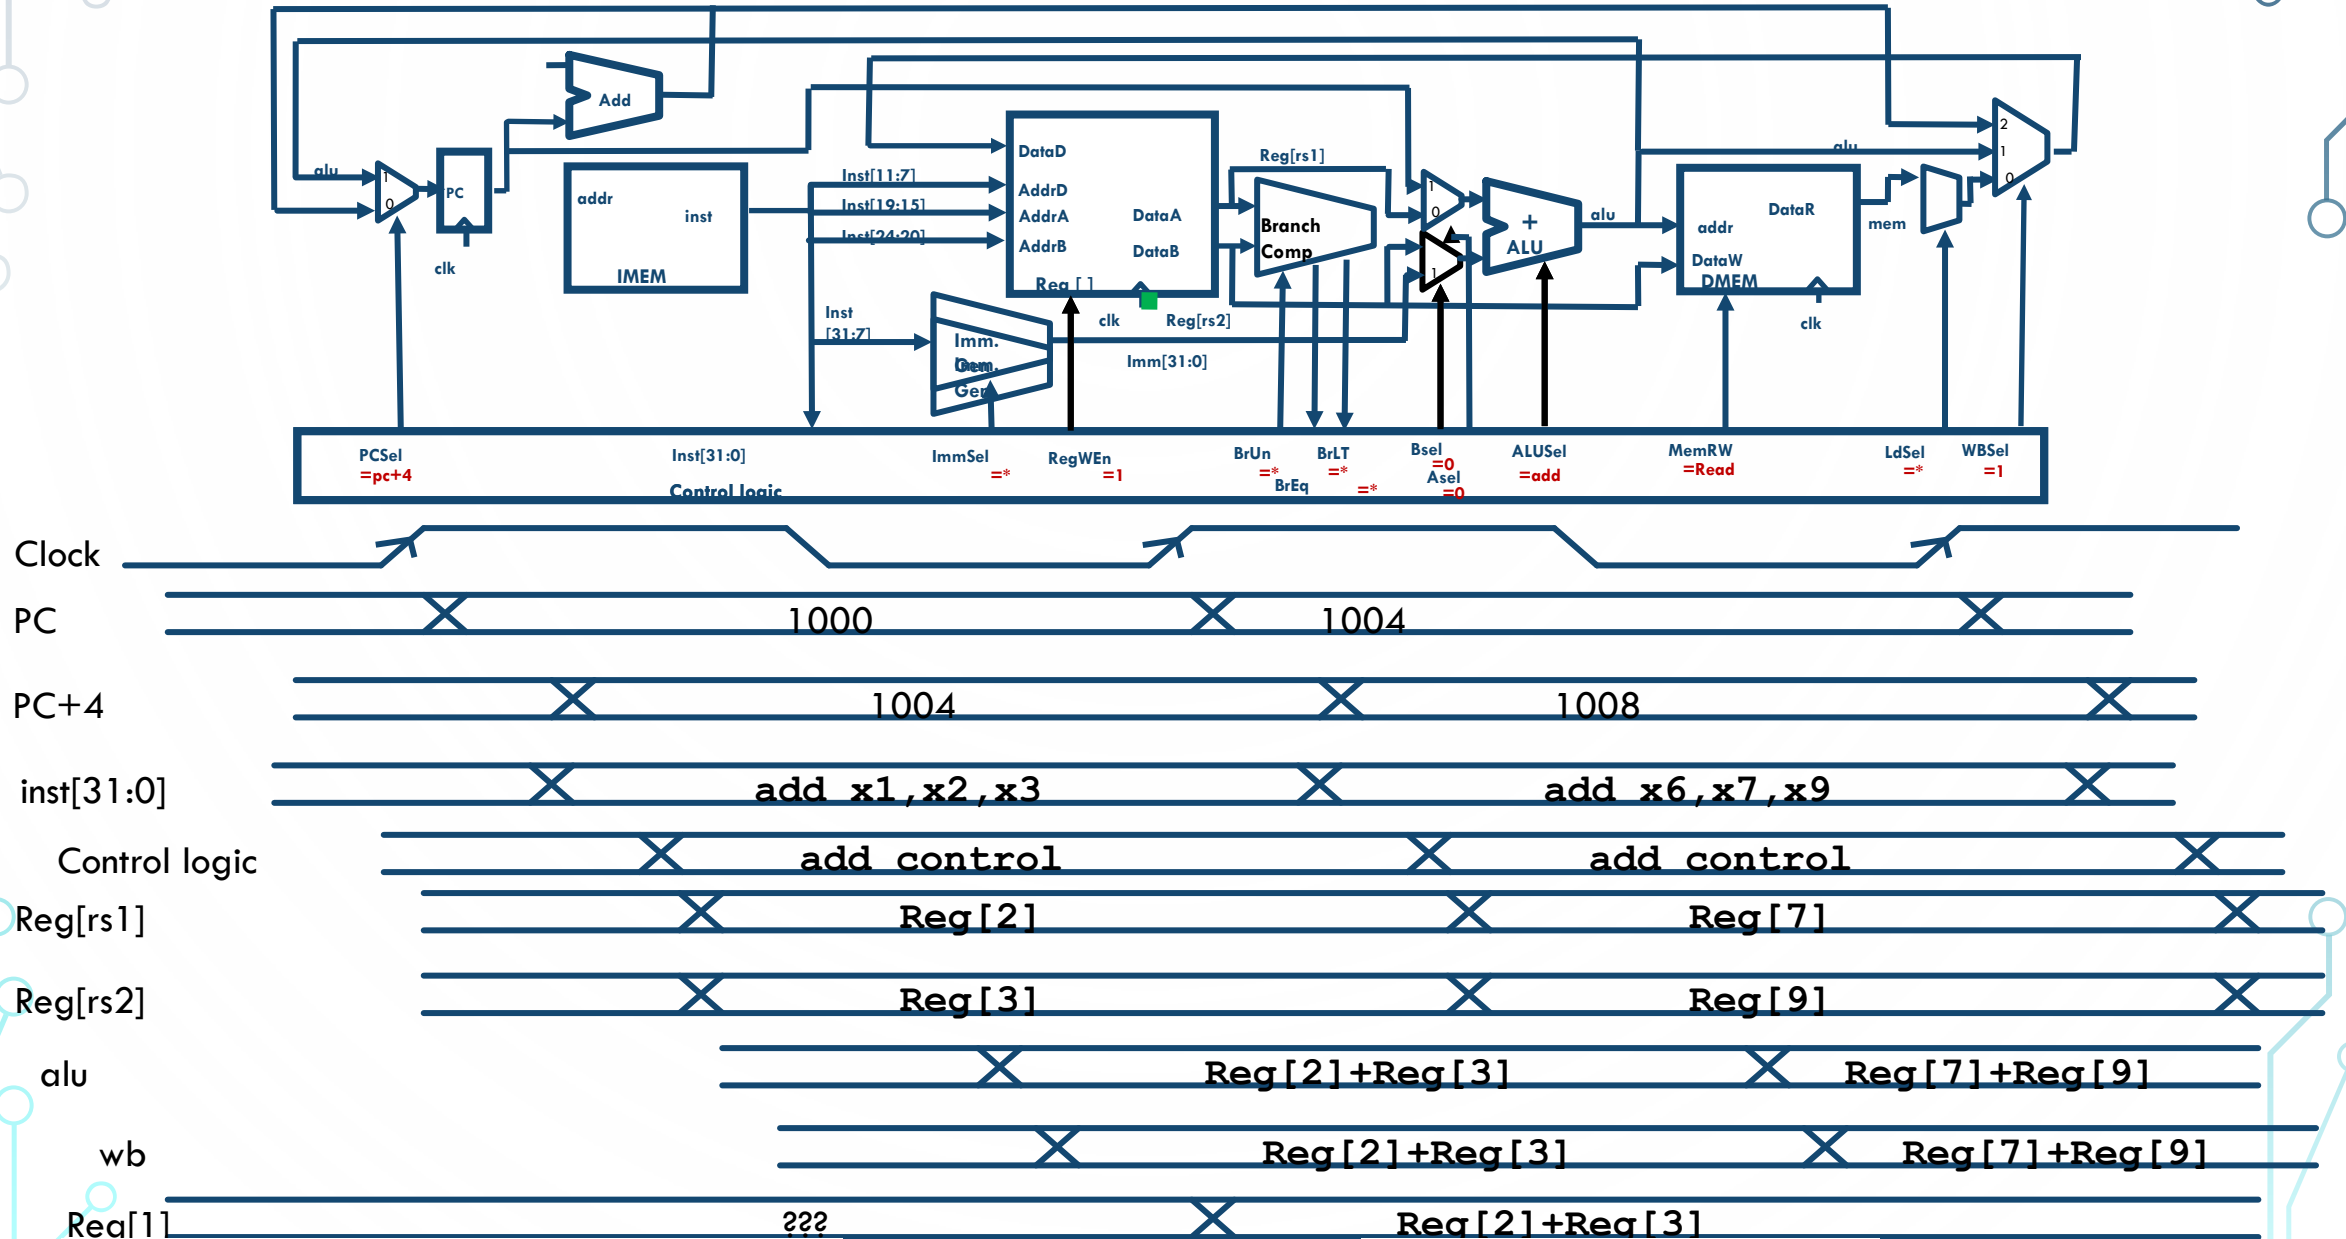
RISC-V Control Logic

[illegible]

Example: add



add Execution



Control Logic Truth Table

Inst[31:0]	BrEq	BrLT	PCSel	ImmSel	BrUn	ASel	BSel	ALUSel	MemRW	RegWEn	WBSel
add	*	*	+4	*	*	Reg	Reg	Add	Read	1	ALU
sub	*	*	+4	*	*	Reg	Reg	Sub	Read	1	ALU
(R-R Op)	*	*	+4	*	*	Reg	Reg	(Op)	Read	1	ALU
addi	*	*	+4	I	*	Reg	Imm	Add	Read	1	ALU
lw	*	*	+4	I	*	Reg	Imm	Add	Read	1	Mem
sw	*	*	+4	S	*	Reg	Imm	Add	Write	0	*
beq	0	*	+4	B	*	PC	Imm	Add	Read	0	*
beq	1	*	ALU	B	*	PC	Imm	Add	Read	0	*
bne	0	*	ALU	B	*	PC	Imm	Add	Read	0	*
bne	1	*	+4	B	*	PC	Imm	Add	Read	0	*
blt	*	1	ALU	B	0	PC	Imm	Add	Read	0	*
bltu	*	1	ALU	B	1	PC	Imm	Add	Read	0	*
jalr	*	*	ALU	I	*	Reg	Imm	Add	Read	1	PC+4
jal	*	*	ALU	J	*	PC	Imm	Add	Read	1	PC+4
auipc	*	*	+4	U	*	PC	Imm	Add	Read	1	ALU

RV32I, a nine-bit ISA!

imm[31:12]					rd	01101	11	LUI
imm[31:12]					rd	00101	11	AUIPC
imm[20 10:1 11 19:12]					rd	11011	11	JAL
imm[11:0]		rs1		000	rd	11001	11	JALR
imm[12 10:5]	rs2	rs1		000	imm[4:1 11]	11000	11	BEQ
imm[12 10:5]	rs2	rs1		001	imm[4:1 11]	11000	11	BNE
imm[12 10:5]	rs2	rs1		100	imm[4:1 11]	11000	11	BLT
imm[12 10:5]	rs2	rs1		101	imm[4:1 11]	11000	11	BGE
imm[12 10:5]	rs2	rs1		110	imm[4:1 11]	11000	11	BLTU
imm[12 10:5]	rs2	rs1		111	imm[4:1 11]	11000	11	BGEU
imm[11:0]		rs1		000	rd	00000	11	LB
imm[11:0]		rs1		001	rd	00000	11	LH
imm[11:0]		rs1		010	rd	00000	11	LW
imm[11:0]		rs1		100	rd	00000	11	LBU
imm[11:0]		rs1		101	rd	00000	11	LHU
imm[11:5]	rs2	rs1		000	imm[4:0]	01000	11	SB
imm[11:5]	rs2	rs1		001	imm[4:0]	01000	11	SH
imm[11:5]	rs2	rs1		010	imm[4:0]	01000	11	SW
imm[11:0]		rs1		000	rd	00100	11	ADDI
imm[11:0]		rs1		010	rd	00100	11	SLTI
imm[11:0]		rs1		011	rd	00100	11	SLTIU
imm[11:0]		rs1		100	rd	00100	11	XORI
imm[11:0]		rs1		110	rd	00100	11	ORI
imm[11:0]		rs1		111	rd	00100	11	ANDI
0000000	shamt	rs1		001	rd	00100	11	SLLI
0000000	shamt	rs1		101	rd	00100	11	SRLI
0100000	shamt	rs1		101	rd	00100	11	SRAI
0000000	rs2	rs1		000	rd	01100	11	ADD
0100000	rs2	rs1		000	rd	01100	11	SUB
0000000	rs2	rs1		001	rd	01100	11	SLL
0000000	rs2	rs1		010	rd	01100	11	SLT
0000000	rs2	rs1		011	rd	01100	11	SLTU
0000000	rs2	rs1		100	rd	01100	11	XOR
0000000	rs2	rs1		101	rd	01100	11	SRL
0100000	rs2	rs1		101	rd	01100	11	SRA
0000000	rs2	rs1		110	rd	01100	11	OR
0000000	rs2	rs1		111	rd	01100	11	AND

inst[30]

inst[14:12]

inst[6:2]

Instruction type encoded by using only 9 bits: inst[30],inst[14:12], inst[6:2]

Control Realization Options



- ROM

- “Read-Only Memory”
- Regular structure
- Can be easily reprogrammed
 - fix errors
 - add instructions

- Combinatorial Logic

- Start from a truth table
- More compact, faster
- Use synthesis tools

inst[30],
inst[14:12],
inst[6:2]



Control bits

Combinational Logic Control

- Decoder is typically hierarchical
 - First decode opcode, and figure out instruction type
 - E.g. branches are $\text{Inst}[6:2] = 11000$
 - Then determine the actual instruction
 - $\text{Inst}[30] + \text{Inst}[14:12]$
- Modularity helps simplify and speed up logic
 - Narrows problem space for logic synthesis

Combinational Logic Control

- Simple example: BrUn

inst[14:12] inst[6:2]

imm[12:10:5]	rs2	rs1	000	imm[4:1:11]	1100011	BEQ
imm[12:10:5]	rs2	rs1	001	imm[4:1:11]	1100011	BNE
imm[12:10:5]	rs2	rs1	100	imm[4:1:11]	1100011	BLT
imm[12:10:5]	rs2	rs1	101	imm[4:1:11]	1100011	BGE
imm[12:10:5]	rs2	rs1	110	imm[4:1:11]	1100011	BLTU
imm[12:10:5]	rs2	rs1	111	imm[4:1:11]	1100011	BGEU

inst[14:13]

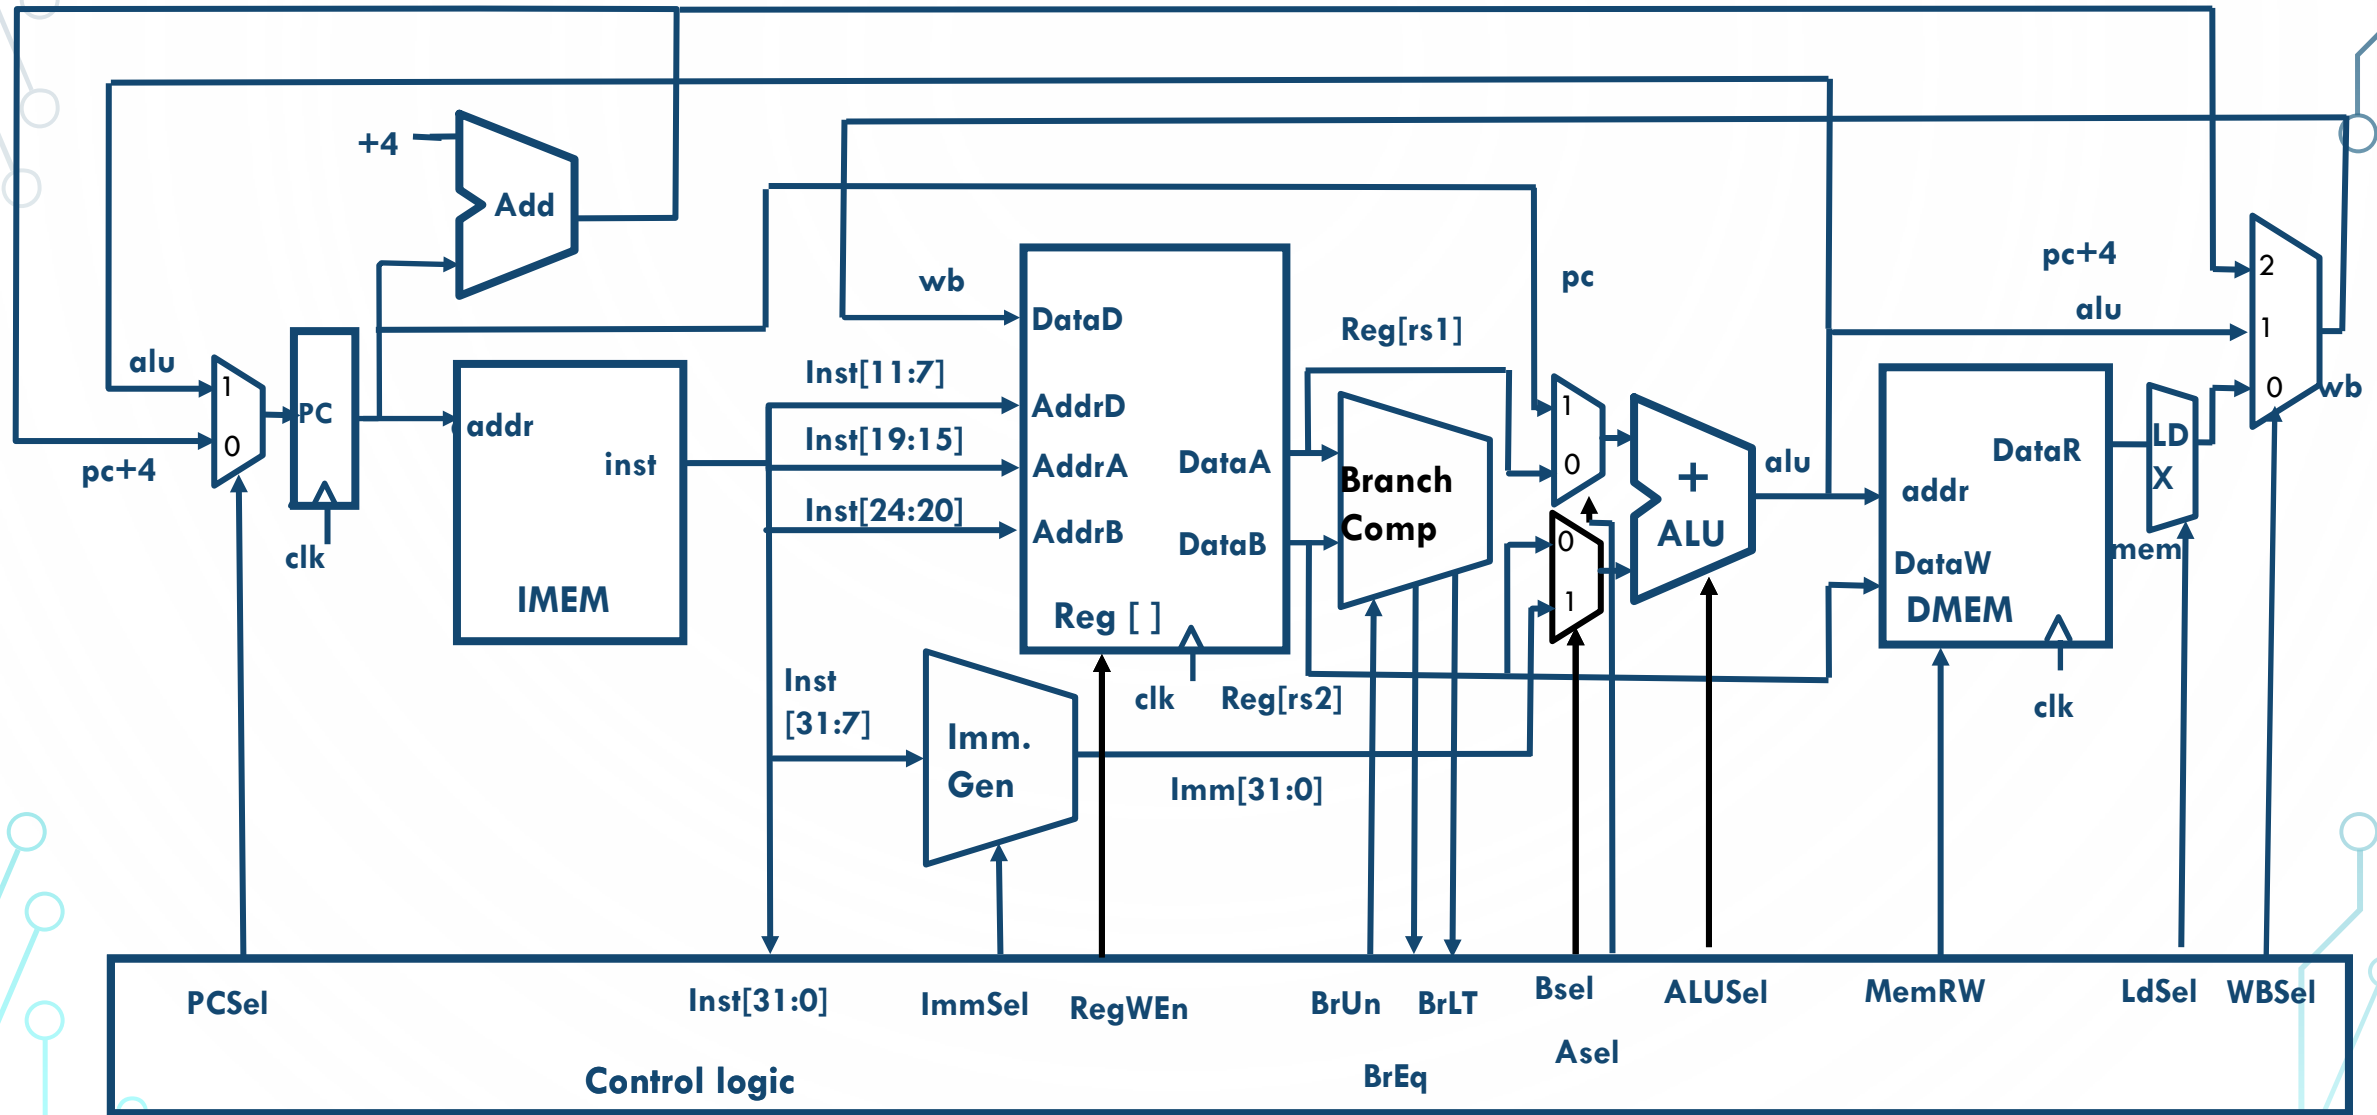
	00	01	11	10
inst[12]				
0				
1				

- How to decode whether BrUn is 1?

- BrUn = Inst [13] • Branch

- Branch = Inst[6] • Inst[5] • !Inst[4] • !Inst[3] • !Inst[2]

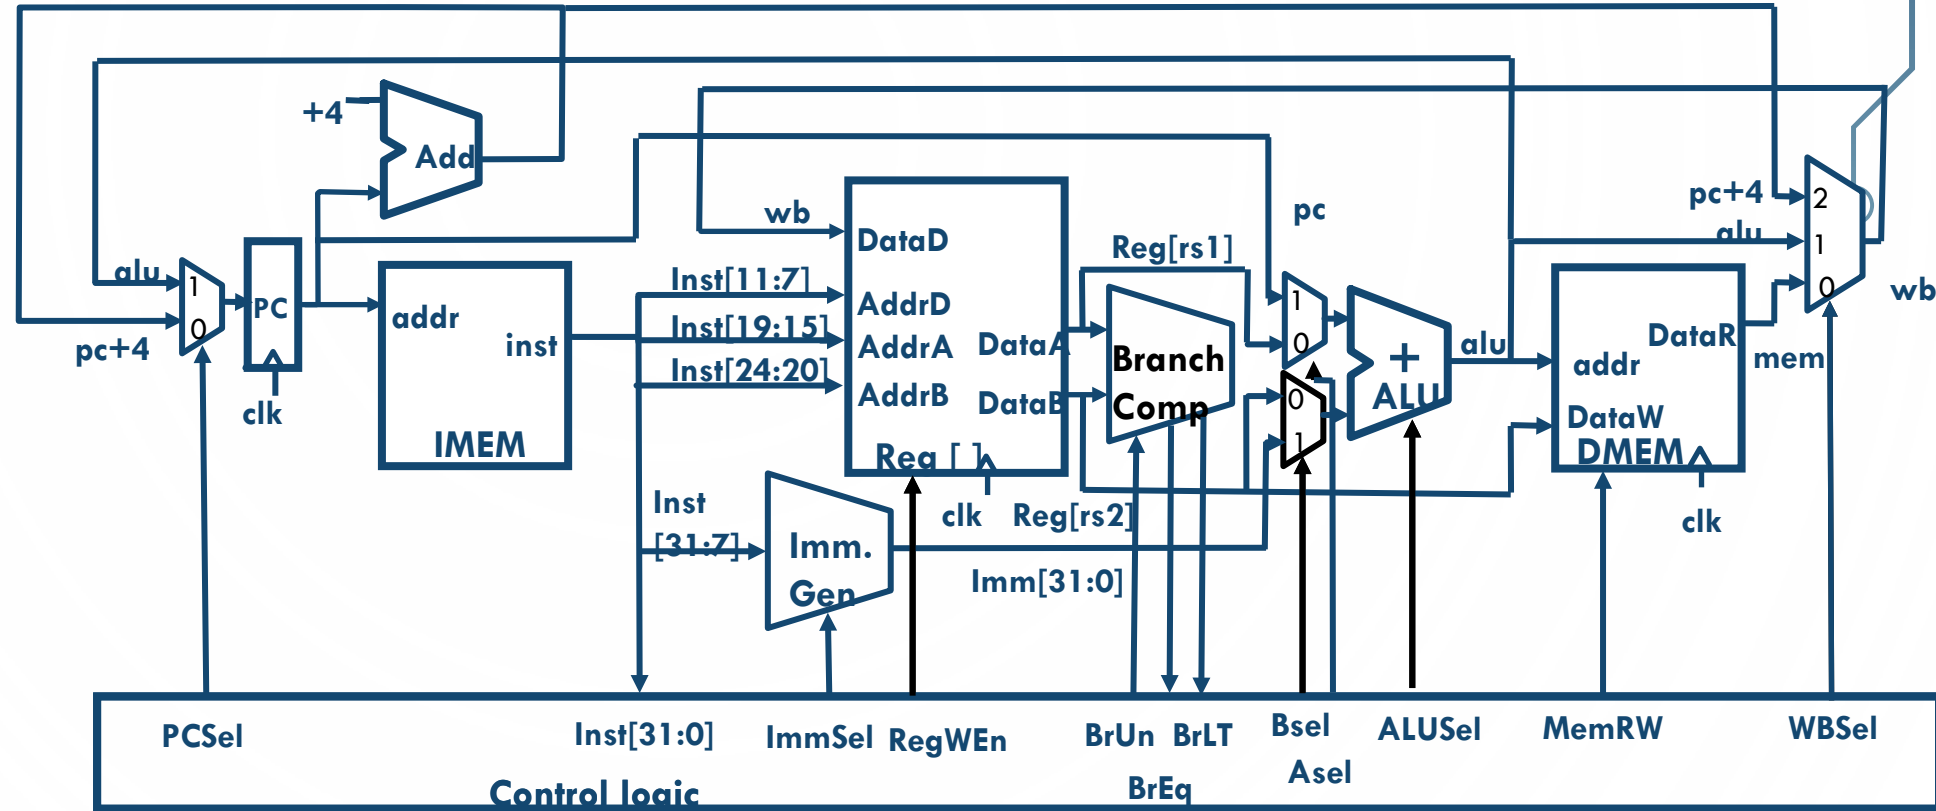
Complete RV32I Datapath with Control



Peer Instruction(s): Critical Path yellkey: crime

Critical path for a addi

$$R[rd] = R[rs1] + imm$$

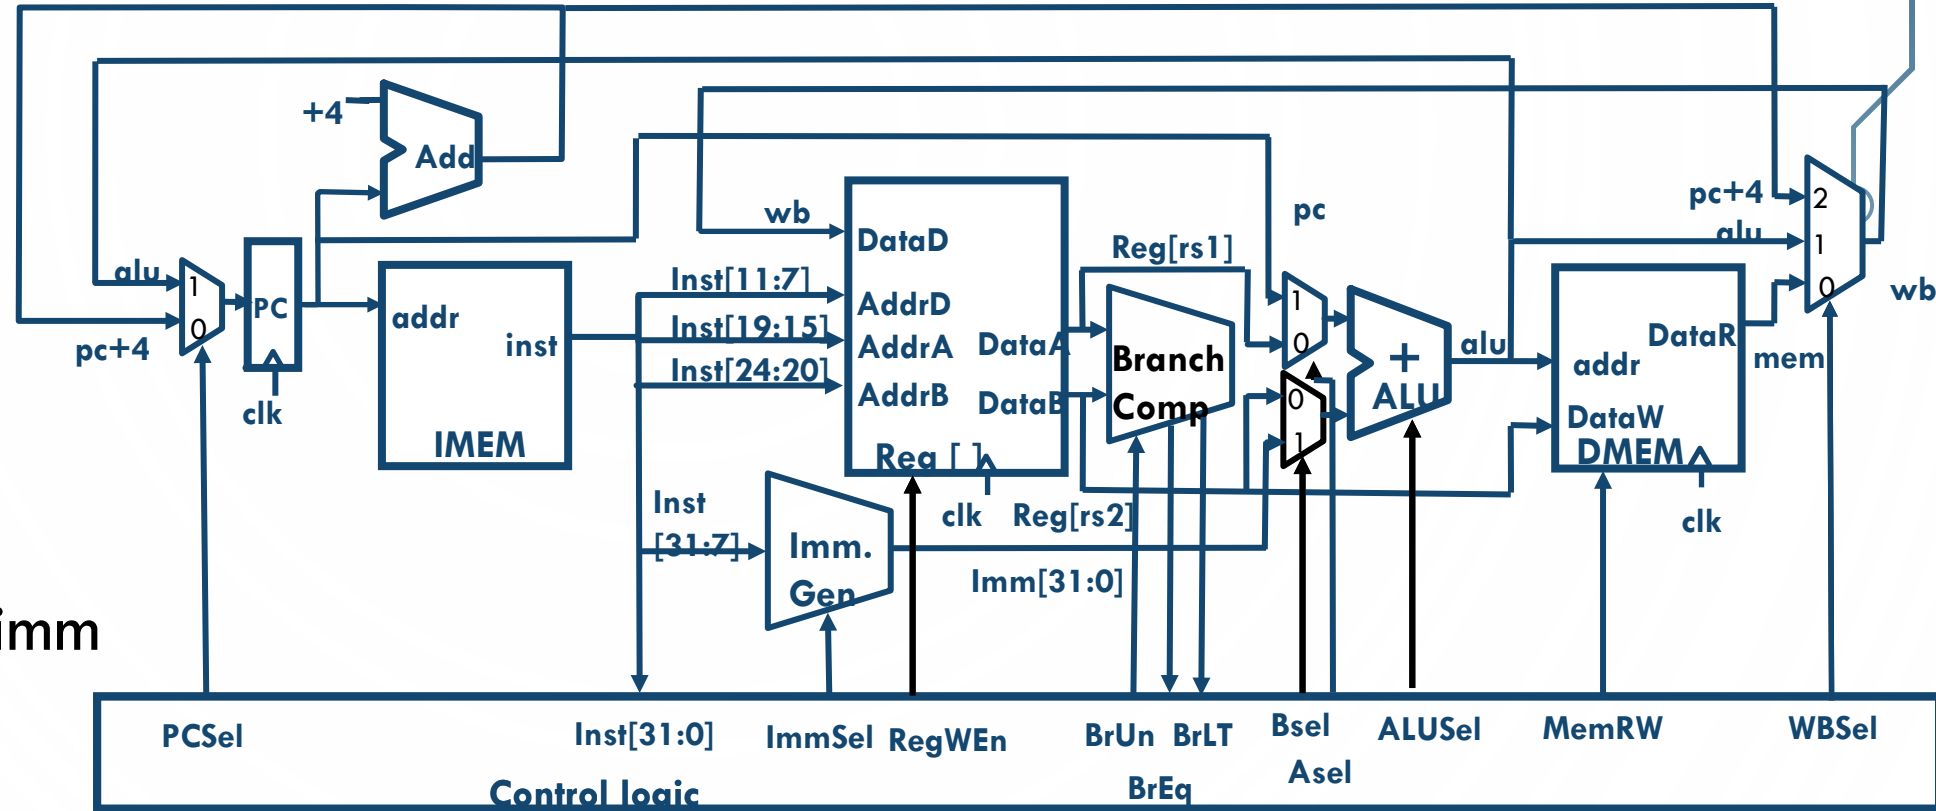


- 1) $t_{\text{clk-q}} + t_{\text{Add}} + t_{\text{IMEM}} + t_{\text{Reg}} + t_{\text{BComp}} + t_{\text{ALU}} + t_{\text{DMEM}} + t_{\text{mux}} + t_{\text{Setup}}$
- 2) $t_{\text{clk-q}} + t_{\text{IMEM}} + \max\{t_{\text{Reg}}, t_{\text{Imm}}\} + t_{\text{ALU}} + 2t_{\text{mux}} + t_{\text{Setup}}$
- 3) $t_{\text{clk-q}} + t_{\text{IMEM}} + \max\{t_{\text{Reg}}, t_{\text{Imm}}\} + t_{\text{ALU}} + 3t_{\text{mux}} + t_{\text{DMEM}} + t_{\text{Setup}}$
- 4) None of the above

Peer Instruction(s): Critical Path yellkey: physical

Critical path for a addi

$$R[rd] = R[rs1] + imm$$



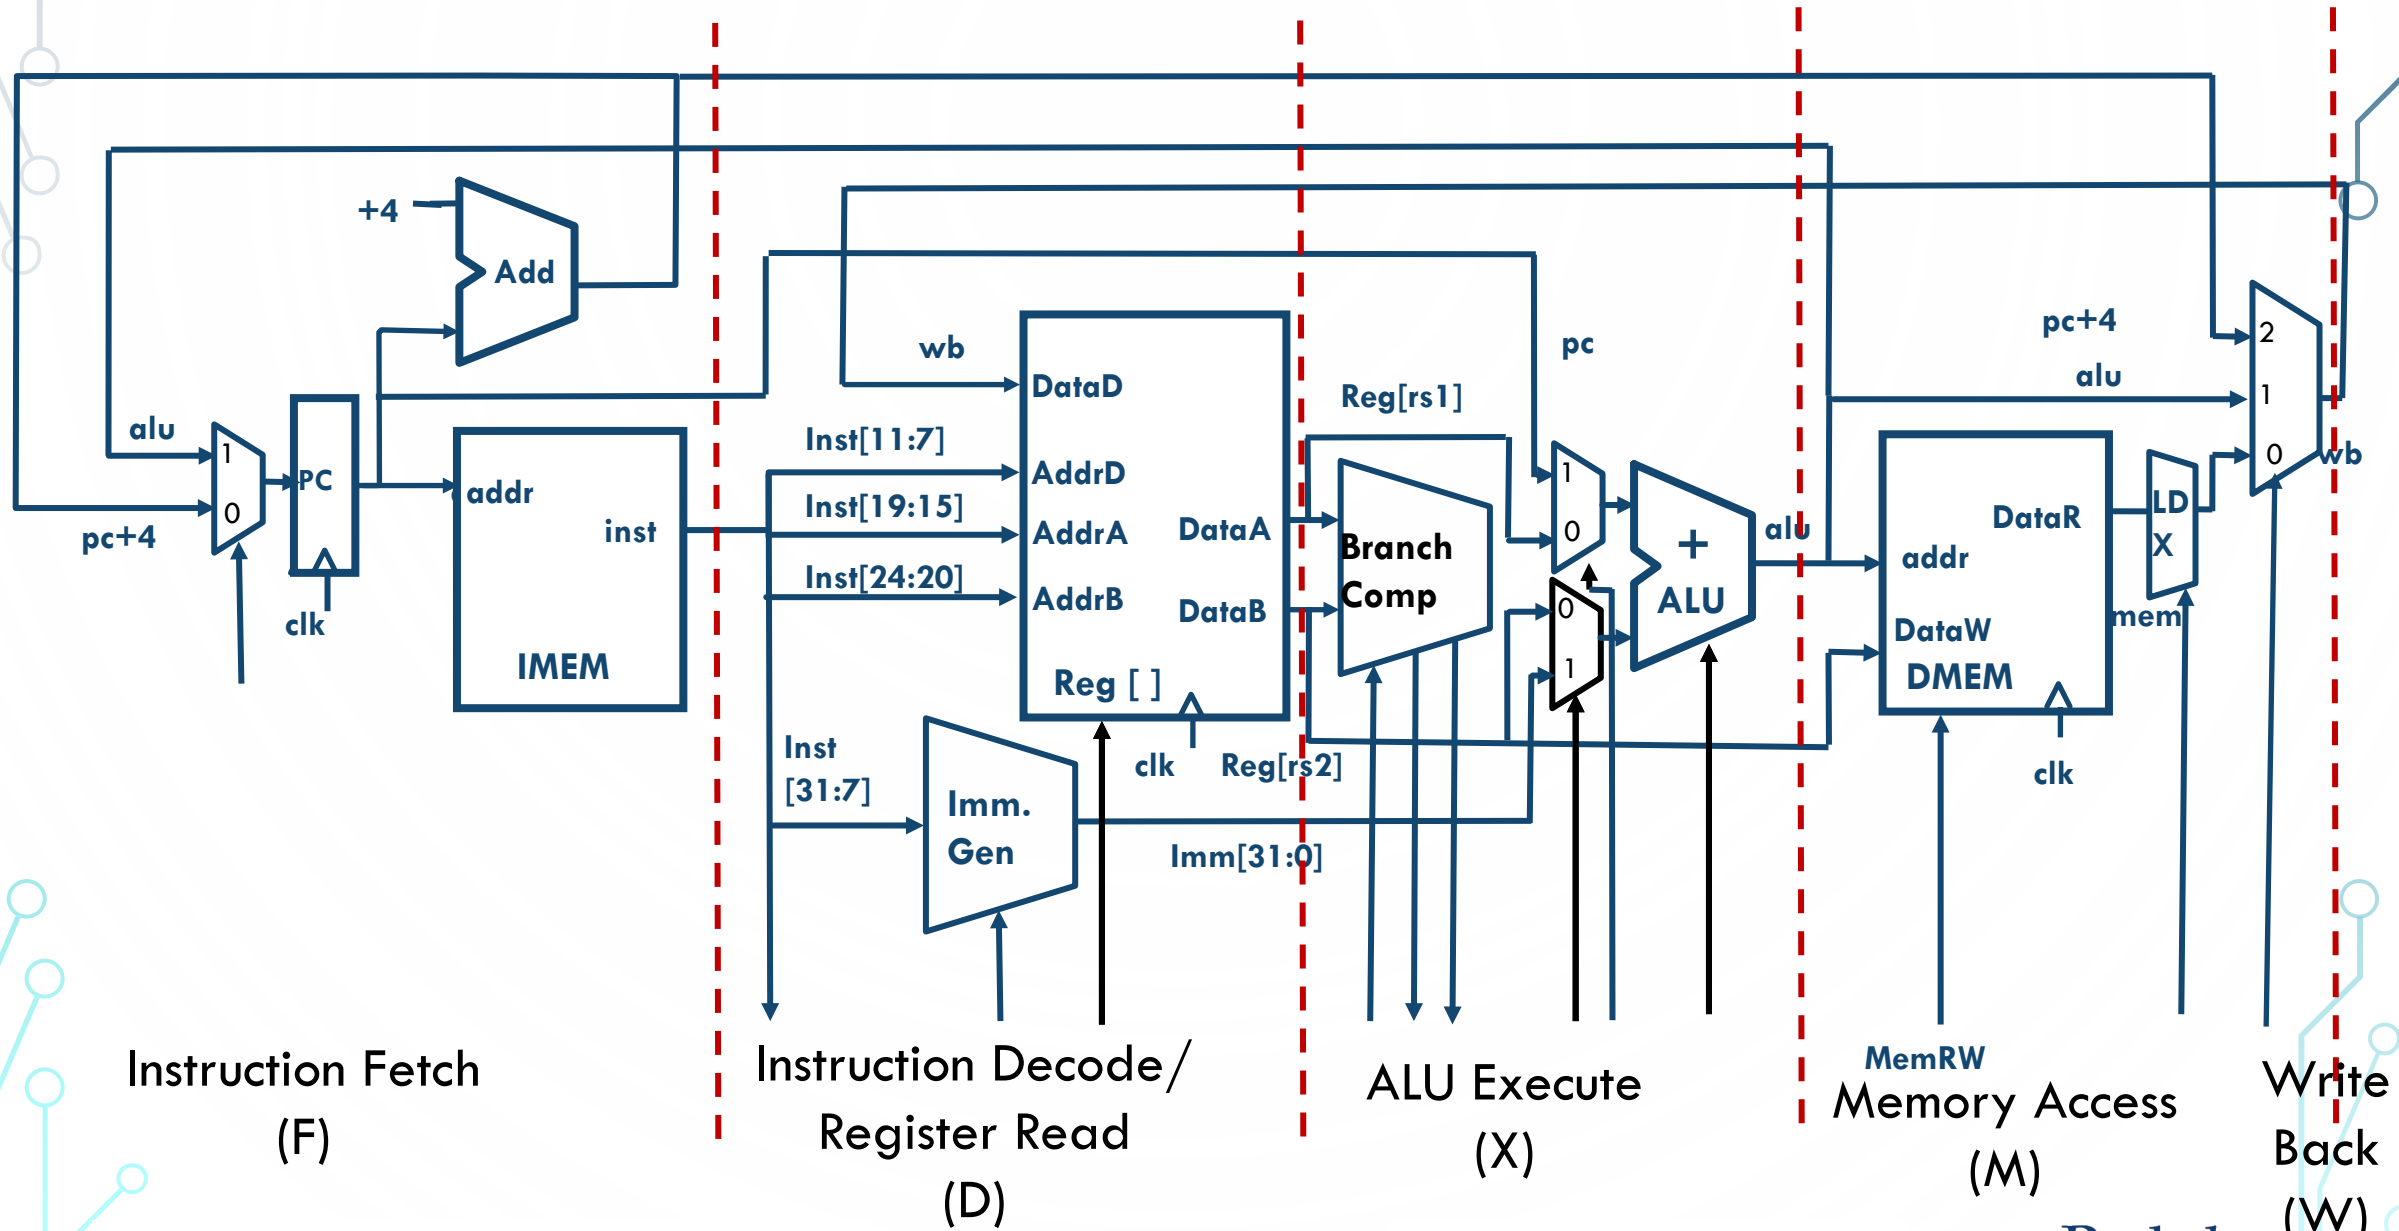
- 1) $t_{\text{clk-q}} + t_{\text{Add}} + t_{\text{IMEM}} + t_{\text{Reg}} + t_{\text{BComp}} + t_{\text{ALU}} + t_{\text{DMEM}} + t_{\text{mux}} + t_{\text{Setup}}$
- 2) $t_{\text{clk-q}} + t_{\text{IMEM}} + \max\{t_{\text{Reg}}, t_{\text{Imm}}\} + t_{\text{ALU}} + 2t_{\text{mux}} + t_{\text{Setup}}$
- 3) $t_{\text{clk-q}} + t_{\text{IMEM}} + \max\{t_{\text{Reg}}, t_{\text{Imm}}\} + t_{\text{ALU}} + 3t_{\text{mux}} + t_{\text{DMEM}} + t_{\text{Setup}}$
- 4) None of the above



Pipelining

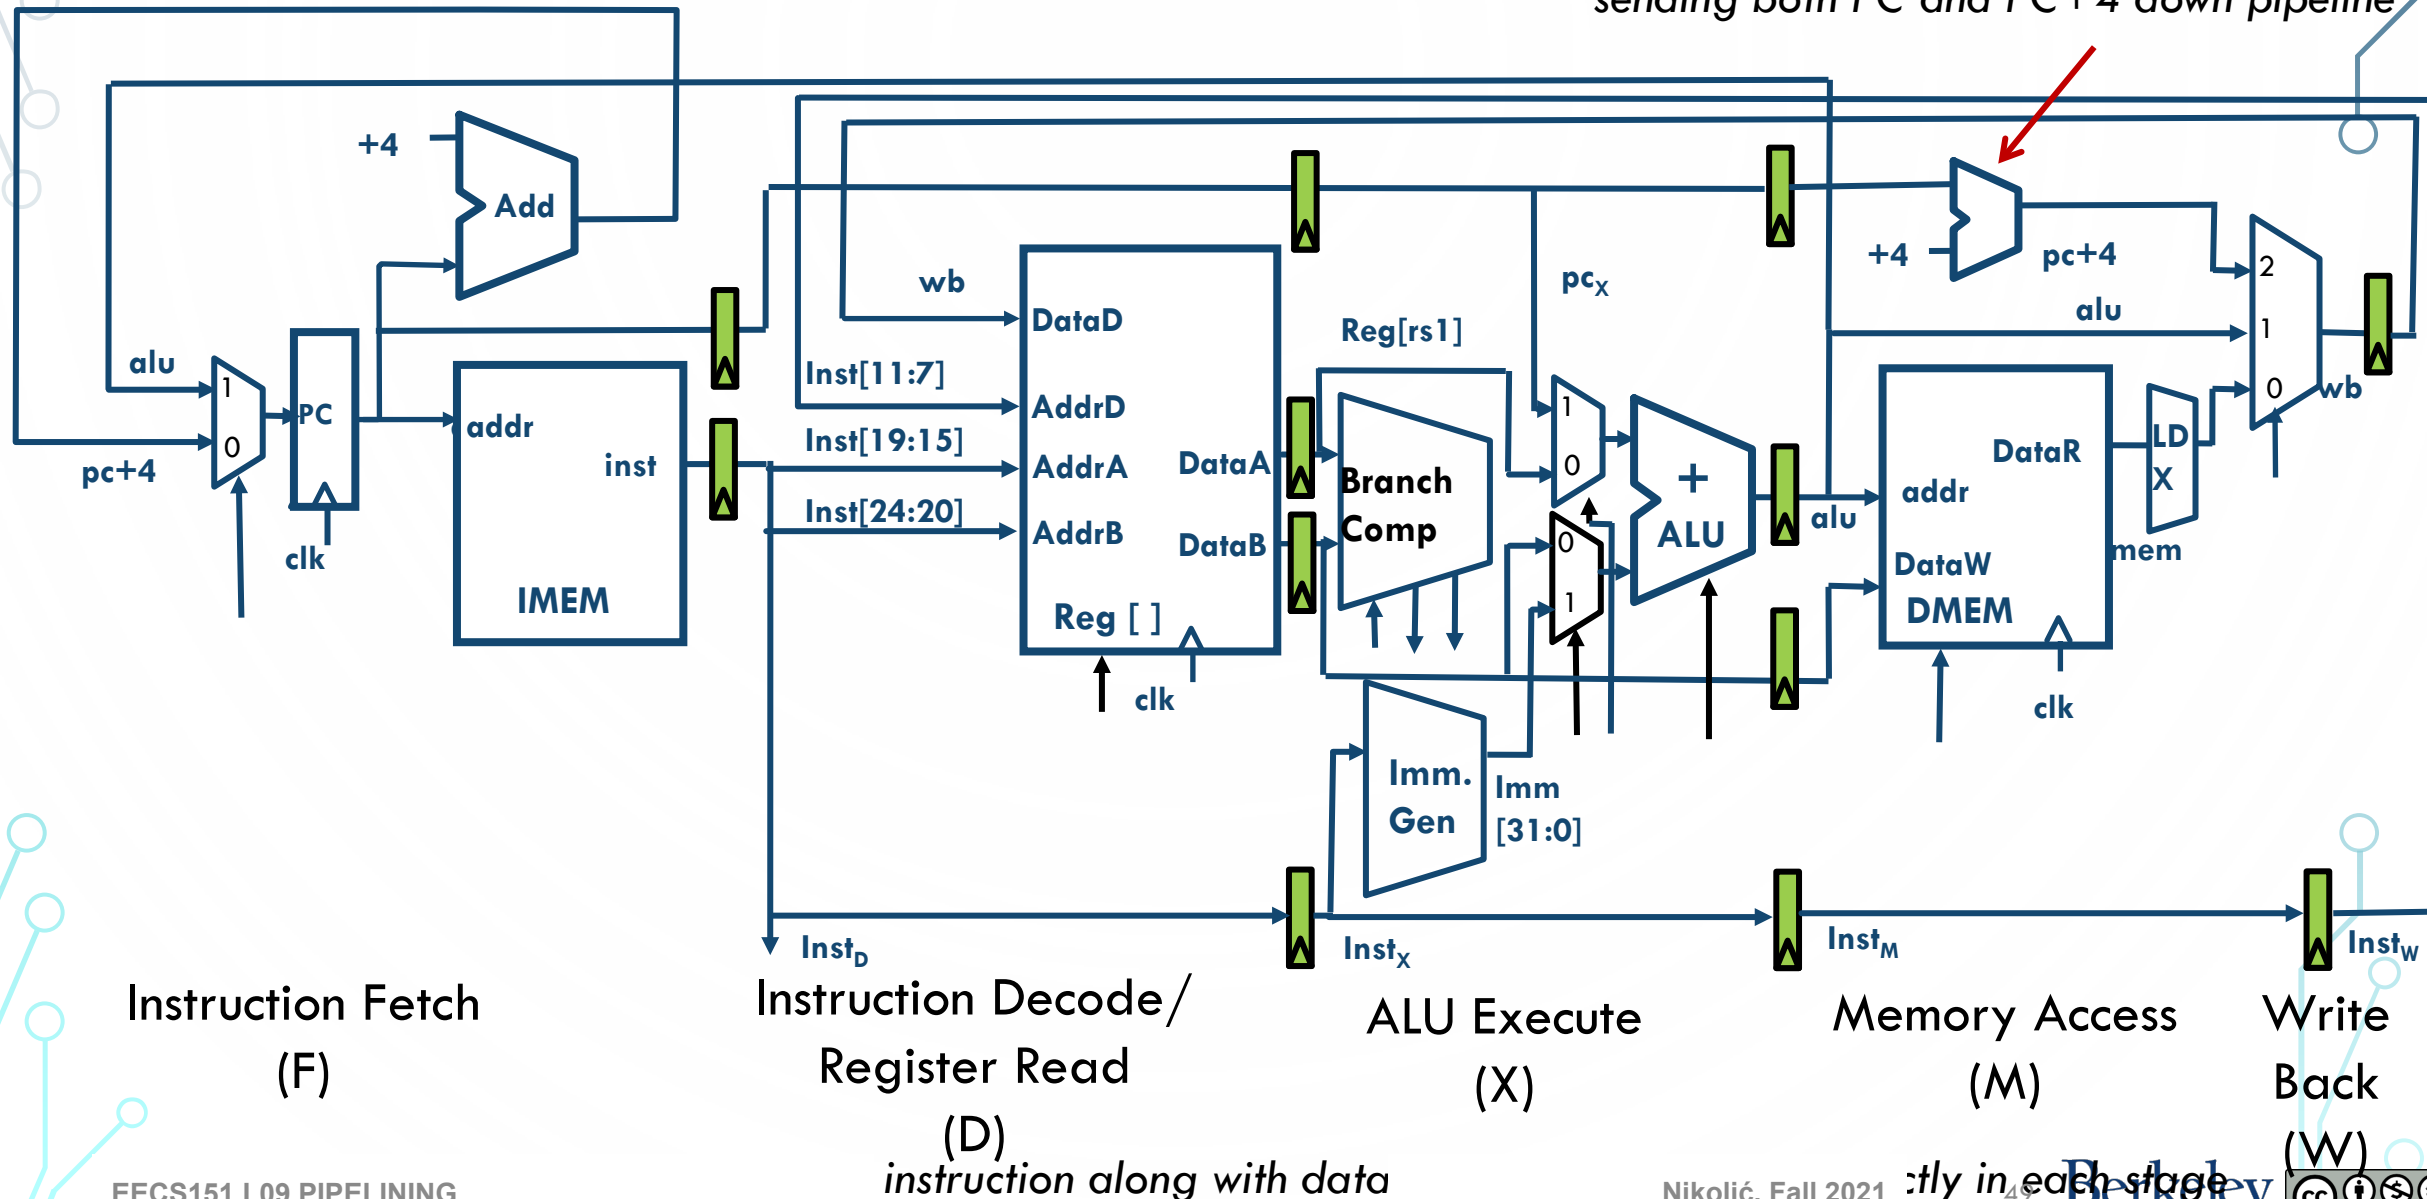
EECS151 L09 PIPELINING

Complete RV32I Datapath with Control



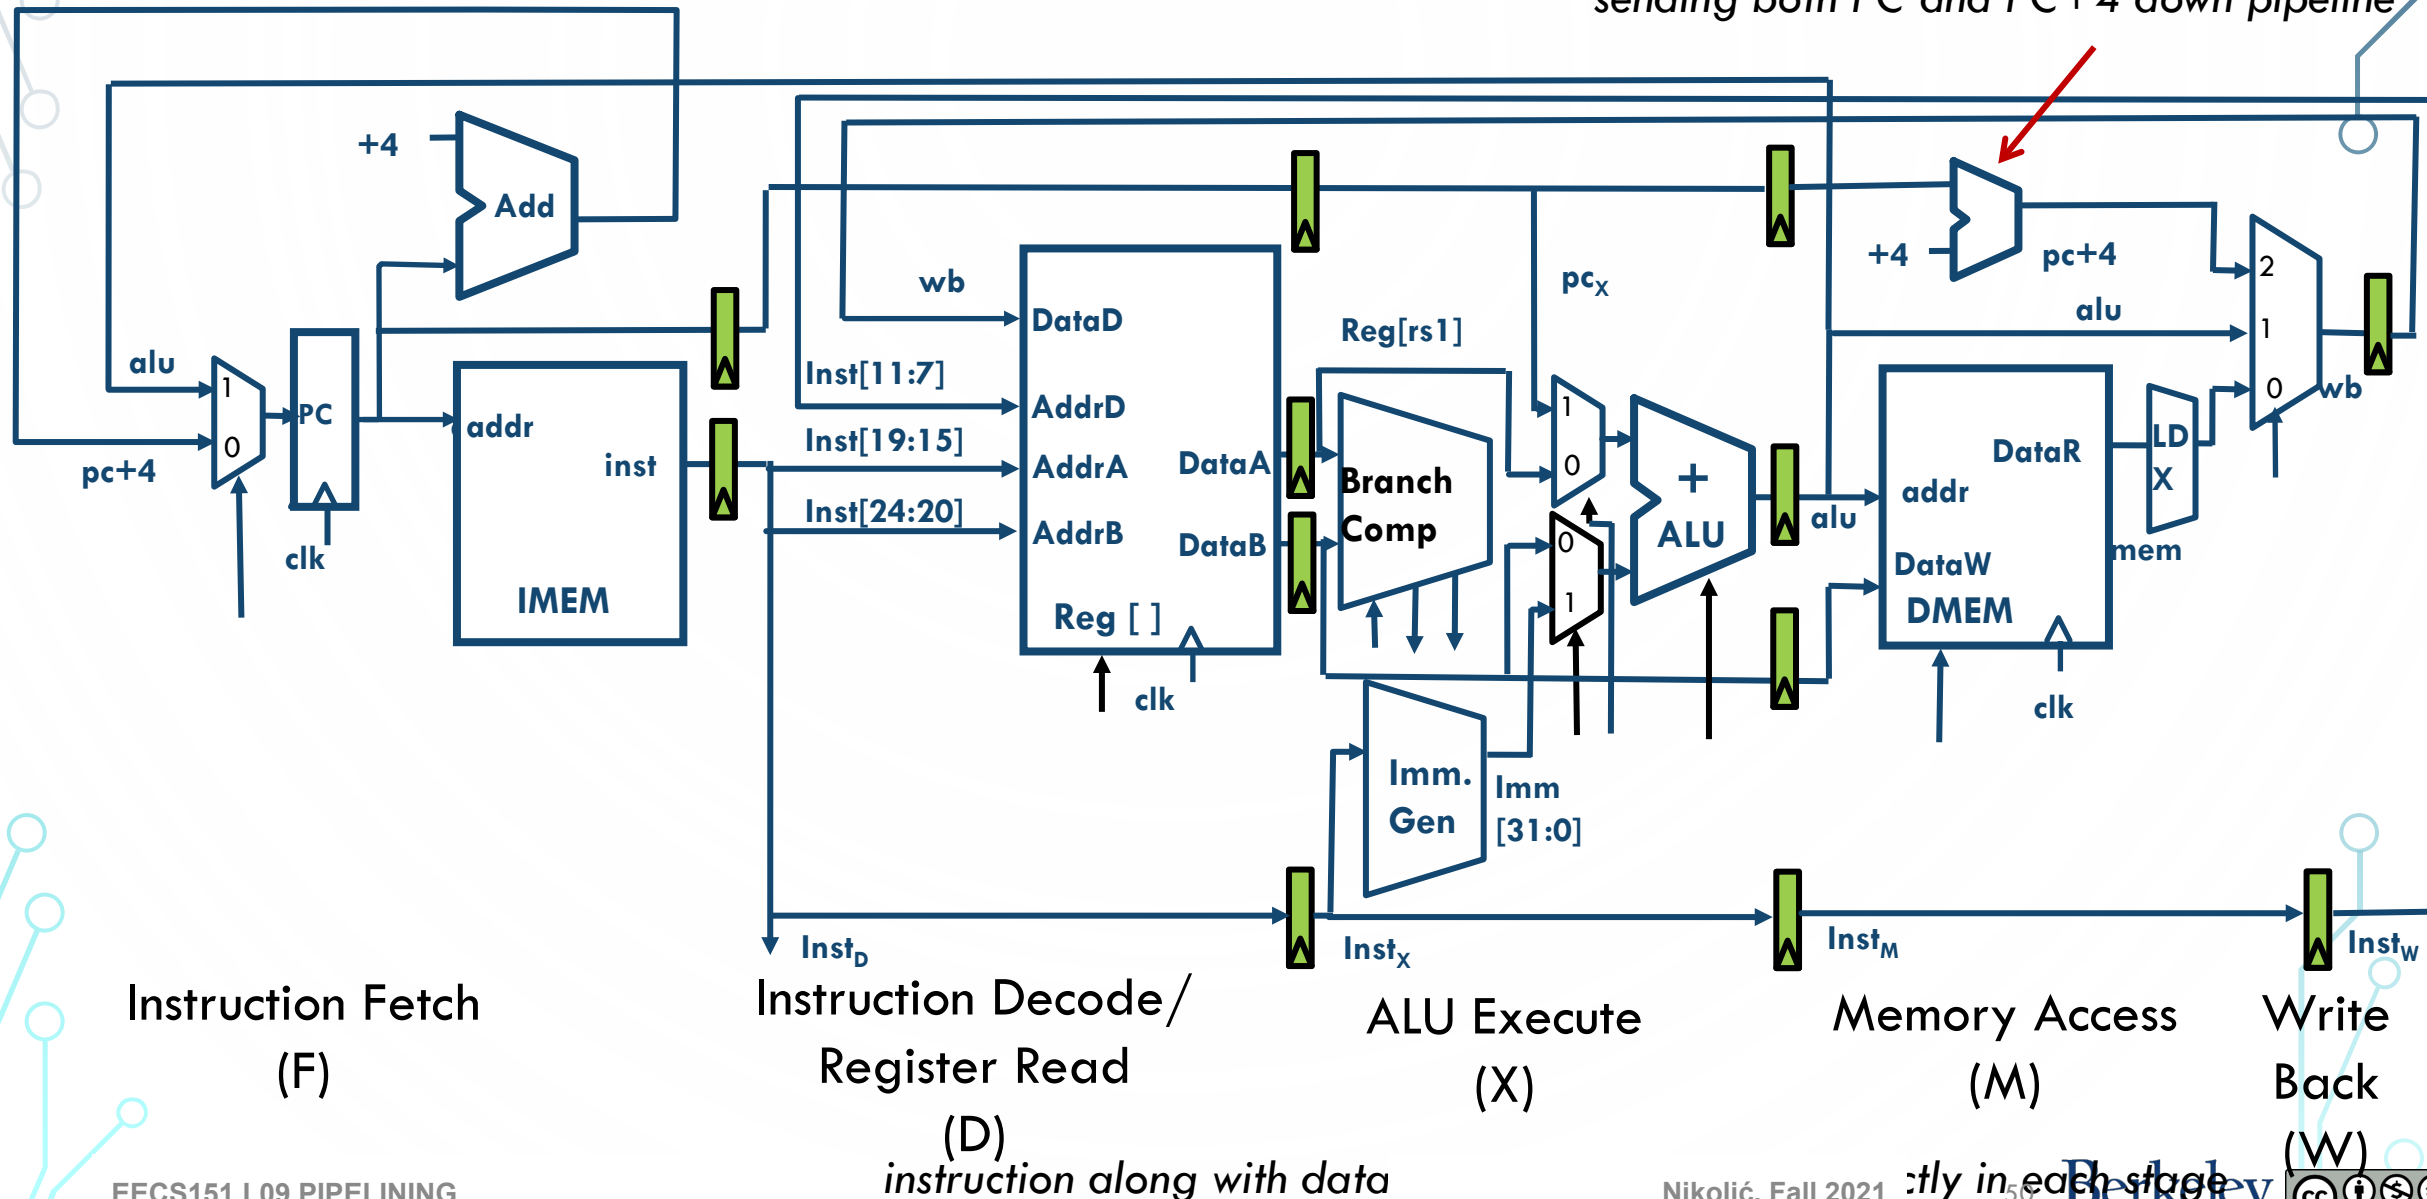
Pipelining RV32I Datapath

Recalculate $PC+4$ in M stage to avoid sending both PC and $PC+4$ down pipeline

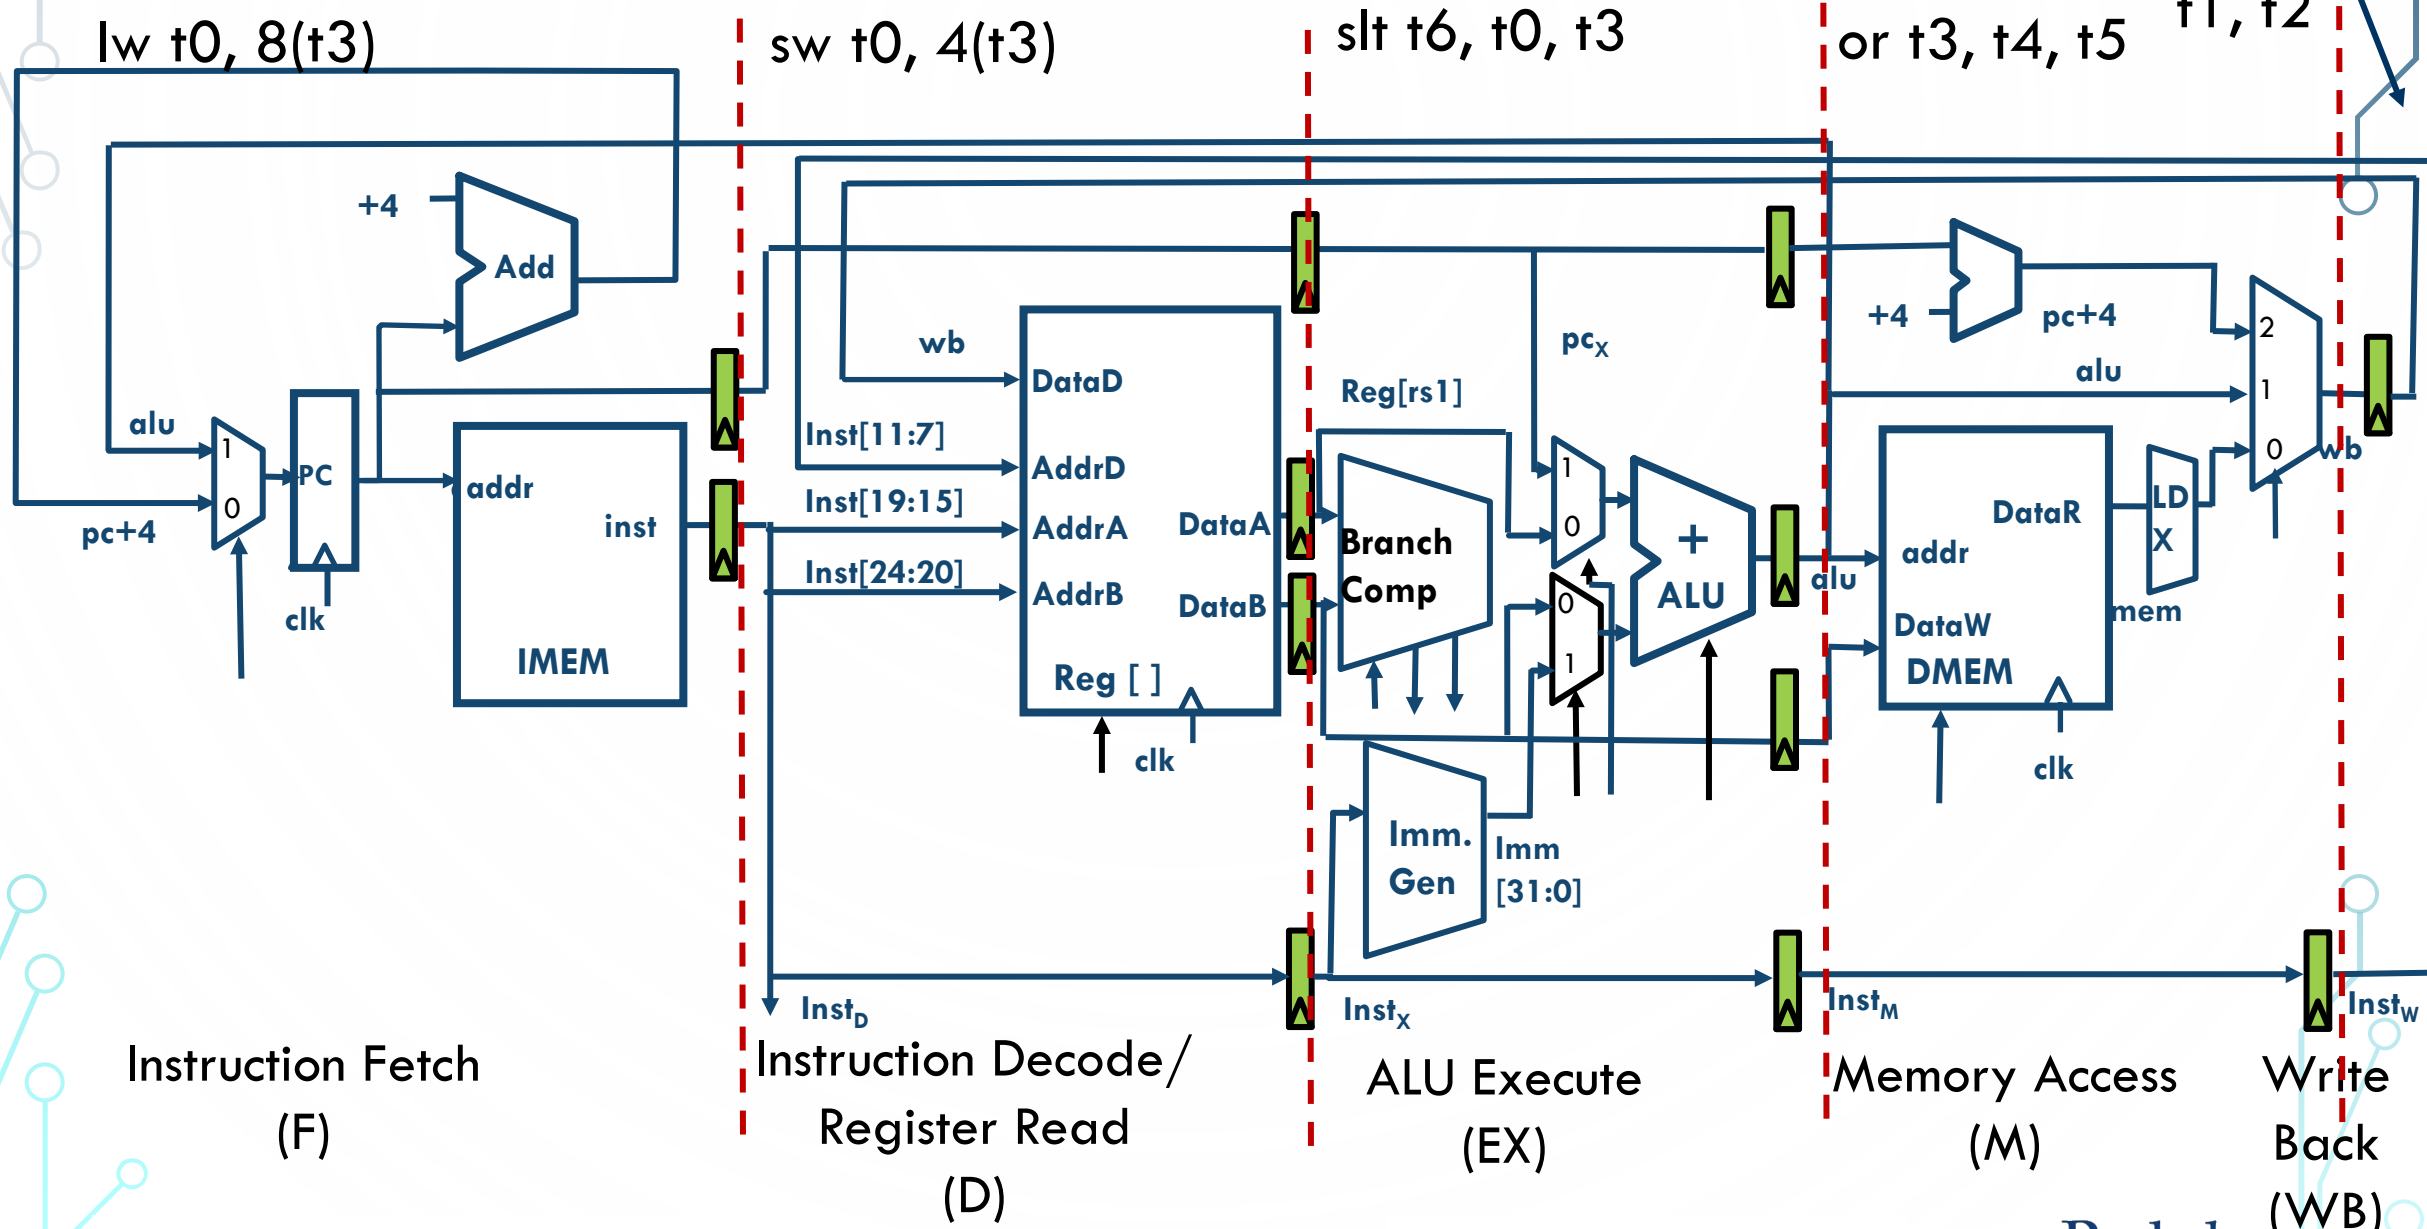


Pipelining RV32I Datapath

Recalculate $PC+4$ in M stage to avoid sending both PC and $PC+4$ down pipeline




Different Instructions in Flight



Pipelined Control

- Control signals derived from instruction
 - As in single-cycle implementation
 - Information is stored in pipeline registers for use by later stages



Summary

- RISC-V ISA
 - Completed the datapath with B-, J-, U-instructions
- Control
 - Can be implemented as a ROM while prototyping
 - Synthesized as custom logic
- Pipelining to increase throughput
 - 5-stage pipeline example