

inst.eecs.berkeley.edu/~eecs151

EECS151 : Introduction to Digital Design and ICs

Lecture 9 – RISC-V ISA, Pipelining

Bora Nikolić



August 20, 2021, Tom's hardware
**Tesla Packs 50 Billion Transistors Onto D1 Dojo Chip
 Designed to Conquer Artificial Intelligence Training**

D1 delivers 362 TeraFLOPs of power.
 Called the D1, the chip resembles a part of the Dojo supercomputer used to train AI models inside Tesla HQ, which are later deployed in various applications. The D1 chip is a product of TSMC's manufacturing efforts, forged in a 7nm semiconductor node. Packing over 50 billion transistors, the chip boasts a huge die size of 645mm^2 .




Image credit: Dennis Hong / Twitter

EECS151 L09 PIPELINING

Nikolić, Fall 2021



1

Review

- RISC-V ISA
 - Open, with increasing adoption
- RISC-V processor
 - A large state machine
 - Datapath + control
 - Reviewed R-, I-, S-format instructions and corresponding datapath elements

EECS151 L09 PIPELINING

Nikolić, Fall 2021



2



RISC-V B-Format Instructions

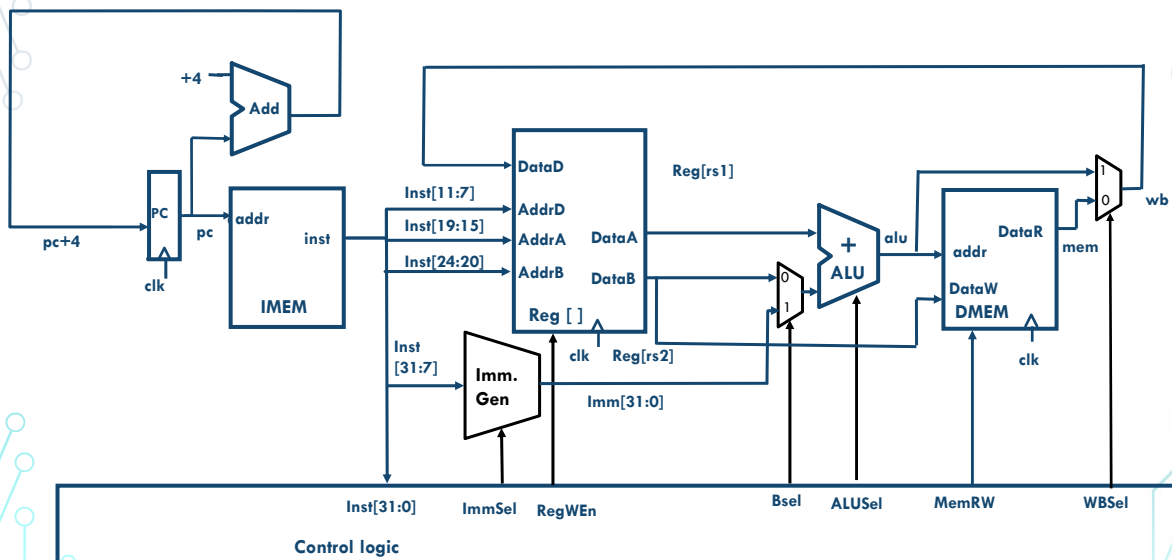
EECS151 L09 PIPELINING

Nikolić, Fall 2021

 3 Berkeley
UNIVERSITY OF CALIFORNIA

3

Datapath So Far (R-, I-, S Instruction Types)



EECS151 L09 PIPELINING

Nikolić, Fall 2021

 4 Berkeley
UNIVERSITY OF CALIFORNIA

4

B-Format - RISC-V Conditional Branches


- E.g., `BEQ x1, x2, Label`
- Branches read two registers but don't write a register (similar to stores)
- How to encode label, i.e., where to branch to?

RISC-V Feature, $n \times 16$ -bit instructions

- Extensions to RISC-V base ISA support 16-bit compressed instructions and also variable-length instructions that are multiples of 16-bits in length
- To enable this, RISC-V scales the branch offset by 2 bytes even when there are no 16-bit instructions
- Reduces branch reach by half and means that $\frac{1}{2}$ of possible targets will be errors on RISC-V processors that only support 32-bit instructions (as used in this class)
- RISC-V conditional branches can only reach $\pm 2^{10} \times 32$ -bit instructions on either side of PC

RISC-V Branch Immediates

- 12-bit immediate encodes PC-relative offset of -4096 to +4094 bytes in multiples of 2 bytes
- RISC-V approach: keep 11 immediate bits in fixed position in output value, and rotate LSB of S-format to be bit 12 of B-format



Only one bit changes position between S and B, so only need a single-bit 2-way mux

EECS151 L09 PIPELINING

Nikolić, Fall 2021

7

7

RISC-V Immediate Encoding

Instruction encodings, inst[31:0]

31	30	25	24	20	19	15	14	12	11	8	7	6	0	
funct7			rs2		rs1		funct3		rd			opcode		R-type
			imm[11:0]		rs1		funct3		rd			opcode		I-type
			imm[11:5]		rs2		rs1		funct3		imm[4:0]		opcode	S-type
			imm[12 10:5]		rs2		rs1		funct3		imm[4:1 11]		opcode	B-type

32-bit immediates produced, imm[31:0]

31	25	24	12	11	10	5	4	1	0	
-inst[31]-					inst[30:25]	inst[24:21]		inst[20]		I-imm.
-inst[31]-					inst[30:25]	inst[11:8]		inst[7]		S-imm.
-inst[31]-					inst[7]	inst[30:25]	inst[11:8]		0	B-imm.

Upper bits sign-extended from inst[31]
always

Only bit 7 of instruction changes role in
immediate between S and B

EECS151 L09 PIPELINING

Nikolić, Fall 2021

8

8

Branch Example, complete encoding

beq **x19,x10,** **offset = 16 bytes**

13-bit immediate, imm[12:0], with value 16

0000000010000

imm[0] discarded,
always zero

imm[12]

imm[11]

0	000000	01010	10011	000	1000	0	1100011
---	--------	-------	-------	-----	------	---	---------

imm[10:5] rs2=10 rs1=19 BEQ imm[4:1] BRANCH

EECS151 L09 PIPELINING

Nikolić, Fall 2021

9 Berkeley UNIVERSITY OF CALIFORNIA

9

Implementing Branches

31	30	25	24	20	19	15	14	12	11	8	7	6	0
imm[12]	imm[10:5]	rs2	rs1	funct3	imm[4:1]	imm[11]	opcode						
1	6	5	5	3	4	1	7						
offset[12:10:5]		rs2	rs1	funct3	offset[4:1:11]		BRANCH						

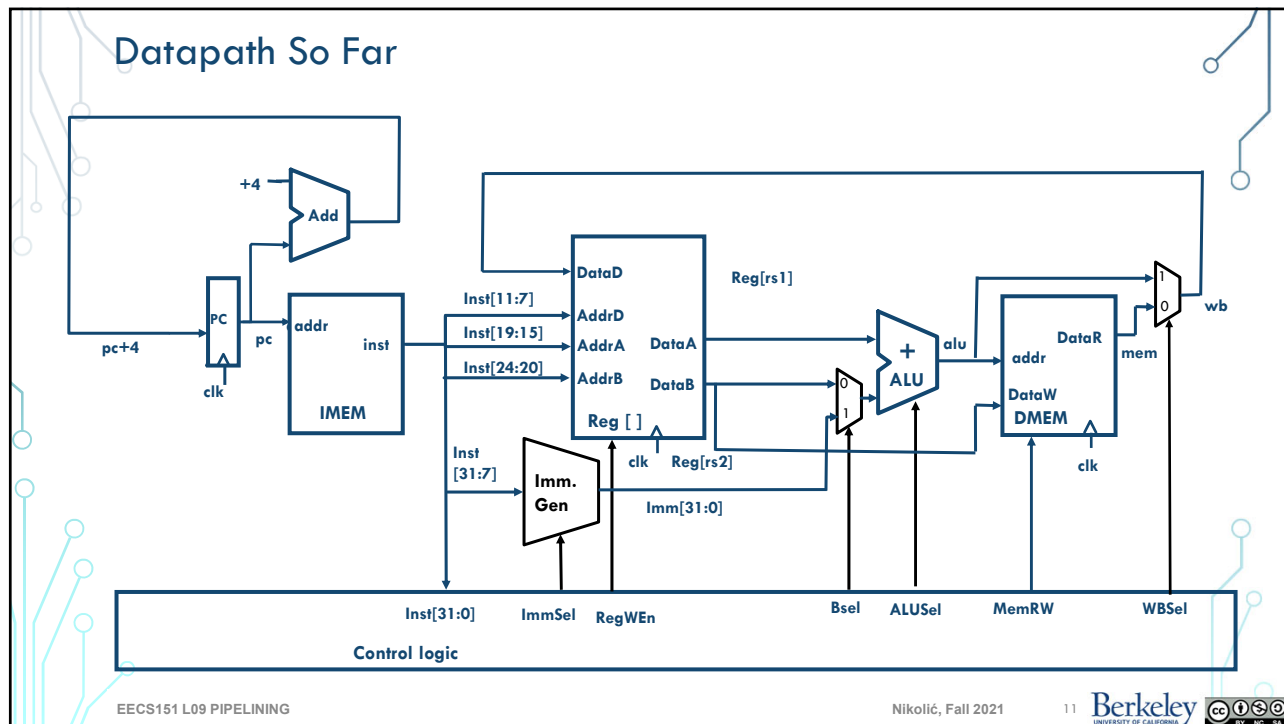
- B-format is mostly same as S-format, with two register sources (rs1/rs2) and a 12-bit immediate
- But now immediate represents values -4096 to +4094 in 2-byte increments
- The 12 immediate bits encode even 13-bit signed byte offsets (lowest bit of offset is always zero, so no need to store it)

EECS151 L09 PIPELINING

Nikolić, Fall 2021

10 Berkeley UNIVERSITY OF CALIFORNIA

10



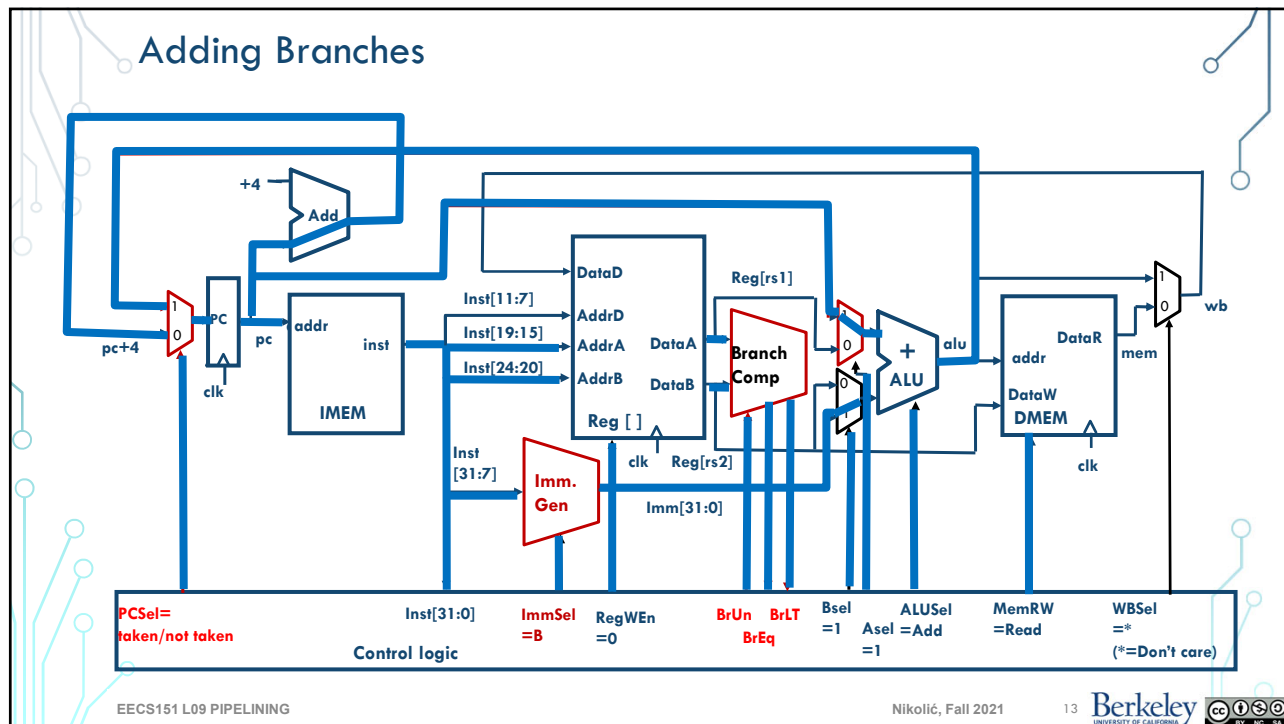
11

To Add Branches

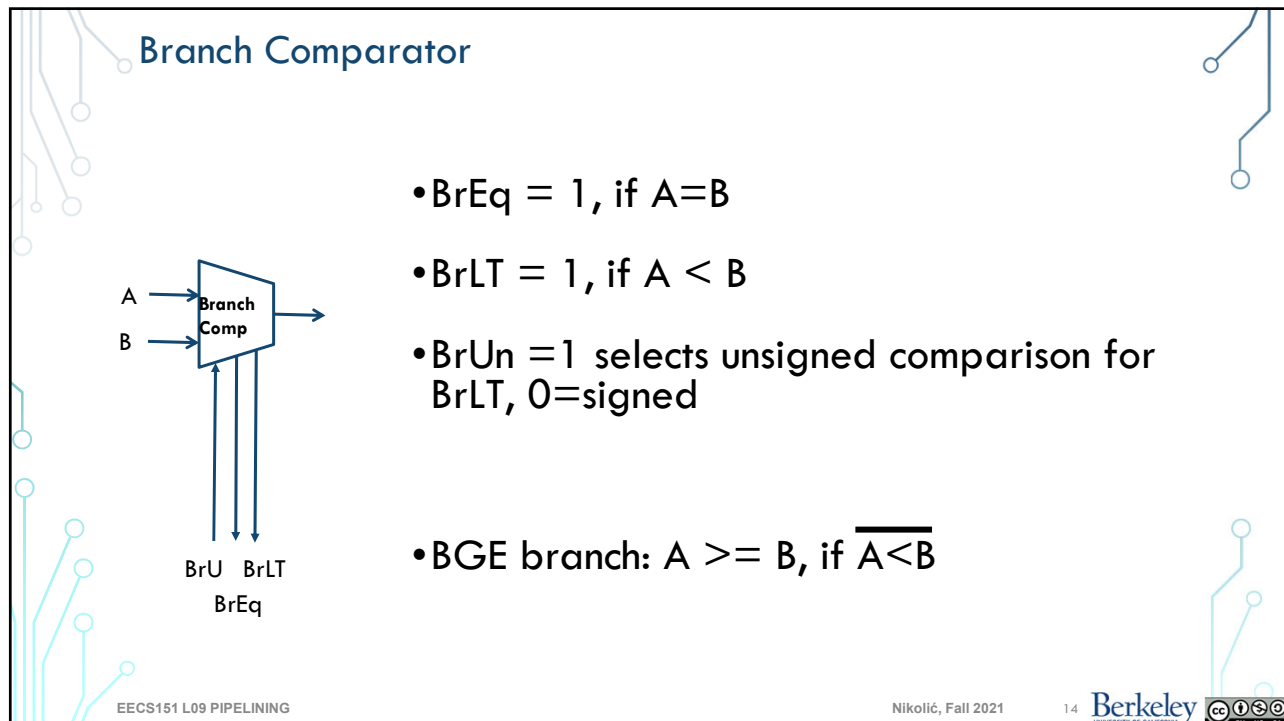
- Different change to the state:
 - $PC = \begin{cases} PC + 4, & \text{branch not taken} \\ PC + \text{immediate}, & \text{branch taken} \end{cases}$
- Six branch instructions: **BEQ, BNE, BLT, BGE, BLTU, BGEU**
- Need to compute **PC + immediate** and to compare values of **rs1** and **rs2**
 - Need another add/sub unit

EECS151 L09 PIPELINING Nikolić, Fall 2021 12 Berkeley

12



13



14

All RISC-V Branch Instructions

imm[12 10:5]	rs2	rs1	000	imm[4:1 11]	1100011	BEQ
imm[12 10:5]	rs2	rs1	001	imm[4:1 11]	1100011	BNE
imm[12 10:5]	rs2	rs1	100	imm[4:1 11]	1100011	BLT
imm[12 10:5]	rs2	rs1	101	imm[4:1 11]	1100011	BGE
imm[12 10:5]	rs2	rs1	110	imm[4:1 11]	1100011	BLTU
imm[12 10:5]	rs2	rs1	111	imm[4:1 11]	1100011	BGEU

Administrivia

- Homework 4 is due next Monday
 - No new homework this week
 - Homework 5 will be posted next week, due after the midterm
- Lab 5 this week
 - No lab next week
 - Lab 6 (last) after the midterm
- Midterm 1 on October 7, 7-8:30pm



RISC-V J-Format Instructions

EECS151 L09 PIPELINING


Nikolić, Fall 2021

17

Berkeley
UNIVERSITY OF CALIFORNIA

17

J-Format for Jump Instructions



- JAL saves PC+4 in register rd (the return address)
 - Assembler “j” jump is pseudo-instruction, uses JAL but sets rd=x0 to discard return address
- Set PC = PC + offset (PC-relative jump)
- Target somewhere within $\pm 2^{19}$ locations, 2 bytes apart
 - $\pm 2^{18}$ 32-bit instructions
- Immediate encoding optimized similarly to branch instruction to reduce hardware cost

EECS151 L09 PIPELINING


Nikolić, Fall 2021

18

Berkeley
UNIVERSITY OF CALIFORNIA

18

JALR Instruction (I-Format)



- JALR rd, rs, immediate
 - Writes PC+4 to rd (return address)
 - Sets PC = rs1 + immediate (and sets the LSB to 0)
 - Uses same immediates as arithmetic and loads
 - **no** multiplication by 2 bytes
 - In contrast to branches and JAL

Otherwise, we would have yet another new encoding


EECS151 L09 PIPELINING

Nikolić, Fall 2021

19 Berkeley UNIVERSITY OF CALIFORNIA

19

Let's Add JALR (I-Format)




- JALR rd, rs, immediate
- Two changes to the state
 - Writes PC+4 to rd (return address)
 - Sets PC = rs + immediate
 - Uses same immediates as arithmetic and loads
 - **no** multiplication by 2 bytes
 - LSB is ignored

EECS151 L09 PIPELINING

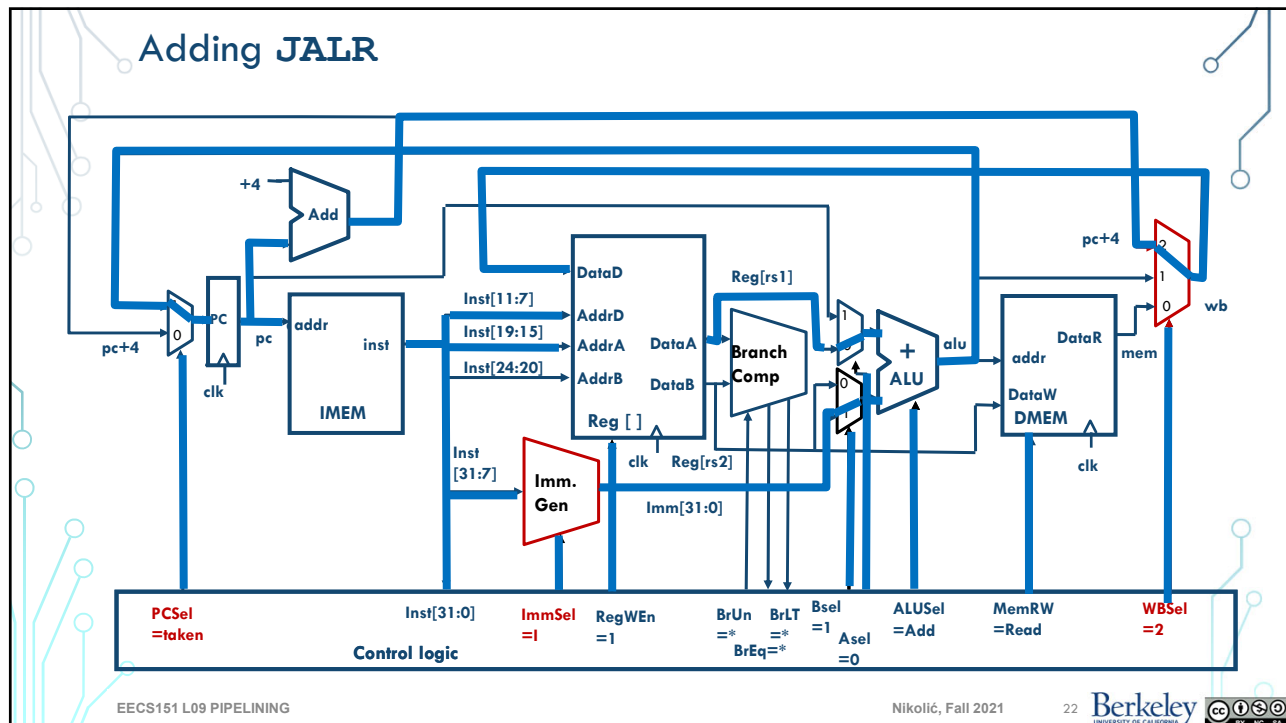
Nikolić, Fall 2021

20 Berkeley UNIVERSITY OF CALIFORNIA

20




21



22

Adding JAL



- JAL saves PC+4 in register rd (the return address)
- Set PC = PC + offset (PC-relative jump)
- Target somewhere within $\pm 2^{19}$ locations, 2 bytes apart
 - $\pm 2^{18}$ 32-bit instructions
- Immediate encoding optimized similarly to branch instruction to reduce hardware cost

EECS151 L09 PIPELINING

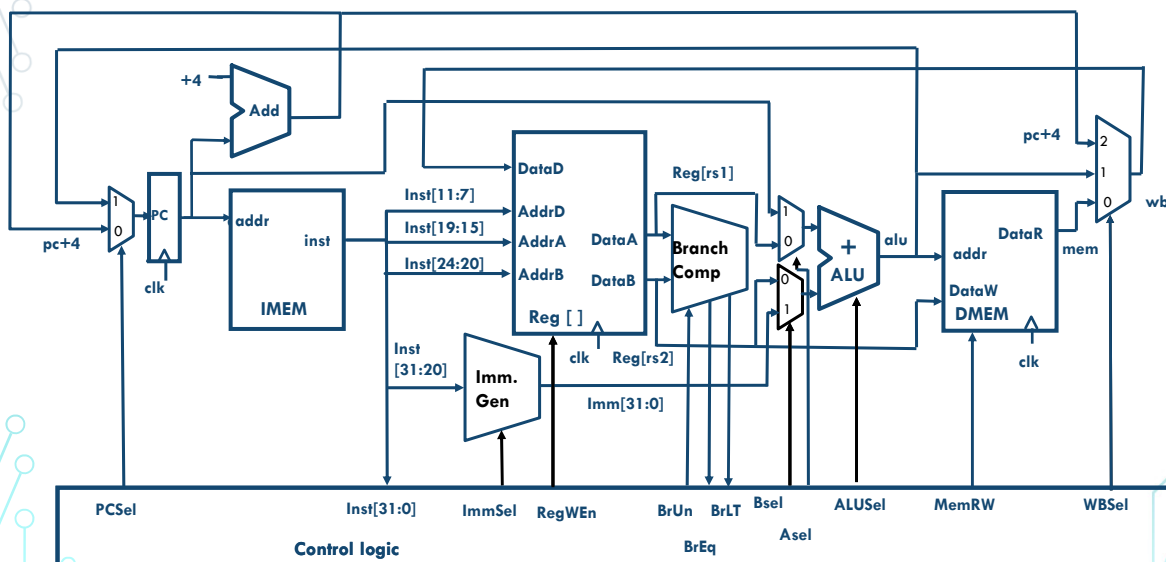
Nikolić, Fall 2021

23



23

Datapath with JALR



EECS151 L09 PIPELINING

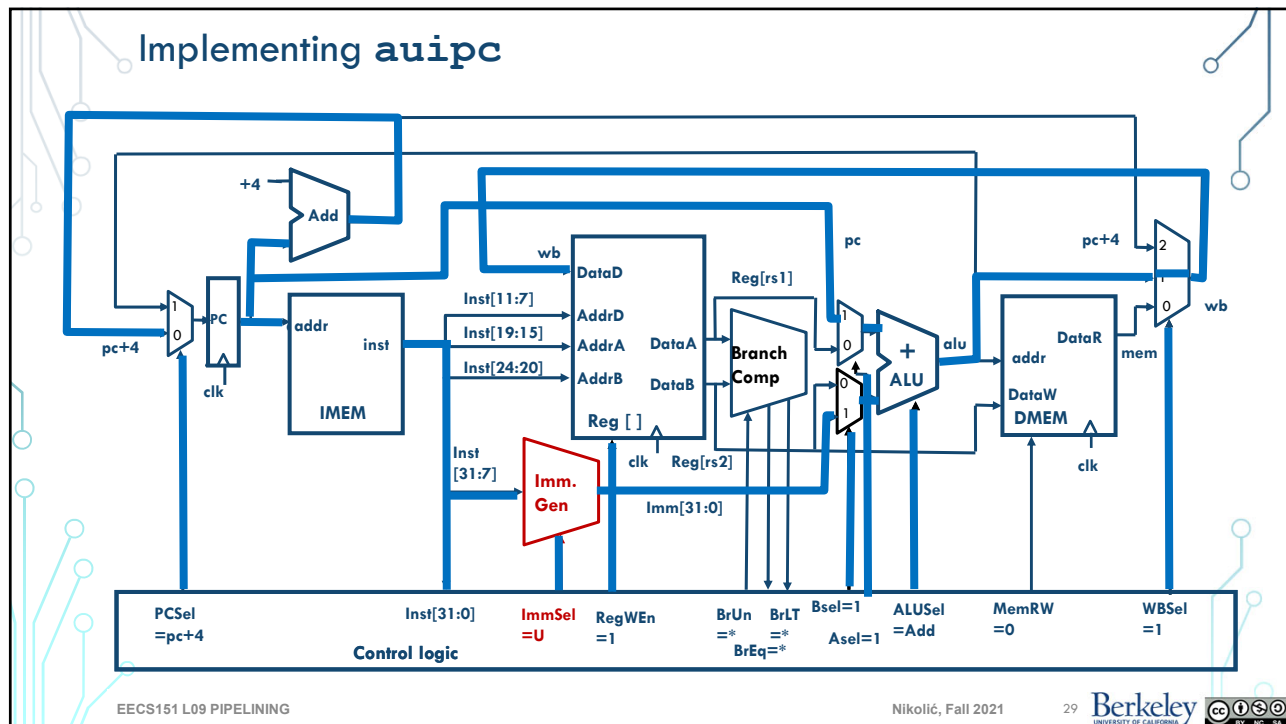
Nikolić, Fall 2021

24




24





29



30

Recap: Complete RV32I ISA

Open		Reference Card					
Base Integer Instructions: RV32I							
Category	Name	Fmt	RV32I Base	Category	Name	Fmt	RV32I Base
Shifts	Shift Left Logical	R	SLL rd,rs1,rs2	Loads	Load Byte	I	LB rd,rs1,imm
	Shift Left Log. Imm.	I	SLLI rd,rs1,shamt		Load Halfword	I	LH rd,rs1,imm
	Shift Right Logical	R	SRL rd,rs1,rs2		Load Byte Unsigned	I	LBU rd,rs1,imm
	Shift Right Log. Imm.	I	SRLI rd,rs1,shamt		Load Half Unsigned	I	LHU rd,rs1,imm
	Shift Right Arithmetic	R	SRA rd,rs1,rs2		Load Word	I	LW rd,rs1,imm
	Shift Right Arith. Imm.	I	SRAI rd,rs1,shamt	Stores	Store Byte	S	SB rs1,rs2,imm
Arithmetic	ADD	R	ADD rd,rs1,rs2		Store Halfword	S	SH rs1,rs2,imm
	ADD Immediate	I	ADDI rd,rs1,imm		Store Word	S	SW rs1,rs2,imm
	SUBtract	R	SUB rd,rs1,rs2	Branches	Branch =	B	BEQ rs1,rs2,imm
	Load Upper Imm	U	LUI rd,imm		Branch ≠	B	BNE rs1,rs2,imm
	Add Upper Imm to PC	U	AUIPC rd,imm		Branch <	B	BLT rs1,rs2,imm
Logical	XOR	R	XOR rd,rs1,rs2		Branch ≥	B	BGE rs1,rs2,imm
	XOR Immediate	I	XORI rd,rs1,imm		Branch < Unsigned	B	BLTU rs1,rs2,imm
	OR	R	OR rd,rs1,rs2		Branch ≥ Unsigned	B	BGEU rs1,rs2,imm
	OR Immediate	I	ORI rd,rs1,imm	Jump & Link	J&L	J	JAL rd,imm
	AND	R	AND rd,rs1,rs2		Jump & Link Register	I	JALR rd,rs1,imm
	AND Immediate	I	ANDI rd,rs1,imm				
Compare	Set <	R	SLT rd,rs1,rs2	Synch	Synch thread	I	FENCE
	Set < Immediate	I	SLTI rd,rs1,imm				
	Set < Unsigned	R	SLTU rd,rs1,rs2	Environment	CALL	I	ECALL
	Set < Imm Unsigned	I	SLTIU rd,rs1,imm		BREAK	I	EBREAK

- 40 instructions are enough to run any C program

EECS151 L09 PIPELINING

Nikolić, Fall 2021

31



31

Summary of RISC-V Instruction Formats

31	30	25	24	21	20	19	15	14	12	11	8	7	6	0			
funct7				rs2		rs1		funct3		rd		opcode		R-type			
imm[11:0]						rs1		funct3		rd		opcode		I-type			
imm[11:5]				rs2		rs1		funct3		imm[4:0]				opcode		S-type	
imm[12 10:5]				rs2		rs1		funct3		imm[4:1 11]				opcode		B-type	
imm[31:12]										rd		opcode				U-type	
imm[20 10:1 11]]						imm[19:12]				rd		opcode				J-type	

EECS151 L09 PIPELINING

Nikolić, Fall 2021


32



32

Control and Status Registers (CSRs)

- 4096 CSRs in a separate address space



- **csrrw** reads the old value of CSR, zero-extends and writes to **rd**
- Initial value of **rs1** is written to CSR
- Pseudo-instructions: **csrr**, **csrw** (**csrrw** **x0**, **csr**, **rs1**)

EECS151 L09 PIPELINING

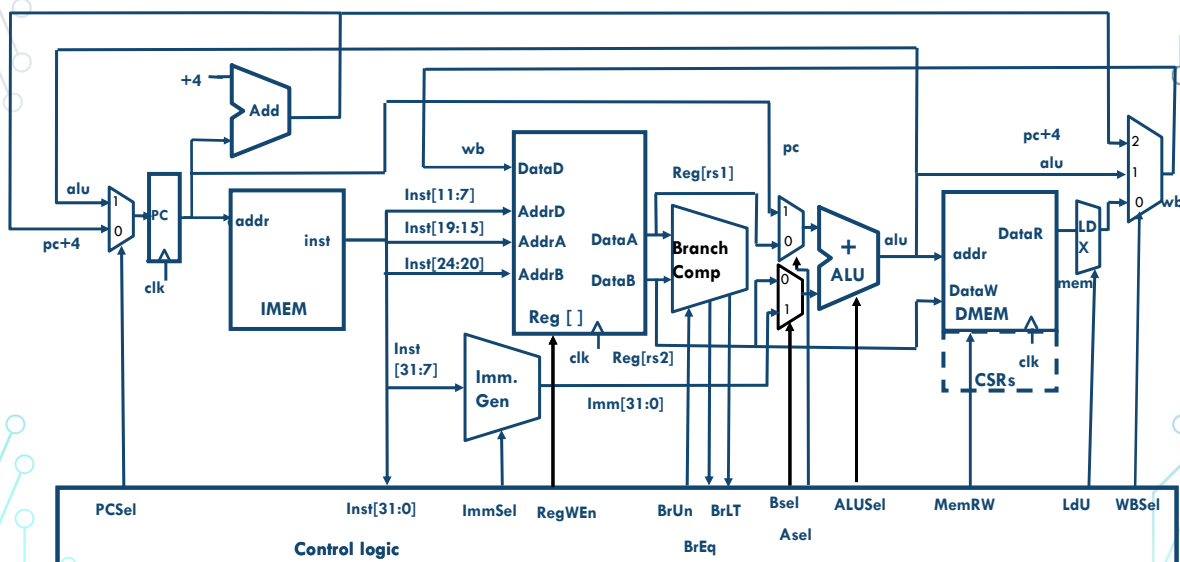
Nikolić, Fall 2021

33



33

Add the **csrrw**!



EECS151 L09 PIPELINING

Nikolić, Fall 2021

34



34



RISC-V Control Logic

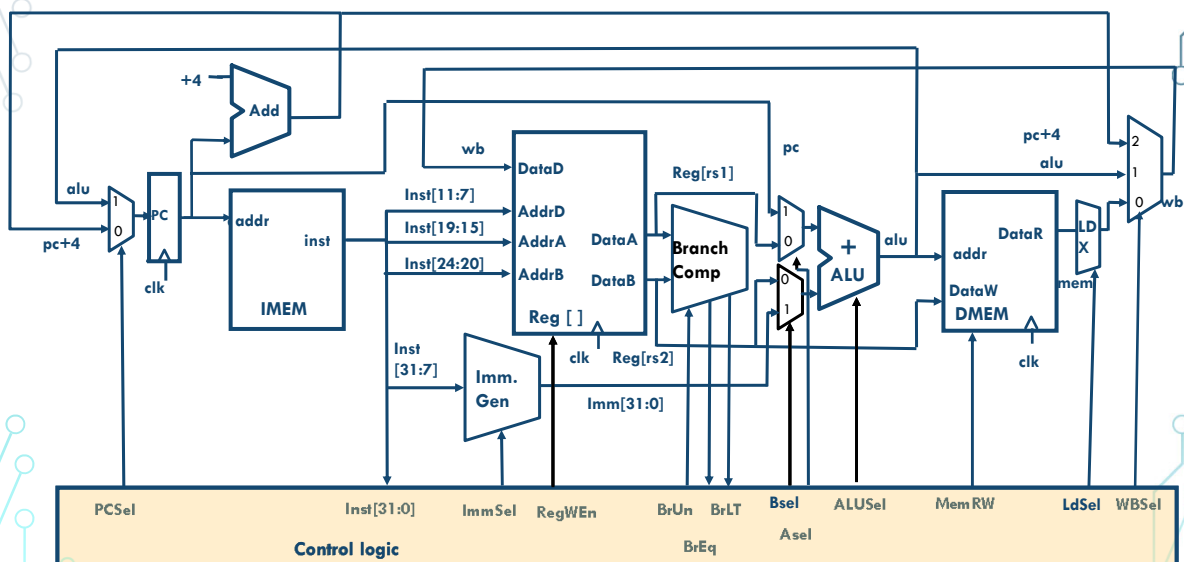
EECS151 L09 PIPELINING

Nikolić, Fall 2021

35 Berkeley UNIVERSITY OF CALIFORNIA

35

Complete RV32I Datapath with Control

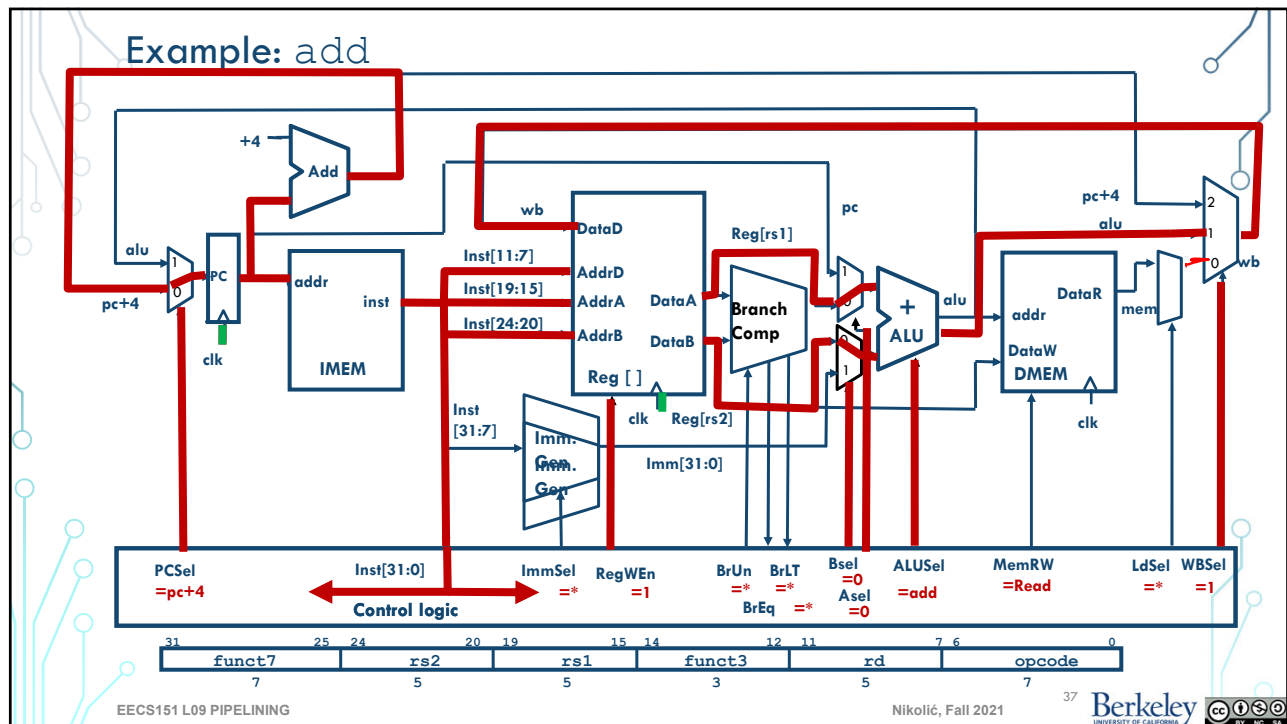


EECS151 L09 PIPELINING

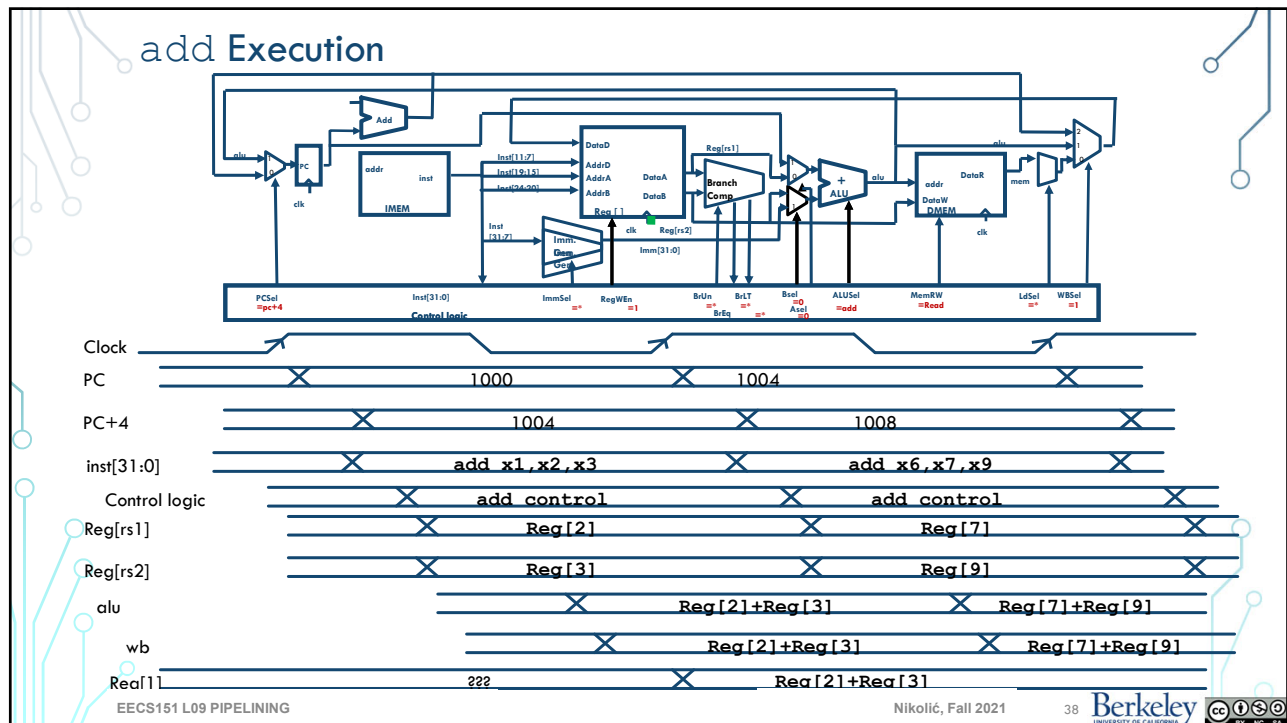
Nikolić, Fall 2021

36 Berkeley UNIVERSITY OF CALIFORNIA

36



37



38

Control Logic Truth Table

Inst[31:0]	BrEq	BrLT	PCSel	ImmSel	BrUn	ASel	BSel	ALUSel	MemRW	RegWEn	WBSel
add	*	*	+4	*	*	Reg	Reg	Add	Read	1	ALU
sub	*	*	+4	*	*	Reg	Reg	Sub	Read	1	ALU
(R-R Op)	*	*	+4	*	*	Reg	Reg	(Op)	Read	1	ALU
addi	*	*	+4	I	*	Reg	Imm	Add	Read	1	ALU
lw	*	*	+4	I	*	Reg	Imm	Add	Read	1	Mem
sw	*	*	+4	S	*	Reg	Imm	Add	Write	0	*
beq	0	*	+4	B	*	PC	Imm	Add	Read	0	*
beq	1	*	ALU	B	*	PC	Imm	Add	Read	0	*
bne	0	*	ALU	B	*	PC	Imm	Add	Read	0	*
bne	1	*	+4	B	*	PC	Imm	Add	Read	0	*
blt	*	1	ALU	B	0	PC	Imm	Add	Read	0	*
bltu	*	1	ALU	B	1	PC	Imm	Add	Read	0	*
jalc	*	*	ALU	I	*	Reg	Imm	Add	Read	1	PC+4
jal	*	*	ALU	J	*	PC	Imm	Add	Read	1	PC+4
auipc	*	*	+4	U	*	PC	Imm	Add	Read	1	ALU

EECS151 L09 PIPELINING

Nikolić, Fall 2021

39

Berkeley



39

RV32I, a nine-bit ISA!

imm[31:12]				rd	011011	LUI
imm[31:12]				rd	001011	AUIPC
imm[20:10]	imm[19:12]			rd	110111	JAL
imm[11:0]				rd	110011	JALR
imm[12:10:5]	rs2	rs1	000	imm[4:1]	110001	BEQ
imm[12:10:5]	rs2	rs1	001	imm[4:1]	110001	BNE
imm[12:10:5]	rs2	rs1	100	imm[4:1]	110001	BLT
imm[12:10:5]	rs2	rs1	101	imm[4:1]	110001	BGE
imm[12:10:5]	rs2	rs1	110	imm[4:1]	110001	BLTU
imm[12:10:5]	rs2	rs1	111	imm[4:1]	110001	BGEU
imm[11:0]		rs1	000	rd	000001	LB
imm[11:0]		rs1	001	rd	000001	LH
imm[11:0]		rs1	010	rd	000001	LW
imm[11:0]		rs1	100	rd	000001	LBU
imm[11:0]		rs1	101	rd	000001	LHU
imm[11:5]	rs2	rs1	000	imm[4:0]	010001	SB
imm[11:5]	rs2	rs1	001	imm[4:0]	010001	SH
imm[11:5]	rs2	rs1	010	imm[4:0]	010001	SW
imm[11:0]		rs1	000	rd	001001	ADDI
imm[11:0]		rs1	010	rd	001001	SLTI
imm[11:0]		rs1	011	rd	001001	SLTIU
imm[11:0]		rs1	100	rd	001001	XORI
imm[11:0]		rs1	110	rd	001001	ORI
imm[11:0]		rs1	111	rd	001001	ANDI
0000000	shamt	rs1	001	rd	001001	SLLI
0000000	shamt	rs1	101	rd	001001	SRLI
0100000	shamt	rs1	101	rd	001001	SRAI
0000000	rs2	rs1	000	rd	011001	ADD
0100000	rs2	rs1	000	rd	011001	SUB
0000000	rs2	rs1	001	rd	011001	SLL
0000000	rs2	rs1	010	rd	011001	SLT
0000000	rs2	rs1	011	rd	011001	SLTU
0000000	rs2	rs1	100	rd	011001	XOR
0000000	rs2	rs1	101	rd	011001	SRL
0100000	rs2	rs1	101	rd	011001	SRA
0000000	rs2	rs1	110	rd	011001	OR
0000000	rs2	rs1	111	rd	011001	AND

inst[30]

inst[14:12]

inst[6:2]

Instruction type encoded by using only 9 bits: inst[30], inst[14:12], inst[6:2]

Nikolić, Fall 2021

40

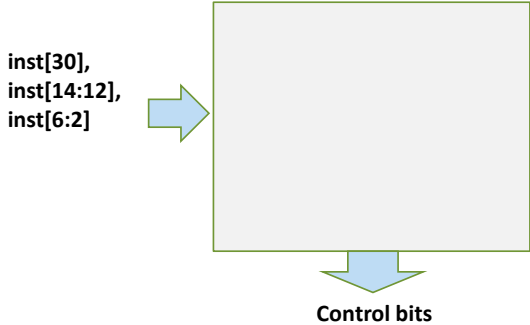
Berkeley



40

Control Realization Options

- ROM
 - “Read-Only Memory”
 - Regular structure
 - Can be easily reprogrammed
 - fix errors
 - add instructions
- Combinatorial Logic
 - Start from a truth table
 - More compact, faster
 - Use synthesis tools



The diagram illustrates a control realization block. On the left, three instruction fields are listed: `inst[30]`, `inst[14:12]`, and `inst[6:2]`. A blue arrow points from these fields into a light gray rectangular block. From the bottom of this block, another blue arrow points downwards to the text "Control bits".

EECS151 L09 PIPELINING

Nikolić, Fall 2021

41 Berkeley UNIVERSITY OF CALIFORNIA

41

Combinational Logic Control

- Decoder is typically hierarchical
 - First decode opcode, and figure out instruction type
 - E.g. branches are $\text{Inst}[6:2] = 11000$
 - Then determine the actual instruction
 - $\text{Inst}[30] + \text{Inst}[14:12]$
- Modularity helps simplify and speed up logic
 - Narrows problem space for logic synthesis

EECS151 L09 PIPELINING

Nikolić, Fall 2021

42 Berkeley UNIVERSITY OF CALIFORNIA

42

Combinational Logic Control

- Simple example: BrUn

		inst[14:12]		inst[6:2]		
imm[12:10:5]	rs2	rs1	000	imm[4:1:11]	1100011	BEQ
imm[12:10:5]	rs2	rs1	001	imm[4:1:11]	1100011	BNE
imm[12:10:5]	rs2	rs1	100	imm[4:1:11]	1100011	BLT
imm[12:10:5]	rs2	rs1	101	imm[4:1:11]	1100011	BGE
imm[12:10:5]	rs2	rs1	110	imm[4:1:11]	1100011	BLTU
imm[12:10:5]	rs2	rs1	111	imm[4:1:11]	1100011	BGEU

		inst[14:13]			
inst[12]		00	01	11	10
0					
1					

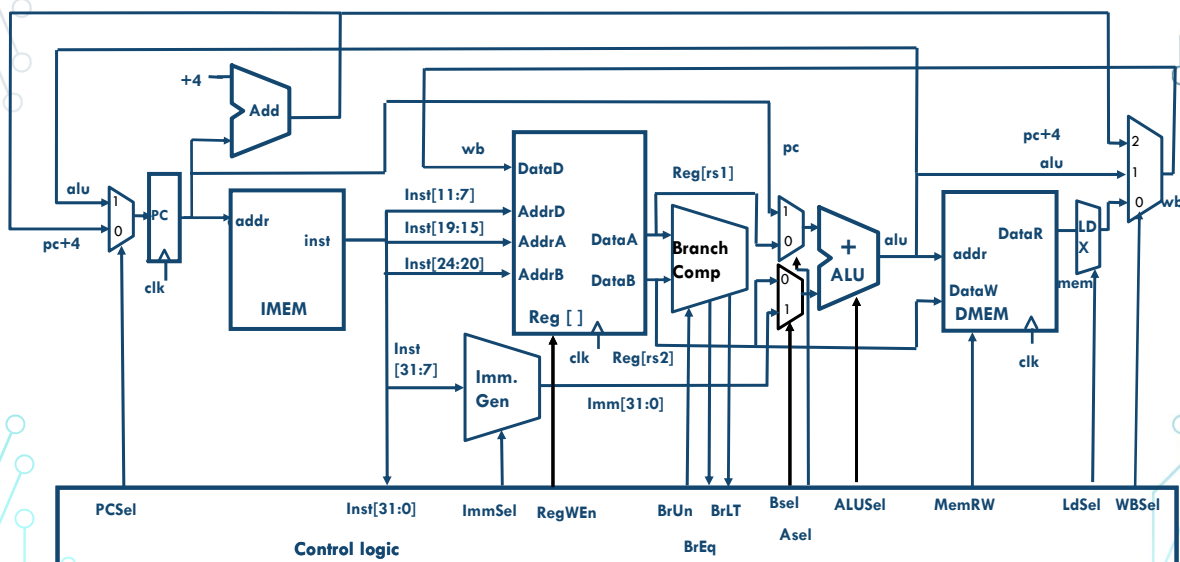
- How to decode whether BrUn is 1?

• BrUn = Inst [13] • Branch

• Branch = Inst[6] • Inst[5] • !Inst[4] • !Inst[3] • !Inst[2]

43

Complete RV32I Datapath with Control



44

Peer Instruction(s): Critical Path yellkey: crime

Critical path for a addi

$R[rd] = R[rs1] + imm$

Control logic

- 1) $t_{clk-q} + t_{Add} + t_{IMEM} + t_{Reg} + t_{BComp} + t_{ALU} + t_{DMEM} + t_{mux} + t_{Setup}$
- 2) $t_{clk-q} + t_{IMEM} + \max\{t_{Reg}, t_{Imm}\} + t_{ALU} + 2t_{mux} + t_{Setup}$
- 3) $t_{clk-q} + t_{IMEM} + \max\{t_{Reg}, t_{Imm}\} + t_{ALU} + 3t_{mux} + t_{DMEM} + t_{Setup}$
- 4) None of the above

EECS151 L09 PIPELINING Nikolić, Fall 2021 45 Berkeley

45

Peer Instruction(s): Critical Path yellkey: physical

Critical path for a addi

$R[rd] = R[rs1] + imm$

Control logic

- 1) $t_{clk-q} + t_{Add} + t_{IMEM} + t_{Reg} + t_{BComp} + t_{ALU} + t_{DMEM} + t_{mux} + t_{Setup}$
- 2) $t_{clk-q} + t_{IMEM} + \max\{t_{Reg}, t_{Imm}\} + t_{ALU} + 2t_{mux} + t_{Setup}$
- 3) $t_{clk-q} + t_{IMEM} + \max\{t_{Reg}, t_{Imm}\} + t_{ALU} + 3t_{mux} + t_{DMEM} + t_{Setup}$
- 4) None of the above

EECS151 L09 PIPELINING Nikolić, Fall 2021 46 Berkeley

46



Pipelining

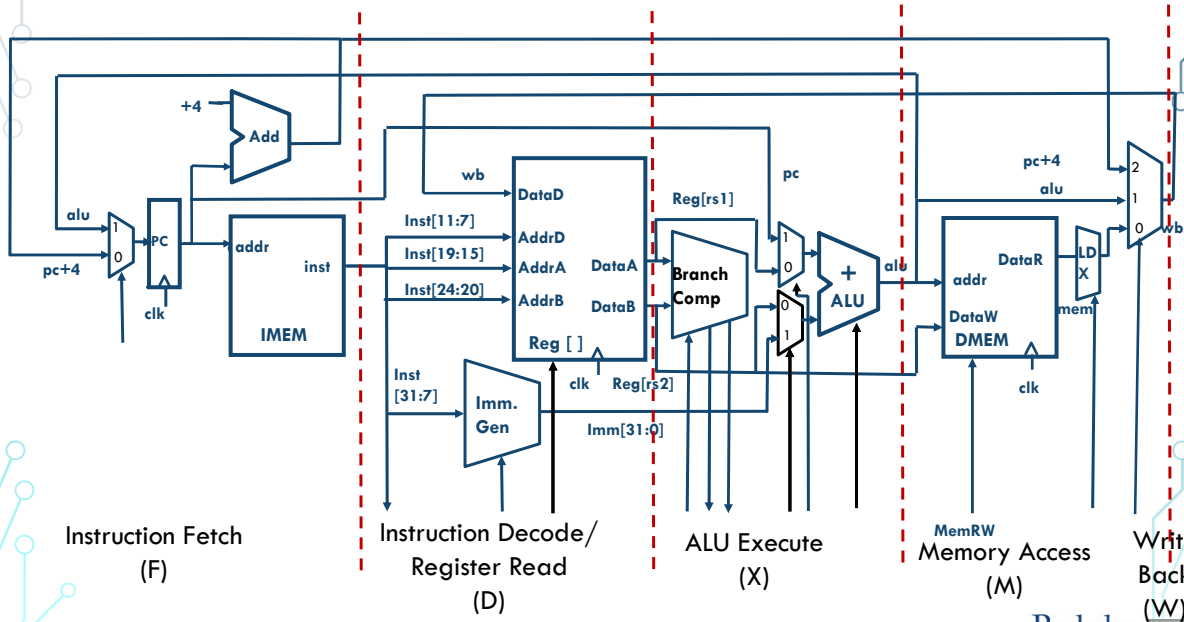
EECS151 L09 PIPELINING

Nikolić, Fall 2021

47

47

Complete RV32I Datapath with Control

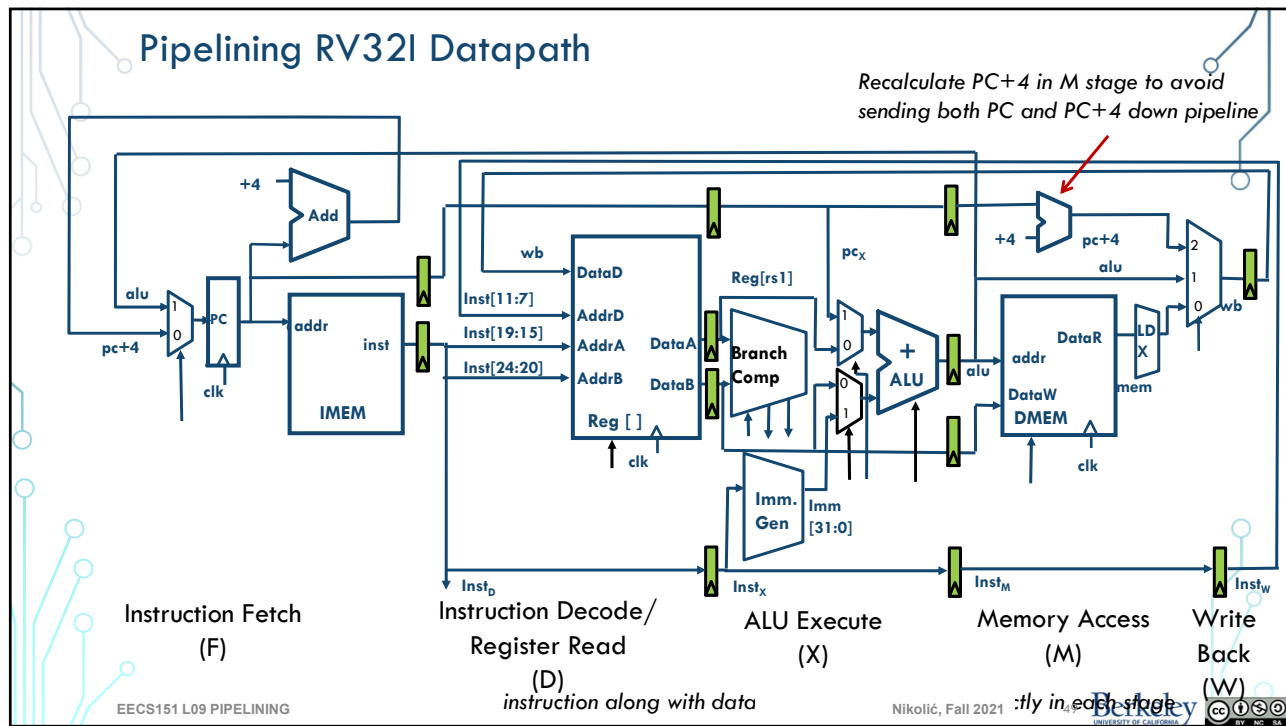


EECS151 L09 PIPELINING

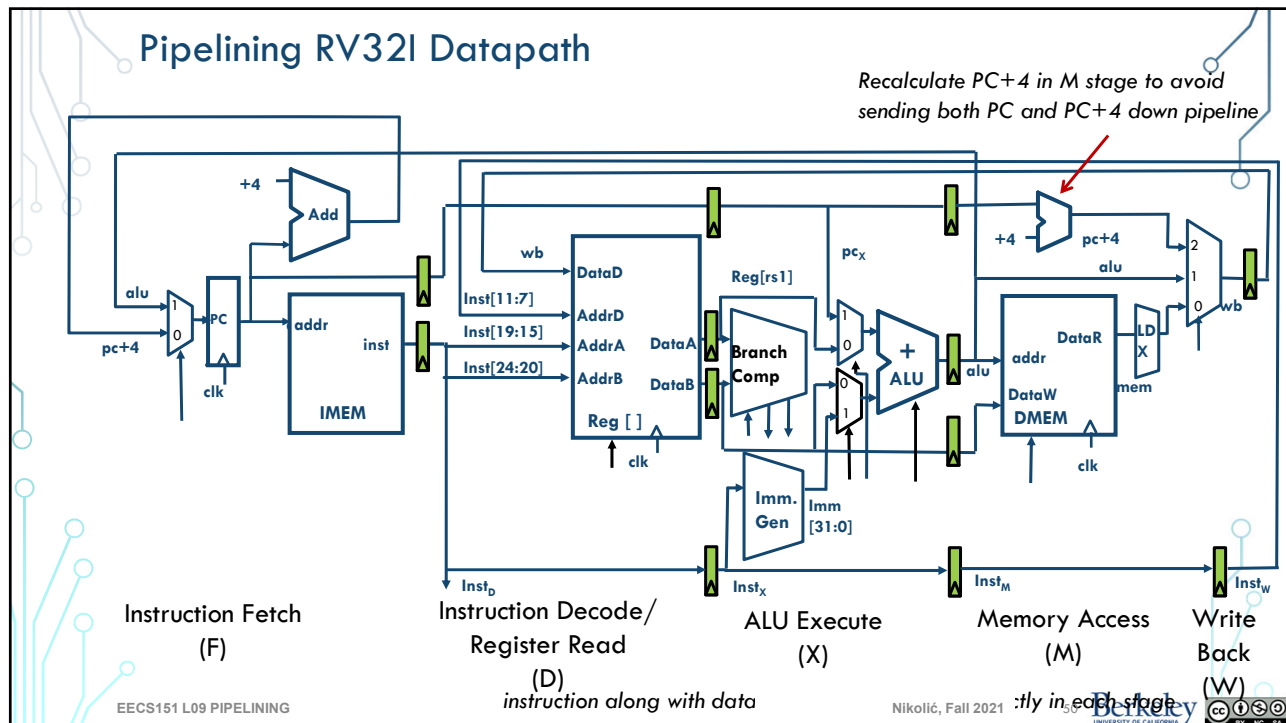
Nikolić, Fall 2021

48

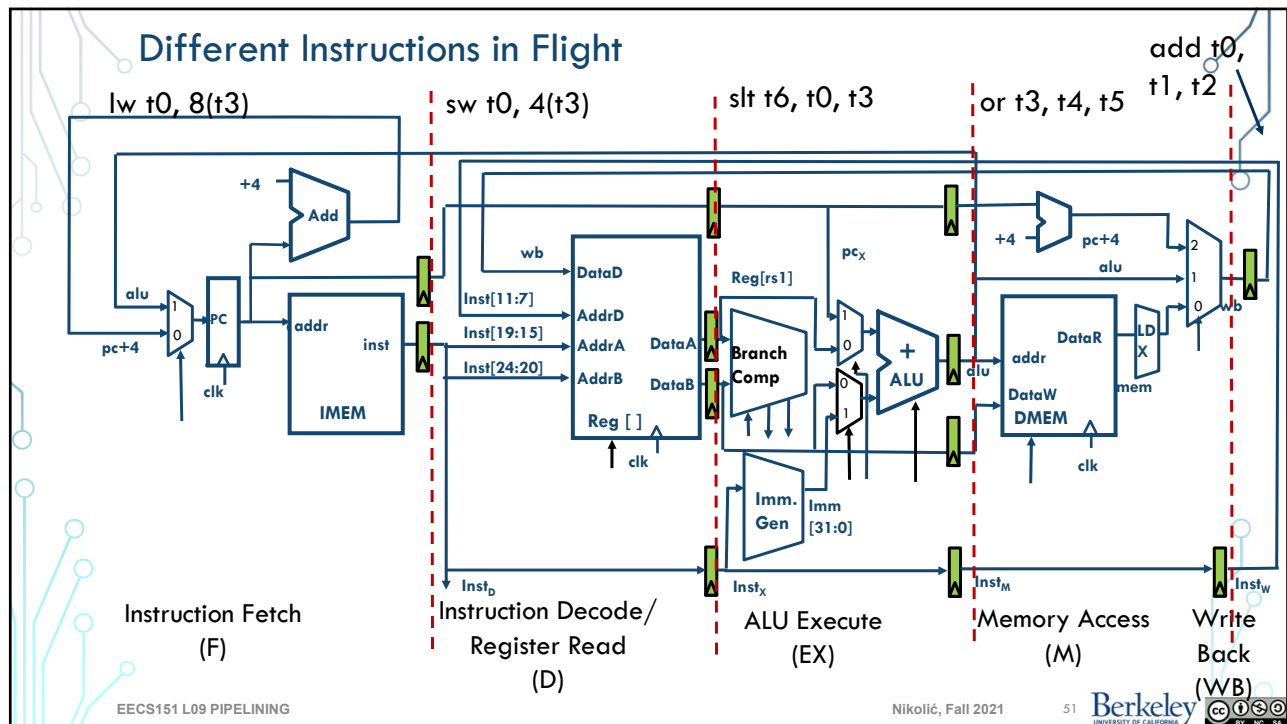
48



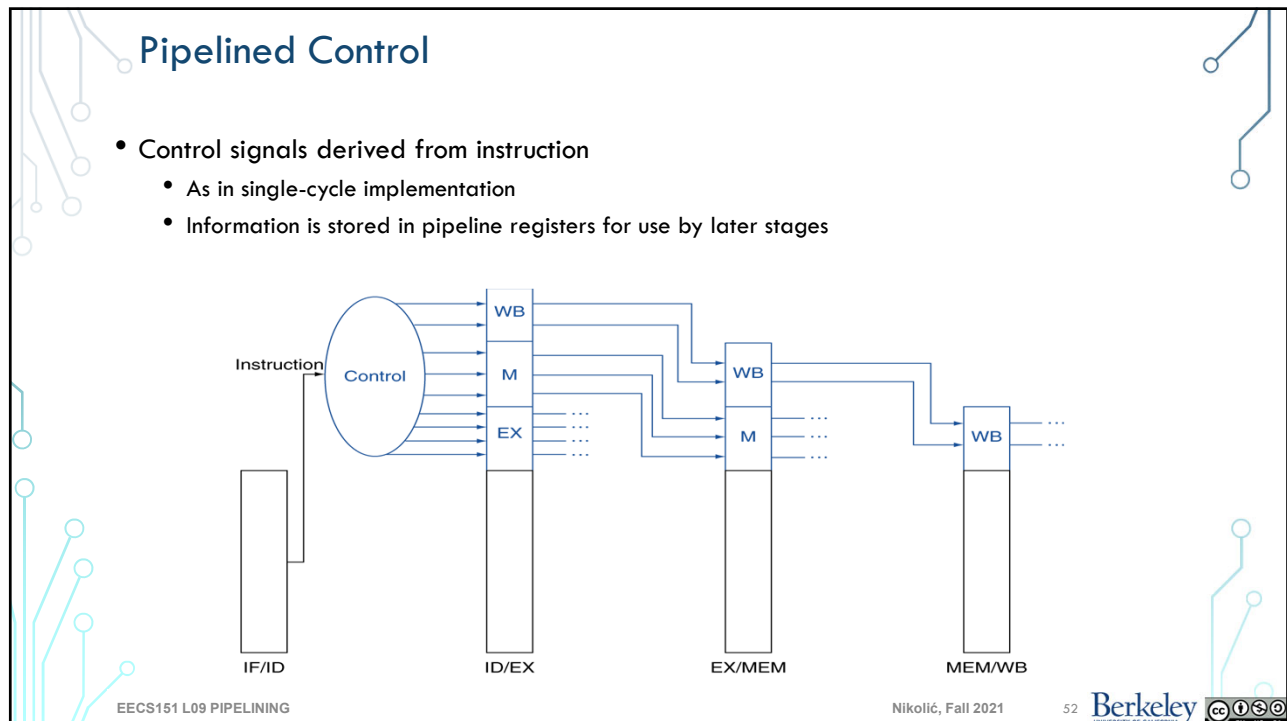
49



50



51



52

Summary

- RISC-V ISA
 - Completed the datapath with B-, J-, U-instructions
- Control
 - Can be implemented as a ROM while prototyping
 - Synthesized as custom logic
- Pipelining to increase throughput
 - 5-stage pipeline example