

EECS151 : Introduction to Digital Design and ICs

Lecture 9 – RISC-V ISA, Pipelining

Bora Nikolić

August 20, 2021, Tom's hardware
**Tesla Packs 50 Billion Transistors Onto D1 Dojo Chip
 Designed to Conquer Artificial Intelligence Training**

D1 delivers 362 TeraFLOPs of power.
 Called the D1, the chip resembles a part of the Dojo supercomputer used to train AI models inside Tesla HQ, which are later deployed in various applications. The D1 chip is a product of TSMC's manufacturing efforts, forged in a 7nm semiconductor node. Packing over 50 billion transistors, the chip boasts a huge die size of 645mm².




Image credit: Dennis Hong / Twitter



EECS151 L09 PIPELINING

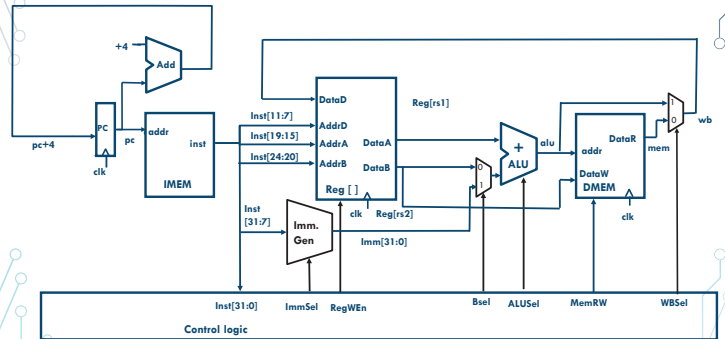
Nikolić, Fall 2021

Berkeley

Review

- RISC-V ISA
 - Open, with increasing adoption
- RISC-V processor
 - A large state machine
 - Datapath + control
 - Reviewed R-, I-, S-format instructions and corresponding datapath elements

Datapath So Far (R-, I-, S Instruction Types)

RISC-V
B-Format Instructions

B-Format - RISC-V Conditional Branches


- E.g., BEQ x1, x2, Label
- Branches read two registers but don't write a register (similar to stores)
- How to encode label, i.e., where to branch to?

RISC-V Feature, n×16-bit instructions

- Extensions to RISC-V base ISA support 16-bit compressed instructions and also variable-length instructions that are multiples of 16-bits in length
- To enable this, RISC-V scales the branch offset by 2 bytes even when there are no 16-bit instructions
- Reduces branch reach by half and means that 1/2 of possible targets will be errors on RISC-V processors that only support 32-bit instructions (as used in this class)
- RISC-V conditional branches can only reach $\pm 2^{10} \times 32$ -bit instructions on either side of PC

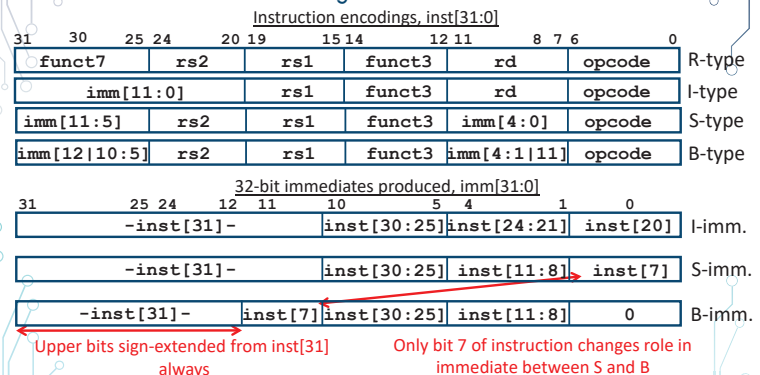
RISC-V Branch Immediates

- 12-bit immediate encodes PC-relative offset of -4096 to +4094 bytes in multiples of 2 bytes
- RISC-V approach: keep 11 immediate bits in fixed position in output value, and rotate LSB of S-format to be bit 12 of B-format



Only one bit changes position between S and B, so only need a single-bit 2-way mux

RISC-V Immediate Encoding



Upper bits sign-extended from inst[31] always


Only bit 7 of instruction changes role in immediate between S and B

Branch Example, complete encoding


beq x19,x10, offset = 16 bytes

13-bit immediate, imm[12:0], with value 16

imm[0] discarded, always zero

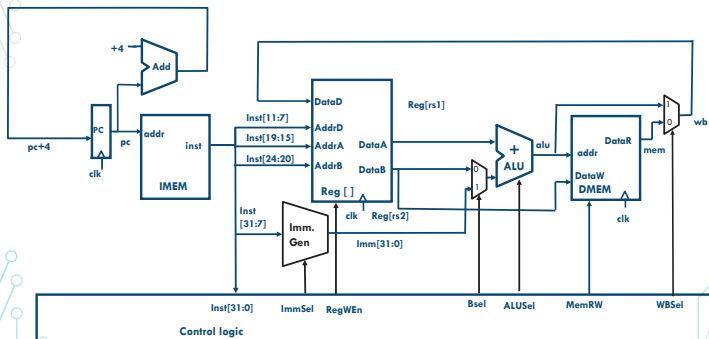


Implementing Branches



- B-format is mostly same as S-format, with two register sources (rs1/rs2) and a 12-bit immediate
- But now immediate represents values -4096 to +4094 in 2-byte increments
- The 12 immediate bits encode even 13-bit signed byte offsets (lowest bit of offset is always zero, so no need to store it)

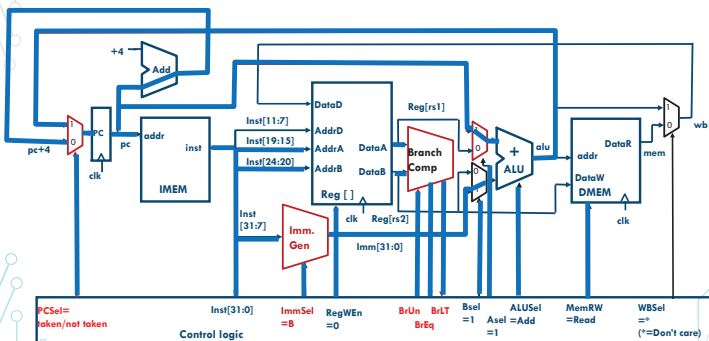
Datapath So Far




To Add Branches

- Different change to the state:
 - $PC = \begin{cases} PC + 4, & \text{branch not taken} \\ PC + \text{immediate}, & \text{branch taken} \end{cases}$
- Six branch instructions: BEQ, BNE, BLT, BGE, BLTU, BGEU
- Need to compute $PC + \text{immediate}$ and to compare values of rs1 and rs2
 - Need another add/sub unit

Adding Branches



Branch Comparator



- BrEq = 1, if A=B
- BrLT = 1, if A < B
- BrUn = 1 selects unsigned comparison for BrLT, 0=signed
- BGE branch: $A \geq B$, if $\overline{A < B}$

All RISC-V Branch Instructions

| | | | | | | |
|--------------|-----|-----|-----|-------------|---------|------|
| imm[12:10:5] | rs2 | rs1 | 000 | imm[4:1:11] | 1100011 | BEQ |
| imm[12:10:5] | rs2 | rs1 | 001 | imm[4:1:11] | 1100011 | BNE |
| imm[12:10:5] | rs2 | rs1 | 100 | imm[4:1:11] | 1100011 | BLT |
| imm[12:10:5] | rs2 | rs1 | 101 | imm[4:1:11] | 1100011 | BGE |
| imm[12:10:5] | rs2 | rs1 | 110 | imm[4:1:11] | 1100011 | BLTU |
| imm[12:10:5] | rs2 | rs1 | 111 | imm[4:1:11] | 1100011 | BGEU |


Administrivia

- Homework 4 is due next Monday
 - No new homework this week
 - Homework 5 will be posted next week, due after the midterm
- Lab 5 this week
 - No lab next week
 - Lab 6 (last) after the midterm
- Midterm 1 on October 7, 7-8:30pm




RISC-V J-Format Instructions

J-Format for Jump Instructions



- JAL saves PC+4 in register rd (the return address)
 - Assembler "j" jump is pseudo-instruction, uses JAL but sets rd=x0 to discard return address
- Set PC = PC + offset (PC-relative jump)
- Target somewhere within $\pm 2^{19}$ locations, 2 bytes apart
 - $\pm 2^{18}$ 32-bit instructions
- Immediate encoding optimized similarly to branch instruction to reduce hardware cost


JALR Instruction (I-Format)



- JALR rd, rs, immediate
 - Writes PC+4 to rd (return address)
 - Sets PC = rs1 + immediate (and sets the LSB to 0)
 - Uses same immediates as arithmetic and loads
 - no** multiplication by 2 bytes
 - In contrast to branches and JAL

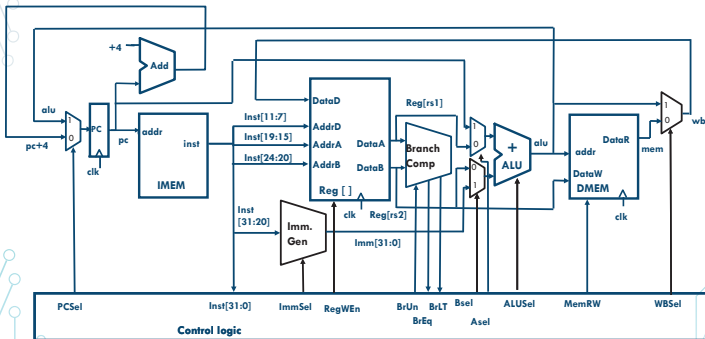
Otherwise, we would have yet another new encoding

Let's Add JALR (I-Format)

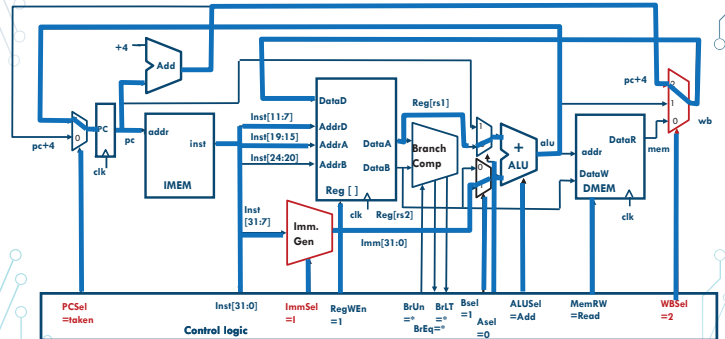


- JALR rd, rs, immediate
- Two changes to the state
 - Writes PC+4 to rd (return address)
 - Sets PC = rs + immediate
 - Uses same immediates as arithmetic and loads
 - no** multiplication by 2 bytes
 - LSB is ignored


Datapath So Far, with Branches



Adding JALR

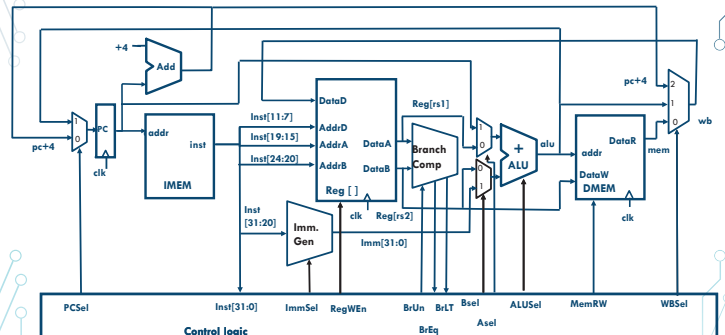


Adding JAL

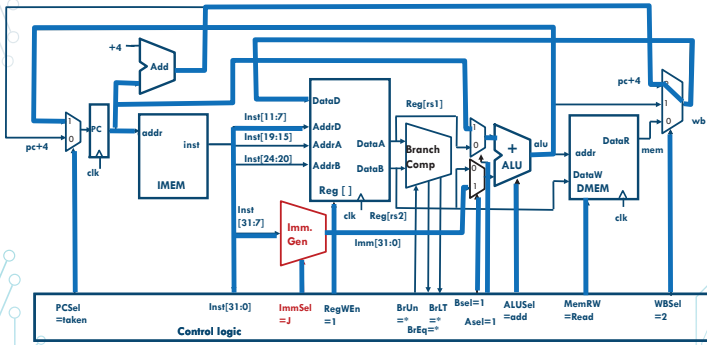


- JAL saves PC+4 in register rd (the return address)
- Set PC = PC + offset (PC-relative jump)
- Target somewhere within $\pm 2^{19}$ locations, 2 bytes apart
 - $\pm 2^{18}$ 32-bit instructions
- Immediate encoding optimized similarly to branch instruction to reduce hardware cost

Datapath with JALR



Adding JAL



EECS151 L09 PIPELINING

Nikolić, Fall 2021


25

Berkeley



RISC-V U-Format Instructions

U-Format for “Upper Immediate” Instructions



- Has 20-bit immediate in upper 20 bits of 32-bit instruction word
- One destination register, rd
- Used for two instructions
 - lui – Load Upper Immediate
 - auipc – Add Upper Immediate to PC

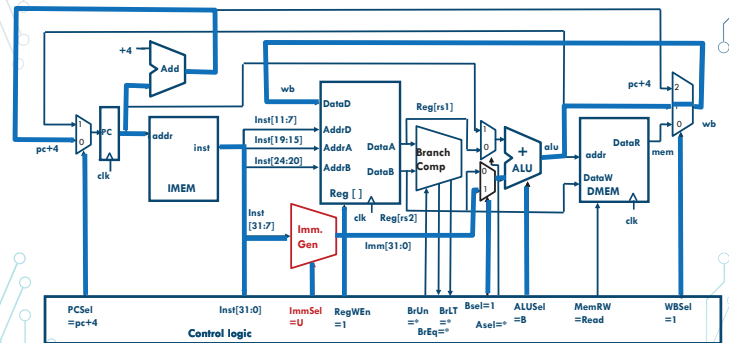
EECS151 L09 PIPELINING

Nikolić, Fall 2021

27

Berkeley

Implementing lui



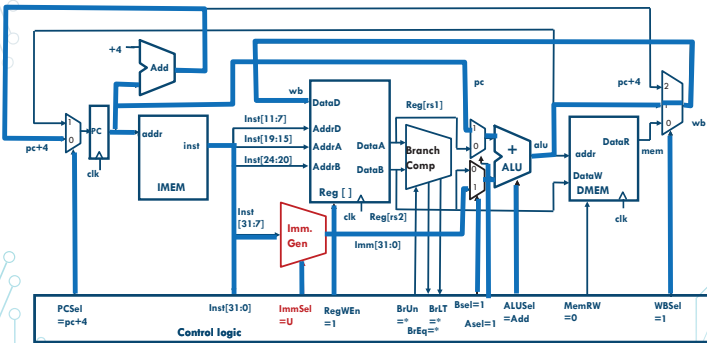
EECS151 L09 PIPELINING

Nikolić, Fall 2021

28

Berkeley

Implementing auipc



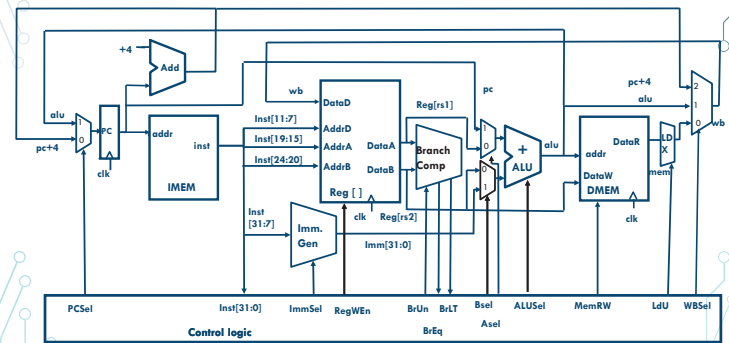
EECS151 L09 PIPELINING

Nikolić, Fall 2021

29

Berkeley

Complete RV32I Datapath!



EECS151 L09 PIPELINING

Nikolić, Fall 2021

30

Berkeley

Recap: Complete RV32I ISA

| Open Reference Card | | | |
|----------------------------------|-----------------------------------|----------------------------------|-------------------------|
| Category | Name | Base Integer Instructions: RV32I | RV32I Base |
| Shifts | Shift Left Logical | R SLL rd,rs1,rs2 | Load Byte |
| | Shift Left Logical Immediate | I SLLI rd,rs1,shamt | Load Halfword |
| | Shift Right Logical | R SRL rd,rs1,rs2 | Load Byte Unaligned |
| | Shift Right Logical Immediate | I SRLI rd,rs1,shamt | Load Halfword Unaligned |
| | Shift Right Arithmetic | R SRA rd,rs1,rs2 | Load Word |
| Shift Right Arithmetic Immediate | Shift Right Arithmetic Immediate | I SRAI rd,rs1,shamt | Store Byte |
| | Shift Right Arithmetic Immediate | I SRAI rd,rs1,shamt | Store Halfword |
| Arithmetic | ADD | R ADD rd,rs1,rs2 | Store Word |
| | ADD Immediate | I ADDI rd,rs1,imm | Store Word |
| Logical | XOR | R XOR rd,rs1,rs2 | Branch |
| | XOR Immediate | I XORI rd,rs1,imm | Branch |
| Compare | Set Less Than | R SLT rd,rs1,rs2 | Branch |
| | Set Less Than Immediate | I SLTI rd,rs1,imm | Branch |
| Environment | Set Less Than Unaligned | R SLTU rd,rs1,rs2 | Branch |
| | Set Less Than Unaligned Immediate | I SLTIU rd,rs1,imm | Branch |

- 40 instructions are enough to run any C program

EECS151 L09 PIPELINING

Nikolić, Fall 2021

31

Berkeley

Summary of RISC-V Instruction Formats

| 31 | 30 | 25 | 24 | 21 | 20 | 19 | 15 | 14 | 12 | 11 | 8 | 7 | 6 | 0 | | |
|------------------|----|----|----|----|----|-----|----|------------|----|--------|---|-------------|--------|--------|--|--------|
| funct7 | | | | | | rs2 | | rs1 | | funct3 | | rd | | opcode | | R-type |
| imm[11:0] | | | | | | | | rs1 | | funct3 | | rd | | opcode | | I-type |
| imm[11:5] | | | | | | rs2 | | rs1 | | funct3 | | imm[4:0] | | opcode | | S-type |
| imm[12:10:5] | | | | | | rs2 | | rs1 | | funct3 | | imm[4:1:11] | | opcode | | B-type |
| imm[31:12] | | | | | | | | | | rd | | | opcode | | | U-type |
| imm[20:10:1:11]] | | | | | | | | imm[19:12] | | | | rd | | opcode | | J-type |

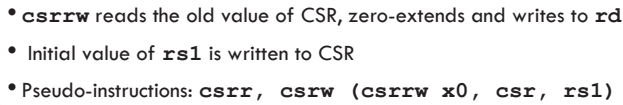
EECS151 L09 PIPELINING

Nikolić, Fall 2021

32

Berkeley

- 4096 CSRs in a separate address space

[illegible][illegible]


The diagram illustrates a 5-stage EECs151 L09 pipeline. The top part shows a block diagram of the pipeline with stages: PCOut, Instruction, Decode, Execute, and WriteBack. Each stage has a multiplexer and a register. The bottom part shows a timing diagram for the pipeline. The diagram includes signals for Clock, PC, PC+4, inst[31:0], Control logic, Reg[r1], Reg[r2], alu, wb, and Real1. The timing diagram shows the data flow and control signals over time, with labels for instructions 1000, 1004, and 1008. The Real1 signal is shown at the bottom, with labels for Reg[2]+Reg[3] and Reg[7]+Reg[9].

| Inst[31:0] | BrEq | BrLt | PCSel | ImmSel | BrUn | ASel | BSel | ALUSel | MemRW | RegWen | WBSel |
|------------|------|------|-------|--------|------|------|------|--------|-------|--------|-------|
| add | * | * | +4 | * | * | Reg | Reg | Add | Read | 1 | ALU |
| sub | * | * | +4 | * | * | Reg | Reg | Sub | Read | 1 | ALU |
| (R-R Op) | * | * | +4 | * | * | Reg | Reg | (Op) | Read | 1 | ALU |
| addi | * | * | +4 | I | * | Reg | Imm | Add | Read | 1 | ALU |
| lw | * | * | +4 | I | * | Reg | Imm | Add | Read | 1 | Mem |
| sw | * | * | +4 | S | * | Reg | Imm | Add | Write | 0 | * |
| beq | 0 | * | +4 | B | * | PC | Imm | Add | Read | 0 | * |
| bne | 1 | * | ALU | B | * | PC | Imm | Add | Read | 0 | * |
| bne | 0 | * | ALU | B | * | PC | Imm | Add | Read | 0 | * |
| bne | 1 | * | +4 | B | * | PC | Imm | Add | Read | 0 | * |
| blt | * | 1 | ALU | B | 0 | PC | Imm | Add | Read | 0 | * |
| bltu | * | 1 | ALU | B | 1 | PC | Imm | Add | Read | 0 | * |
| jalr | * | * | ALU | I | * | Reg | Imm | Add | Read | 1 | PC+4 |
| jal | * | * | ALU | J | * | PC | Imm | Add | Read | 1 | PC+4 |
| auipc | * | * | +4 | I | * | PC | Imm | Arif | Read | 1 | ALU |

[illegible]

Control Realization Options

- ROM
 - “Read-Only Memory”
 - Regular structure
 - Can be easily reprogrammed
 - fix errors
 - add instructions
- Combinatorial Logic
 - Start from a truth table
 - More compact, faster
 - Use synthesis tools



Combinational Logic Control

- Decoder is typically hierarchical
 - First decode opcode, and figure out instruction type
 - E.g. branches are $\text{Inst}[6:2] = 11000$
 - Then determine the actual instruction
 - $\text{Inst}[30] + \text{Inst}[14:12]$
- Modularity helps simplify and speed up logic
 - Narrows problem space for logic synthesis

Combinational Logic Control

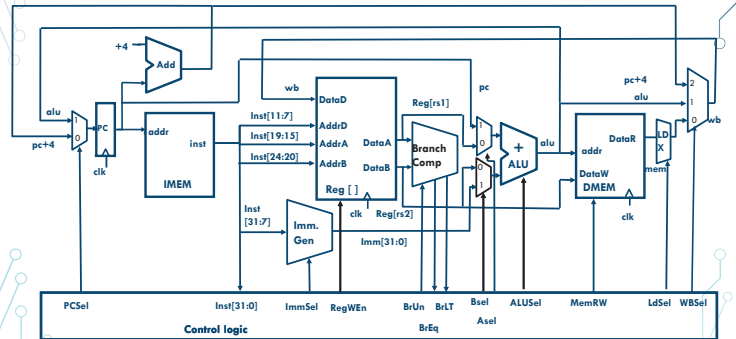
- Simple example: BrUn

| | inst[14:12] | inst[6:2] | inst[14:13] | inst[12] |
|--------------|-------------|-----------|-------------|------------|
| imm[12:10:5] | rs2 | rs1 | 000 | imm[4:1:1] |
| imm[12:10:5] | rs2 | rs1 | 001 | imm[4:1:1] |
| imm[12:10:5] | rs2 | rs1 | 100 | imm[4:1:1] |
| imm[12:10:5] | rs2 | rs1 | 101 | imm[4:1:1] |
| imm[12:10:5] | rs2 | rs1 | 110 | imm[4:1:1] |
| imm[12:10:5] | rs2 | rs1 | 111 | imm[4:1:1] |

- How to decode whether BrUn is 1?

- $\text{BrUn} = \text{Inst}[13] \bullet \text{Branch}$
- $\text{Branch} = \text{Inst}[6] \bullet \text{Inst}[5] \bullet \neg \text{Inst}[4] \bullet \neg \text{Inst}[3] \bullet \neg \text{Inst}[2]$

Complete RV32I Datapath with Control



Peer Instruction(s): Critical Path yellkey: crime

Critical path for a addi

$$R[rd] = R[rs1] + \text{imm}$$

- 1) $t_{\text{clk-q}} + t_{\text{Add}} + t_{\text{IMEM}} + t_{\text{Reg}} + t_{\text{BComp}} + t_{\text{ALU}} + t_{\text{DMEM}} + t_{\text{mux}} + t_{\text{Setup}}$
- 2) $t_{\text{clk-q}} + t_{\text{IMEM}} + \max\{t_{\text{Reg}}, t_{\text{Imm}}\} + t_{\text{ALU}} + 2t_{\text{mux}} + t_{\text{Setup}}$
- 3) $t_{\text{clk-q}} + t_{\text{IMEM}} + \max\{t_{\text{Reg}}, t_{\text{Imm}}\} + t_{\text{ALU}} + 3t_{\text{mux}} + t_{\text{Setup}}$
- 4) None of the above

Peer Instruction(s): Critical Path yellkey: physical

Critical path for a addi

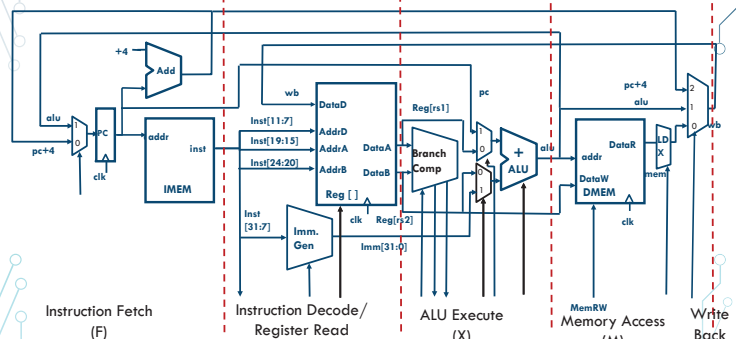
$$R[rd] = R[rs1] + \text{imm}$$

- 1) $t_{\text{clk-q}} + t_{\text{Add}} + t_{\text{IMEM}} + t_{\text{Reg}} + t_{\text{BComp}} + t_{\text{ALU}} + t_{\text{DMEM}} + t_{\text{mux}} + t_{\text{Setup}}$
- 2) $t_{\text{clk-q}} + t_{\text{IMEM}} + \max\{t_{\text{Reg}}, t_{\text{Imm}}\} + t_{\text{ALU}} + 2t_{\text{mux}} + t_{\text{Setup}}$
- 3) $t_{\text{clk-q}} + t_{\text{IMEM}} + \max\{t_{\text{Reg}}, t_{\text{Imm}}\} + t_{\text{ALU}} + 3t_{\text{mux}} + t_{\text{Setup}}$
- 4) None of the above

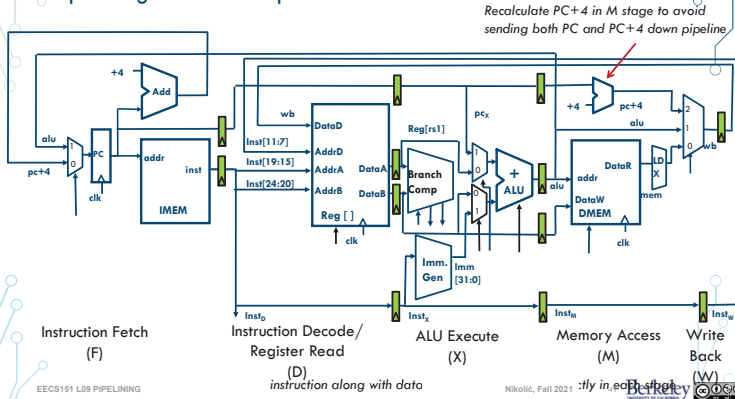
Pipelining



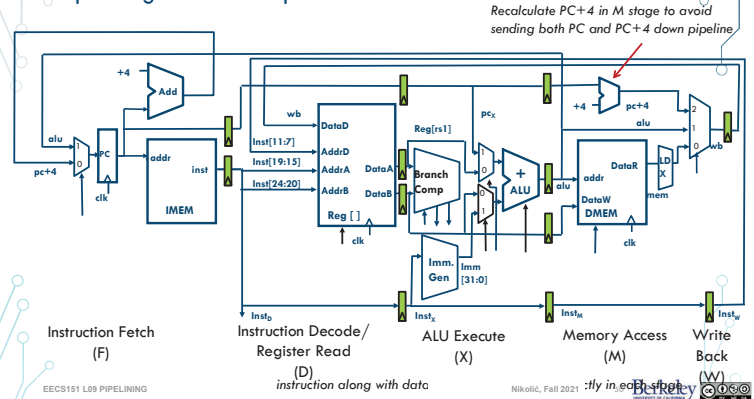
Complete RV32I Datapath with Control



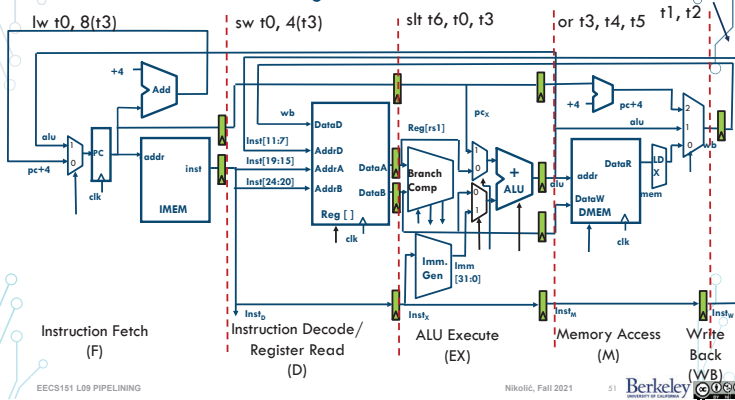
Pipelining RV32I Datapath



Pipelining RV32I Datapath




Different Instructions in Flight



Pipelined Control

- Control signals derived from instruction
 - As in single-cycle implementation
 - Information is stored in pipeline registers for use by later stages



Summary

- RISC-V ISA
 - Completed the datapath with B-, J-, U-instructions
- Control
 - Can be implemented as a ROM while prototyping
 - Synthesized as custom logic
- Pipelining to increase throughput
 - 5-stage pipeline example