

EECS151 : Introduction to Digital Design and ICs

Lecture 10 – Pipelining, FPGAs

Bora Nikolić

Stalking the Elusive Computer Bug

From at least the time of Thomas Edison, U.S. engineers have used the word "bug" to refer to flaws in the systems they developed. This short word conveniently covered a multitude of possible problems. It also suggested that difficulties were small and could be easily corrected. IBM engineers who installed the ASAC Mark I at Harvard University in 1944 taught the phrase to the staff there. Grace Murray Hopper used the word with particular enthusiasm in documents relating to her work. In 1947, when technicians building the Mark II computer at Harvard discovered a moth in one of the relays, they saved it as the first actual case of a bug being found. In the early 1950s, the terms "bug" and "debug," as applied to computers and computer programs, began to appear not only in computer documentation but even in the popular press.

Peggy Aldrich Kidwell,
IEEE Annals of the History of Computing, 1998.

EECS151 L10 FPGA



Grace Murray Hopper
Logbook of the Mark II for 9/9/1947

Nikolić, Fall 2021

Berkeley

Review

- RISC-V ISA
 - Completed the datapath with B-, J-, U-instructions
- Control
 - Can be implemented as a ROM while prototyping
 - Synthesized as custom logic
- Pipelining to increase throughput
 - 5-stage pipeline example

EECS151 L10 FPGA

Nikolić, Fall 2021

Berkeley

Pipelining

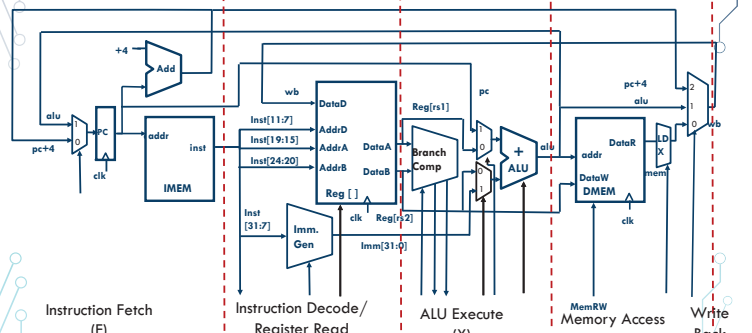


EECS151 L10 FPGA

Nikolić, Fall 2021

Berkeley

Complete RV32I Datapath with Control



Instruction Fetch (F)

Instruction Decode/ Register Read (D)

ALU Execute (X)

Memory Access (M)

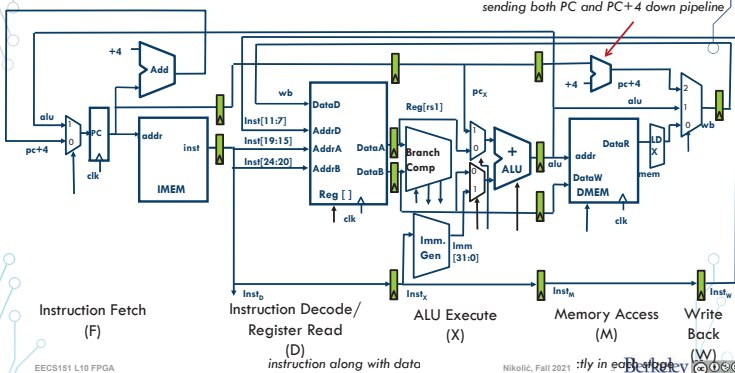
Write Back (W)

EECS151 L10 FPGA

Nikolić, Fall 2021

Berkeley

Pipelining RV32I Datapath



Instruction Fetch (F)

Instruction Decode/ Register Read (D)

ALU Execute (X)

Memory Access (M)

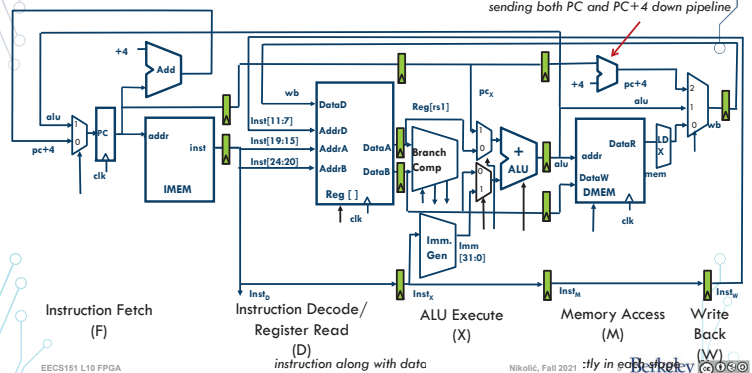
Write Back (W)

EECS151 L10 FPGA

Nikolić, Fall 2021

Berkeley

Pipelining RV32I Datapath



Instruction Fetch (F)

Instruction Decode/ Register Read (D)

ALU Execute (X)

Memory Access (M)

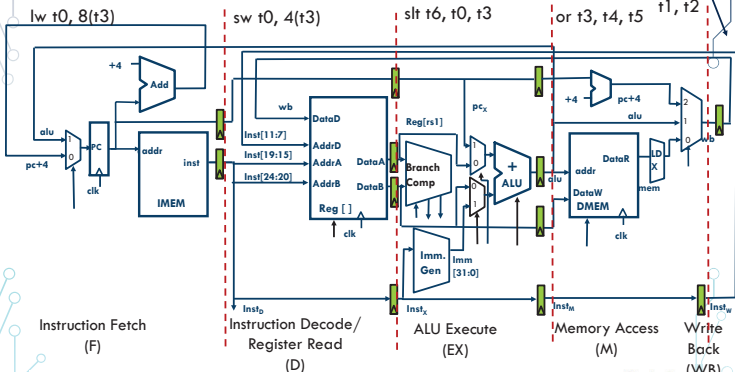
Write Back (W)

EECS151 L10 FPGA

Nikolić, Fall 2021

Berkeley

Different Instructions in Flight



Instruction Fetch (F)

Instruction Decode/ Register Read (D)

ALU Execute (EX)

Memory Access (M)

Write Back (WB)


EECS151 L10 FPGA

Nikolić, Fall 2021

Berkeley

Pipelined Control

- Control signals derived from instruction
 - As in single-cycle implementation
 - Information is stored in pipeline registers for use by later stages



EECS151 L10 FPGA

Nikolić, Fall 2021

Berkeley

Pipeline Hazards



Pipelining Hazards

A hazard is a situation that prevents starting the next instruction in the next clock cycle

1) Structural hazard

- A required resource is busy (e.g. needed in multiple stages)


2) Data hazard

- Data dependency between instructions
- Need to wait for previous instruction to complete its data read/write

3) Control hazard

- Flow of execution depends on previous instruction


Control Hazards



Observation

- If branch not taken, then instructions fetched sequentially after branch are correct
- If branch or jump taken, then need to flush incorrect instructions from pipeline by converting to NOPs


Kill Instructions after Branch if Taken



Reducing Branch Penalties

- Every taken branch in simple pipeline costs 2 'dead' cycles
- To improve performance, use "branch prediction" to guess which way branch will go earlier in pipeline
- Only flush pipeline if branch prediction was incorrect

Branch Prediction



Quiz: Hazards

- How many data hazards exist in the following sequence (assuming a 5-stage pipeline)?



```
add x3, x1, x2
or x5, x3, x4
add x2, x5, x3
lw x6, x2, 12
sw x1, x6, 36
```

FPGAs: Overview



Field Programmable Gate Arrays (FPGAs)

- An integrated circuit designed to be configured by a customer or a designer after manufacturing, i.e., field programmable.




- The FPGA configuration is generally specified using a hardware description language, similar to that used for ASICs.

- Two dominant FPGA makers:

- Xilinx
- Altera (now Intel)

FPGA Overview

- Basic idea:
 - Two dimensional array of logic blocks and flip-flops with means for the user to configure:
 - The function of each block
 - The interconnection between blocks
- Configurable Logic Blocks (CLBs)
 - FPGA's Functional Units
- Reconfigurable Interconnect
 - Connecting CLBs together




State-of-the-art Xilinx FPGAs

	45nm	28nm	20nm	16nm
	SPARTAN [®]	VIRTEX [®] KINTEX [®] ARTIX [®] SPARTAN [®]	KINTEX [®]	VIRTEX [®] KINTEX [®]
				Virtex Ultra-scale
Device Name	VU1P	VU1P	VU1P	VU1P
System Logic Cells (K)	862	1,314	1,724	2,586
System Logic Cells (K)	862	1,314	1,724	2,586
CLB Flip-Flops (K)	788	1,201	1,576	2,364
CLB LUTs (K)	394	601	788	1,182
Max. Dist. RAM (KB)	12.0	18.3	24.1	36.1
Total Block RAM (KB)	25.3	36.0	50.6	75.9
UltraRAM (KB)	90.0	132.2	180.0	270.0
HBM DRAM (GB)	—	—	—	—
HBM AXI Interfaces	—	—	—	—
Clock Mgmt. Tiles (CMTs)	30	30	30	30
DSP Slices	2,280	3,474	4,560	6,840
Peak INT8 DSP (TOP/s)	7.1	10.8	14.2	21.3
PCIe Gen3 x16	2	4	4	6
PCIe Gen3 x16 Gen4 x8	—	—	—	—
150G Interlink	3	4	6	9
100G Ethernet w/ KRM RS-FEC	3	4	6	9
Max. Single-Ended HP I/Os	510	832	832	832
QTY. 32.75Gb/s Transceivers	40	80	80	120
GTM 18Gb/s/PAM4 Transceivers	—	—	—	—
100G/50G KRM FEC	—	—	—	—
Extended [®]	-1-2-2L-3	-1-2-2L-3	-1-2-2L-3	-1-2-2L-3
Industrial	-1-2	-1-2	-1-2	-1-2

Configurable Logic Blocks (CLBs)

- Basic FPGA functional unit
- Implements both combinational and sequential logic
- Includes:
 - Look-up table
 - Register (Flip-Flop)
 - Multiplexers




FPGA: CLBs



Look-Up Table Implementation

- Implement truth table in small memories
 - SRAM/Latch arrays
 - "Latch" is actually a flip-flop
- n-bit LUT is implemented as a $2^n \times 1$ memory:
 - inputs choose one of 2^n memory locations.
 - memory locations (latches) are normally loaded with values from user's configuration bit stream.
 - Inputs to mux control are the CLB inputs.
- Result is a general purpose "logic gate".
 - n-LUT can implement any function of n inputs!



Look-Up Table Implementation


- An n-LUT is a direct implementation of a function truth-table.
- Each location holds the value of the function corresponding to one input combination.
- LUT size grows exponentially with # of inputs.
 - 64 input LUT requires $2^{64} = 1.84 \times 10^{19}$ bits storage.
 - 4-input ~ 8-input LUT

Example: 4-LUT

INPUTS	
0000	F(0,0,0,0) ← store in 1st latch
0001	F(0,0,0,1) ← store in 2nd latch
0010	F(0,0,1,0)
0011	F(0,0,1,1)
0111	
0100	
0101	
0110	
0111	
1000	
1001	
1010	
1011	
1100	
1101	
1110	
1111	

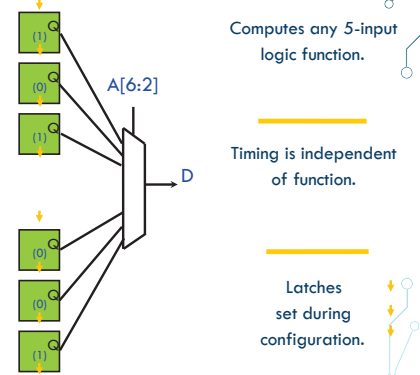

Slices

- Each CLB contains two slices.
- LUTs and registers are split across slices.
 - 4 LUTs and 8 FFs in 7-series.
- Two types of slices:
 - SLICEM: Full slice
 - LUT can be used for logic *and* memory/shift registers.
 - Has wide multiplexers and carry chain
 - SLICE: logic and arithmetic only
 - LUT can only be used for logic (no memory)
 - Has wide multiplexers and carry chain




Constructing a SLICE

5-Input Look-Up Table



Constructing a SLICE


6-input LUT



May be used
as one 6-input LUT
(D6 out)
...
... or as two
5-input LUTs
(D6 and D5)

Combinational
logic
(post configuration)

The Simplest View of A Slice



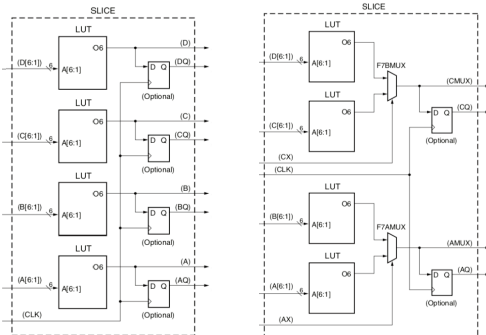
Four 6-LUTs

Four Flip-Flops

Switching fabric may see
combinational and registered
outputs.

An actual Virtex slice adds many small
features to this simplified diagram.
We show them one by one ...

How about 7-input LUT in a slice?

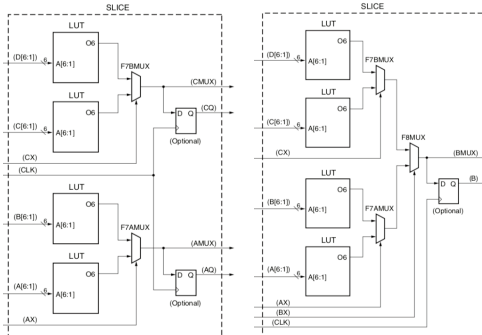


Two 7-LUTs

Extra MUX
(F7AMUX, F7BMUX)

Extra inputs
(AX and CX)

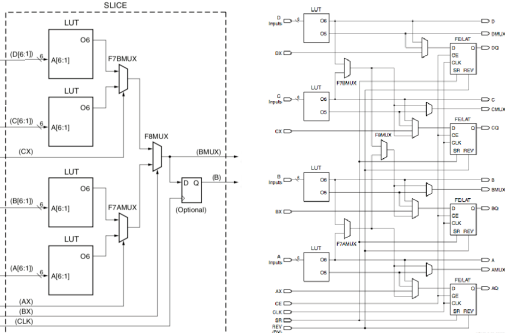
How about 8-input LUT in a slice?



Third MUX
(F8MUX)

Third input
(BX)

Extra MUXes to choose LUT outputs

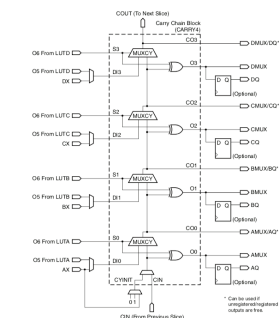



From eight 5-LUTs
... to one 8-LUT.

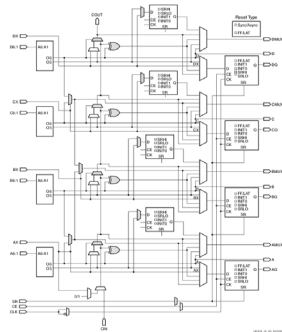
Combinational
or registered outs.

Extra Carry Chain

We can map
ripple-carry addition onto
carry-chain block.



Putting it all together ... a SLICEL.



EECS151 L10 FPGA

The previous slides explain all SLICEL features.

About 50% of the are SLICELs.

The other slices are SLICEMs, and have extra features.

Hlsuic, Fall 2021



Administrivia

- Homework 4 is due next Monday
 - No new homework this week
 - Homework 5 will be posted next week, due after the midterm
- Lab 5 this week
 - No lab next week
 - Lab 6 (last) after the midterm
- Midterm 1 on October 7, 7-8:30pm
 - You will be assigned a classroom

EECS151 L10 FPGA

Hlsuic, Fall 2021



FPGA Interconnect




EECS151 L10 FPGA

Hlsuic, Fall 2021



Configurable Interconnect

- Between rows and columns of CLBs are wiring channels.
- These are programmable. Each wire can be connected in many ways.
- Switch Box:
 - Each interconnection has a transistor switch.
 - Each switch is controlled by 1-bit configuration register.



EECS151 L10 FPGA

Hlsuic, Fall 2021



FPGA Features: BRAMs, DSP, AI



EECS151 L10 FPGA

Hlsuic, Fall 2021




Diverse Resources on FPGA

Colors represent different types of resources:

- Logic
- Block RAM
- DSPs
- Clocking
- I/O
- Serial I/O + PCI

A routing fabric runs throughout the chip to wire everything together.



Virtex-5 Die Photo [Xilinx]

EECS151 L10 FPGA

Hlsuic, Fall 2021



Block RAM

- Block Random Access Memory
- Used for storing large amounts of data:
 - 18Kb or 36Kb
 - Configurable bitwidth
 - 2 read and write ports
- More recently
 - UltraRAM in UltraScale+ devices




Figure 2: Block RAMs Arranged in Columns with Detailed Floorplan of XC3S5000

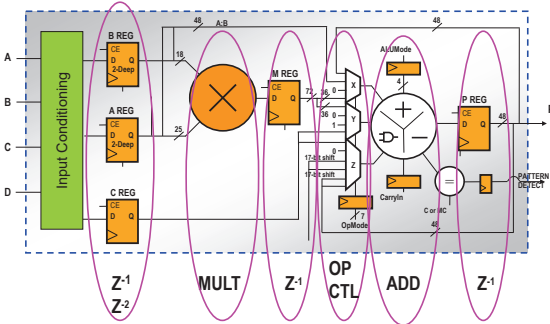
Xilinx Datasheet

EECS151 L10 FPGA

Hlsuic, Fall 2021



DSP Slice



Efficient implementation of multiply, add, bit-wise logical.

Xilinx Resource

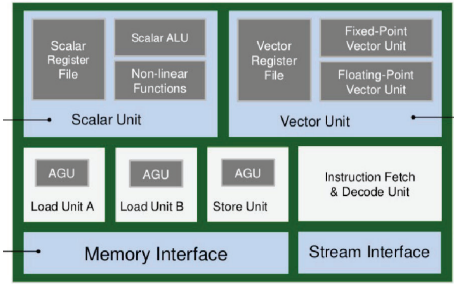
EECS151 L10 FPGA

Hlsuic, Fall 2021



AI Engine

- Versal AI Core



EECS151 L10 FPGA


HB@UC, Fall 2021

Xilinx HotChips'2019





State-of-the-art Xilinx FPGA Platform

- Versal (ACAP: Adaptive Compute Acceleration Platform)



Xilinx HotChips'2019

EECS151 L10 FPGA



HB@UC, Fall 2021



Summary

- Pipelining increases throughput
 - Structural, control and data hazards exist
- FPGAs are widely used for hardware prototyping and accelerating key applications.
- Core FPGA building blocks:
 - Configurable Logic Blocks (CLBs)
 - Slices
 - Look-Up Tables
 - Flip-Flops
 - Carry chain
 - Configurable Interconnect
 - Switch boxes
- Modern FPGA Designs:
 - BRAMs, DSPs, and AI Engines

EECS151 L10 FPGA

HB@UC, Fall 2021

