



EECS151 : Introduction to Digital Design and ICs

Lecture 23 – SRAM

Bora Nikolić



Intel's Process Roadmap to 2025 with 4nm, 3nm, 20A and 18A7
Jan Cunniff, Austin, July 2023
 2020, Intel 10nm SuperFin (10P). Current generation technology in use with Tiger Lake and Intel's Xe-HP discrete graphics solutions (DG1, DG2). The name stays the same.

2021 H2, Intel 7. Previously known as 10nm Enhanced Super Fin or 10E3P. Alder Lake and Sapphire Rapids will now be known as Intel 7nm products, showing a 10-15% performance per watt gain over 10P due to transistor optimizations. Alder Lake is currently in volume production.

2022 H2, Intel 4. Previously known as Intel 7nm, Intel earlier this year stated that its Meteor Lake processor will use a compute tile based on this process node technology, and the 4nm is now back in the air being tested. Intel expects a 20% performance per watt gain over the previous generation, and the technology was more EUV, mostly in the BEOL. Intel's next Xeon Scalable product, Granite Rapids, will also use a compute tile based on Intel 4.

2023 H2, Intel 3. Previously known as Intel 7+. Increased use of EUV and new high density libraries. This is where Intel's strategy becomes more modular – Intel 3 will share some features of Intel 4, but enough will be new enough to describe this as a new full node, in particular new high performance libraries.

2024, Intel 20A. Previously known as Intel 5nm. Moving to double digit naming, with the A standing for Analog, or 15A is equal to 7nm. Few details, but this is where Intel will move from PDKs to the version of GDSII-based (GDS) transistors called PDKs.

2025, Intel 18A. Not listed on the diagram above, but Intel is expecting to have an 18A process in 2025. 18A will be using ASML's latest EUV machines, known as High-NA machines, which are capable of more accurate photolithography. Intel has stated to us that it is ASML's Intel partner which is close to High-NA, and is set to receive the first production model of a High-NA machine.

EECS151 L23 SRAM



Nikolić Fall 2021



Review

- Latches are based on positive feedback
- Clk-Q delay calculated similarly to combinational logic
- Setup, hold defined as D-Clk times that correspond to Clk-Q delay increases
- Flip-flop is typically a latch pair

EECS151 L23 SRAM


Nikolić Fall 2021




SRAM

Random Access Memory Architecture

- Conceptual: Linear array of addresses
 - Each box holds some data
 - Not practical to physically realize
 - millions of 32b/64b words



- Create a 2-D array
 - Decode Row and Column address to get data




EECS151 L23 SRAM

Nikolić Fall 2021



Basic Memory Array


- Core
 - Wordlines to access rows
 - Bitlines to access columns
 - Data multiplexed onto columns
- Decoders
 - Addresses are binary
 - Row/column MUXes are 'one-hot' - only one is active at a time



Nikolić Fall 2021



Basic Static Memory Element



- If D is high, \bar{D} will be driven low
 - Which makes D stay high

- Positive feedback
- Same principle as in latches


EECS151 L23 SRAM

Nikolić Fall 2021




Positive Feedback: Bi-Stability

- As in latches



Writing into a Cross-Coupled Pair




- This is a 5T SRAM cell
 - Access transistor must be able to overpower the feedback; therefore must be large
 - Easier to write a 0, harder to write 1
- Can implement as a transmission gate as well; single-ended 6T cell
- There is a better solution...

EECS151 L23 SRAM

Nikolić Fall 2021



SRAM Cell




Since it is easier to write a 0 through NMOS, write only 0s, but on opposite sides!
When reading, measure the difference

EECS151 L23 SRAM

Nikolić Fall 2021



6-transistor CMOS SRAM Cell



- Wordline (WL) enables read/write access for a row
- Data is written/read differentially through shared BL, \overline{BL}


EECS151 L23 SRAM

Nikolić Fall 2021




SRAM Operation

Write




Hold



SRAM Operation

Read



SRAM read is non-destructive
- Reading the cell should not destroy the stored value

11

EECS151 L23 SRAM

Nikolić Fall 2021




EECS151 L23 SRAM

Nikolić Fall 2021




Sizing SRAM Cell

- Read stability: Cell should not change value during read
 - $Q = 0$: M_5, M_1 both on
 - Voltage divider between M_5, M_1
 - V_Q should stay low, not to flip M_4 - M_3 inverter
 - $(W/L)_1 > (W/L)_5$
- Typically $(W/L)_1 = 1.5 (W/L)_5$
 - In finFETs: $(W/L)_1 = 2(W/L)_5$
- Read speed: Both M_5 and M_1




Sizing SRAM Cell

- Writeability: Cell should be writeable by pulling BL low
 - $Q = 1, M_5, M_2$ both on
 - Voltage divider between M_5, M_2
 - V_Q should pull below the switching point of M_4 - M_3 inverter
 - $(W/L)_5 > (W/L)_2$
- Typically $(W/L)_5 = (W/L)_2$ in planar
 - In finFETs: $(W/L)_5 = 2(W/L)_2$
- 1:2:2 and 1:2:3 sizing



6T High-Current (HC) bitcell
0.049 μm^2 (1:2:2)



Song, ISSCC'16

14

EECS151 L23 SRAM

Nikolić Fall 2021

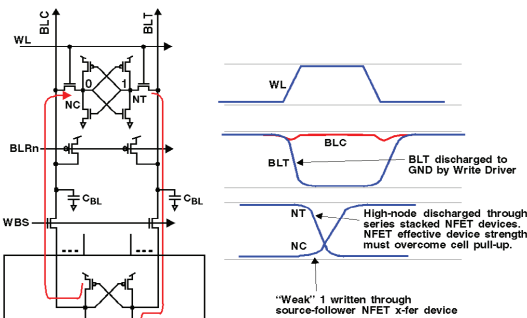


EECS151 L23 SRAM

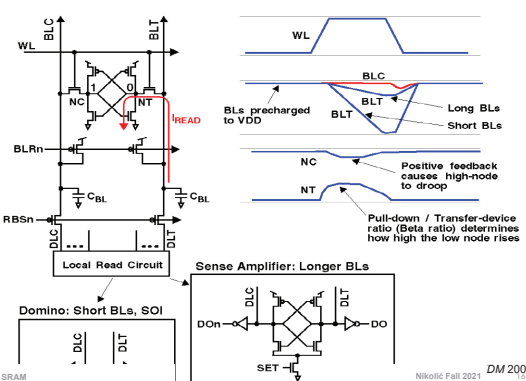
Nikolić Fall 2021



SRAM Column: Write



SRAM Column: Read



EECS151 L23 SRAM

Nikolić Fall 2021





EECS151 L23 SRAM

Nikolić Fall 2021



SRAM Array Layout



Administrivia


- Homework 10 posted on Friday, due 11/22
- No homework during Thanksgiving
- Project checkpoints #3 this week

Multi-Ported Memory


Motivation:

- Consider CPU core register file:
 - 1 read or write per cycle limits processor performance.
 - Complicates pipelining. Difficult for different instructions to simultaneously read or write regfile.
 - Single-issue pipelined CPUs usually needs 2 read ports and 1 write port (2R/1W).
 - Superscalar processors have more (e.g. 6R/3W)

I/O data buffering:




- dual-porting allows both sides to simultaneously access memory at full bandwidth.




Dual-Ported Memory Internals

- Add decoder, another set of read/write logic, bits lines, word lines:




- Example cell: SRAM



- Repeat everything but cross-coupled inverters.
- This scheme extends up to a couple more ports, then need to add additional transistors.

1R/1W 8T SRAM



- Dual-port read/write capability
- Single-cycle read and write, timed appropriately
- Often found in register files, first level (L1) of cache

True or False

www.yellkey.com/pull

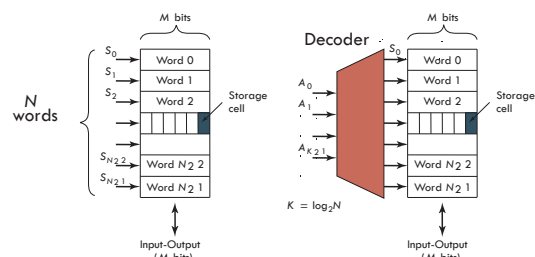
1	F	F	F
2	F	F	T
3	F	T	F
4	F	T	T
5	T	F	F
6	T	F	T
7	T	T	F
8	T	T	T

1. Transistor leakage doesn't affect SRAM read speed
2. One should write into an SRAM cell by pulling BL high
3. One can only write into some cells of a selected WL

Memory Decoders



Decoders



Intuitive architecture for $N \times M$ memory
Too many select signals:
 N words = N select signals

Decoder reduces the number of select signals
 $K = \log_2 N$

Row Decoders

Collection of 2^M complex logic gates
Organized in regular and dense fashion

(N)AND Decoder


$$WL_0 = A_0 A_1 A_2 A_3 A_4 A_5 A_6 A_7 A_8 A_9$$

$$WL_{511} = A_0 A_1 A_2 A_3 A_4 A_5 A_6 A_7 A_8 A_9$$

NOR Decoder


$$WL_0 = A_0 + A_1 + A_2 + A_3 + A_4 + A_5 + A_6 + A_7 + A_8 + A_9$$

$$WL_{511} = A_0 + A_1 + A_2 + A_3 + A_4 + A_5 + A_6 + A_7 + A_8 + A_9$$




Decoder Design Example

- Look at decoder for 256x256 memory block (8KBytes)




Possible Decoder

- 256 8-input AND gates
 - Each built out of a tree of NAND gates and inverters
- Need to drive a lot of capacitance (SRAM cells)
 - What's the best way to do this?




Possible AND8

- Build 8-input NAND gate using 2-input gates and inverters
- Is this the best we can do?
- Is this better than using fewer NAND4 gates?



Problem Setup

- Goal: Build fastest possible decoder with static CMOS logic
- What we know
 - Basically need 256 AND gates, each one of them drives one word line



Problem Setup (1)

- Each wordline has 256 cells connected to it
- $C_{WL} = 256 \cdot C_{cell} + C_{wire}$
 - Ignore wire for now
- Assume that decoder input capacitance is $C_{address} = 4 \cdot C_{cell}$

Problem Setup (2)


- Each address bit drives $2^8/2$ AND gates
 - A0 drives $1/2$ of the gates, A0_b the other $1/2$ of the gates
- Total fanout on each address wire is:

$$F = \Pi B \frac{C_{load}}{C_{in}} = 128 \frac{(256 C_{cell})}{4 C_{cell}} = 2^7 \frac{(2^8 C_{cell})}{2^2 C_{cell}} = 2^{13}$$

Decoder Fan-Out

- F of 2^{13} means that we will want to use more than $\log_4(2^{13}) = 6.5$ stages to implement the AND8
- Need many stages anyway
 - So what is the best way to implement the AND gate?
 - Will see next that it's the one with the most stages and least complicated gates


8-Input AND



$$\begin{array}{lcl} g : 9/2 & 1 & \\ G = 9/2 & & \\ P = 8 & + & 1 \end{array} \quad \begin{array}{lcl} g : 5/2 & 3/2 & \\ G = 15/4 & & \\ P = 4 & + & 2 \end{array} \quad \begin{array}{lcl} g : 3/2 & 3/2 & 3/2 & 1 \\ G = 27/8 & & & \\ P = 2 & + & 2 & + & 2 & + & 1 \end{array}$$


8-Input AND

- Using 2-input NAND gates
 - 8-input gate takes 6 stages
- Total G is $(3/2)^3 \approx 3.4$
- So H is $3.4 * 2^{13}$ – optimal N of ~ 7.4




Decoder So Far

- 256 8-input AND gates
 - Each built out of tree of NAND gates and inverters
- Issue:
 - Every address line has to drive 128 gates (and wire) right away
 - Forces us to add buffers just to drive address inputs




Look Inside Each AND8 Gate



Predecoders


- Use a single gate for each of the shared terms
 - E.g., from $A_0, \bar{A}_0, A_1, \bar{A}_1$, generate four signals: $A_0A_1, \bar{A}_0A_1, A_0\bar{A}_1, \bar{A}_0\bar{A}_1$
- In other words, we are decoding smaller groups of address bits first
 - And using the “predecoded” outputs to do the rest of the decoding

Predecoder and Decoder



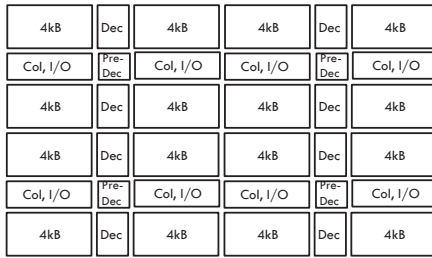
Predecode Options

- Larger predecode usually better:
 - More stages before the long wires
 - Decreases their effect on the circuit
 - Fewer number of long wires switches
 - Lower power
- Easier to fit 2-input gate into cell pitch



Building Larger Arrays

Building Larger Custom Arrays




- Each subarray is 2-8kB
- Hierarchical decoding
- Peripheral overhead is 30-50%
- Delay is wire dominated
- Scratchpads, caches, TLBs

Cascading Memory-Blocks

How to make larger memory blocks out of smaller ones.


Increasing the width. Example: given 1Kx8, want 1Kx16



Cascading Memory-Blocks

How to make larger memory blocks out of smaller ones.


Increasing the depth. Example: given 1Kx8, want 2Kx8



Adding Ports to Primitive Memory Blocks

Adding a read port to a simple dual port (SDP) memory.

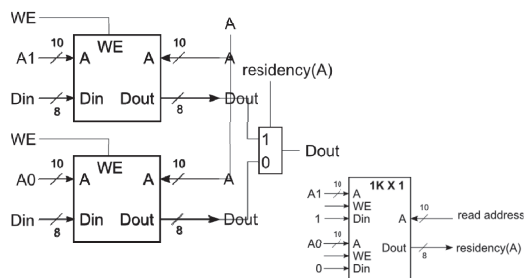
Example: given 1Kx8 SDP, want 1 write & 2 read ports.



Adding Ports to Primitive Memory Blocks

How to add a write port to a simple dual-port memory.

Example: given 1Kx8 SDP, want 1 read & 2 write ports.



Review

- Dense memories are built as arrays of memory elements
 - SRAM is a static memory
- SRAM has unique combination of density, speed, power
- SRAM cells sized for stability and writeability
- SRAM and regfile cells can have multiple R/W ports
- Memory decoding is done hierarchically
 - Wire-limited in large arrays